

Better Code: Concurrency

Sean Parent | Principal Scientist



Better Code

- Regular Type
 - Goal: Implement Complete and Efficient Types
- Algorithms
 - Goal: No Raw Loops
- Data Structures
 - Goal: No Incidental Data Structures
- Runtime Polymorphism
 - Goal: No Inheritance
- Concurrency
 - Goal: No Raw Synchronization Primitives
- ...

Better Code

- Regular Type
 - Goal: Implement Complete and Efficient Types
- Algorithms
 - Goal: No Raw Loops
- Data Structures
 - Goal: No Incidental Data Structures
- Runtime Polymorphism
 - Goal: No Inheritance
- Concurrency
 - Goal: No Raw Synchronization Primitives
- ...



Common Themes

- Manage Relationships
- Understand the Fundamentals
- Code Simply

lightroom.adobe.com

CROP PRESETS **ADJUST**

HISTOGRAM

TREATMENT Color | Black & White

WHITE BALANCE

TONE

Exposure

Contrast

Highlights

Shadows

Whites

Blacks

PRESENCE

Clarity

Vibrance

Saturation

COLOR / B&W

SPLIT TONING

lightroom.adobe.com

✓ Save & Exit Revert to Original Undo Edits CROP PRESETS **ADJUST**

HISTOGRAM

TREATMENT Color | Black & White

WHITE BALANCE

TONE Auto

Exposure

Contrast

Highlights

Shadows

Whites

Blacks

PRESENCE

Clarity

Vibrance

Saturation

COLOR / B&W

SPLIT TONING

Concurrency

- Concurrency: when tasks start, run, and complete in overlapping time periods
- Parallelism: when two or more tasks execute simultaneously
- Why?
 - Enable performance through parallelism
 - Improve interactivity by handling user actions concurrent with processing and IO

Goal: No Raw Synchronization Primitives

What are raw synchronization primitives?

- Synchronization primitives are basic constructs such as:
 - Mutex
 - Atomic
 - Semaphore
 - Memory Fence
 - Condition Variable

Why No Raw Synchronization Primitives?

You Will Likely Get It Wrong

Problems with Locks

```
template <typename T>
class bad_cow {
    struct object_t {
        explicit object_t(const T& x) : data_m(x) { ++count_m; }
        atomic<int> count_m;
        T data_m; };
    object_t* object_m;
public:
    explicit bad_cow(const T& x) : object_m(new object_t(x)) { }
    ~bad_cow() { if (0 == --object_m->count_m) delete object_m; }
    bad_cow(const bad_cow& x) : object_m(x.object_m) { ++object_m->count_m; }

    bad_cow& operator=(const T& x) {
        if (object_m->count_m == 1) object_m->data_m = x;
        else {
            object_t* tmp = new object_t(x);
            --object_m->count_m;
            object_m = tmp;
        }
        return *this;
    }
};
```

Problems with Locks

```
template <typename T>
class bad_cow {
    struct object_t {
        explicit object_t(const T& x) : data_m(x) { ++count_m; }
        atomic<int> count_m;
        T data_m; };
    object_t* object_m;
public:
    explicit bad_cow(const T& x) : object_m(new object_t(x)) { }
    ~bad_cow() { if (0 == --object_m->count_m) delete object_m; }
    bad_cow(const bad_cow& x) : object_m(x.object_m) { ++object_m->count_m; }

    bad_cow& operator=(const T& x) {
        if (object_m->count_m == 1) object_m->data_m = x;
        else {
            object_t* tmp = new object_t(x);
            --object_m->count_m;
            object_m = tmp;
        }
        return *this;
    }
};
```

Problems with Locks

```
template <typename T>
class bad_cow {
    struct object_t {
        explicit object_t(const T& x) : data_m(x) { ++count_m; }
        atomic<int> count_m;
        T data_m; };
    object_t* object_m;
public:
    explicit bad_cow(const T& x) : object_m(new object_t(x)) { }
    ~bad_cow() { if (0 == --object_m->count_m) delete object_m; }
    bad_cow(const bad_cow& x) : object_m(x.object_m) { ++object_m->count_m; }

    bad_cow& operator=(const T& x) {
        if (object_m->count_m == 1) object_m->data_m = x;
        else {
            object_t* tmp = new object_t(x);
            --object_m->count_m;
            object_m = tmp;
        }
        return *this;
    }
};
```

- There is a subtle race condition here:
 - if `count != 1` then the `bad_cow` could also be owned by another thread(s)
 - if the other thread(s) releases the `bad_cow` between these two atomic operations
 - then our `count` will fall to zero and we will leak the object

Problems with Locks

```
template <typename T>
class bad_cow {
    struct object_t {
        explicit object_t(const T& x) : data_m(x) { ++count_m; }
        atomic<int> count_m;
        T data_m; };
    object_t* object_m;
public:
    explicit bad_cow(const T& x) : object_m(new object_t(x)) { }
    ~bad_cow() { if (0 == --object_m->count_m) delete object_m; }
    bad_cow(const bad_cow& x) : object_m(x.object_m) { ++object_m->count_m; }

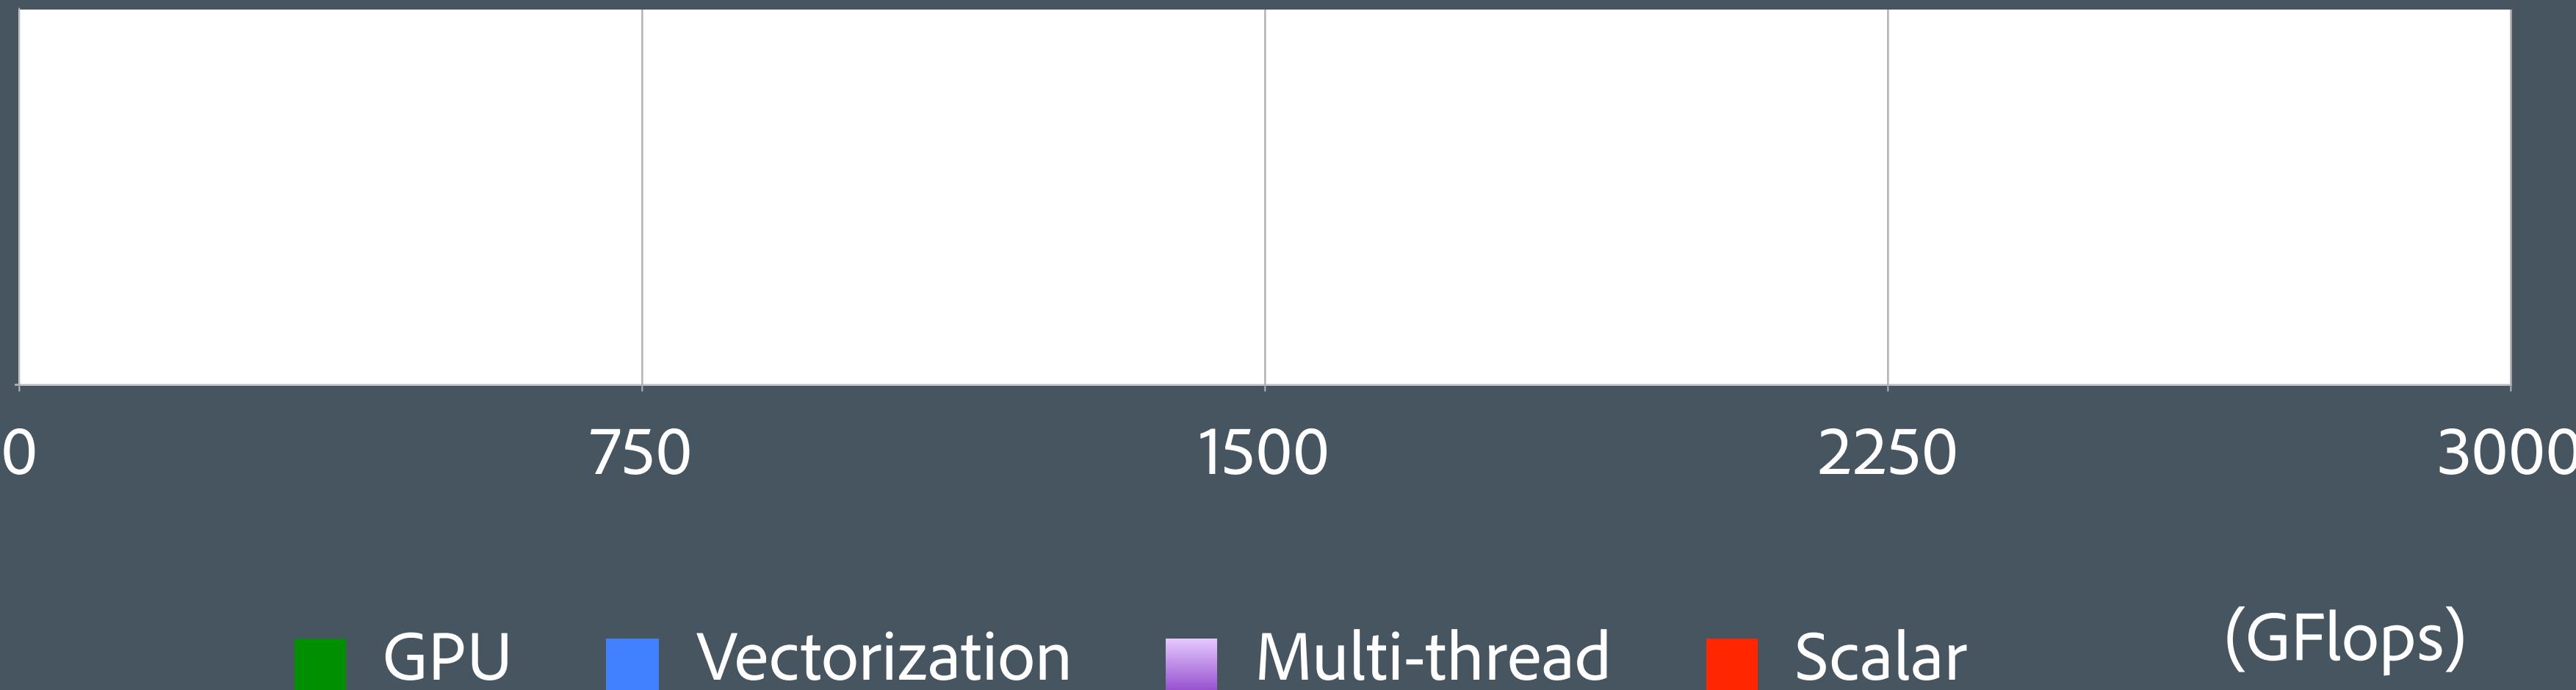
    bad_cow& operator=(const T& x) {
        if (object_m->count_m == 1) object_m->data_m = x;
        else {
            object_t* tmp = new object_t(x);
            if (0 == --object_m->count_m) delete object_m;
            object_m = tmp;
        }
        return *this;
    }
};
```

Why do we want concurrency?

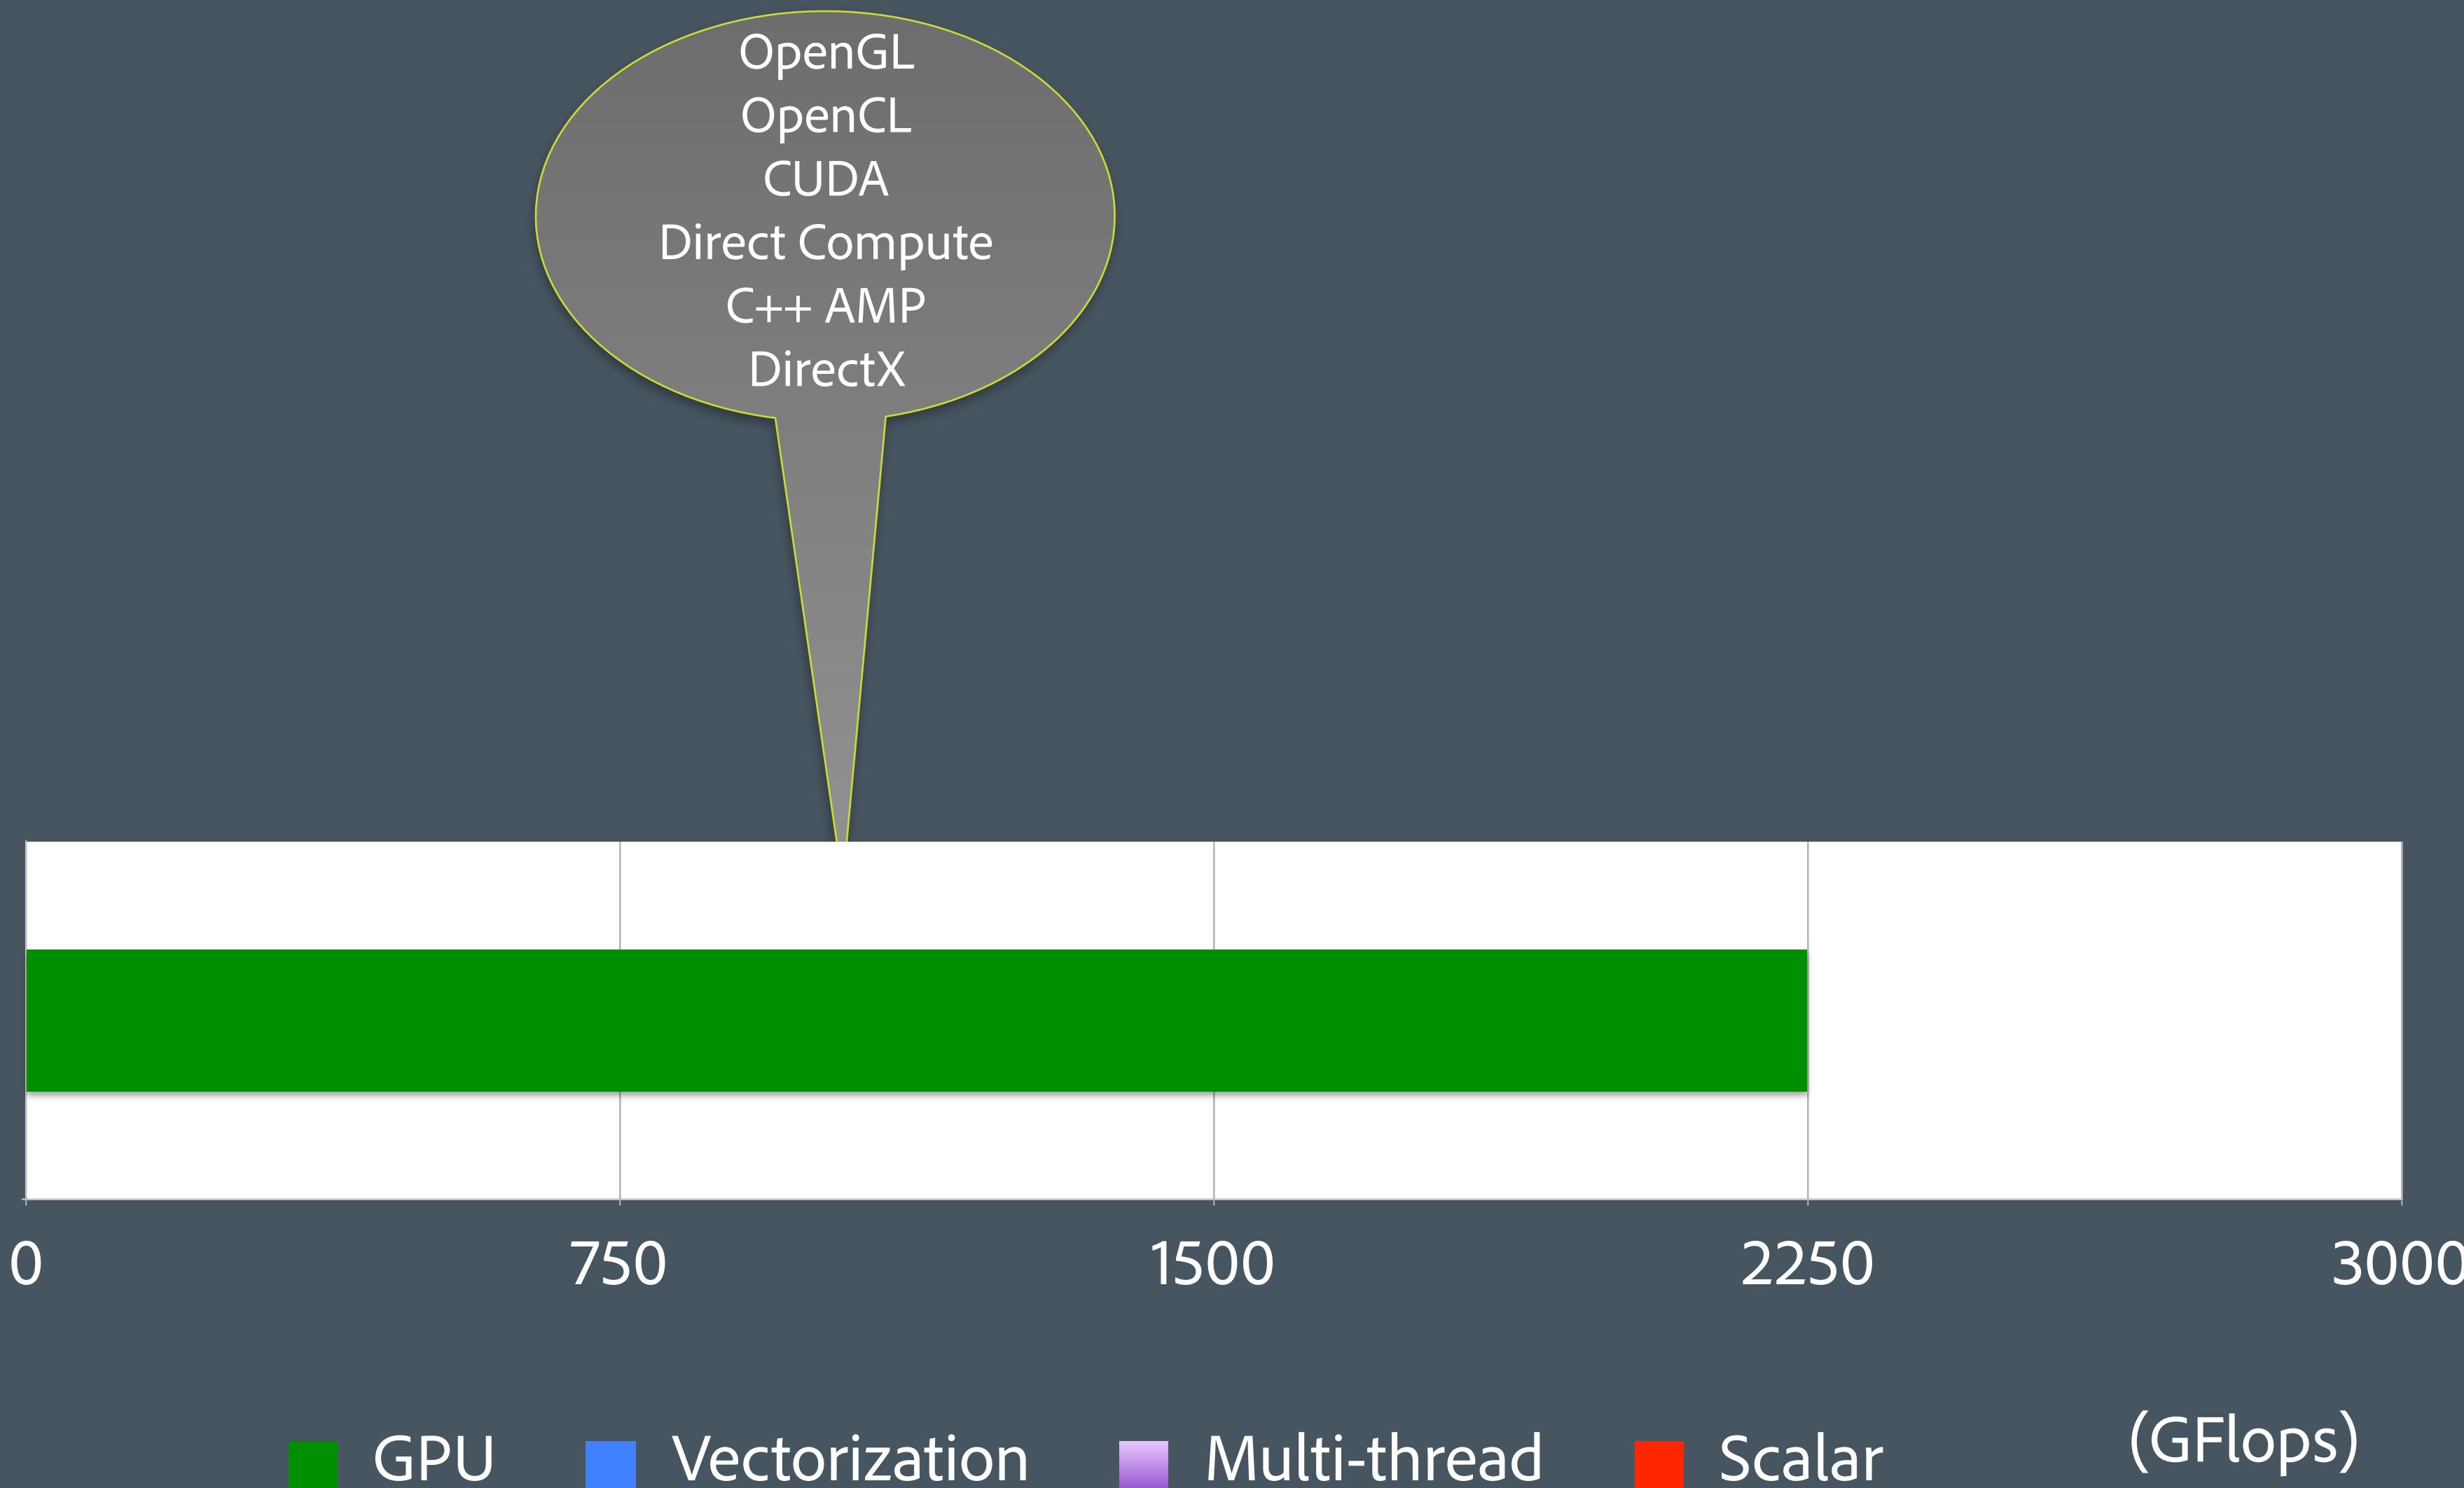
Performance through Parallelism

Desktop Compute Power (8-core 3.5GHz Sandy Bridge + AMD Radeon 6950)

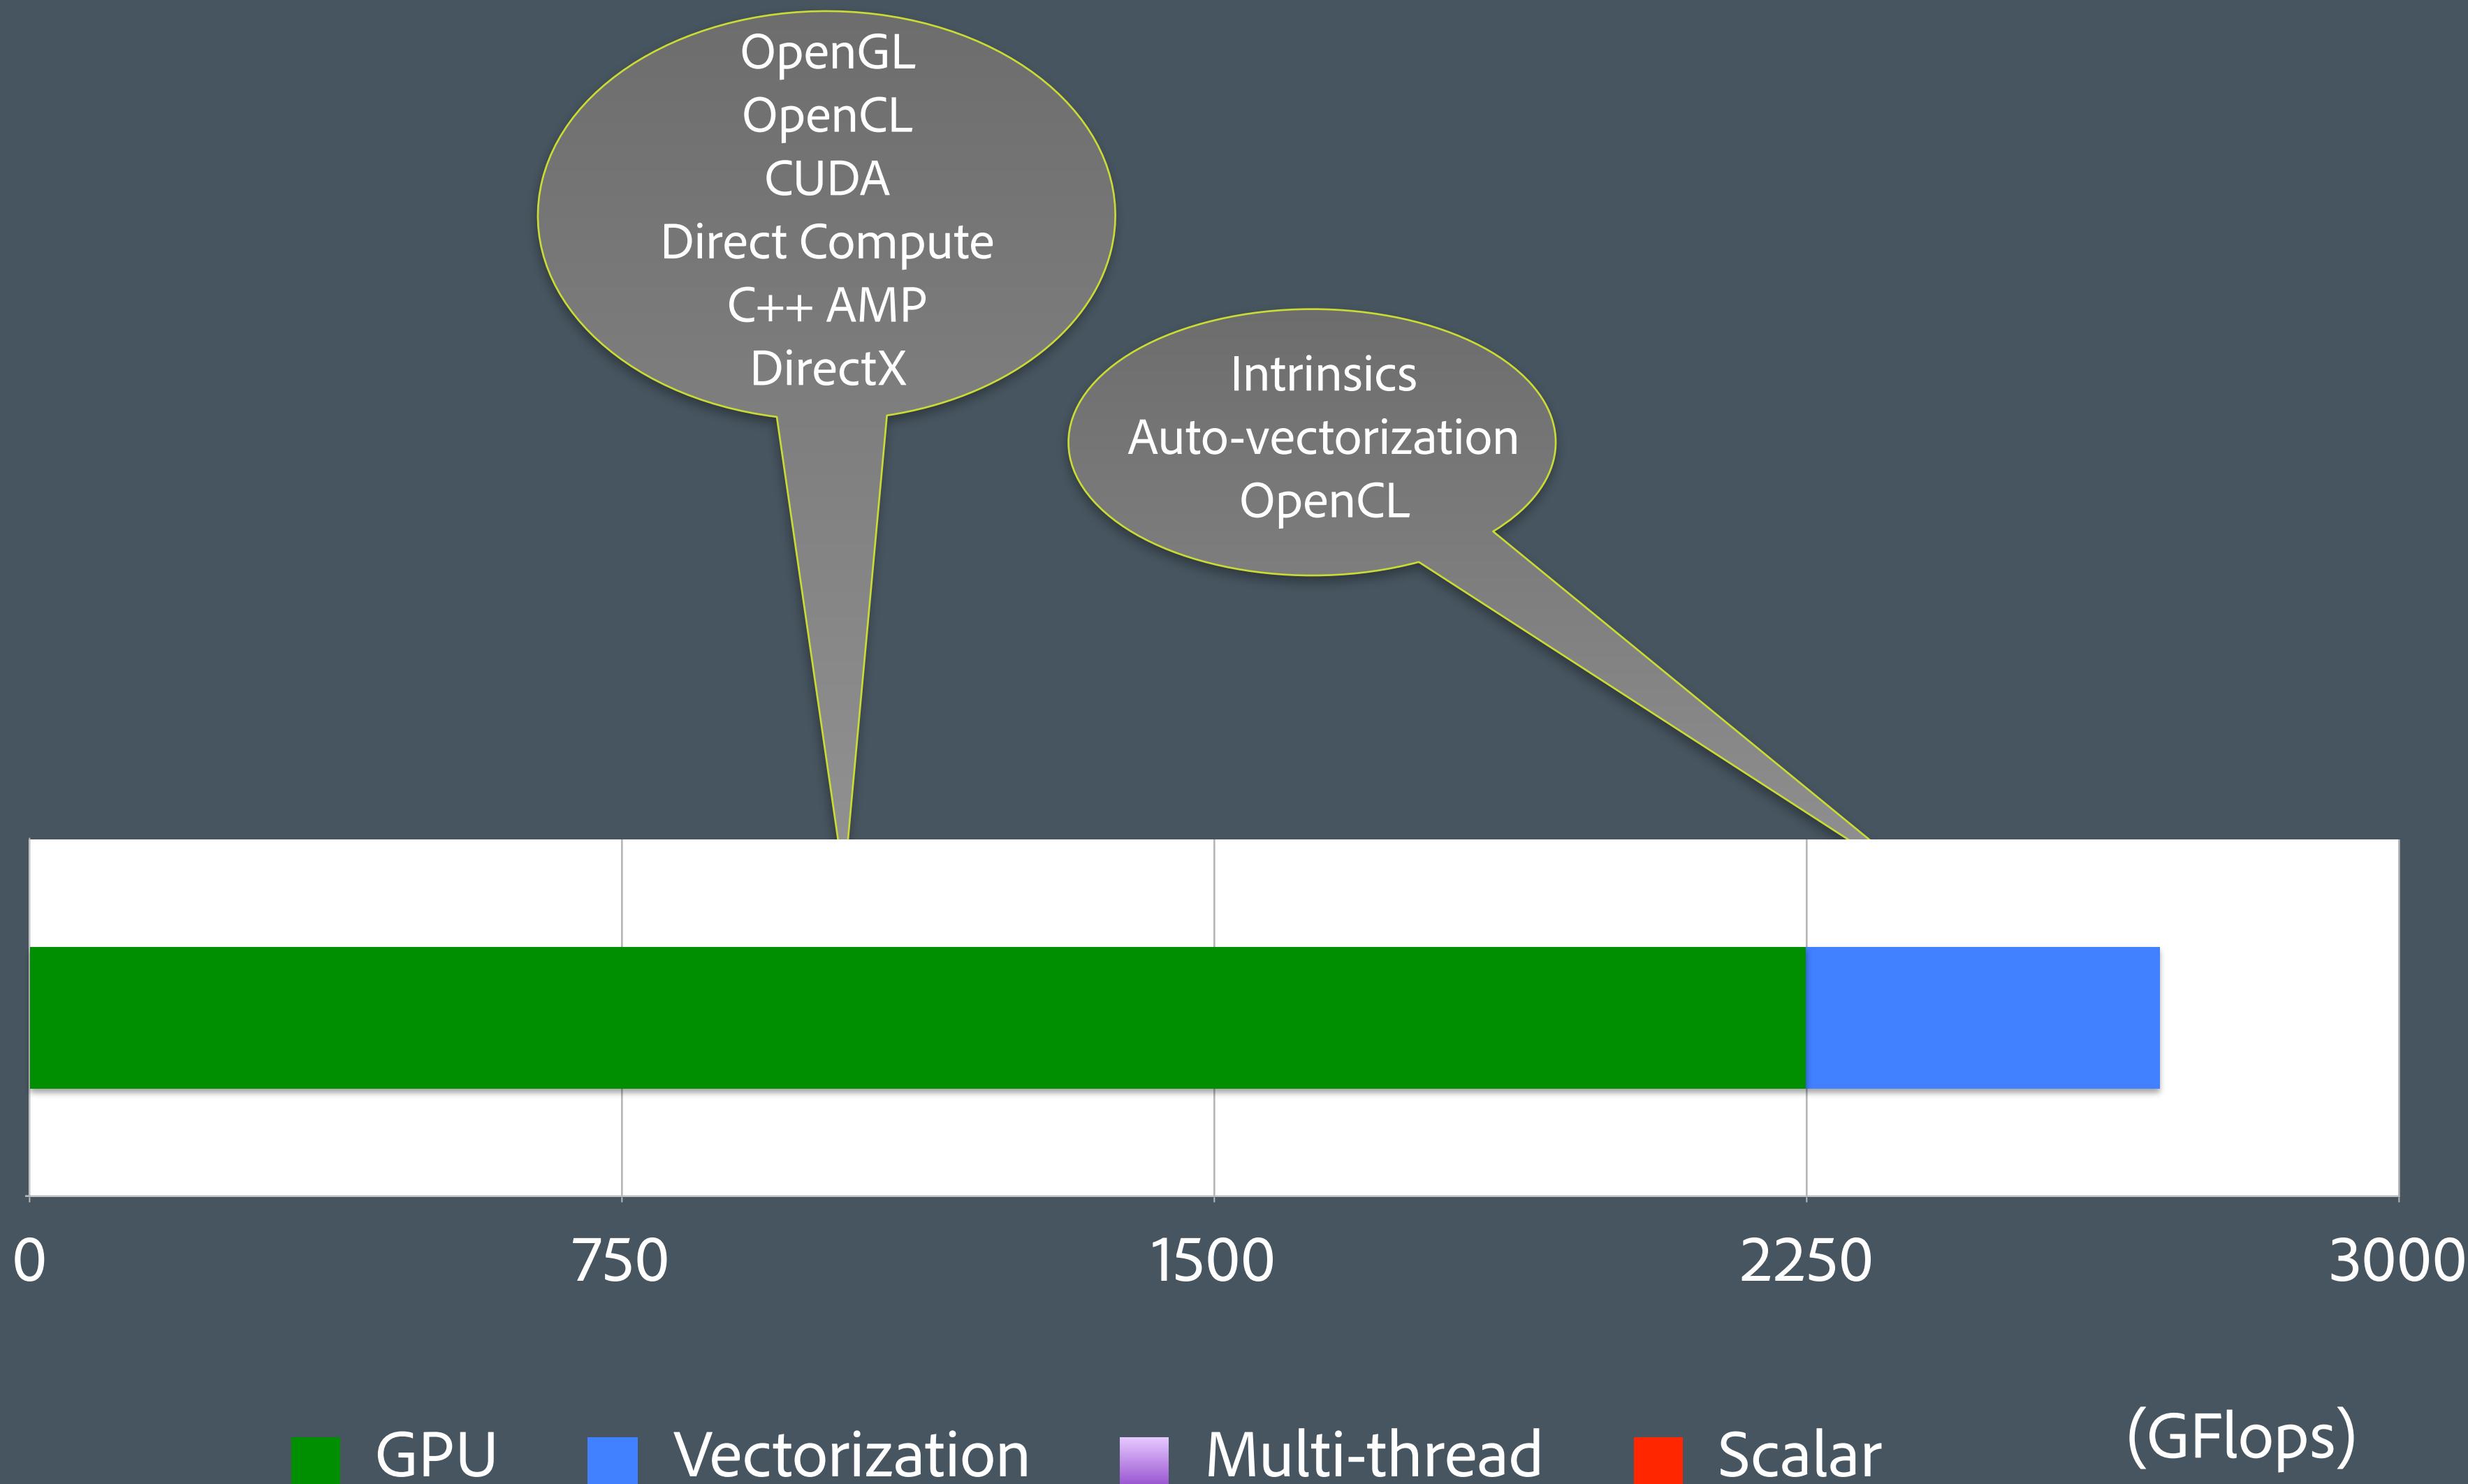
Desktop Compute Power (8-core 3.5GHz Sandy Bridge + AMD Radeon 6950)



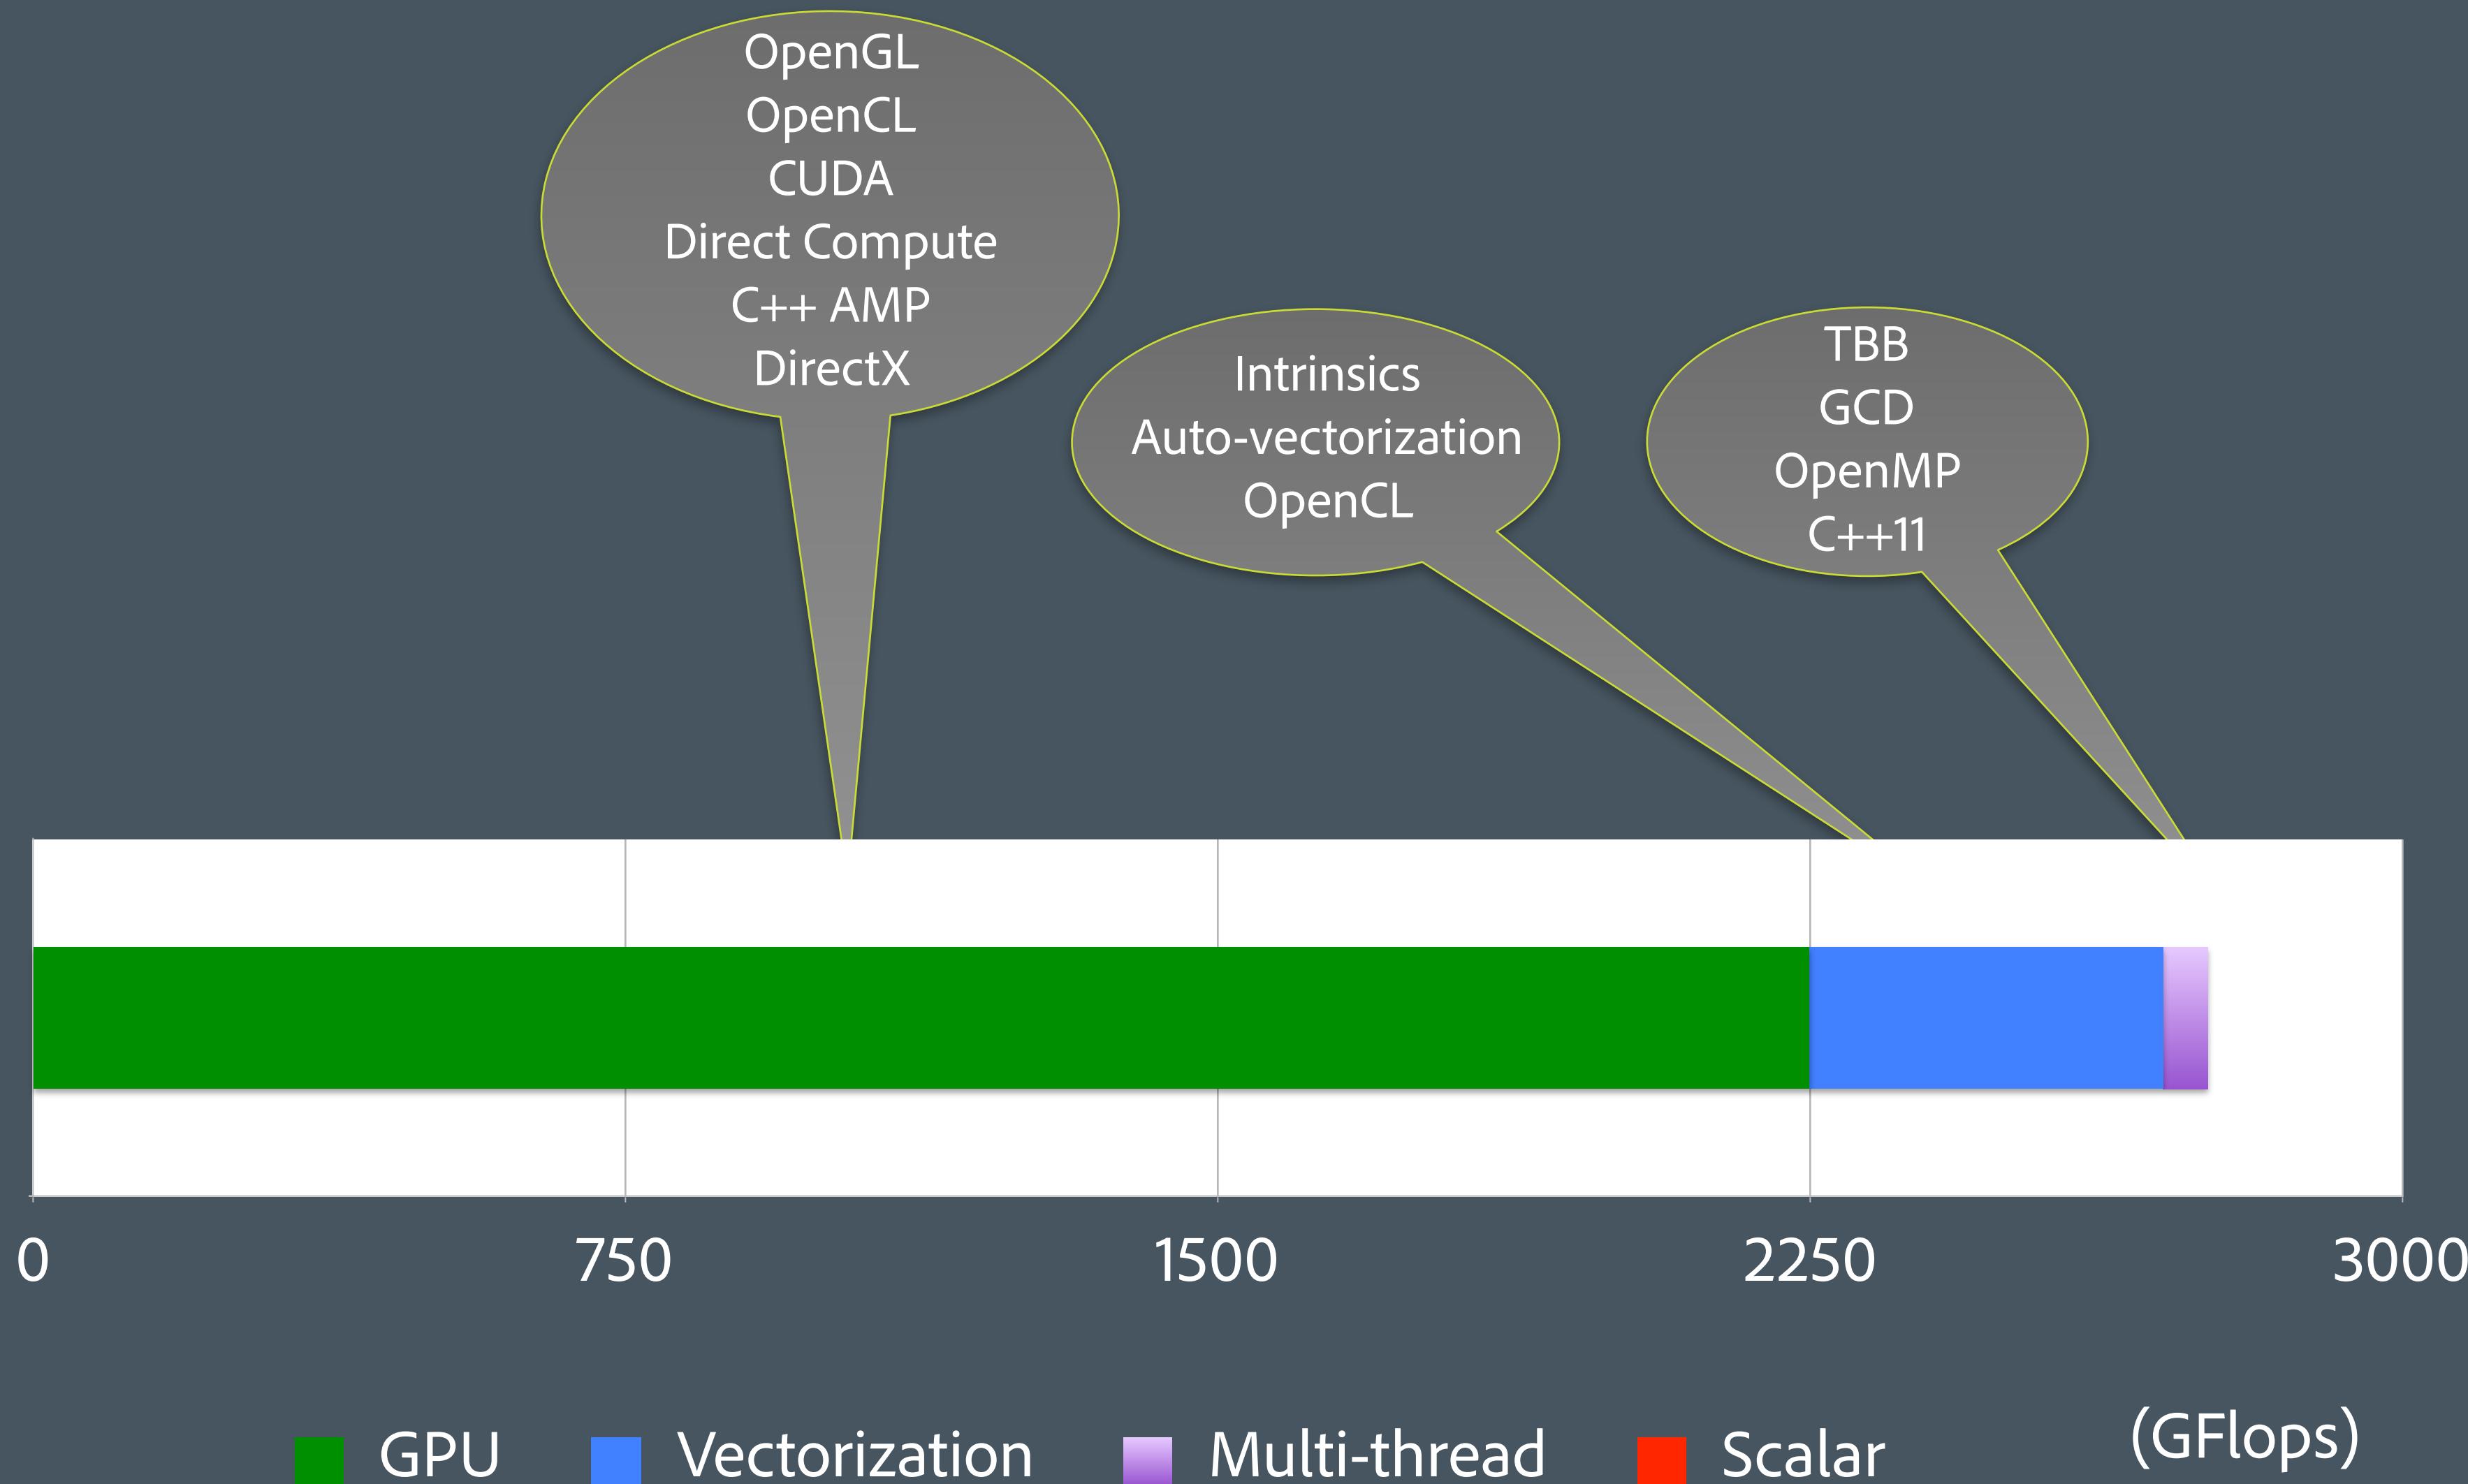
Desktop Compute Power (8-core 3.5GHz Sandy Bridge + AMD Radeon 6950)



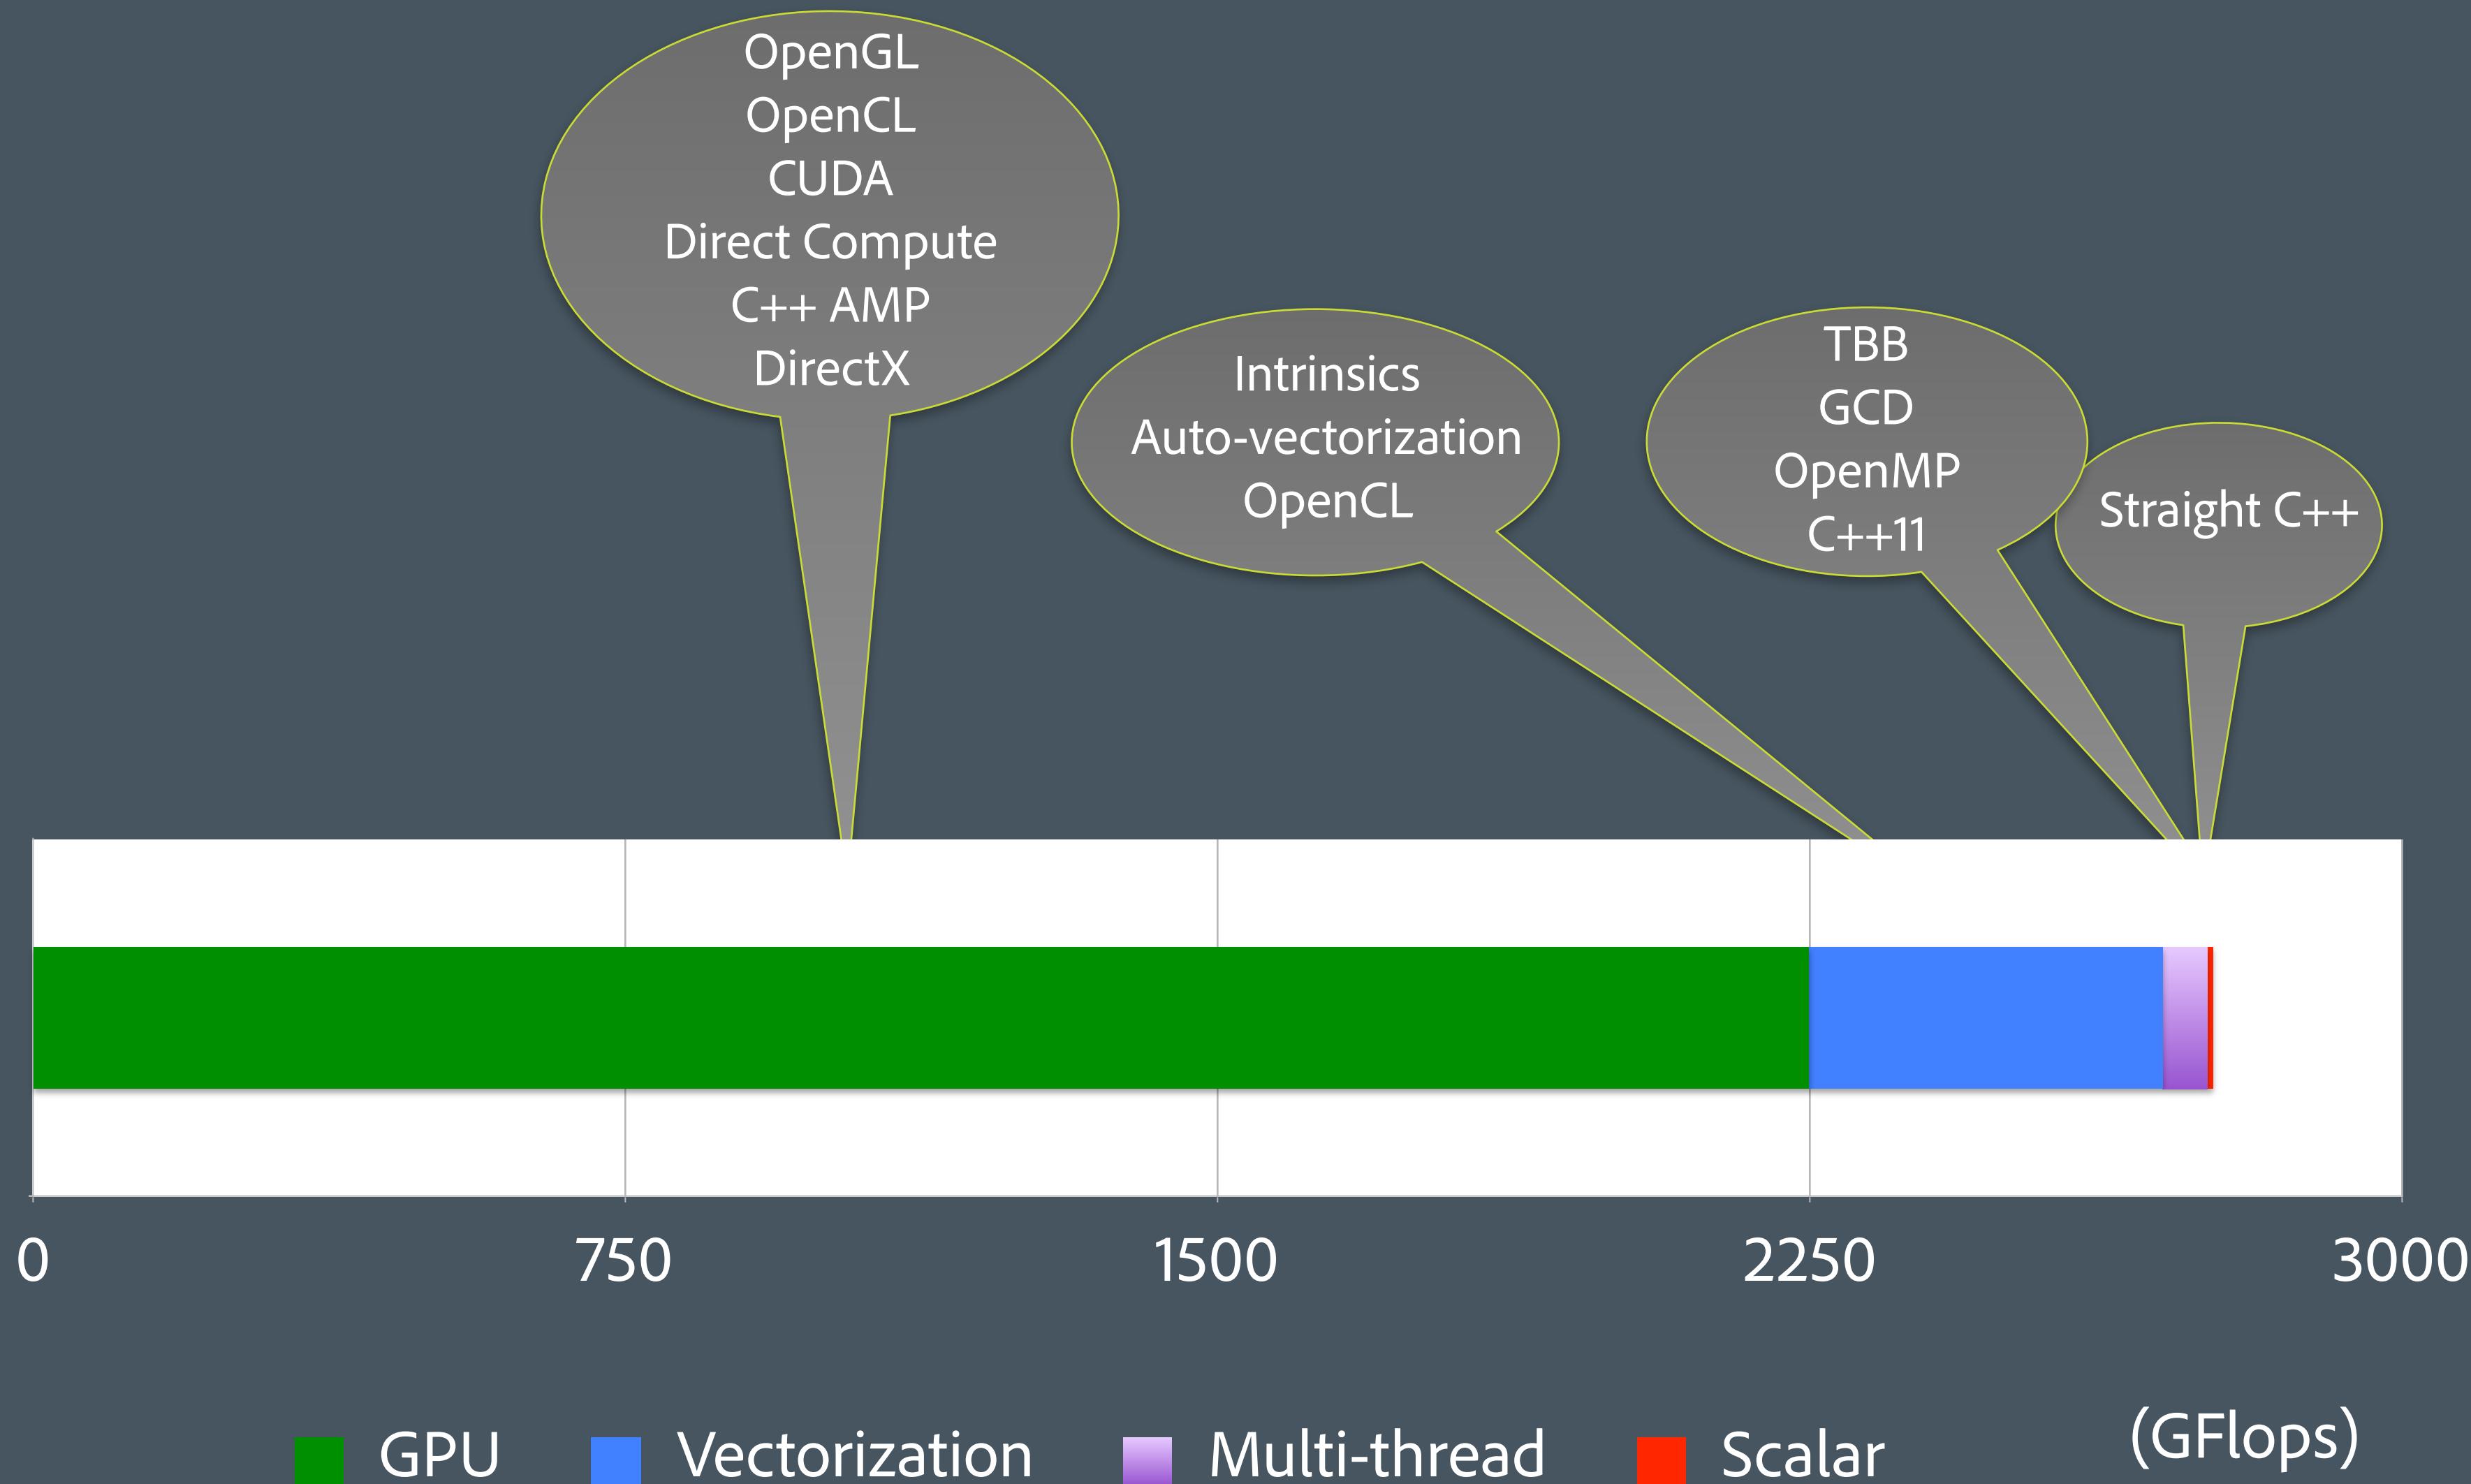
Desktop Compute Power (8-core 3.5GHz Sandy Bridge + AMD Radeon 6950)



Desktop Compute Power (8-core 3.5GHz Sandy Bridge + AMD Radeon 6950)

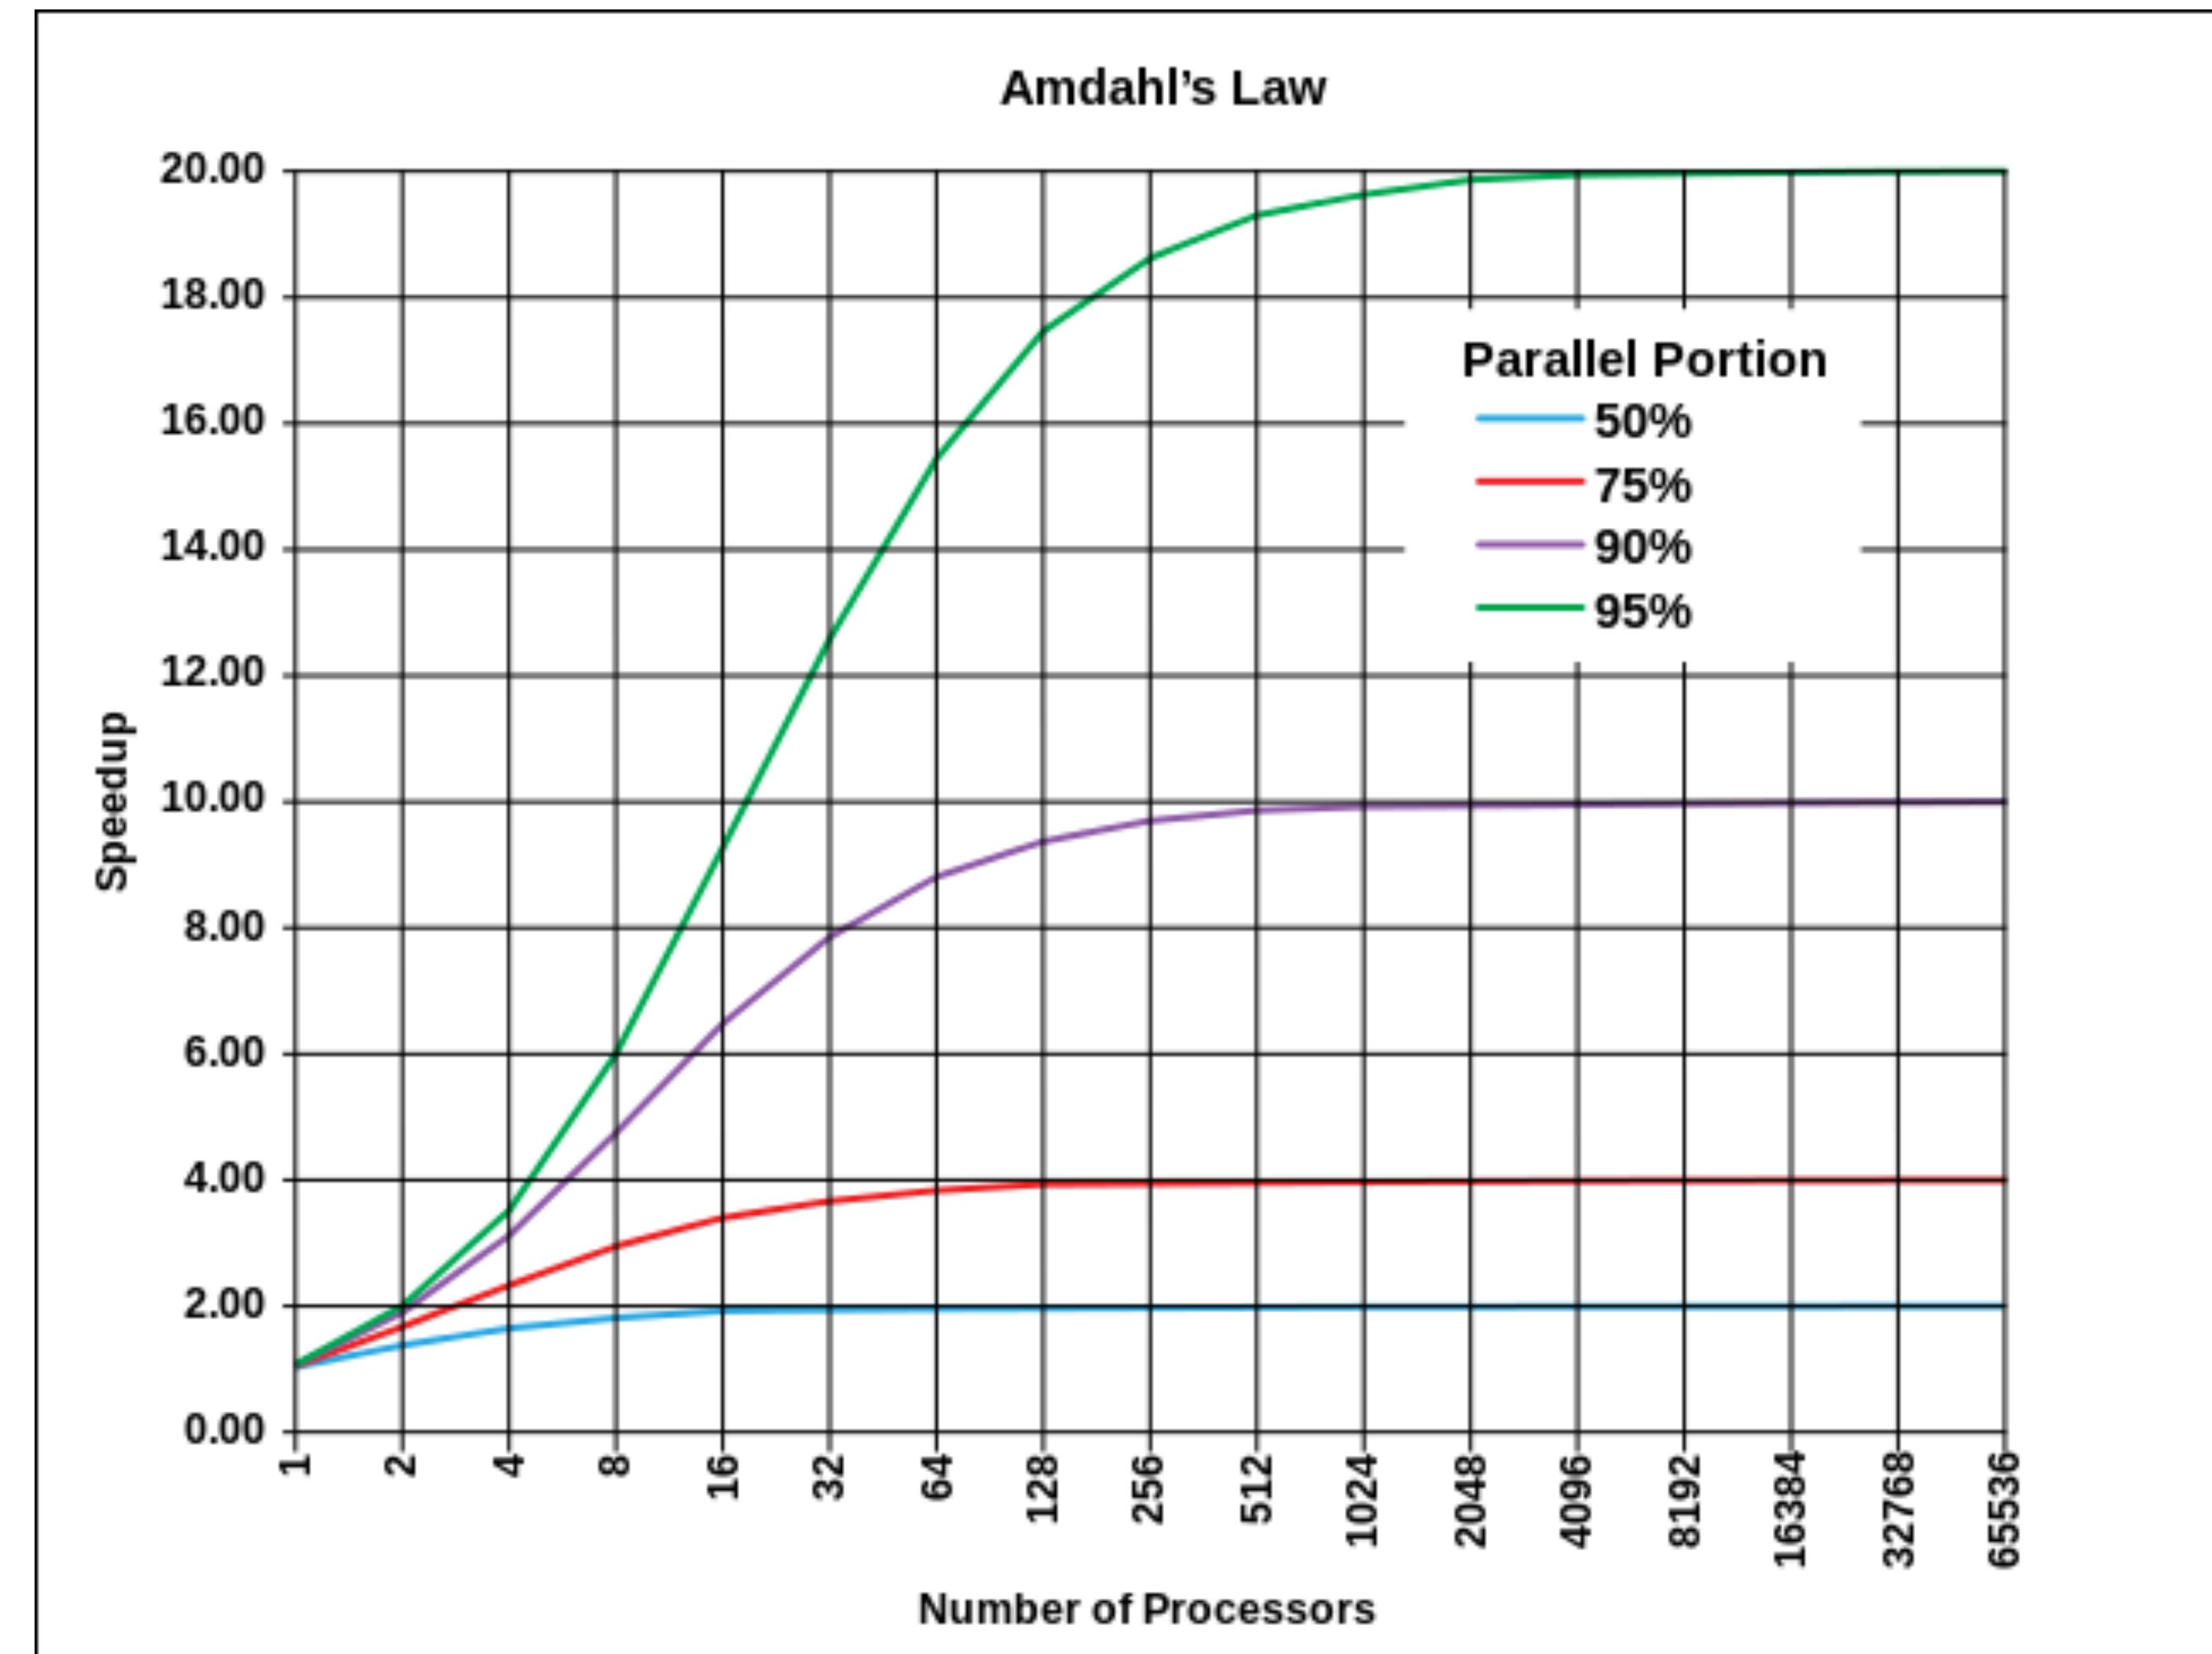


Desktop Compute Power (8-core 3.5GHz Sandy Bridge + AMD Radeon 6950)



Amdahl's Law

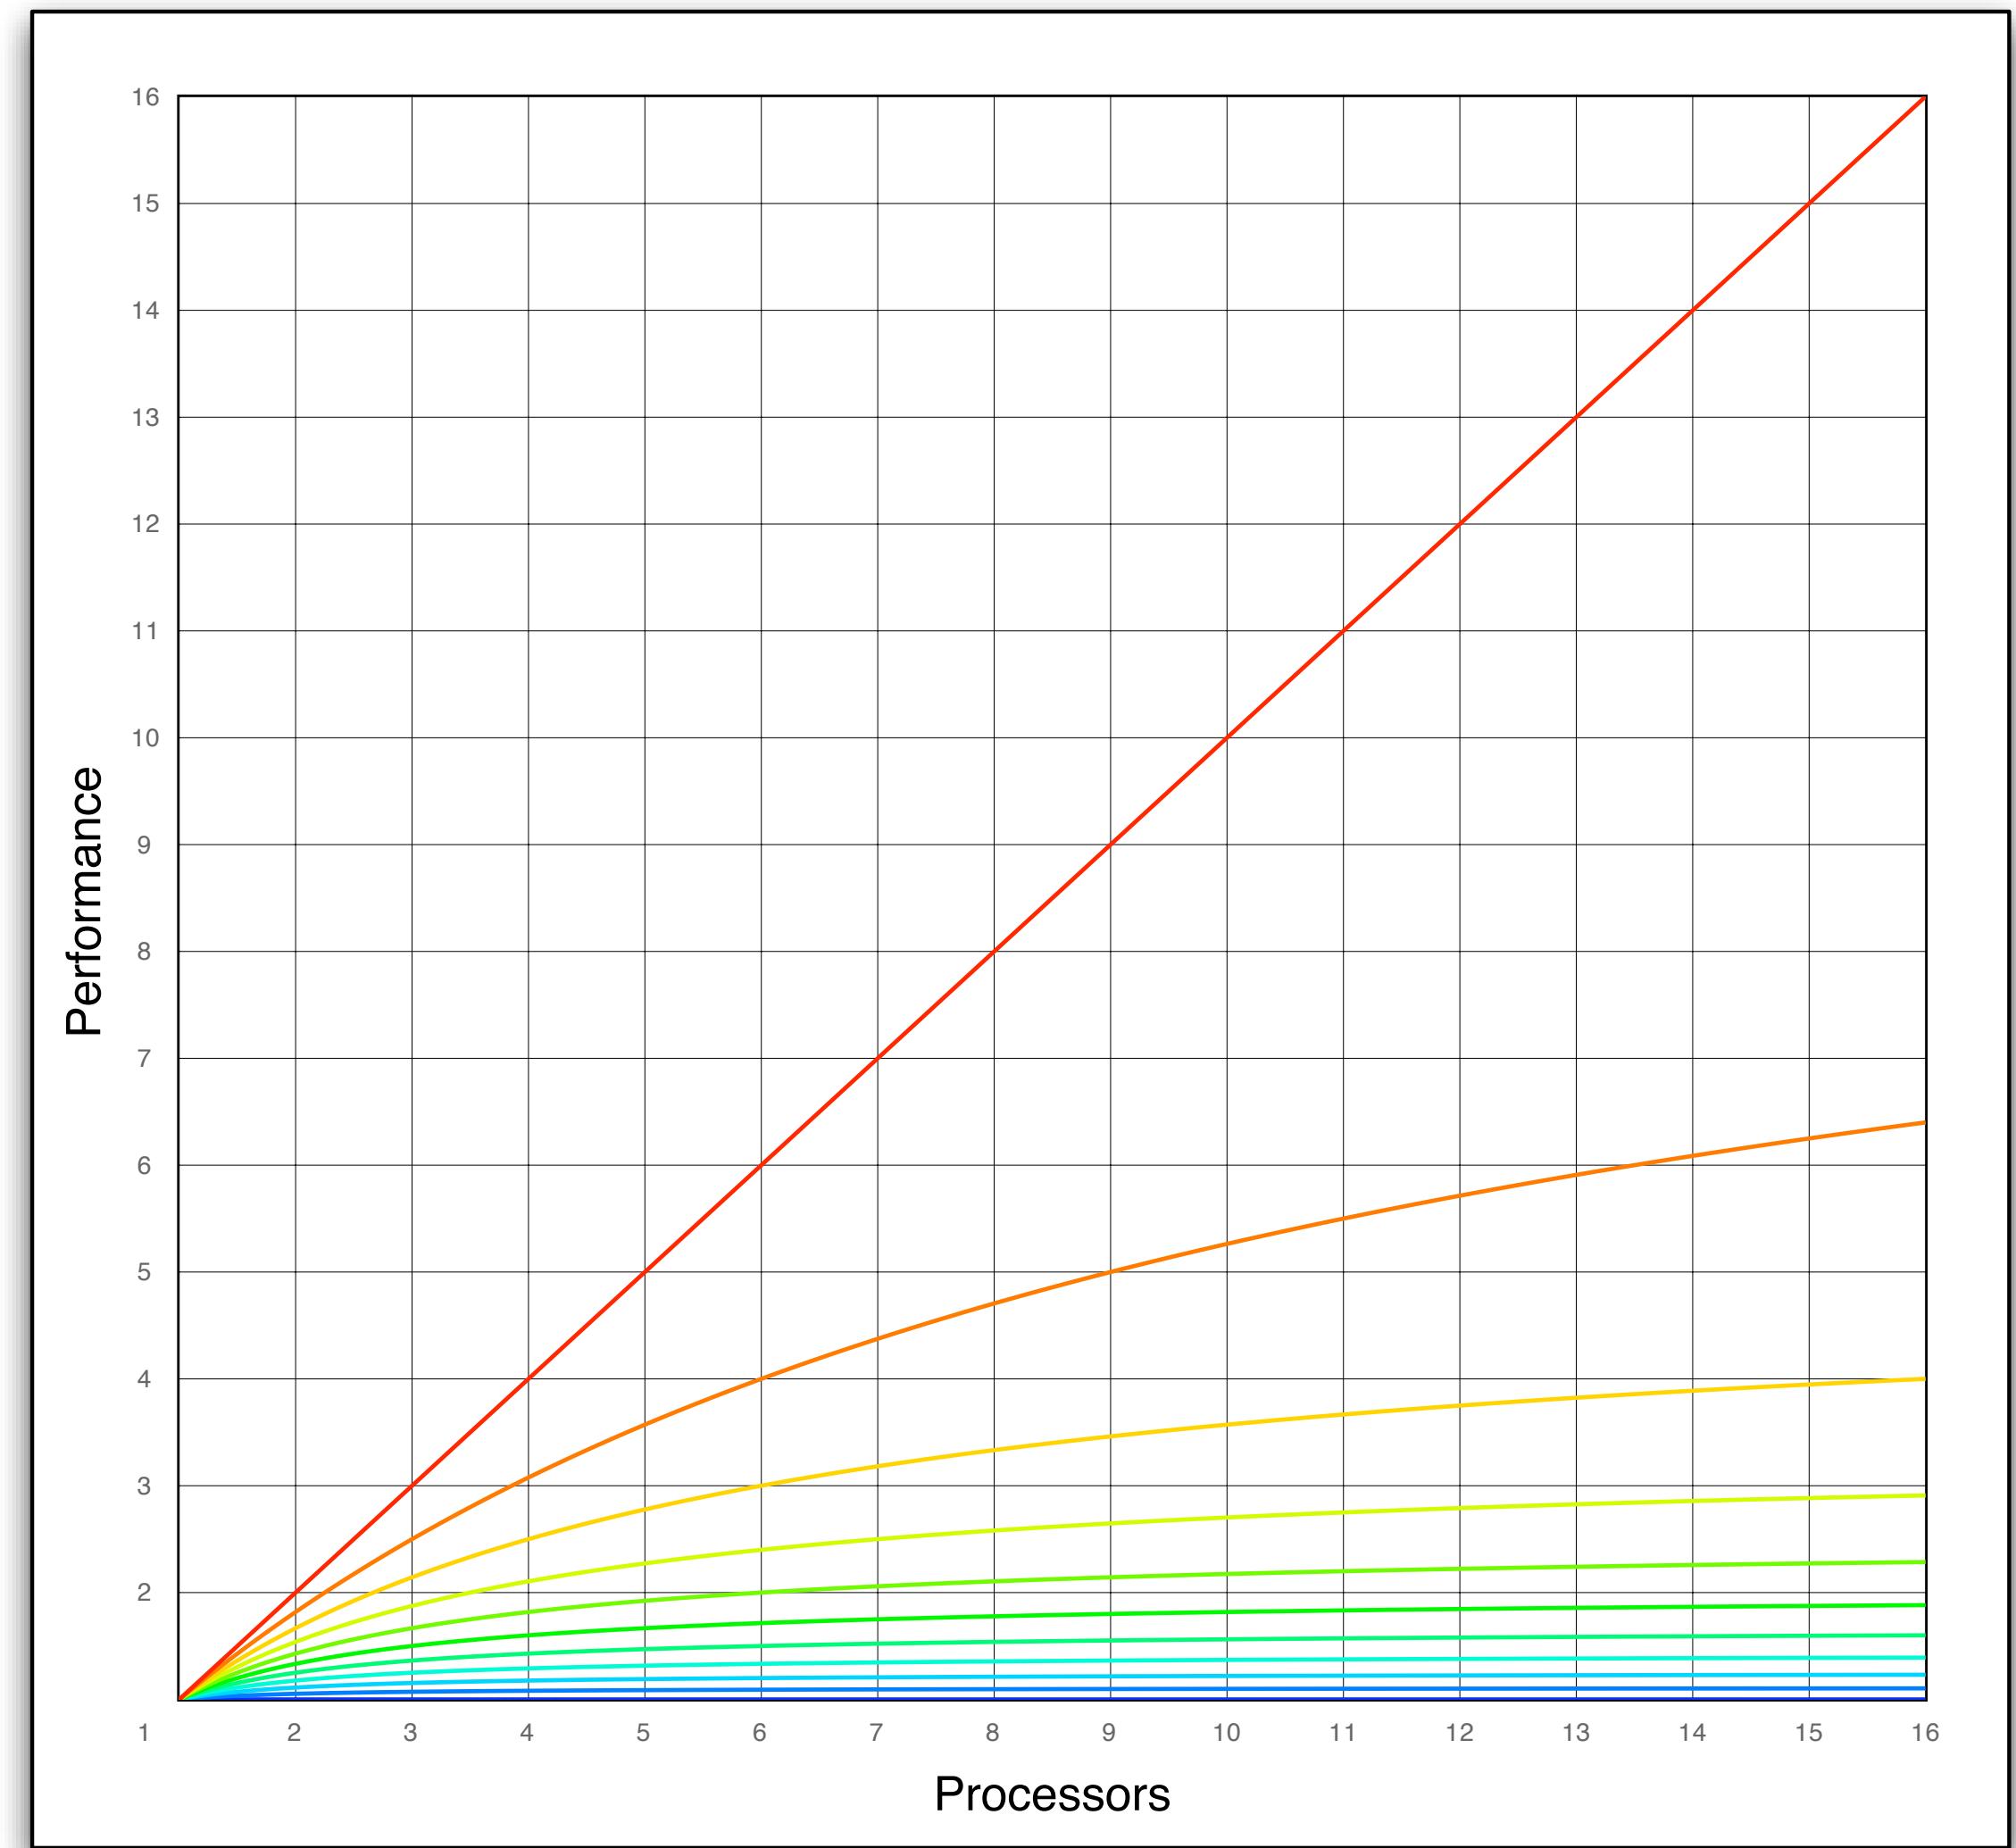
$$S(N) = \frac{1}{(1 - P) + \frac{P}{N}}$$



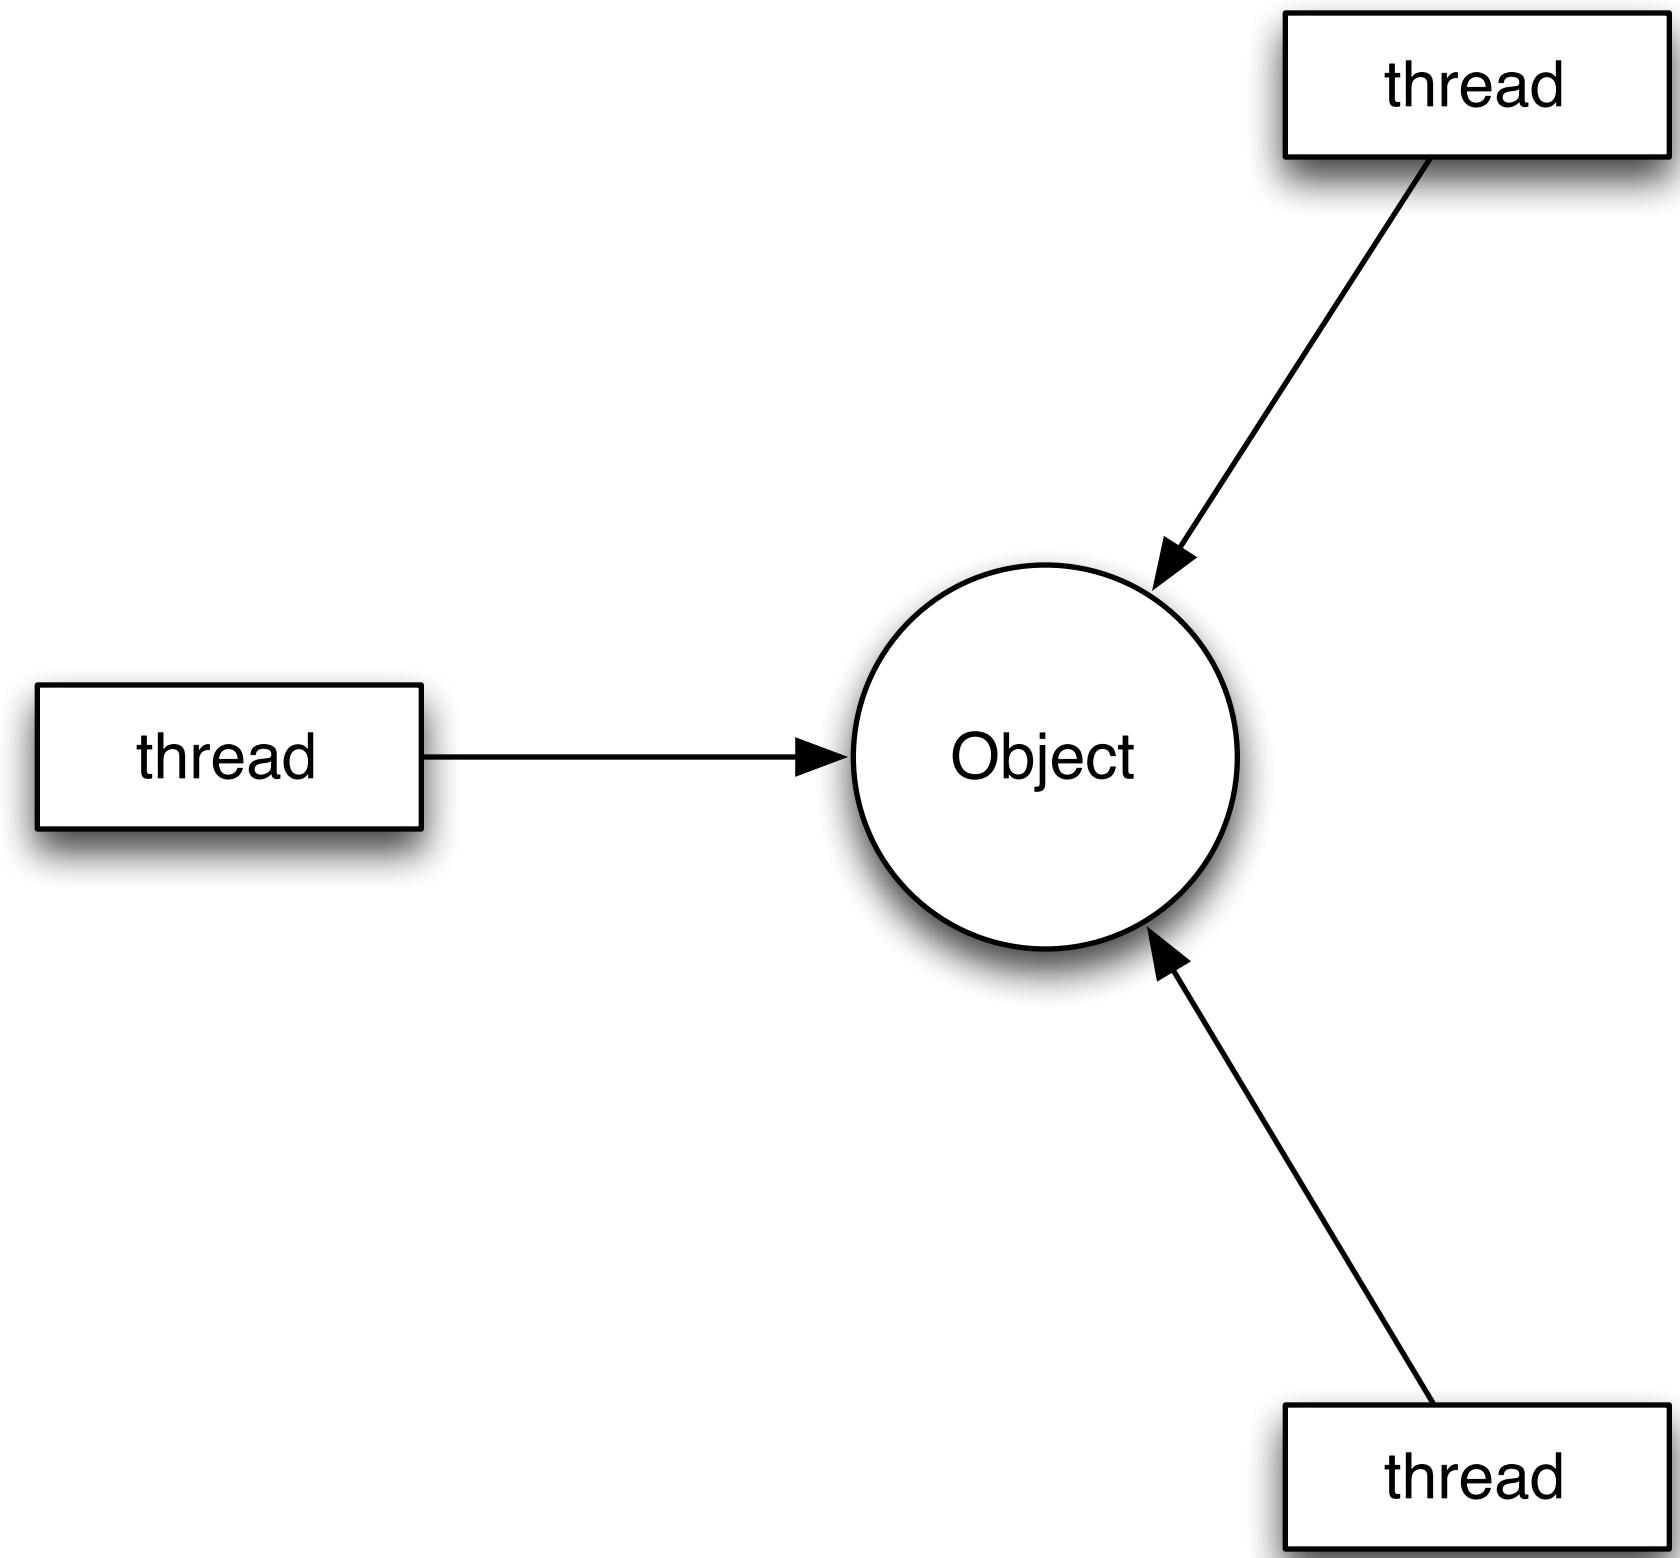
http://en.wikipedia.org/wiki/Amdahl%27s_law

Amdahl's Law

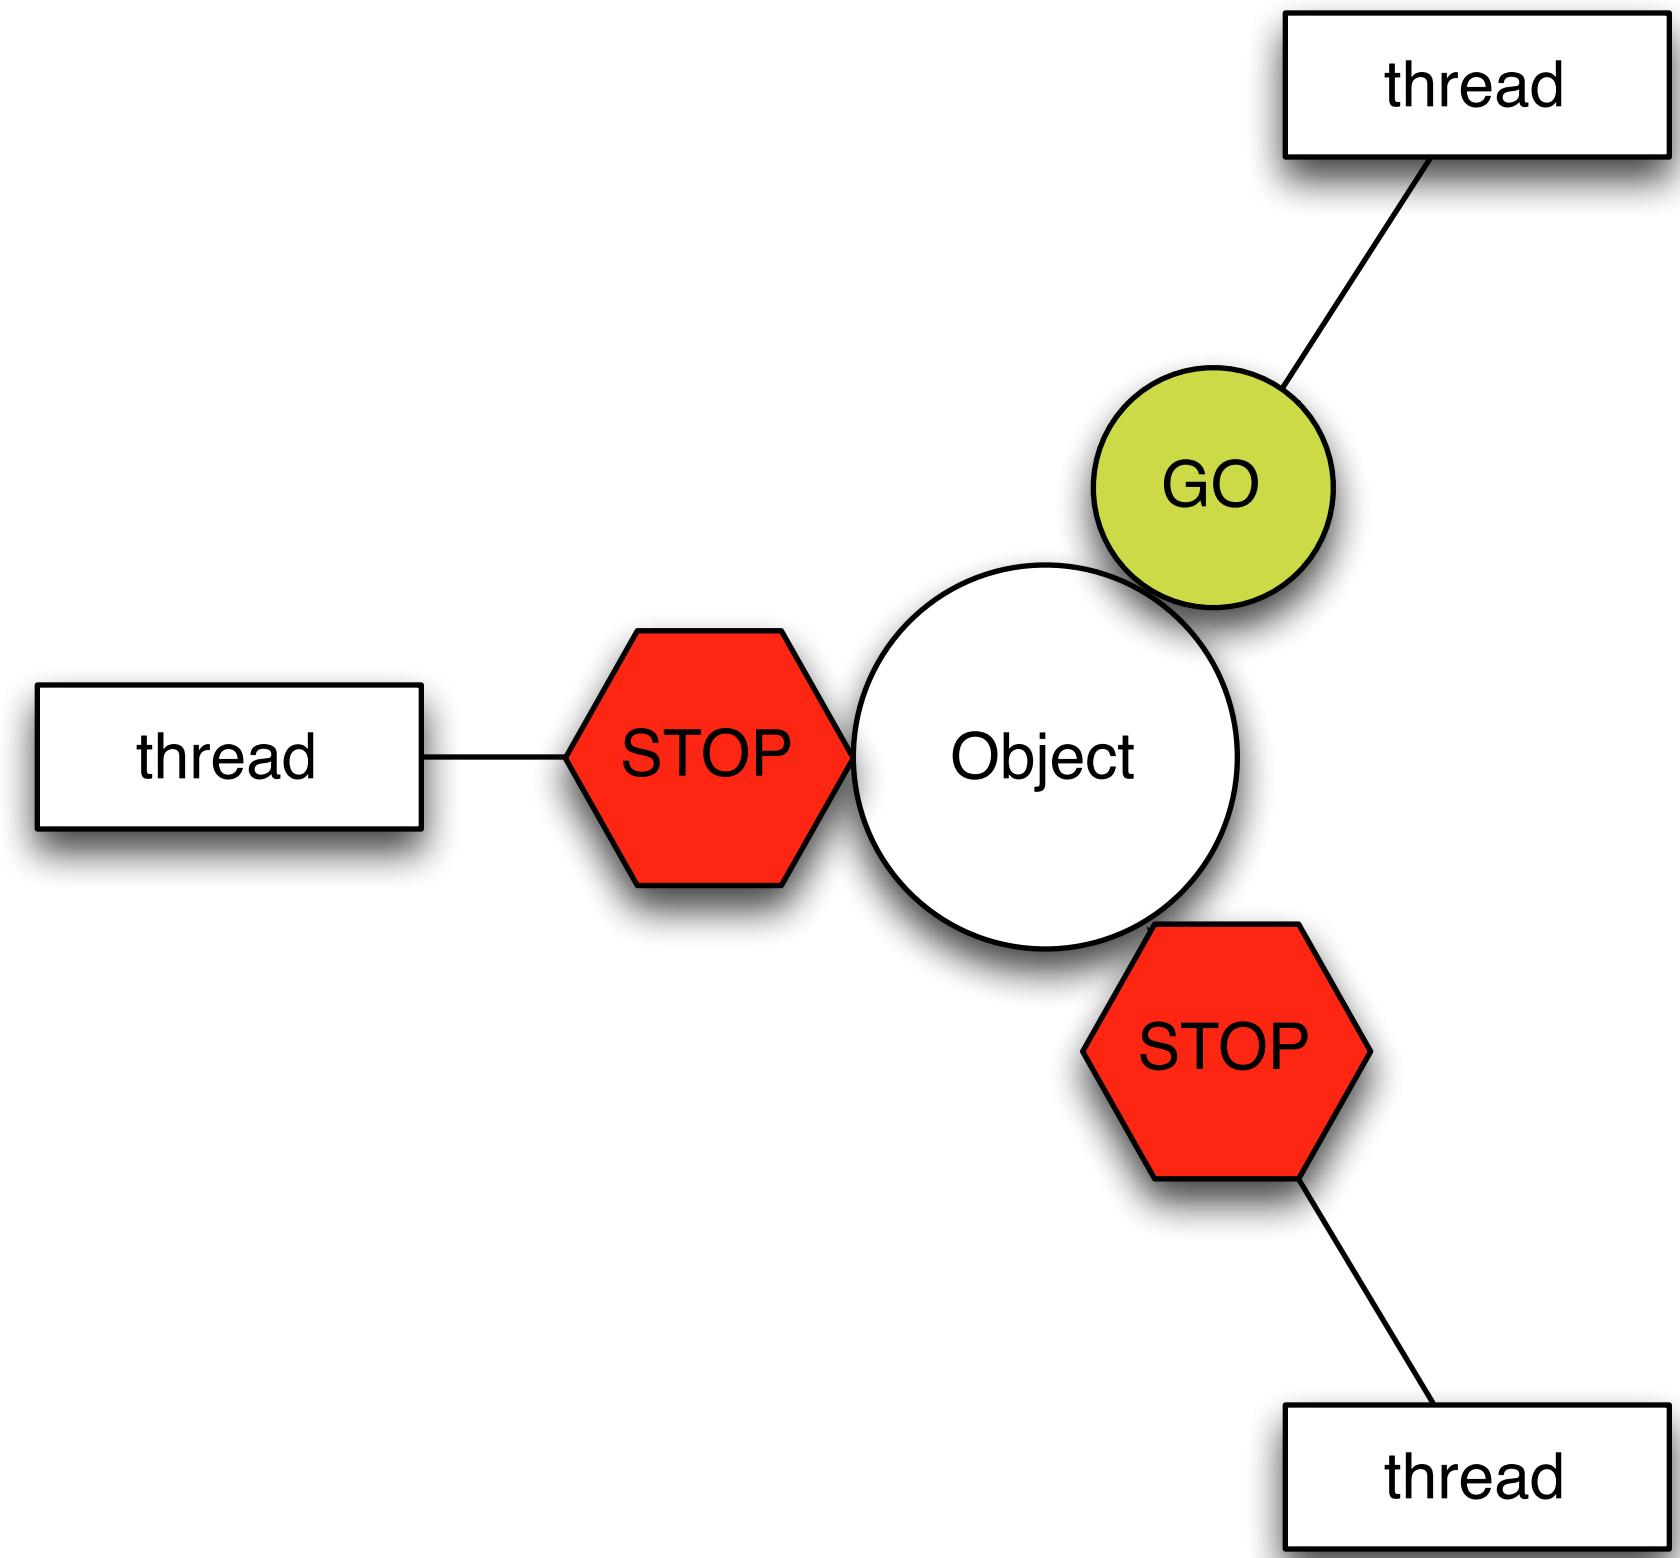
Each line represents 10% more synchronization



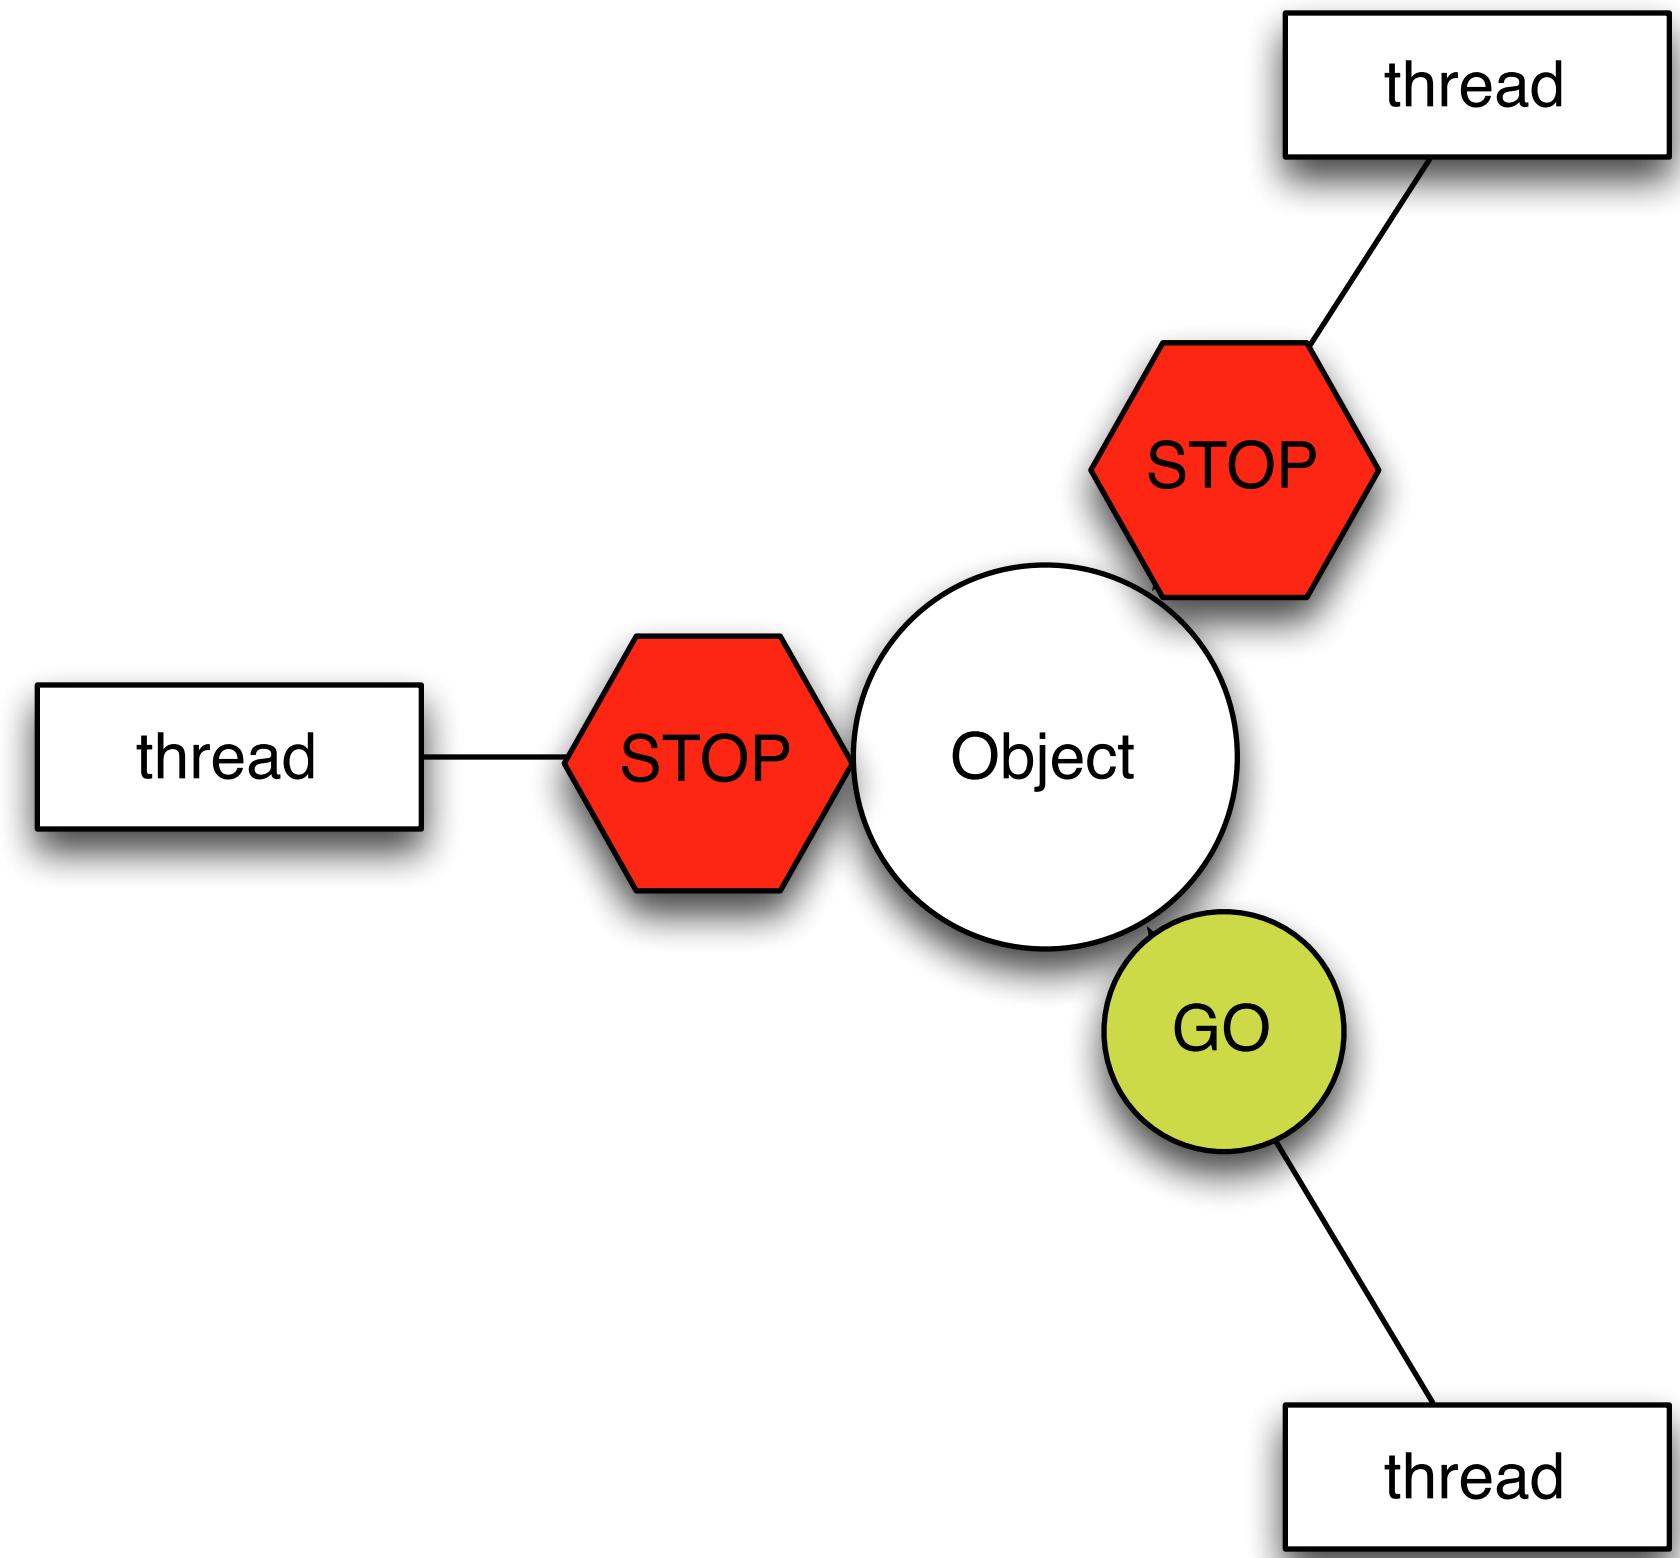
Why No Raw Synchronization Primitives?



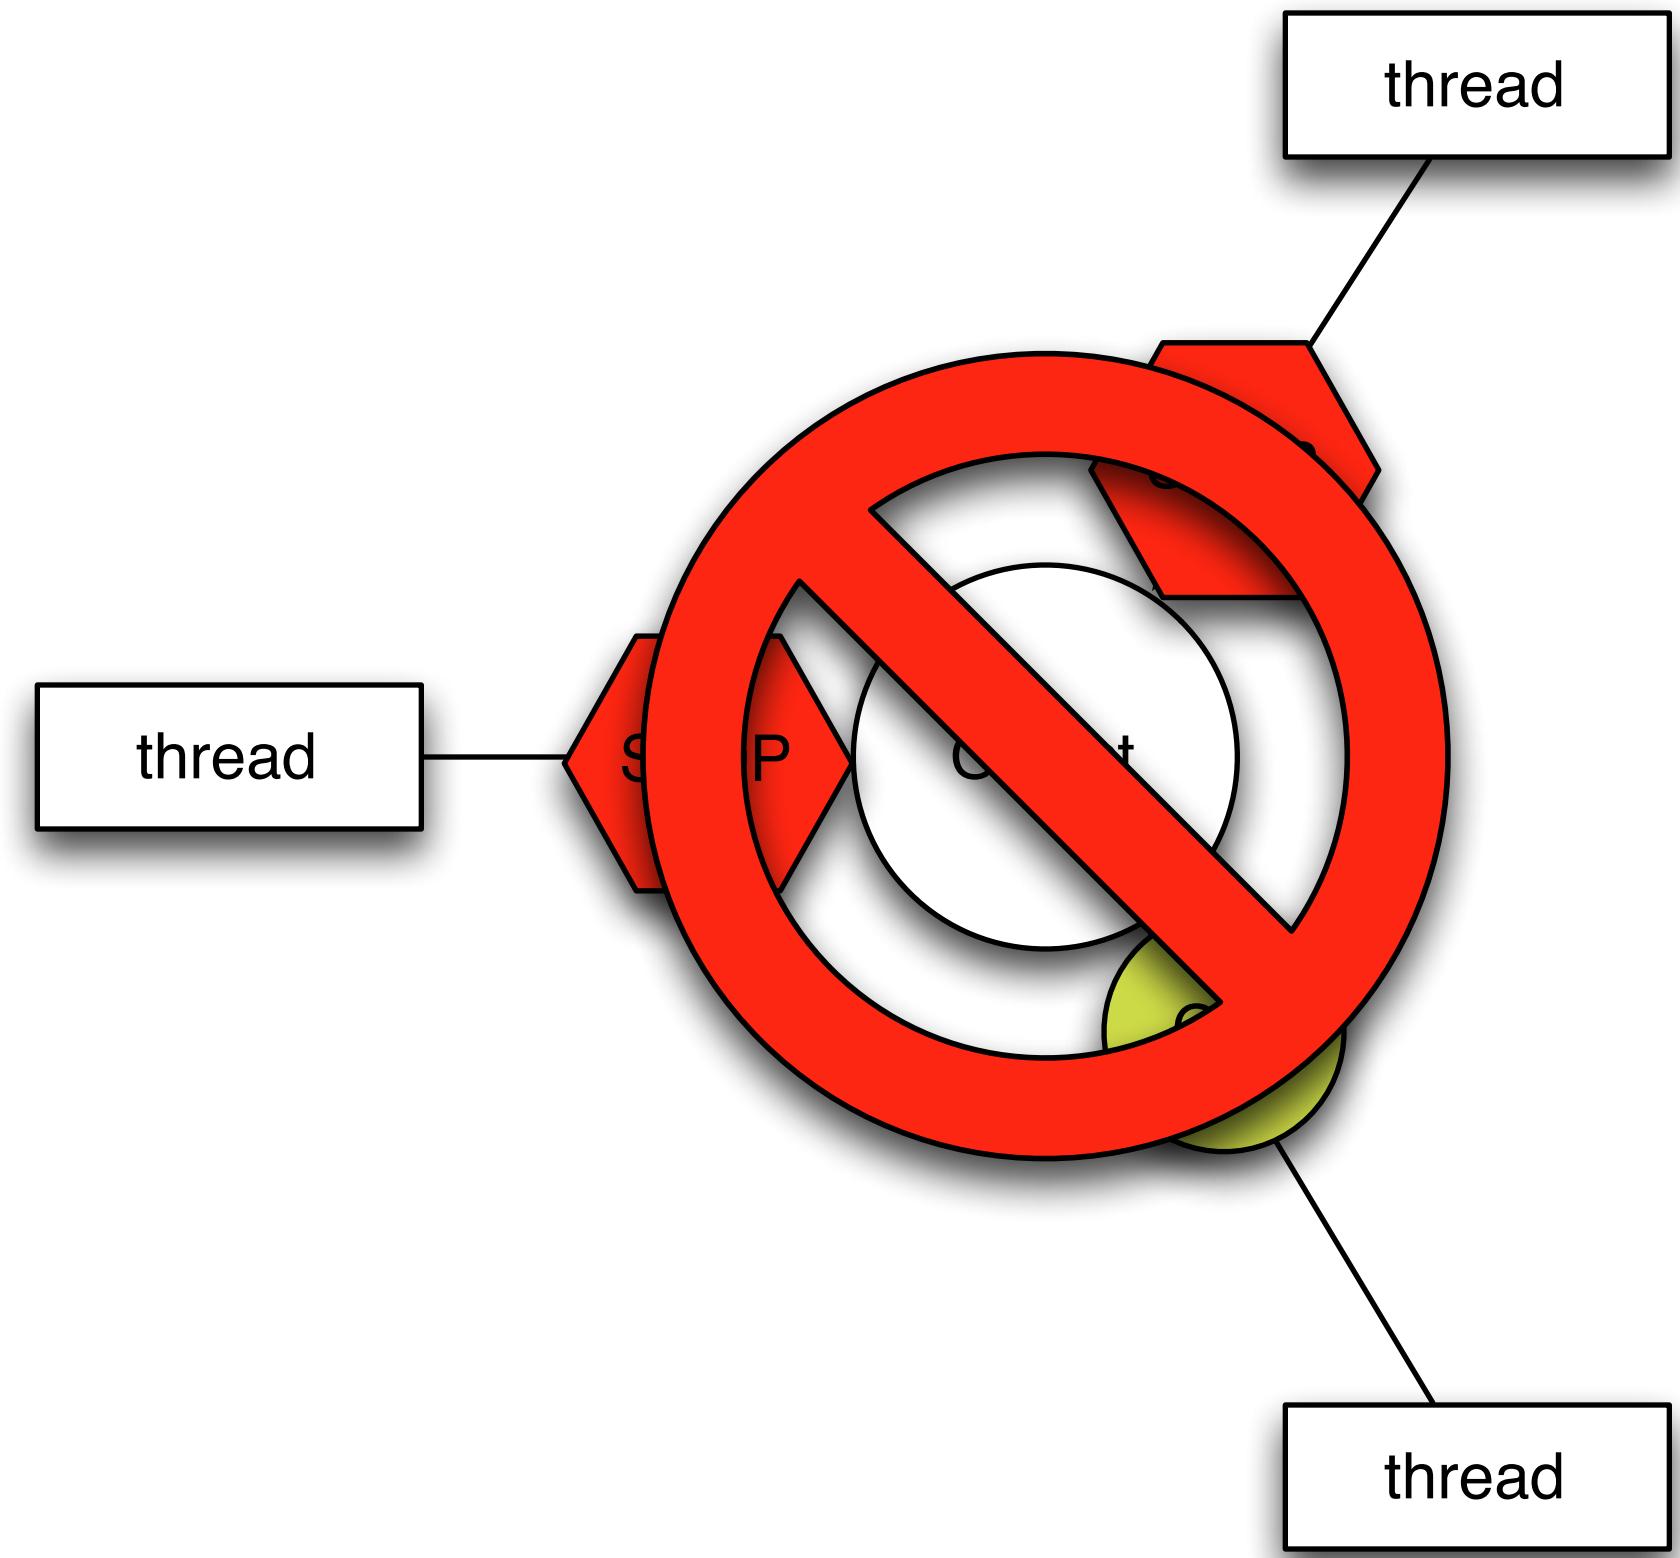
Why No Raw Synchronization Primitives?



Why No Raw Synchronization Primitives?



Why No Raw Synchronization Primitives?



Minimize Locks



Minimize Locks



Threads and Tasks

- Thread: Execution environment consisting of a stack and processor state running in parallel to other threads
- Task: A unit of work, often a function, to be executed on a thread
- Tasks are scheduled on a thread pool to optimize machine utilization

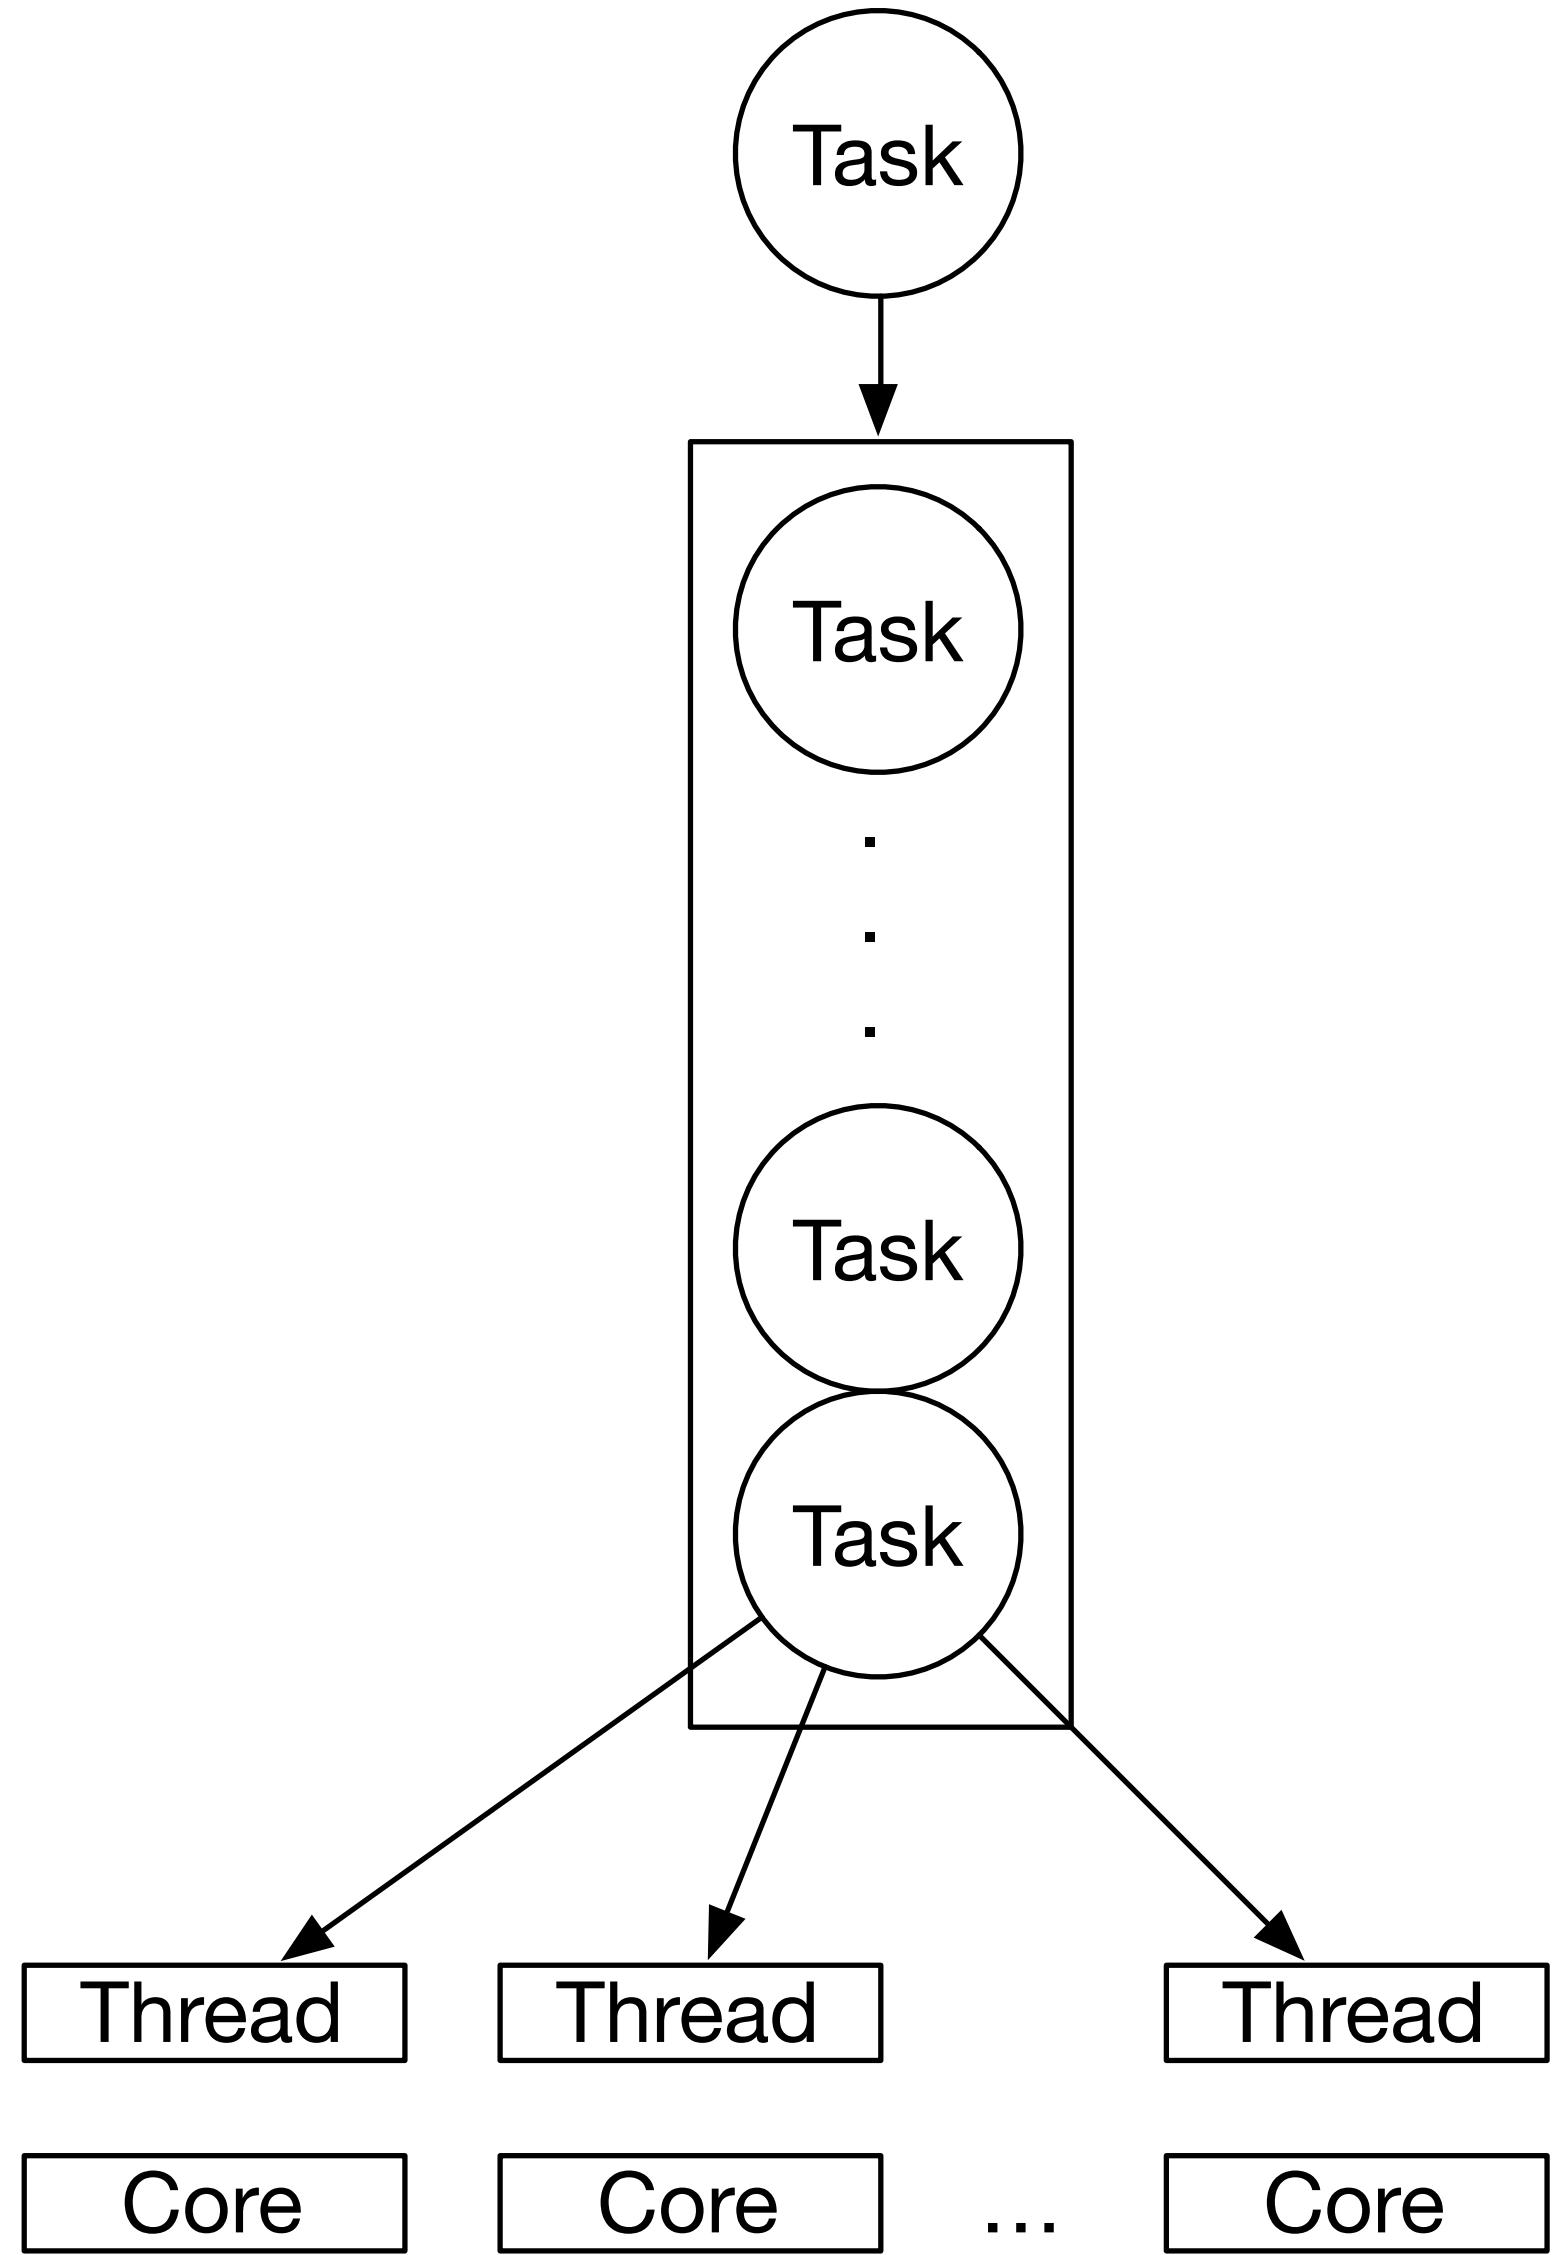
C++14 and Tasks

- C++14 does not (really) have a task system
 - Threads
 - Futures
- It is implementation defined if `std::async()` spins up a thread or executes on a thread pool.

Building a Task System

- Portable Reference Implementation in C++14
- Windows - Window Thread Pool and PPL
- Apple - Grand Central Dispatch (libdispatch)
 - open source, runs on Linux and Android
- Intel TBB - many platforms
 - open source
- HPX - many platforms
 - open source

Building a Task System



<http://docs.oracle.com/cd/E19253-01/816-5137/ggedn/index.html>

Building a Task System

Building a Task System

```
using lock_t = unique_lock<mutex>;
```

Building a Task System

```
using lock_t = unique_lock<mutex>;  
  
class notification_queue {  
    deque<function<void()>> _q;  
    mutex _mutex;  
    condition_variable _ready;
```

Building a Task System

```
using lock_t = unique_lock<mutex>;\n\nclass notification_queue {\n    deque<function<void()>> _q;\n    mutex _mutex;\n    condition_variable _ready;\n\npublic:\n    void pop(function<void()>& x) {\n        lock_t lock{_mutex};\n        while (_q.empty()) _ready.wait(lock);\n        x = move(_q.front());\n        _q.pop_front();\n    }\n}
```

Building a Task System

```
using lock_t = unique_lock<mutex>;\n\nclass notification_queue {\n    deque<function<void()>> _q;\n    mutex _mutex;\n    condition_variable _ready;\n\npublic:\n    void pop(function<void()>& x) {\n        lock_t lock{_mutex};\n        while (_q.empty()) _ready.wait(lock);\n        x = move(_q.front());\n        _q.pop_front();\n    }\n\n    template<typename F>\n    void push(F&& f) {\n        {\n            lock_t lock{_mutex};\n            _q.emplace_back(forward<F>(f));\n        }\n        _ready.notify_one();\n    }\n};
```

Building a Task System

Building a Task System

```
class task_system {  
    const unsigned _count{thread::hardware_concurrency()};  
    vector<thread> _threads;  
    notification_queue _q;
```

Building a Task System

```
class task_system {
    const unsigned _count{thread::hardware_concurrency()};
    vector<thread> _threads;
    notification_queue _q;

    void run(unsigned i) {
        while (true) {
            function<void()> f;
            _q.pop(f);
            f();
        }
    }
}
```

Building a Task System

```
class task_system {
    const unsigned _count{thread::hardware_concurrency()};
    vector<thread> _threads;
    notification_queue _q;

    void run(unsigned i) {
        while (true) {
            function<void()> f;
            _q.pop(f);
            f();
        }
    }

public:
    task_system() {
        for (unsigned n = 0; n != _count; ++n) {
            _threads.emplace_back([&, n]{ run(n); });
        }
    }
}
```

Building a Task System

```
class task_system {
    const unsigned _count{thread::hardware_concurrency()};
    vector<thread> _threads;
    notification_queue _q;

    void run(unsigned i) {
        while (true) {
            function<void()> f;
            _q.pop(f);
            f();
        }
    }

public:
    task_system() {
        for (unsigned n = 0; n != _count; ++n) {
            _threads.emplace_back([&, n]{ run(n); });
        }
    }

    ~task_system() {
        for (auto& e : _threads) e.join();
    }
}
```

Building a Task System

```
class task_system {
    const unsigned _count{thread::hardware_concurrency()};
    vector<thread> _threads;
    notification_queue _q;

    void run(unsigned i) {
        while (true) {
            function<void()> f;
            _q.pop(f);
            f();
        }
    }

public:
    task_system() {
        for (unsigned n = 0; n != _count; ++n) {
            _threads.emplace_back([&, n]{ run(n); });
        }
    }

    ~task_system() {
        for (auto& e : _threads) e.join();
    }

    template <typename F>
    void async_(F&& f) {
        _q.push(forward<F>(f));
    }
};
```

Building a Task System

```
class notification_queue {
    deque<function<void()>> _q;
    bool _done{false};
    mutex _mutex;
    condition_variable _ready;

public:
    void done() {
        {
            unique_lock<mutex> lock{_mutex};
            _done = true;
        }
        _ready.notify_all();
    }

    bool pop(function<void()>& x) {
        lock_t lock{_mutex};
        while (_q.empty() && !_done) _ready.wait(lock);
        if (_q.empty()) return false;
        x = move(_q.front());
        _q.pop_front();
        return true;
    }

template<typename F>
void push(F&& f) {
    lock_t lock{_mutex};
    q.emplace back(forward<F>(f));
}
```

Building a Task System

```
class notification_queue {
    deque<function<void()>> _q;
    bool _done{false};
    mutex _mutex;
    condition_variable _ready;

public:
    void done() {
        {
            unique_lock<mutex> lock{_mutex};
            _done = true;
        }
        _ready.notify_all();
    }

    bool pop(function<void()>& x) {
        lock_t lock{_mutex};
        while (_q.empty() && !_done) _ready.wait(lock);
        if (_q.empty()) return false;
        x = move(_q.front());
        _q.pop_front();
        return true;
    }

template<typename F>
void push(F&& f) {
    lock_t lock{_mutex};
    q.emplace back(forward<F>(f));
}
```

Building a Task System

```
class notification_queue {
    deque<function<void()>> _q;
    bool _done{false};
    mutex _mutex;
    condition_variable _ready;

public:
    void done() {
        {
            unique_lock<mutex> lock{_mutex};
            _done = true;
        }
        _ready.notify_all();
    }

    bool pop(function<void()>& x) {
        lock_t lock{_mutex};
        while (_q.empty() && !_done) _ready.wait(lock);
        if (_q.empty()) return false;
        x = move(_q.front());
        _q.pop_front();
        return true;
    }

    template<typename F>
    void push(F&& f) {
        lock_t lock{_mutex};
        q.emplace back(forward<F>(f));
    }
}
```

Building a Task System

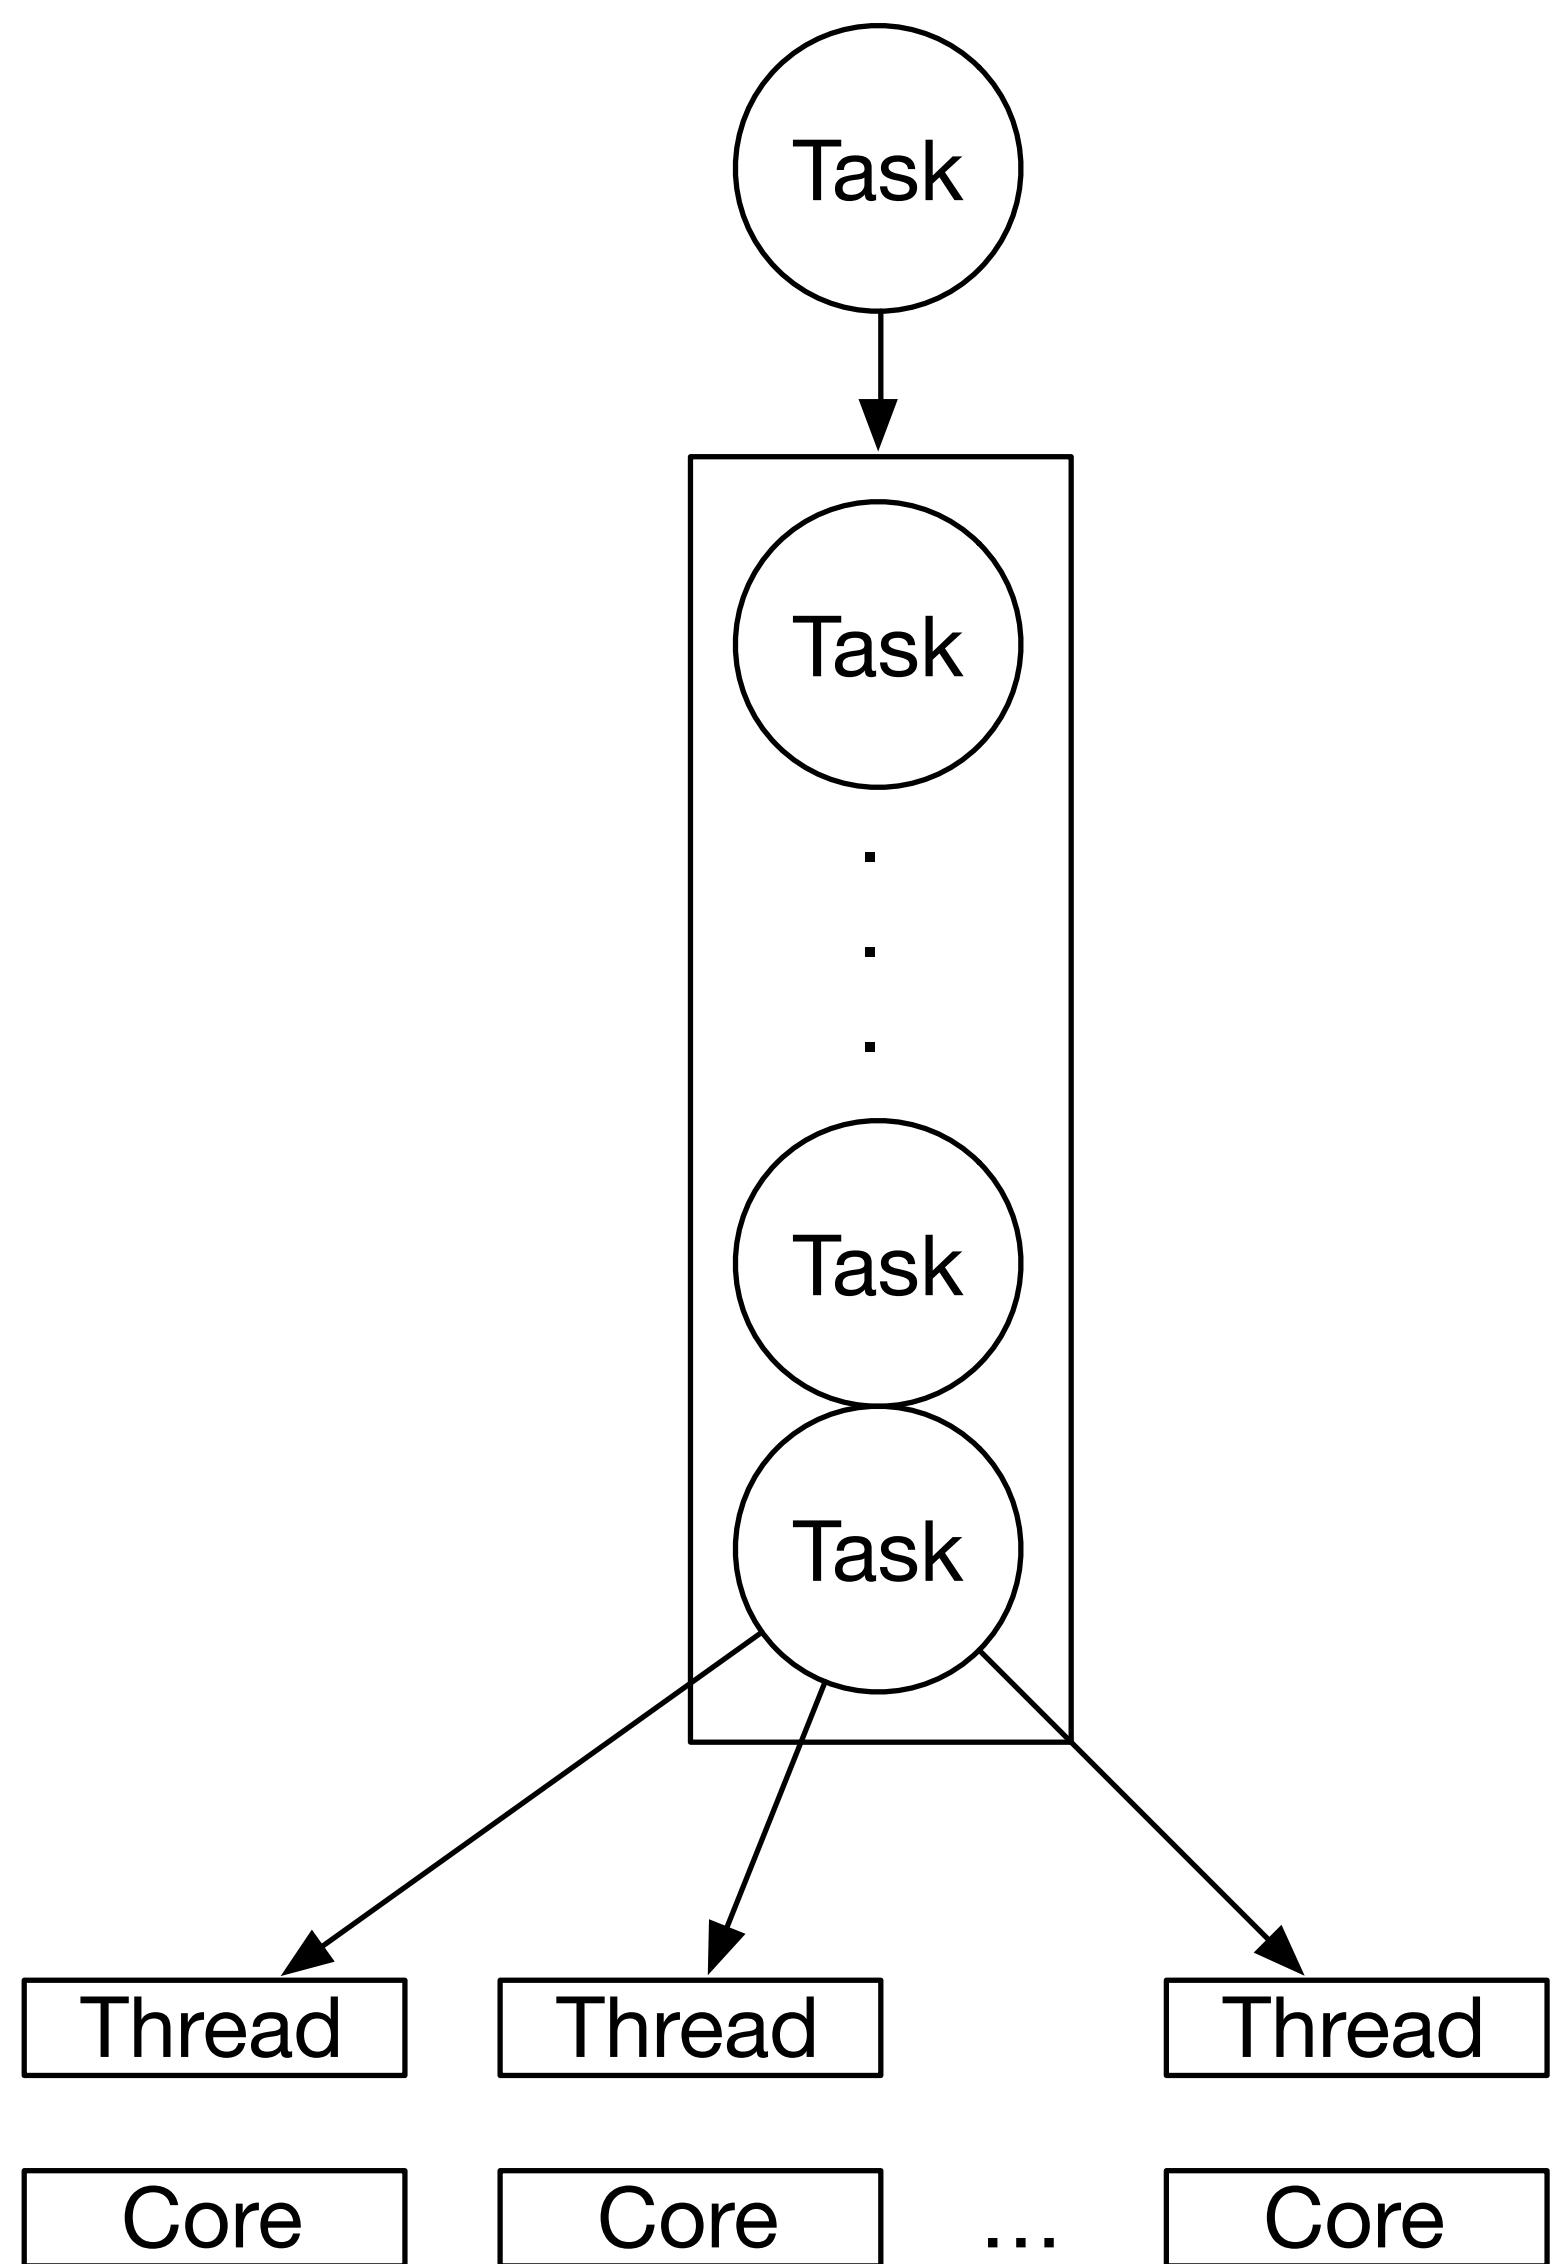
```
class notification_queue {
    deque<function<void()>> _q;
    bool _done{false};
    mutex _mutex;
    condition_variable _ready;

public:
    void done() {
        {
            unique_lock<mutex> lock{_mutex};
            _done = true;
        }
        _ready.notify_all();
    }

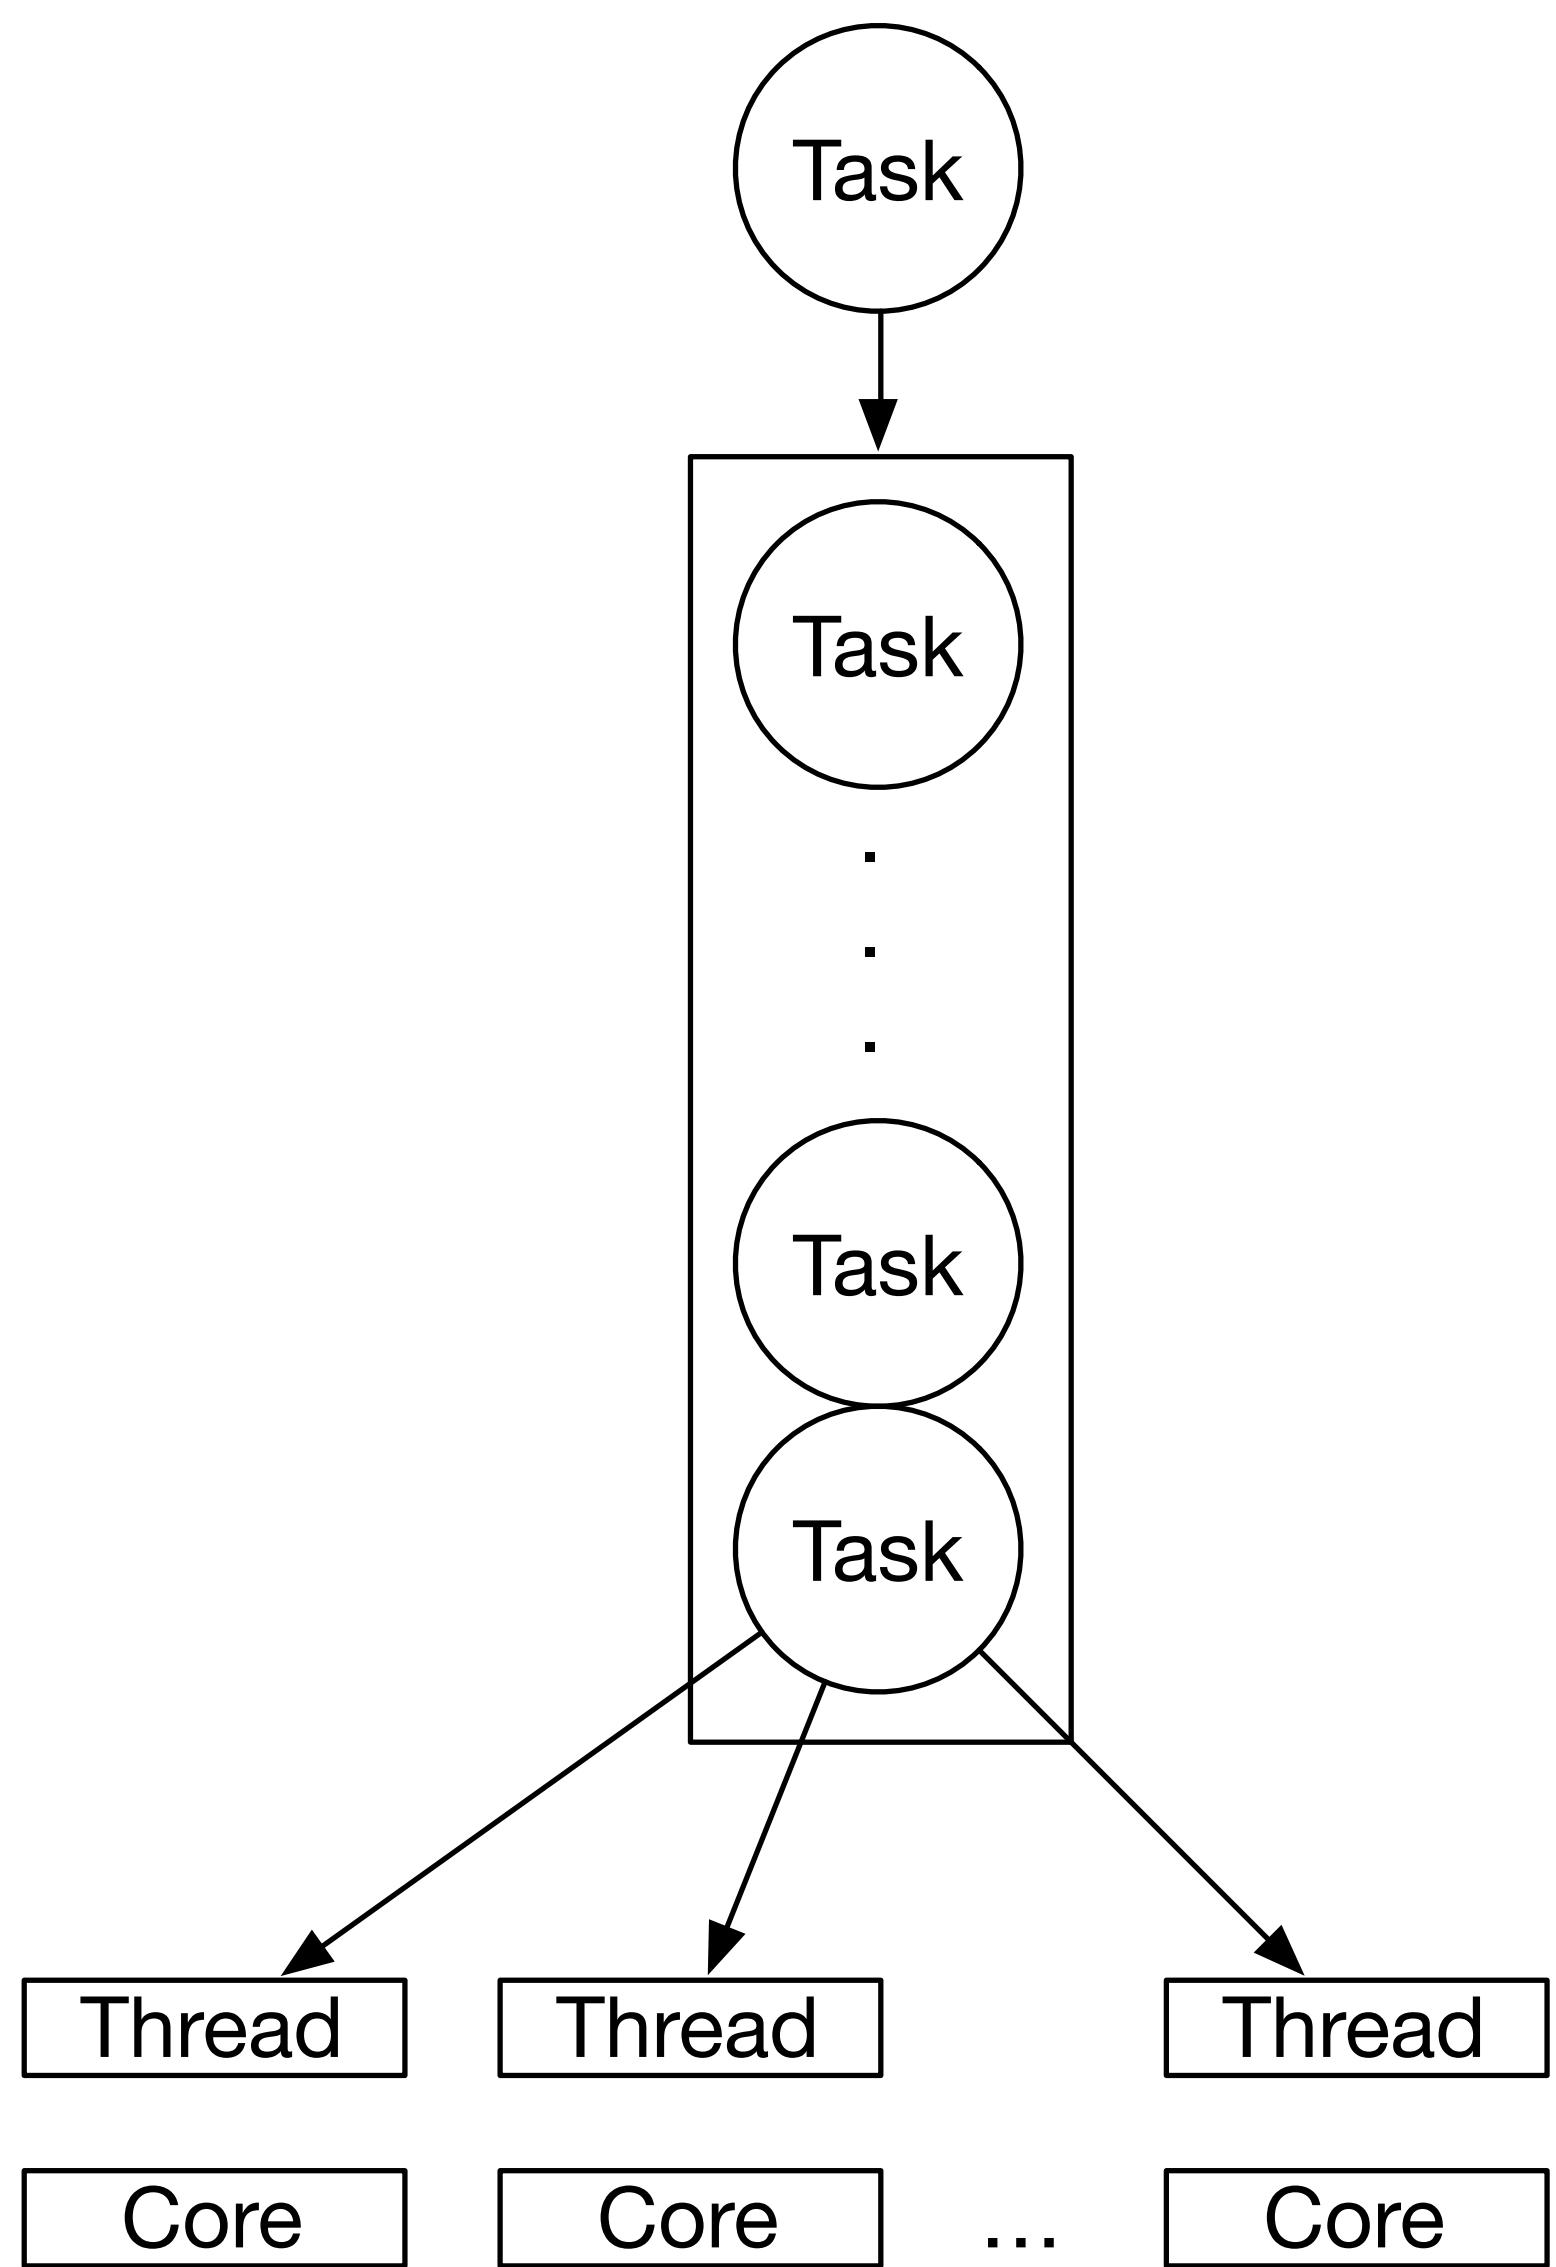
    bool pop(function<void()>& x) {
        lock_t lock{_mutex};
        while (_q.empty() && !_done) _ready.wait(lock);
        if (_q.empty()) return false;
        x = move(_q.front());
        _q.pop_front();
        return true;
    }

template<typename F>
void push(F&& f) {
    lock_t lock{_mutex};
    q.emplace back(forward<F>(f));
}
```

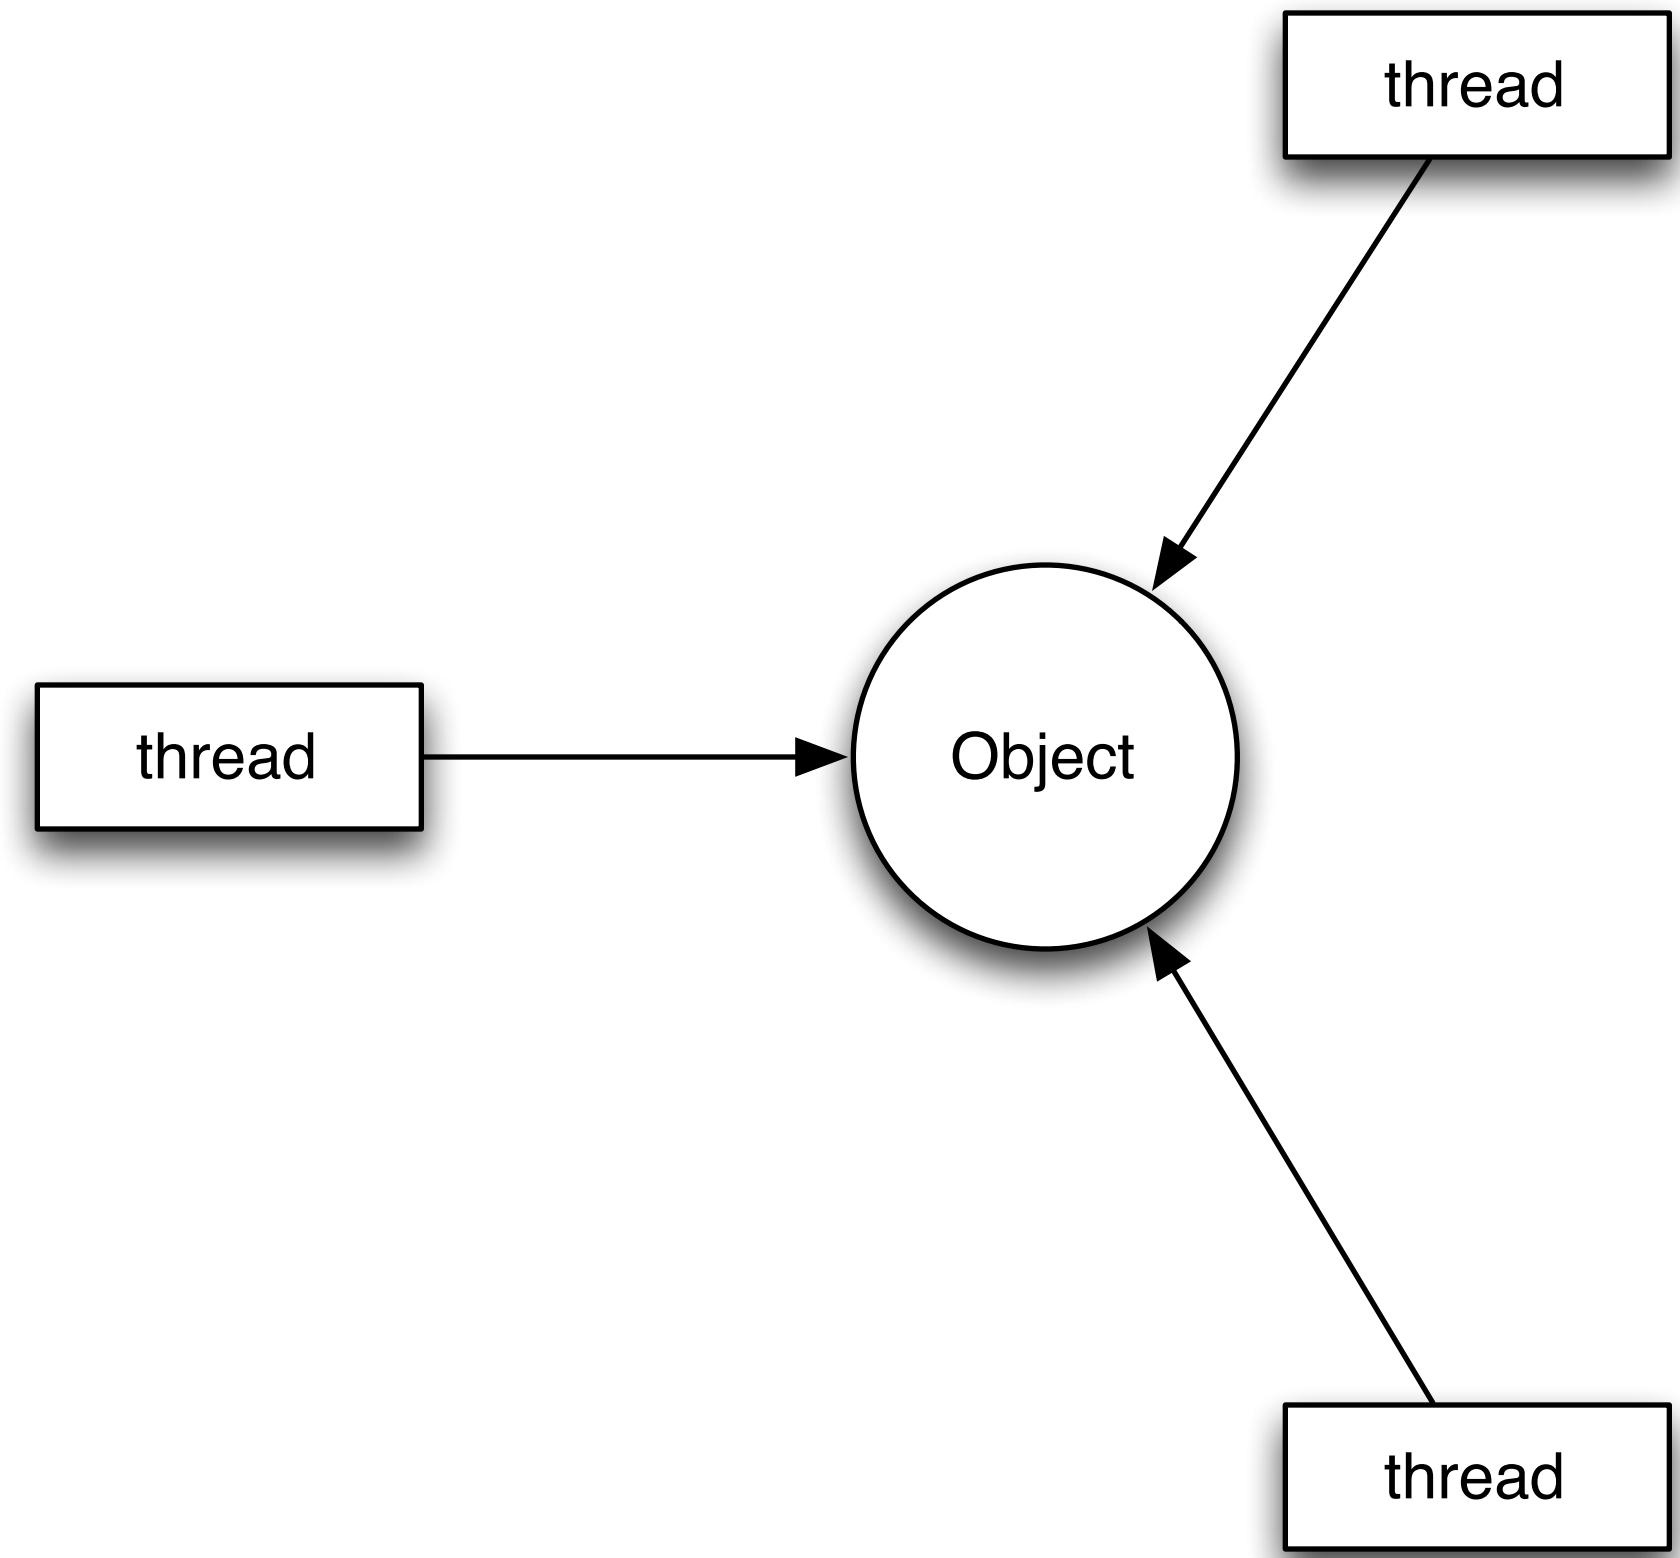
Building a Task System



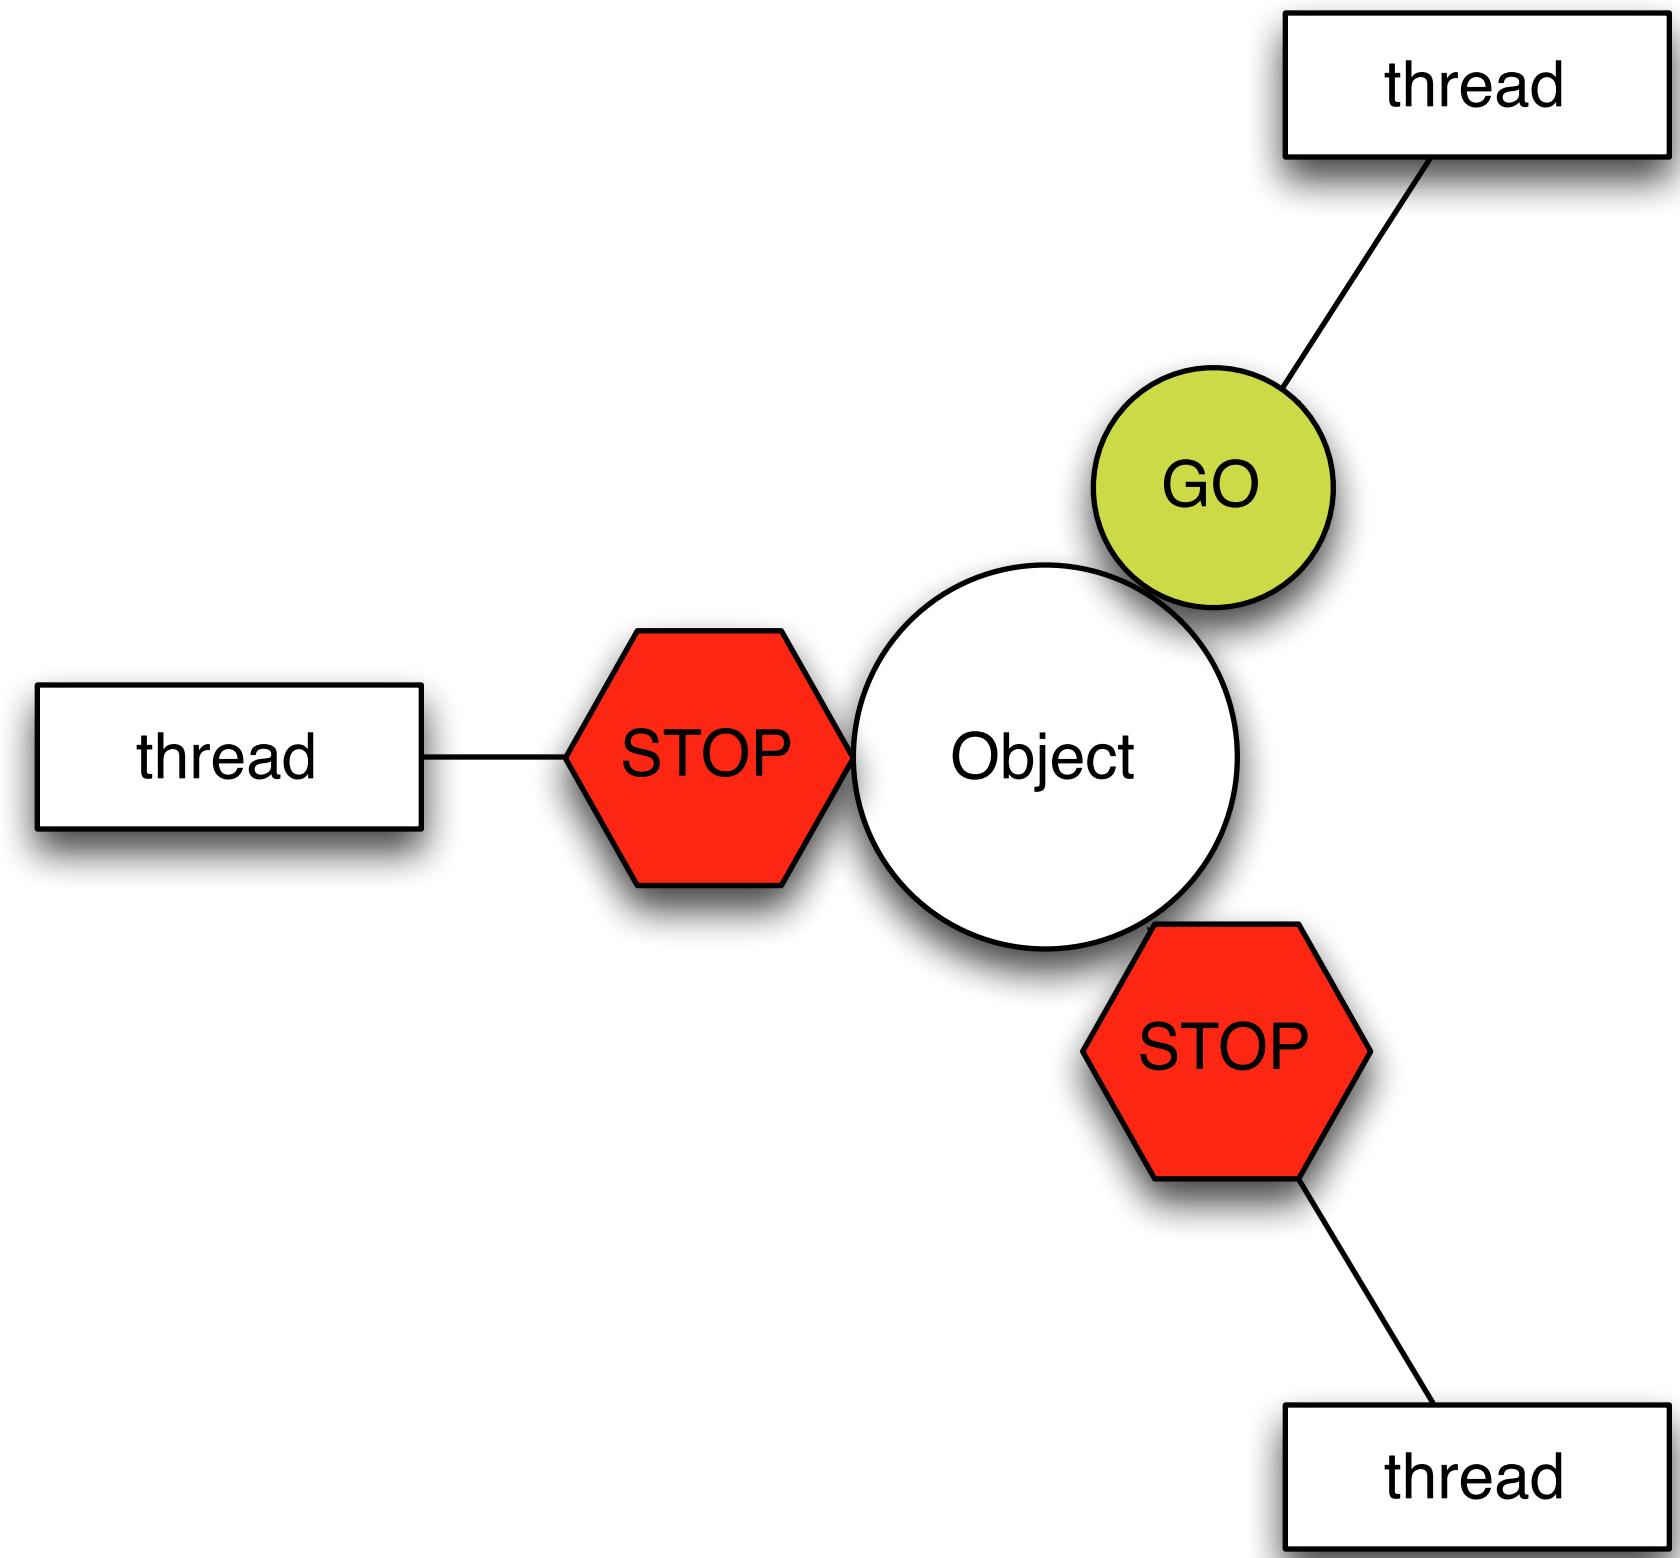
Building a Task System



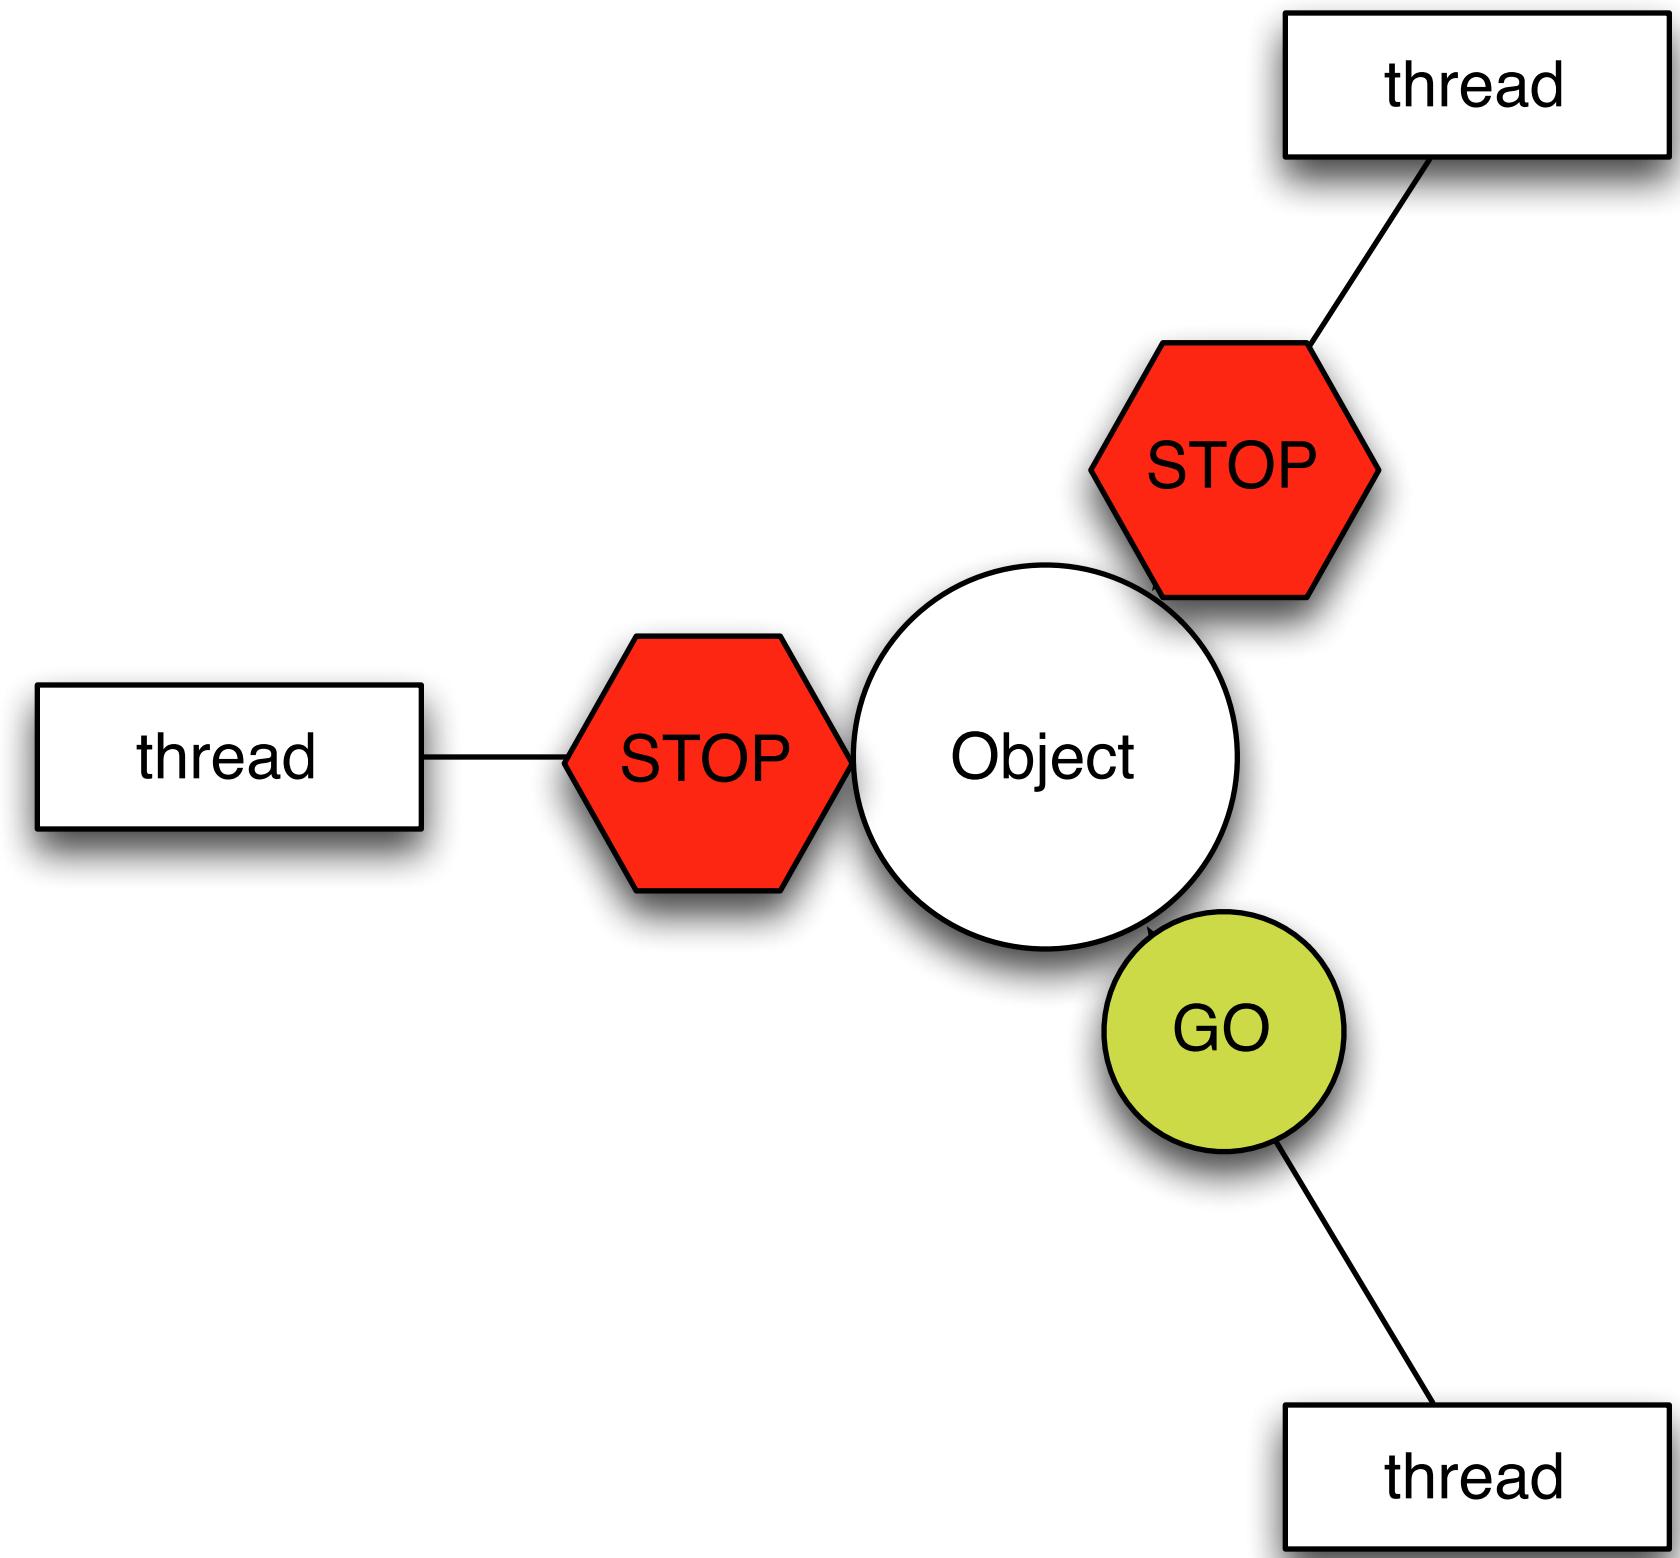
Why No Raw Synchronization Primitives?



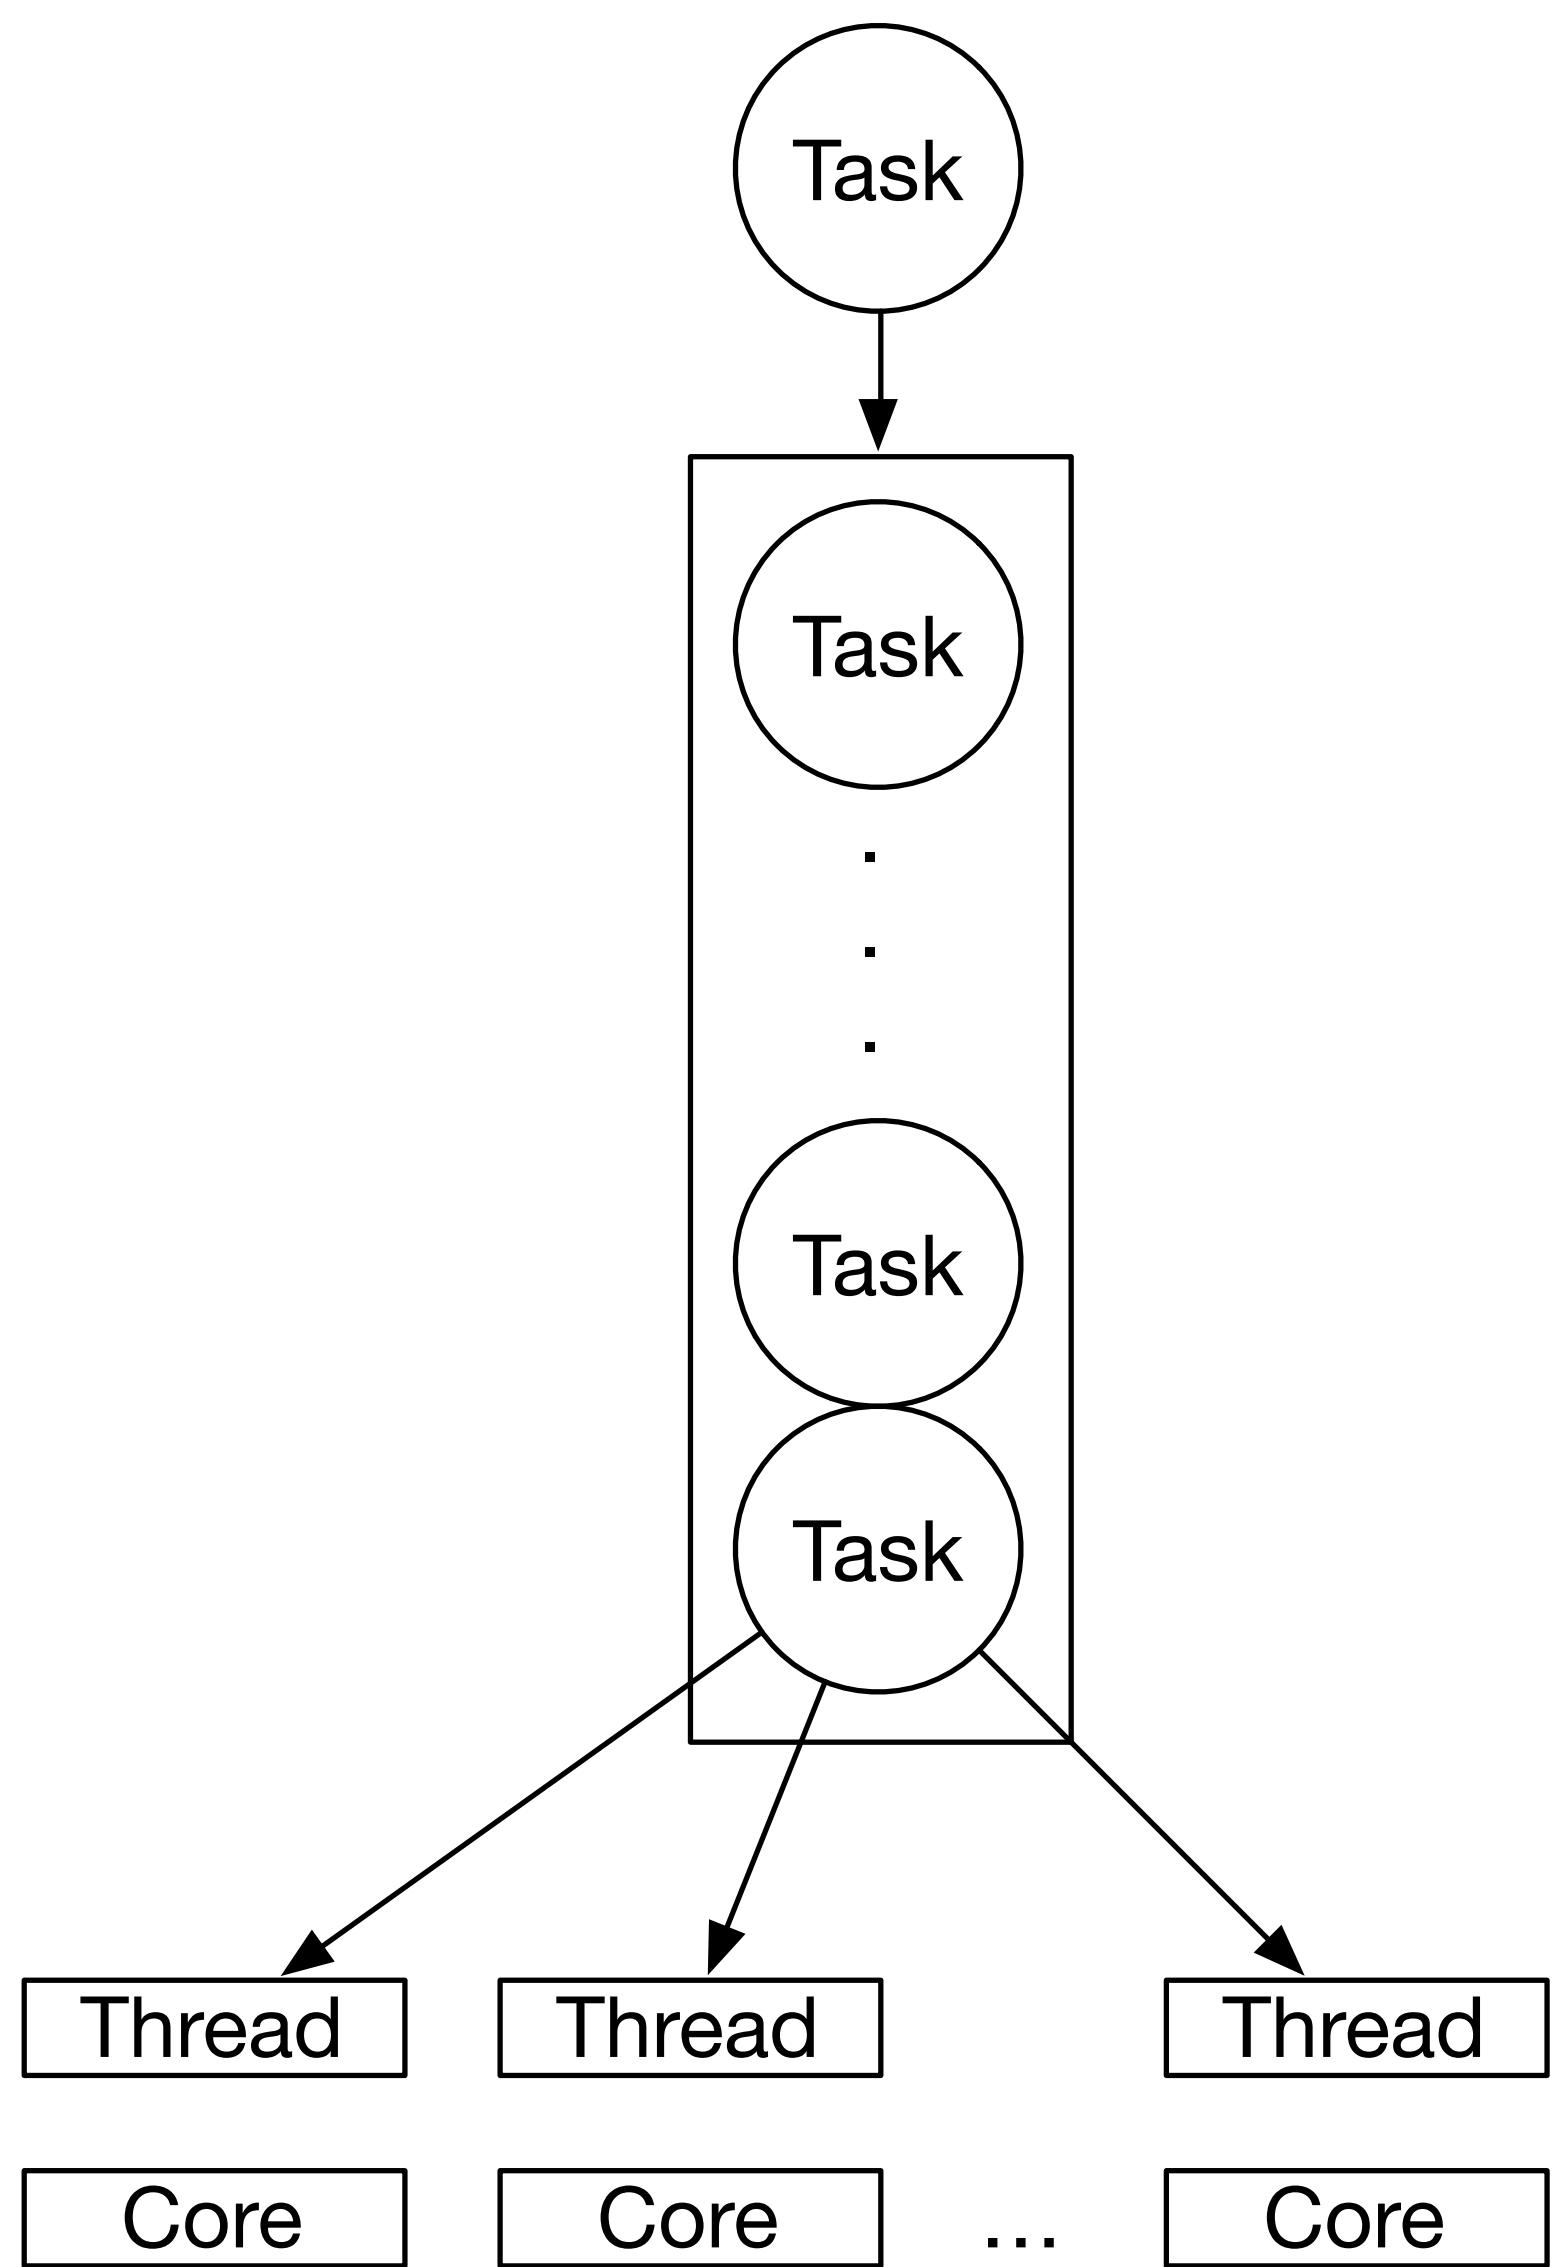
Why No Raw Synchronization Primitives?



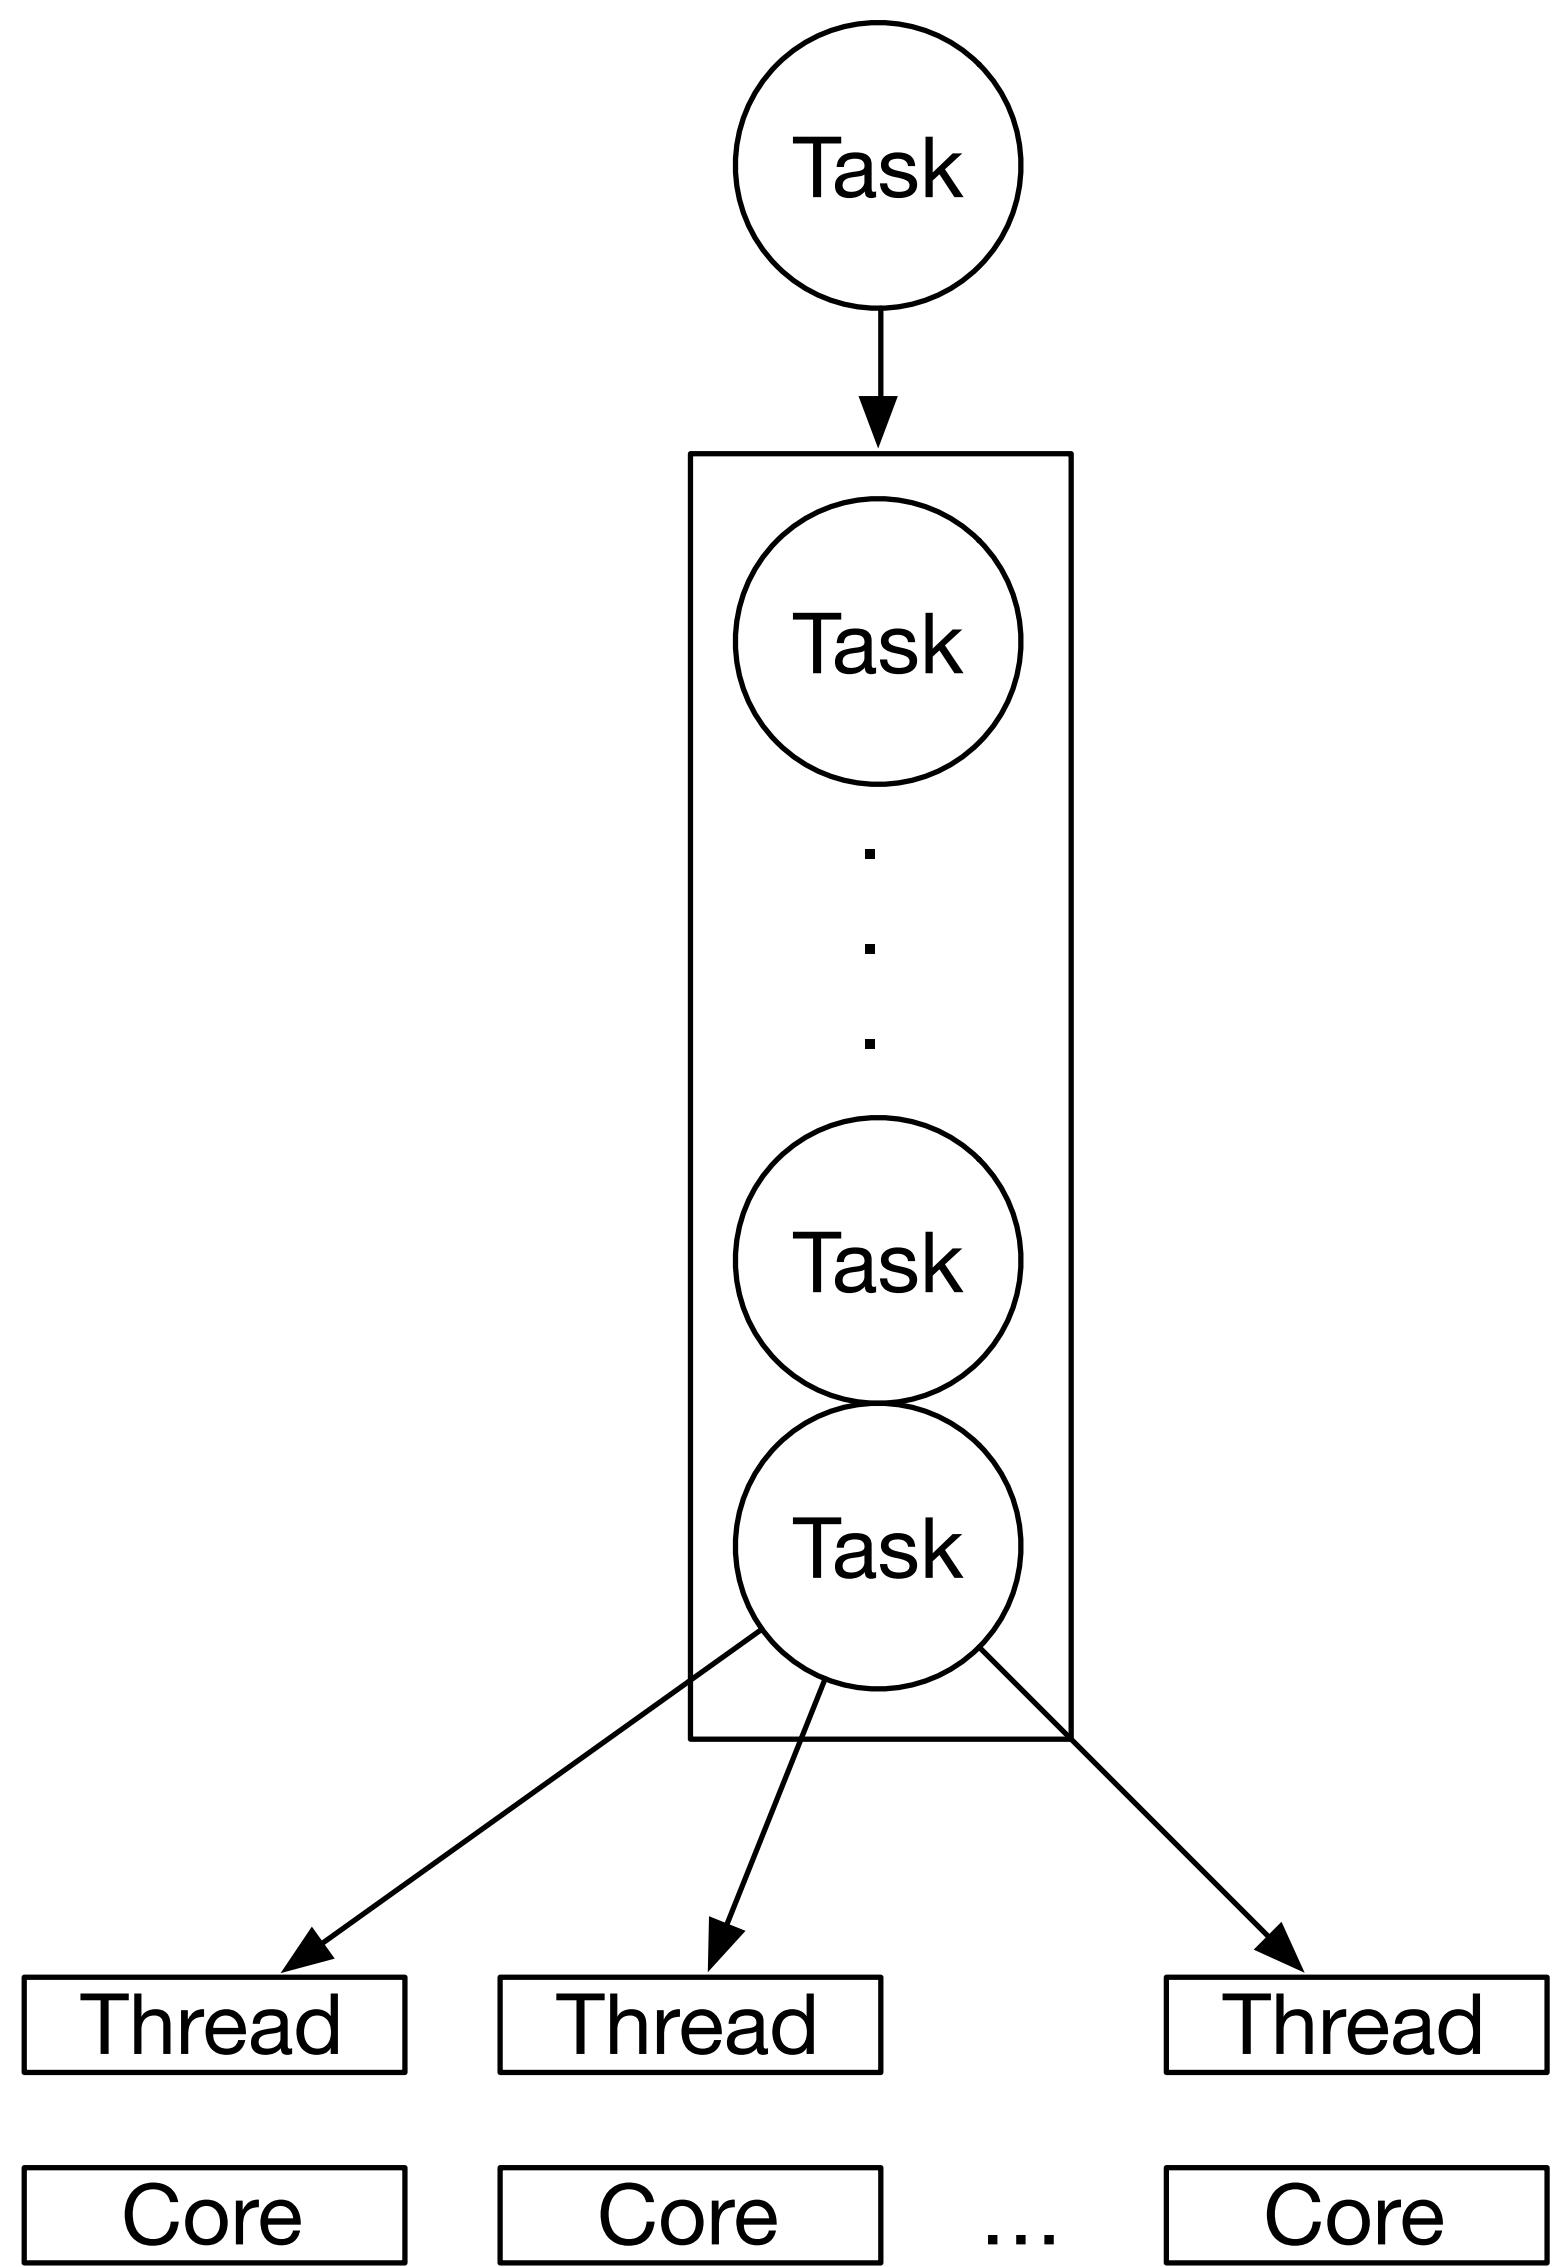
Why No Raw Synchronization Primitives?



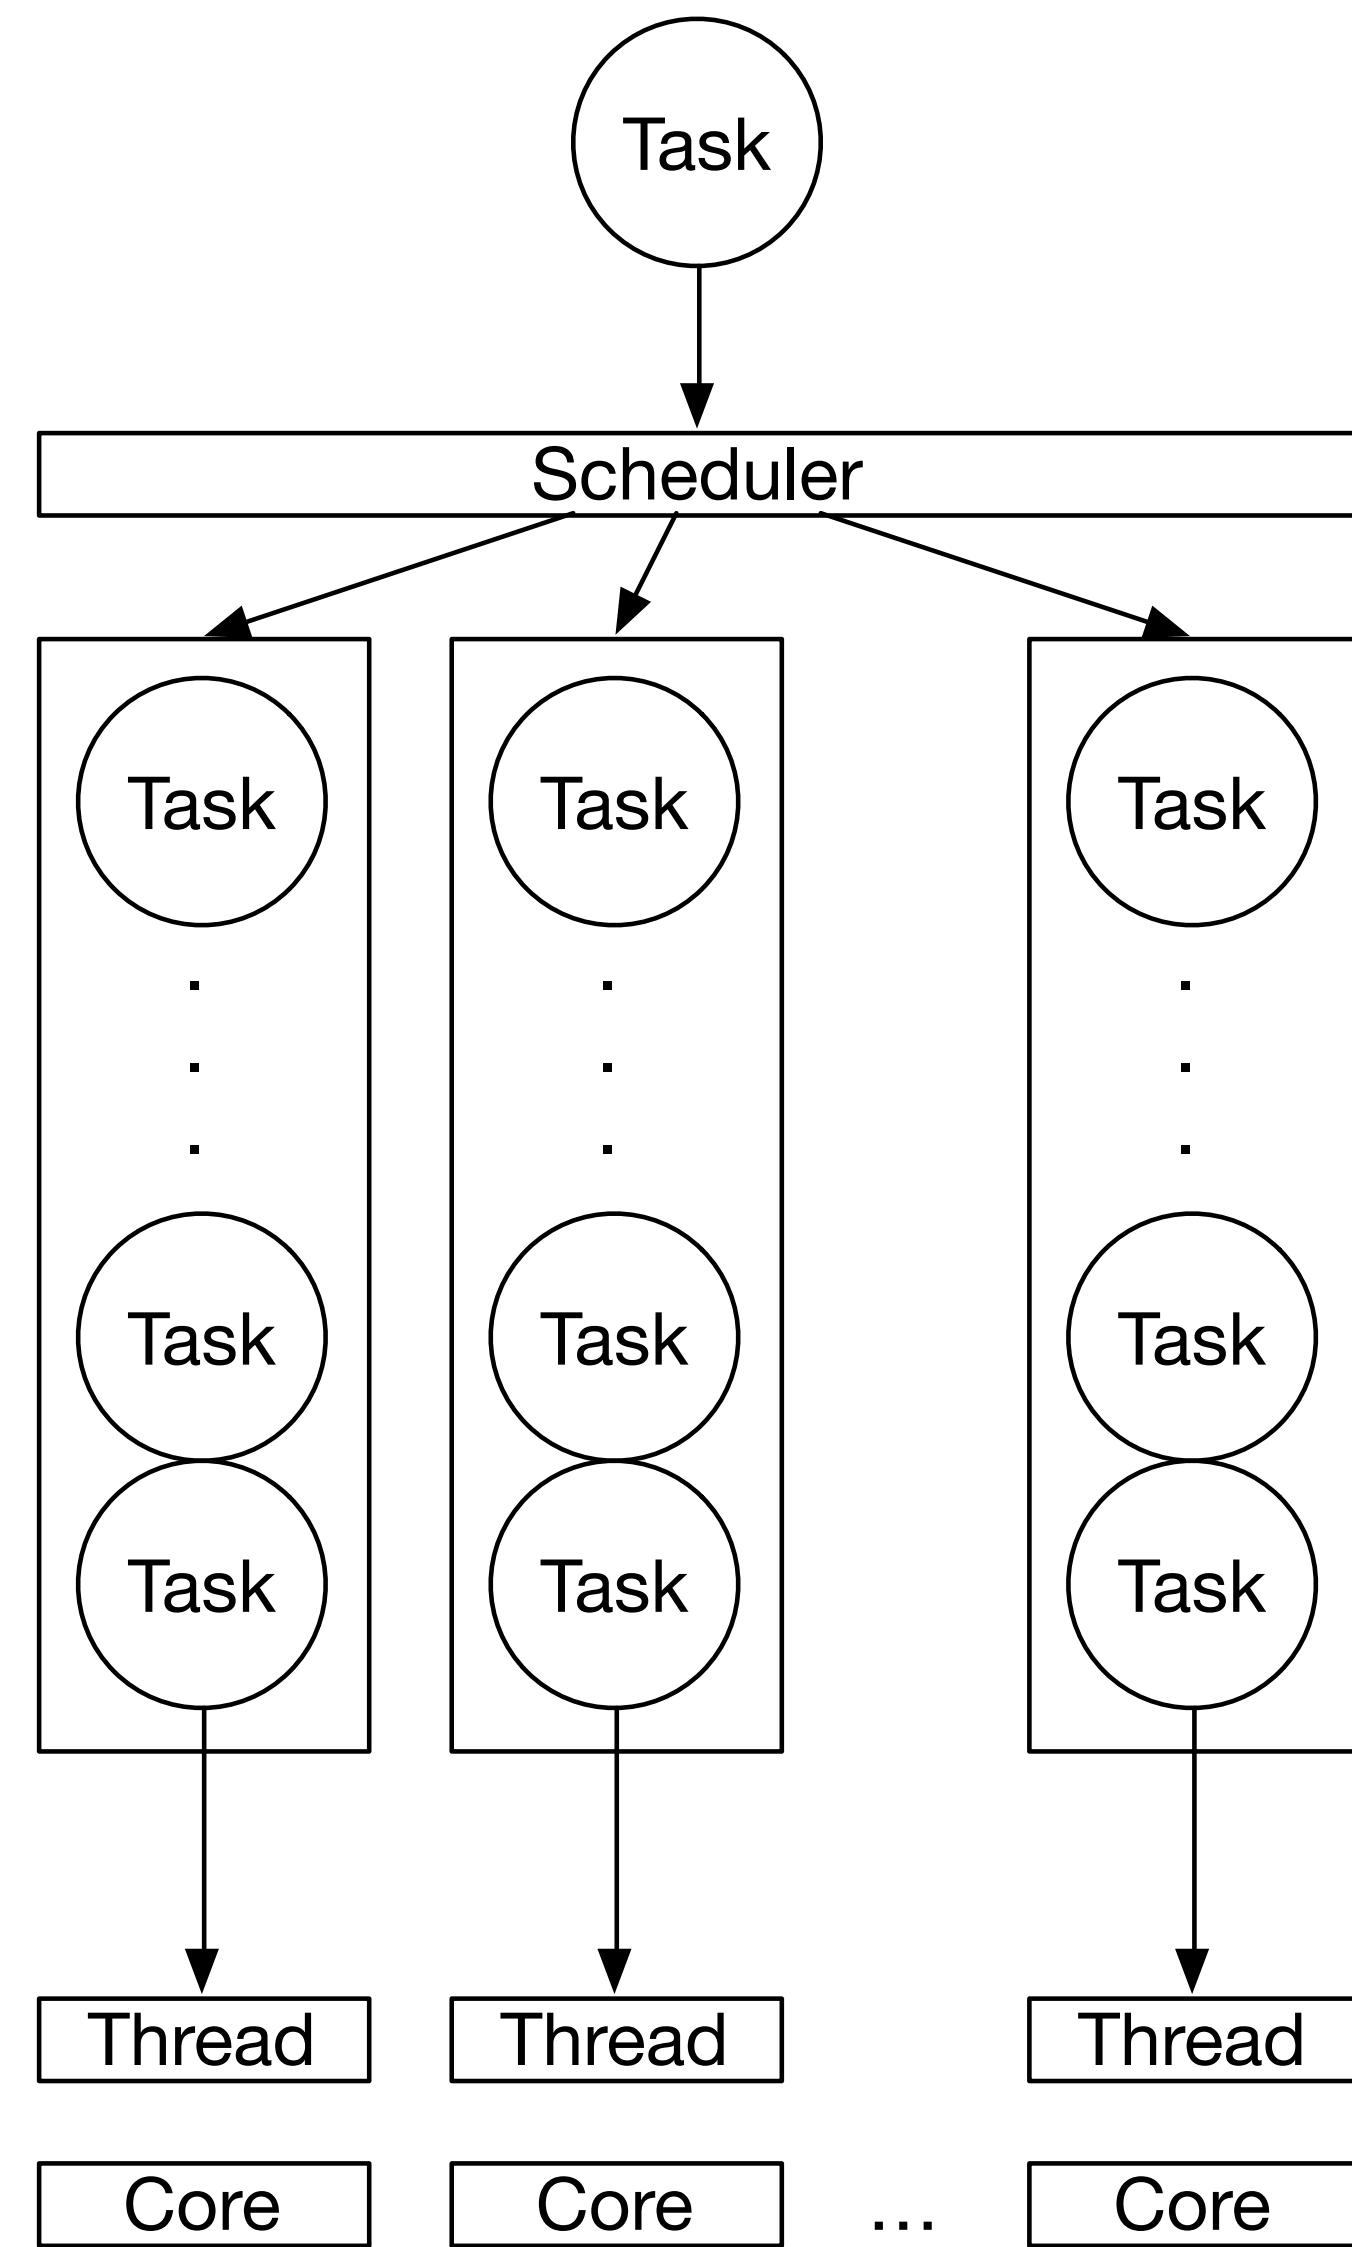
Building a Task System



Building a Task System



Building a Task System



Building a Task System

```
class task_system {
    const unsigned _count{thread::hardware_concurrency()};
    vector<thread> _threads;
    vector<notification_queue> _q{_count};
    atomic<unsigned> _index{0};

    void run(unsigned i) {
        while (true) {
            function<void()> f;
            if (!_q[i].pop(f)) break;
            f();
        }
    }

public:
    task_system() { ... }

    ~task_system() {
        for (auto& e : _q) e.done();
        for (auto& e : _threads) e.join();
    }

    template <typename F>
    void async_(F&& f) {
        auto i = _index++;
        _q[i % _count].push(forward<F>(f));
    }
};
```

Building a Task System

```
class task_system {
    const unsigned _count{thread::hardware_concurrency()};
    vector<thread> _threads;
    vector<notification_queue> _q{_count};
    atomic<unsigned> _index{0};

    void run(unsigned i) {
        while (true) {
            function<void()> f;
            if (!_q[i].pop(f)) break;
            f();
        }
    }

public:
    task_system() { ... }

    ~task_system() {
        for (auto& e : _q) e.done();
        for (auto& e : _threads) e.join();
    }

    template <typename F>
    void async_(F&& f) {
        auto i = _index++;
        _q[i % _count].push(forward<F>(f));
    }
};
```

Building a Task System

```
class task_system {
    const unsigned _count{thread::hardware_concurrency()};
    vector<thread> _threads;
    vector<notification_queue> _q{_count};
    atomic<unsigned> _index{0};

    void run(unsigned i) {
        while (true) {
            function<void()> f;
            if (!_q[i].pop(f)) break;
            f();
        }
    }

public:
    task_system() { ... }

    ~task_system() {
        for (auto& e : _q) e.done();
        for (auto& e : _threads) e.join();
    }

    template <typename F>
    void async_(F&& f) {
        auto i = _index++;
        _q[i % _count].push(forward<F>(f));
    }
};
```

Building a Task System

```
class task_system {
    const unsigned _count{thread::hardware_concurrency()};
    vector<thread> _threads;
    vector<notification_queue> _q{_count};
    atomic<unsigned> _index{0};

    void run(unsigned i) {
        while (true) {
            function<void()> f;
            if (!_q[i].pop(f)) break;
            f();
        }
    }

public:
    task_system() { ... }

    ~task_system() {
        for (auto& e : _q) e.done();
        for (auto& e : _threads) e.join();
    }

    template <typename F>
    void async_(F&& f) {
        auto i = _index++;
        _q[i % _count].push(forward<F>(f));
    }
};
```

Building a Task System

```
class task_system {
    const unsigned _count{thread::hardware_concurrency()};
    vector<thread> _threads;
    vector<notification_queue> _q{_count};
    atomic<unsigned> _index{0};

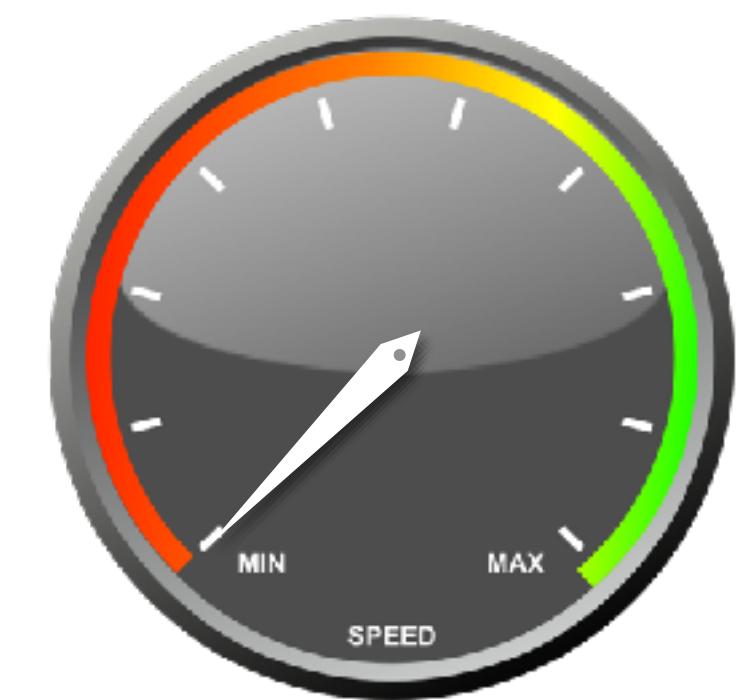
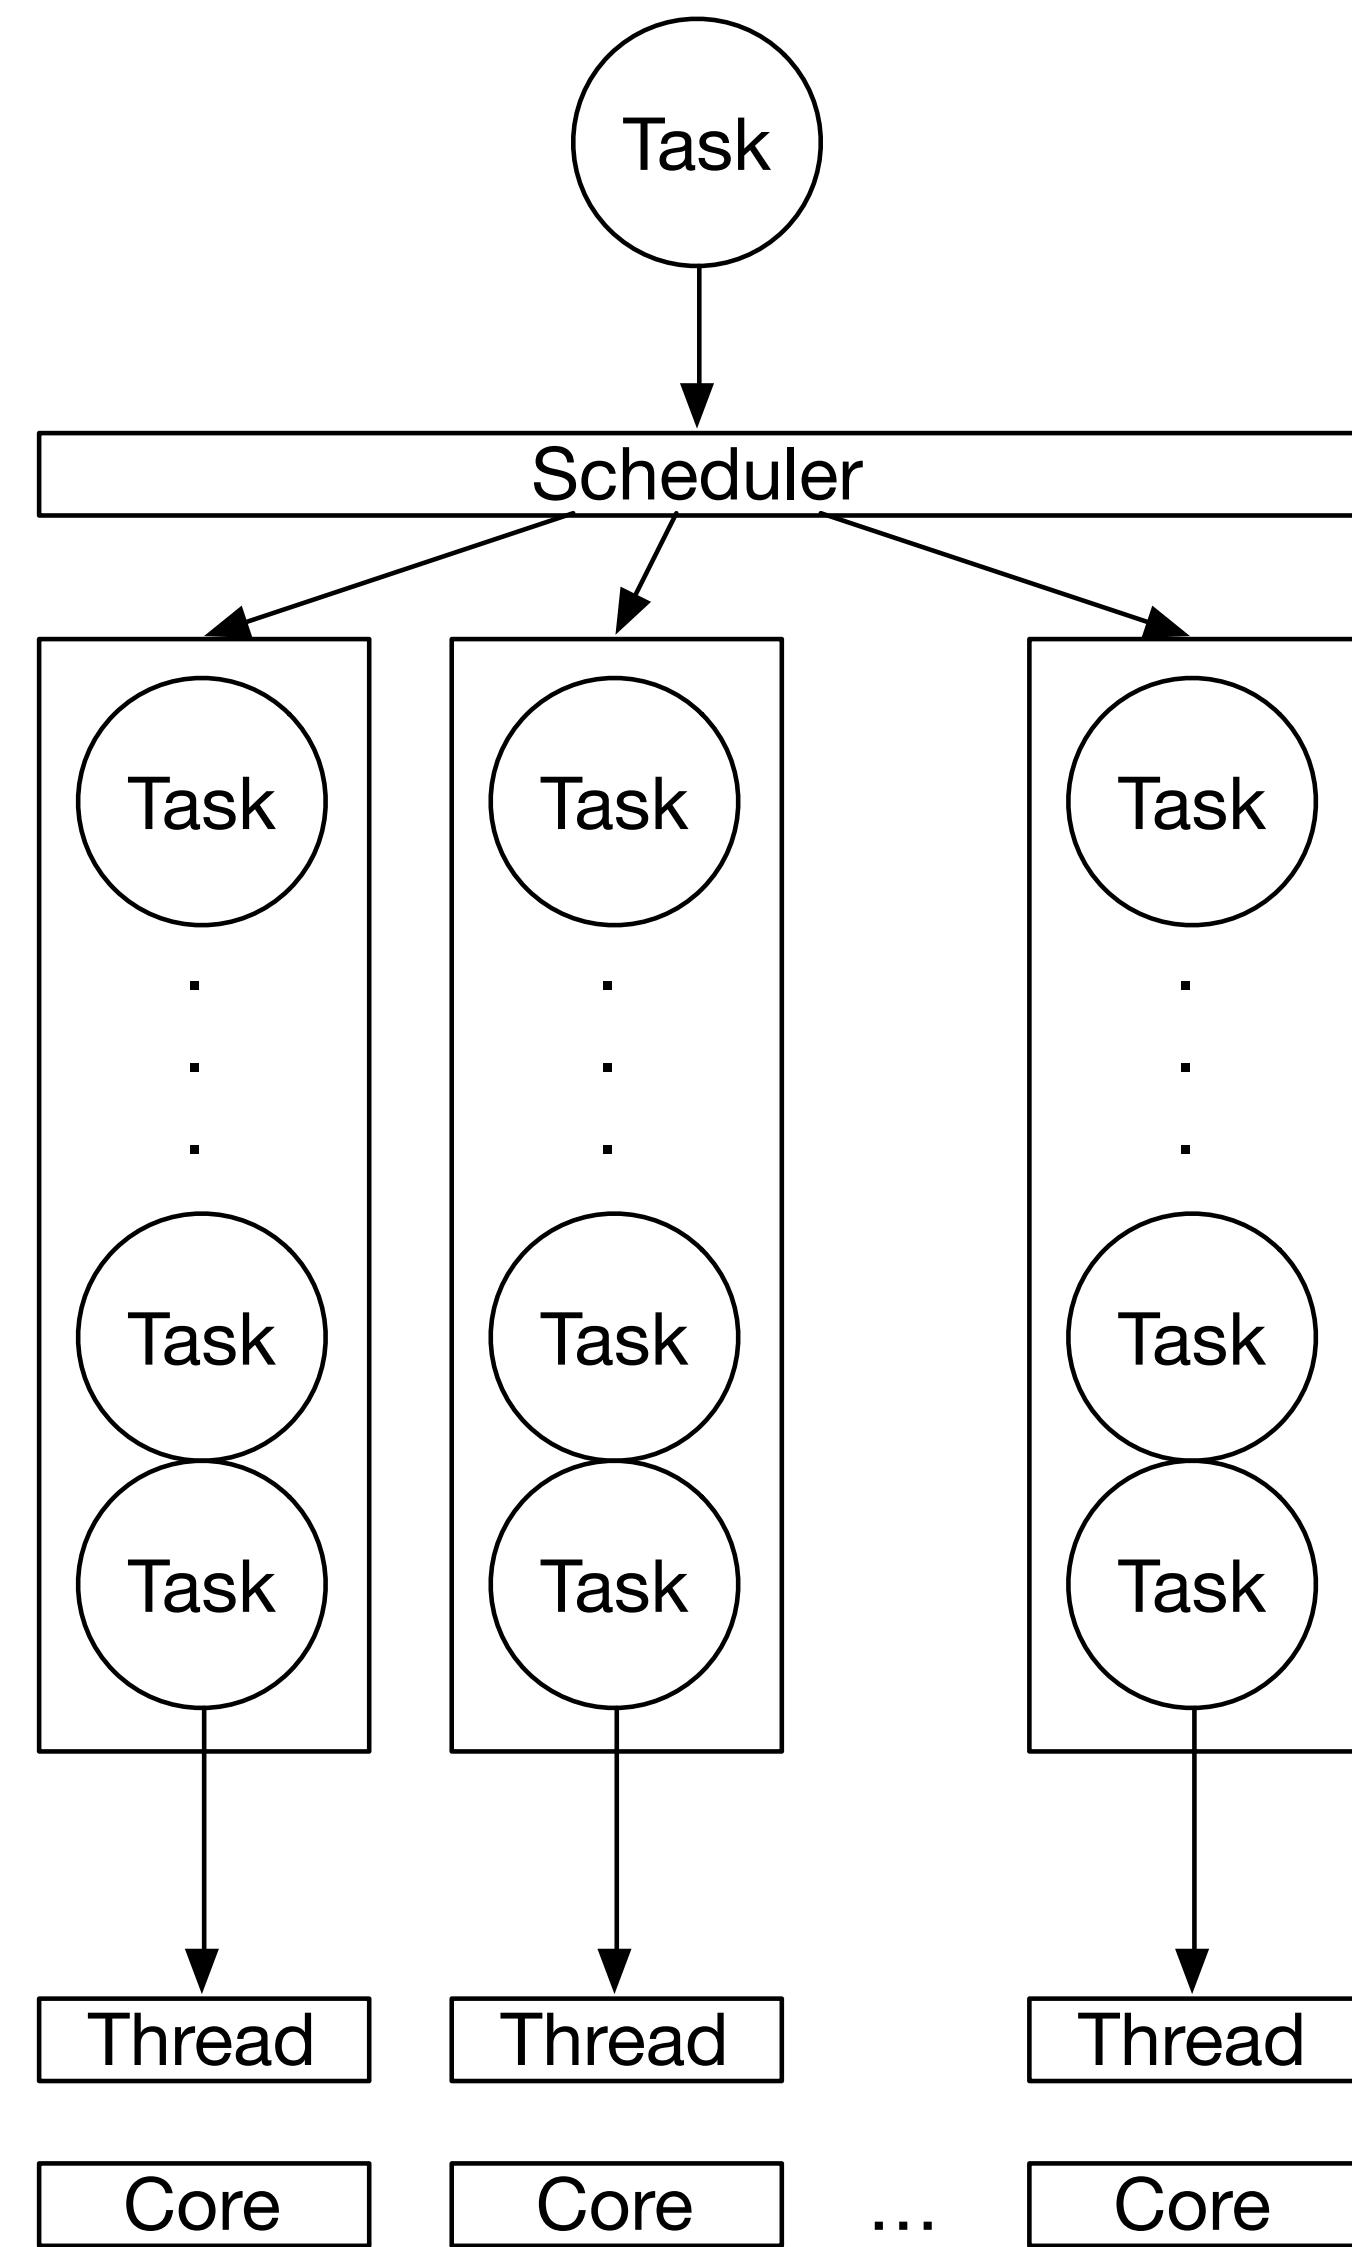
    void run(unsigned i) {
        while (true) {
            function<void()> f;
            if (!_q[i].pop(f)) break;
            f();
        }
    }

public:
    task_system() { ... }

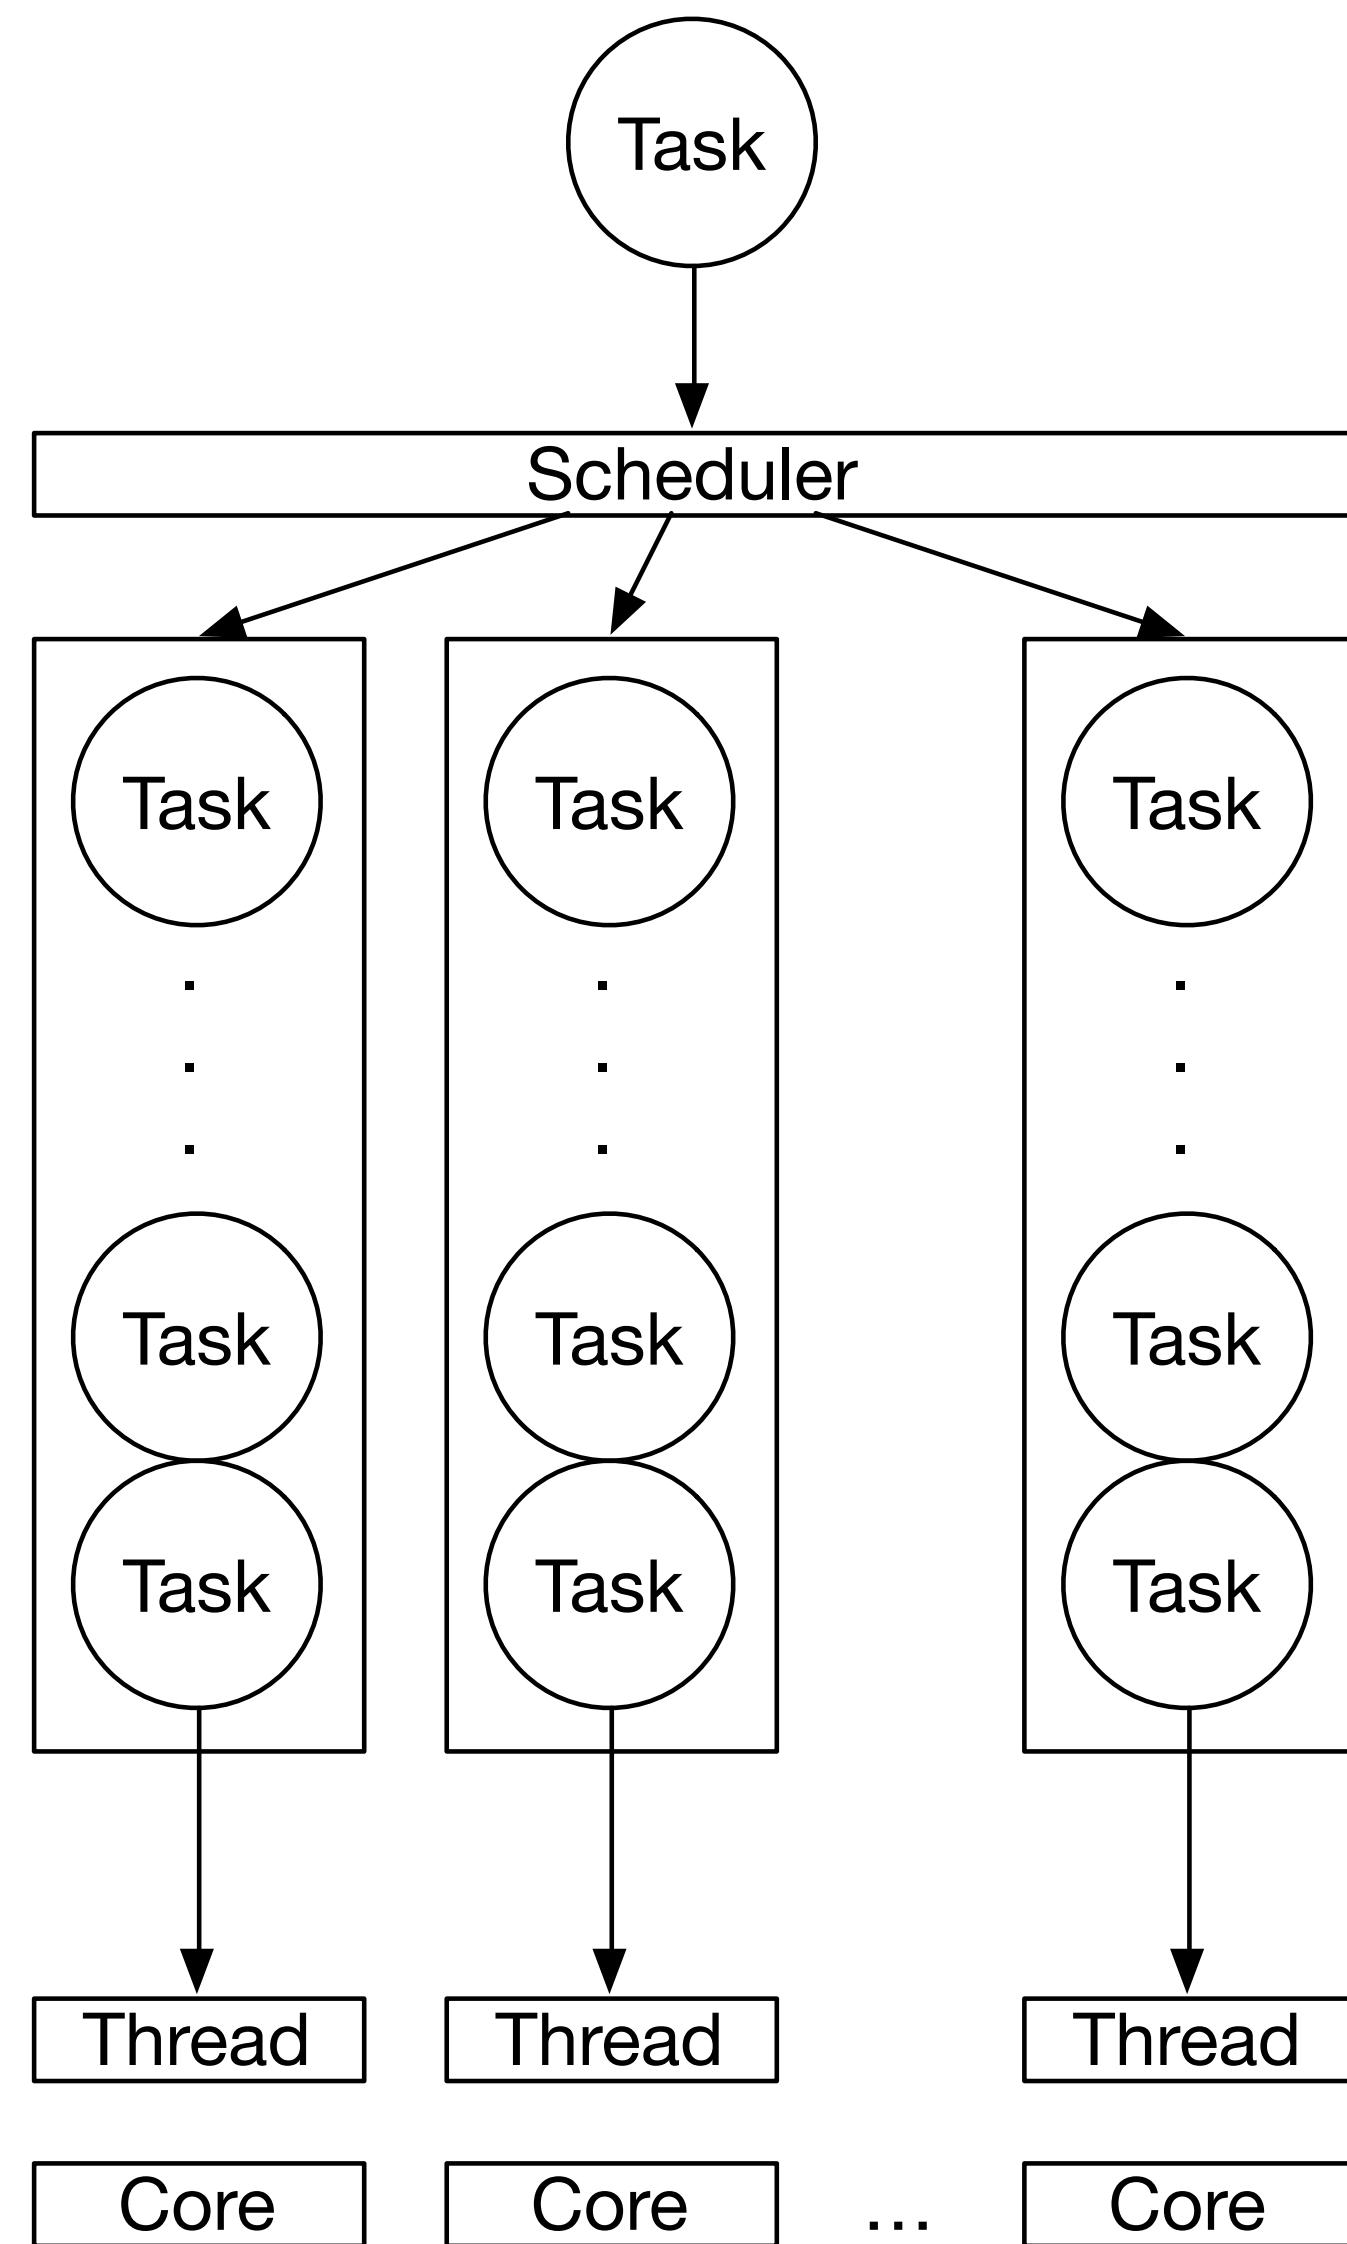
    ~task_system() {
        for (auto& e : _q) e.done();
        for (auto& e : _threads) e.join();
    }

    template <typename F>
    void async_(F&& f) {
        auto i = _index++;
        _q[i % _count].push(forward<F>(f));
    }
};
```

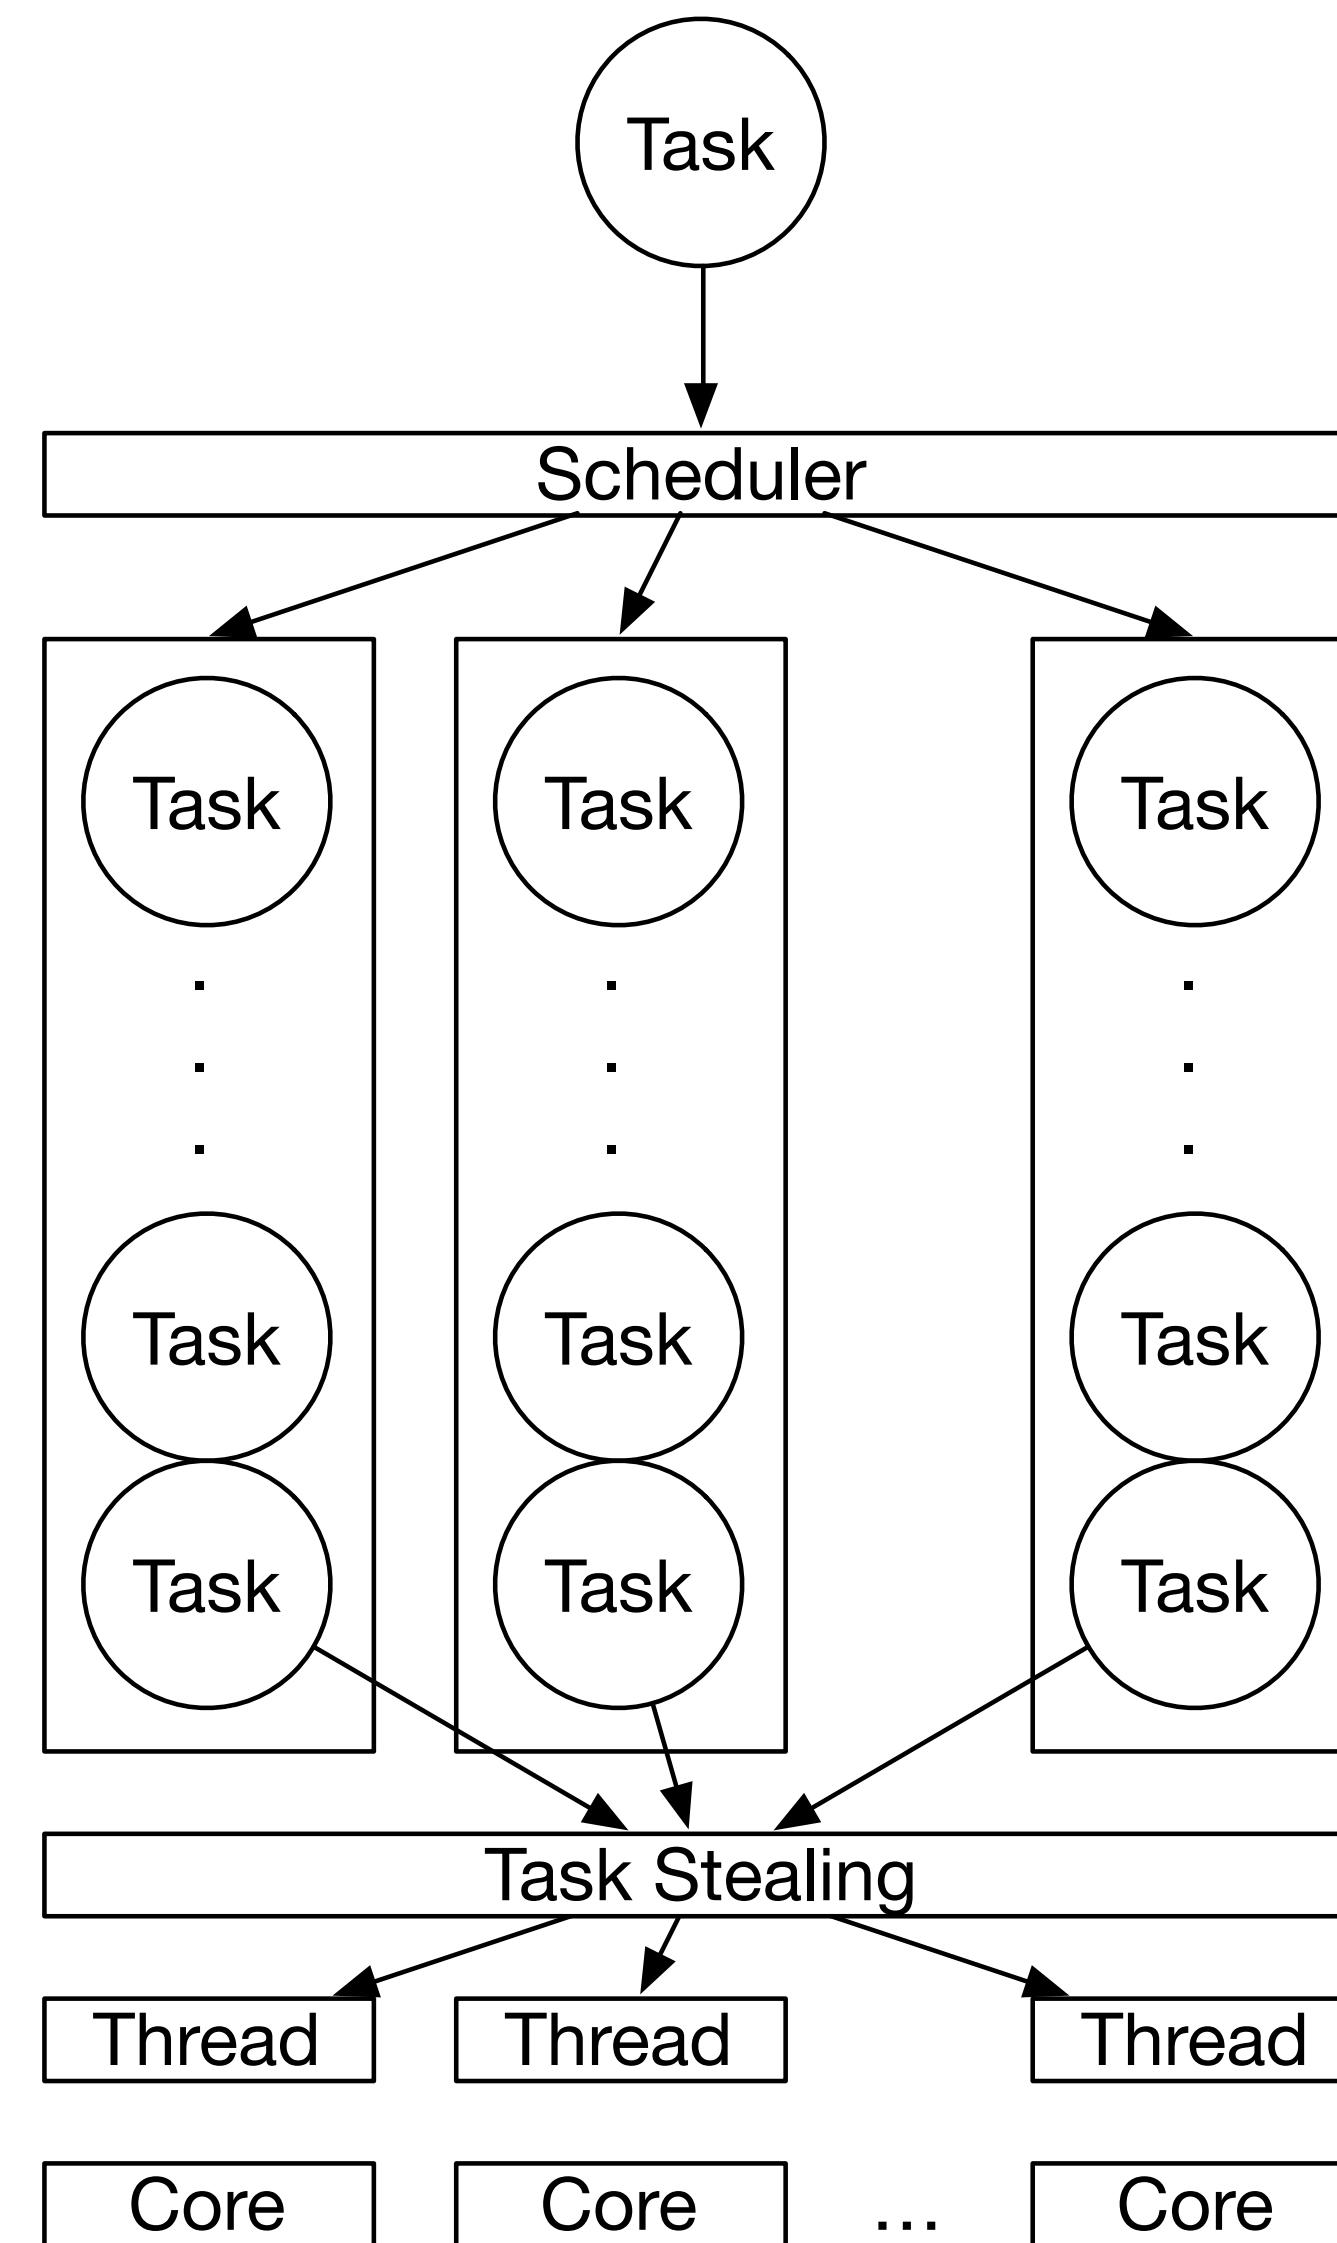
Building a Task System



Building a Task System



Building a Task System



Building a Task System

```
class notification_queue {
    deque<function<void()>> _q;
    bool _done{false};
    mutex _mutex;
    condition_variable _ready;

public:
    bool try_pop(function<void()>& x) {
        lock_t lock{_mutex, try_to_lock};
        if (!lock || !_q.empty()) return false;
        x = move(_q.front());
        _q.pop_front();
        return true;
    }

    template<typename F>
    bool try_push(F&& f) {
        {
            lock_t lock{_mutex, try_to_lock};
            if (!lock) return false;
            _q.emplace_back(forward<F>(f));
        }
        _ready.notify_one();
        return true;
    }

    void done() {
    {
        unique_lock<mutex> lock{_mutex};
```

Building a Task System

```
class notification_queue {
    deque<function<void()>> _q;
    bool _done{false};
    mutex _mutex;
    condition_variable _ready;

public:
    bool try_pop(function<void()>& x) {
        lock_t lock{_mutex, try_to_lock};
        if (!lock || !_q.empty()) return false;
        x = move(_q.front());
        _q.pop_front();
        return true;
    }

    template<typename F>
    bool try_push(F&& f) {
        {
            lock_t lock{_mutex, try_to_lock};
            if (!lock) return false;
            _q.emplace_back(forward<F>(f));
        }
        _ready.notify_one();
        return true;
    }

    void done() {
    {
        unique_lock<mutex> lock{_mutex};
```

Building a Task System

```
class notification_queue {
    deque<function<void()>> _q;
    bool _done{false};
    mutex _mutex;
    condition_variable _ready;

public:
    bool try_pop(function<void()>& x) {
        lock_t lock{_mutex, try_to_lock};
        if (!lock || !_q.empty()) return false;
        x = move(_q.front());
        _q.pop_front();
        return true;
    }

    template<typename F>
    bool try_push(F&& f) {
        {
            lock_t lock{_mutex, try_to_lock};
            if (!lock) return false;
            _q.emplace_back(forward<F>(f));
        }
        _ready.notify_one();
        return true;
    }

    void done() {
    {
        unique_lock<mutex> lock{_mutex};
```

Building a Task System

```
void run(unsigned i) {
    while (true) {
        function<void()> f;

        for (unsigned n = 0; n != _count; ++n) {
            if (_q[(i + n) % _count].try_pop(f)) break;
        }
        if (!f && !_q[i].pop(f)) break;

        f();
    }
}

public:
task_system() { ... }

~task_system() { ... }

template <typename F>
void async_(F&& f) {
    auto i = _index++;

    for (unsigned n = 0; n != _count * K; ++n) {
        if (_q[(i + n) % _count].try_push(forward<F>(f))) return;
    }

    _q[i % _count].push(forward<F>(f));
}
};
```

Building a Task System

```
void run(unsigned i) {
    while (true) {
        function<void()> f;

        for (unsigned n = 0; n != _count; ++n) {
            if (_q[(i + n) % _count].try_pop(f)) break;
        }
        if (!f && !_q[i].pop(f)) break;

        f();
    }
}

public:
task_system() { ... }

~task_system() { ... }

template <typename F>
void async_(F&& f) {
    auto i = _index++;

    for (unsigned n = 0; n != _count * K; ++n) {
        if (_q[(i + n) % _count].try_push(forward<F>(f))) return;
    }

    _q[i % _count].push(forward<F>(f));
}
};
```

Building a Task System

```
void run(unsigned i) {
    while (true) {
        function<void()> f;

        for (unsigned n = 0; n != _count; ++n) {
            if (_q[(i + n) % _count].try_pop(f)) break;
        }
        if (!f && !_q[i].pop(f)) break;

        f();
    }
}

public:
task_system() { ... }

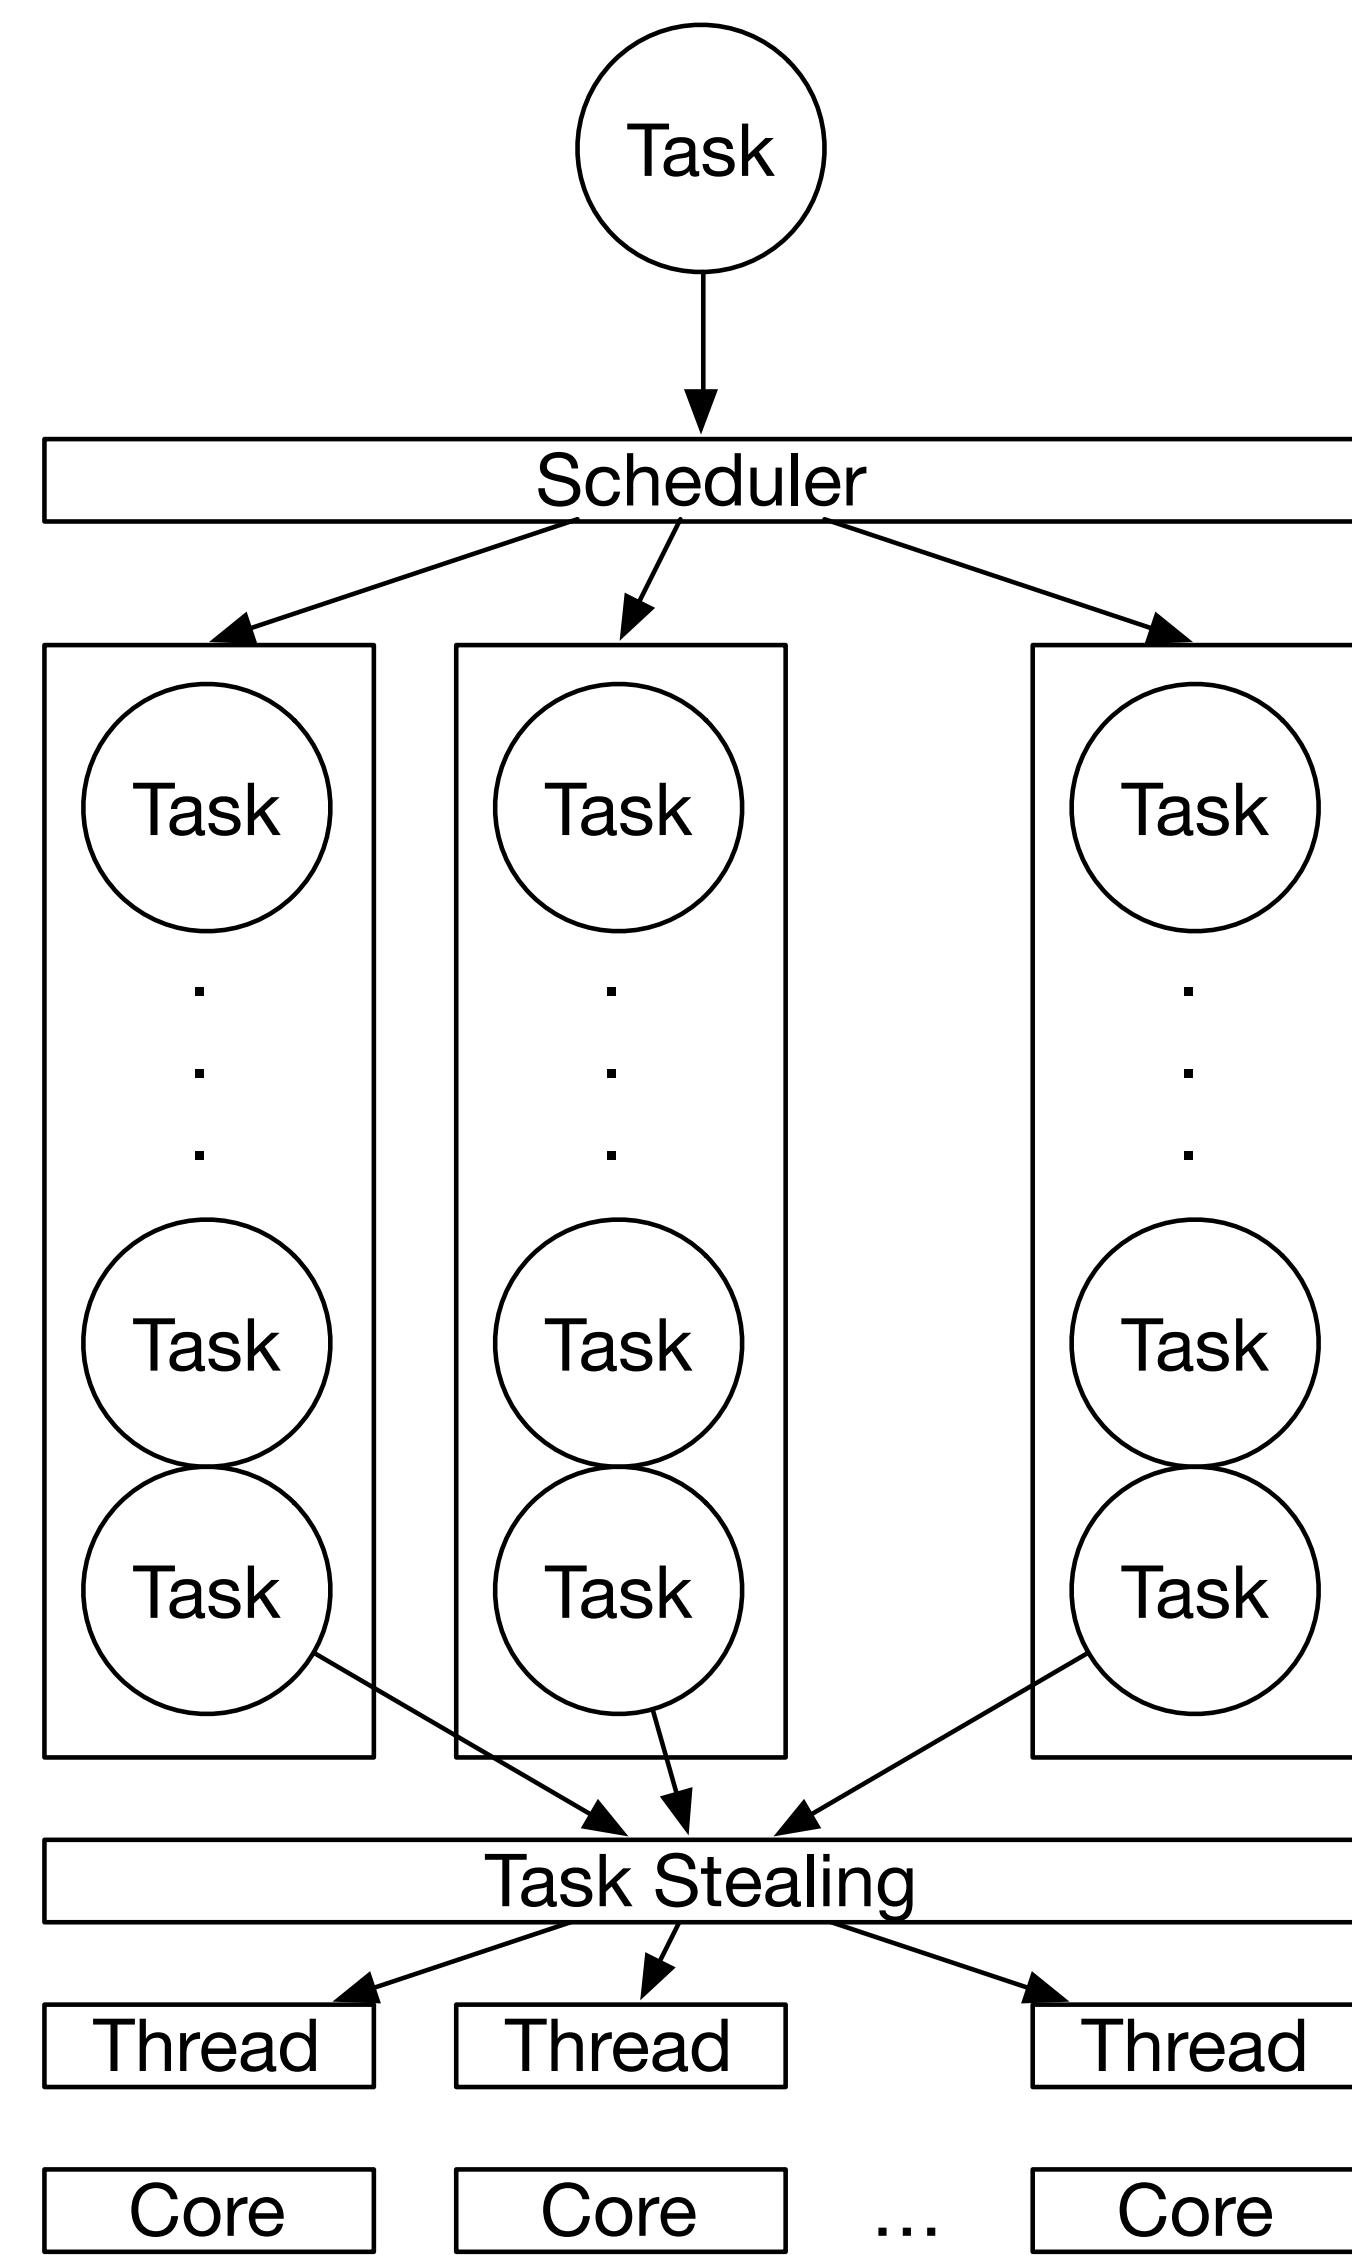
~task_system() { ... }

template <typename F>
void async_(F&& f) {
    auto i = _index++;

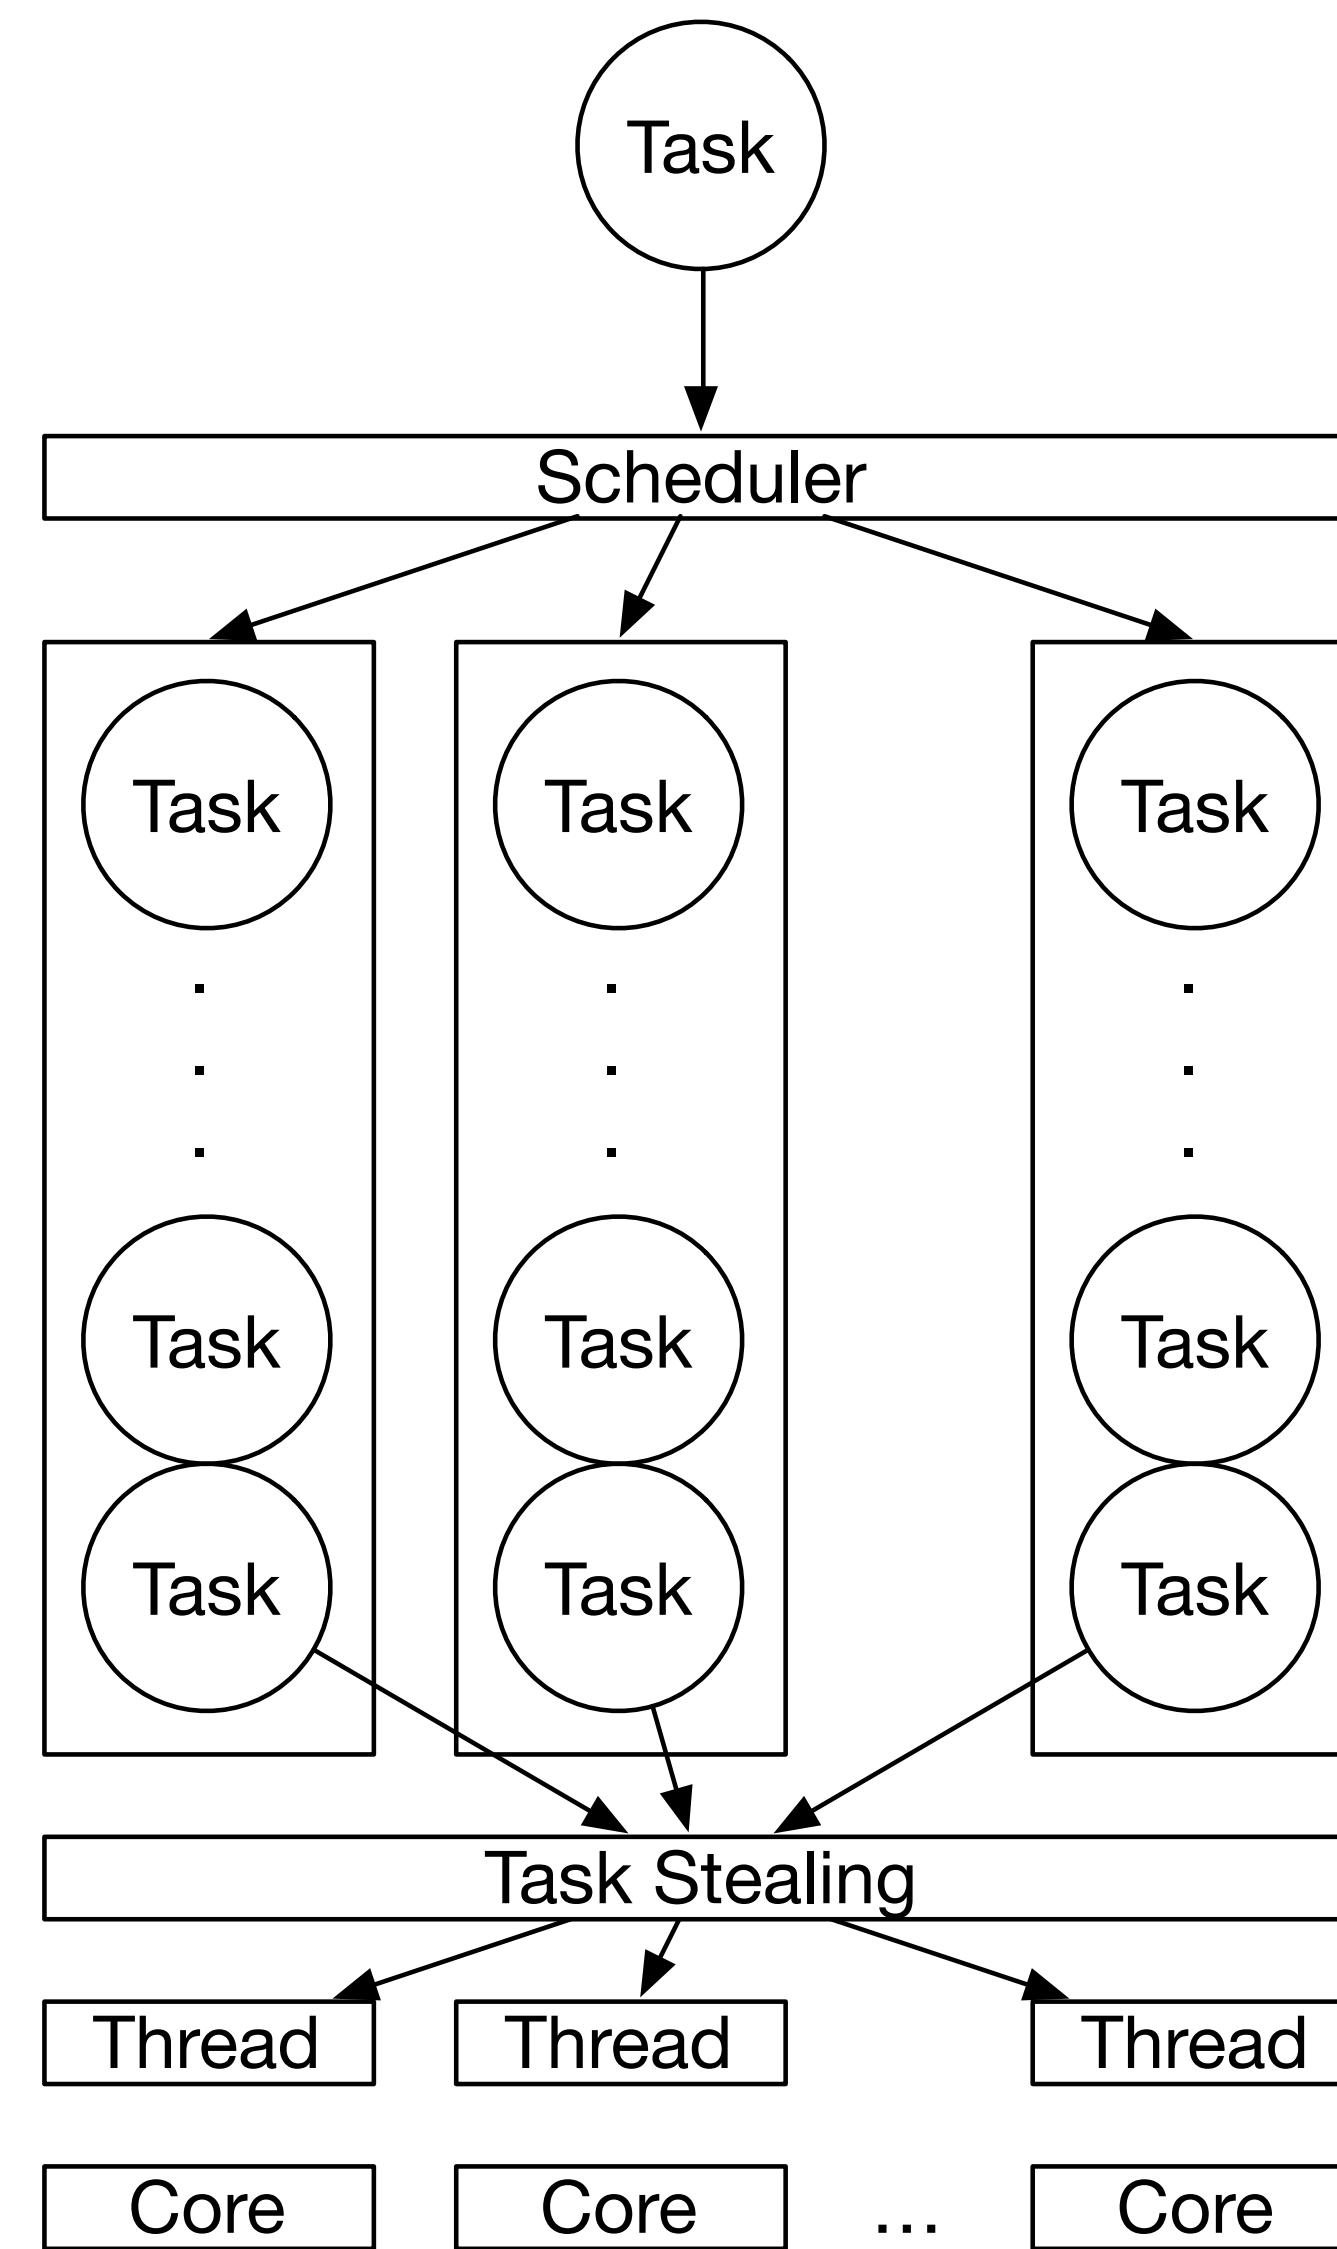
    for (unsigned n = 0; n != _count * K; ++n) {
        if (_q[(i + n) % _count].try_push(forward<F>(f))) return;
    }

    _q[i % _count].push(forward<F>(f));
}
};
```

Building a Task System



Building a Task System



Task System

- Compared to Apple's Grand Central Dispatch (libdispatch)



Task System

- Compared to Apple's Grand Central Dispatch (libdispatch)



C++14 compatible async with libdispatch

```
#include <functional>
#include <future>
#include <type_traits>

#include <dispatch/dispatch.h>

namespace stlab {

template <class Function, class... Args>
auto async(Function&& f, Args&&... args )
{
    using result_type = std::result_of_t<std::decay_t<Function>(std::decay_t<Args>...)>;
    using packaged_type = std::packaged_task<result_type()>;

    auto p = new packaged_type(std::bind(std::forward<Function>(f), std::forward<Args>(args)...));
    auto result = p->get_future();

    dispatch_async_f(dispatch_get_global_queue(DISPATCH_QUEUE_PRIORITY_DEFAULT, 0),
                    p, [](void* f_) {
                        packaged_type* f = static_cast<packaged_type*>(f_);
                        (*f)();
                        delete f;
                    });
}

return result;
}

} // namespace stlab
```

Task System

- Written with ASIO (Boost 1.62.0)

```
class task_system {  
    io_service           _service;  
    vector<thread>     _threads;  
    unique_ptr<io_service::work> _work{make_unique<io_service::work>(_service)};  
  
public:  
    task_system() {  
        for (unsigned n = 0; n != thread::hardware_concurrency(); ++n) {  
            _threads.emplace_back([&]{  
                _service.run();  
            });  
        }  
    }  
  
    ~task_system() {  
        _work.reset();  
        for (auto& e : _threads) e.join();  
    }  
  
    template <typename F>  
    void async_(F&& f) {  
        _service.post(forward<F>(f));  
    }  
};
```



Task System

- Written with ASIO (Boost 1.62.0)

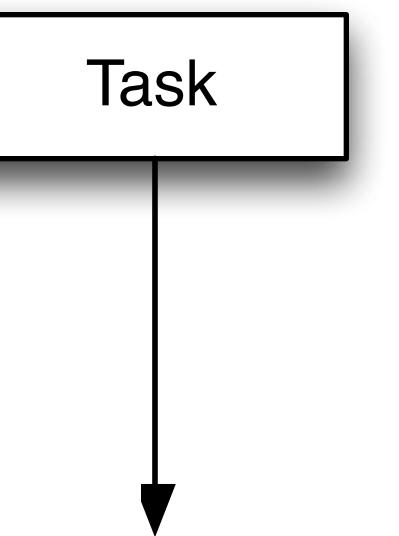
```
class task_system {  
    io_service           _service;  
    vector<thread>     _threads;  
    unique_ptr<io_service::work> _work{make_unique<io_service::work>(_service)};  
  
public:  
    task_system() {  
        for (unsigned n = 0; n != thread::hardware_concurrency(); ++n) {  
            _threads.emplace_back([&]{  
                _service.run();  
            });  
        }  
    }  
  
    ~task_system() {  
        _work.reset();  
        for (auto& e : _threads) e.join();  
    }  
  
    template <typename F>  
    void async_(F&& f) {  
        _service.post(forward<F>(f));  
    }  
};
```



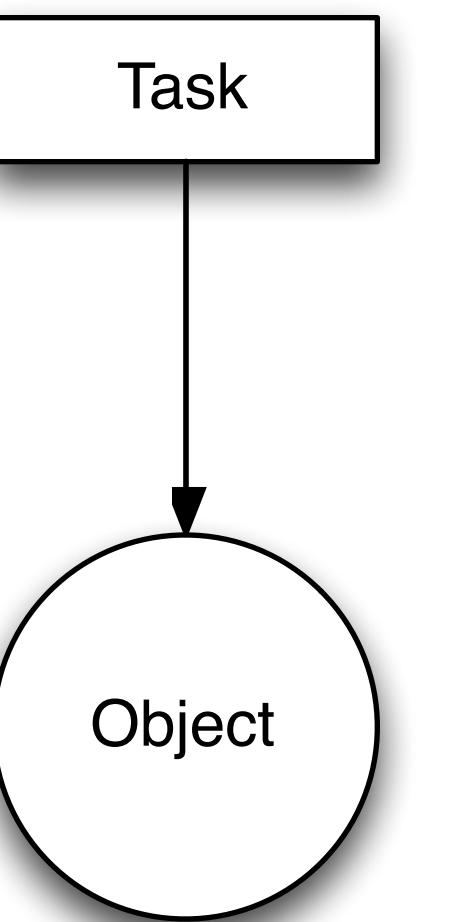
No Raw Synchronization Primitives

Task

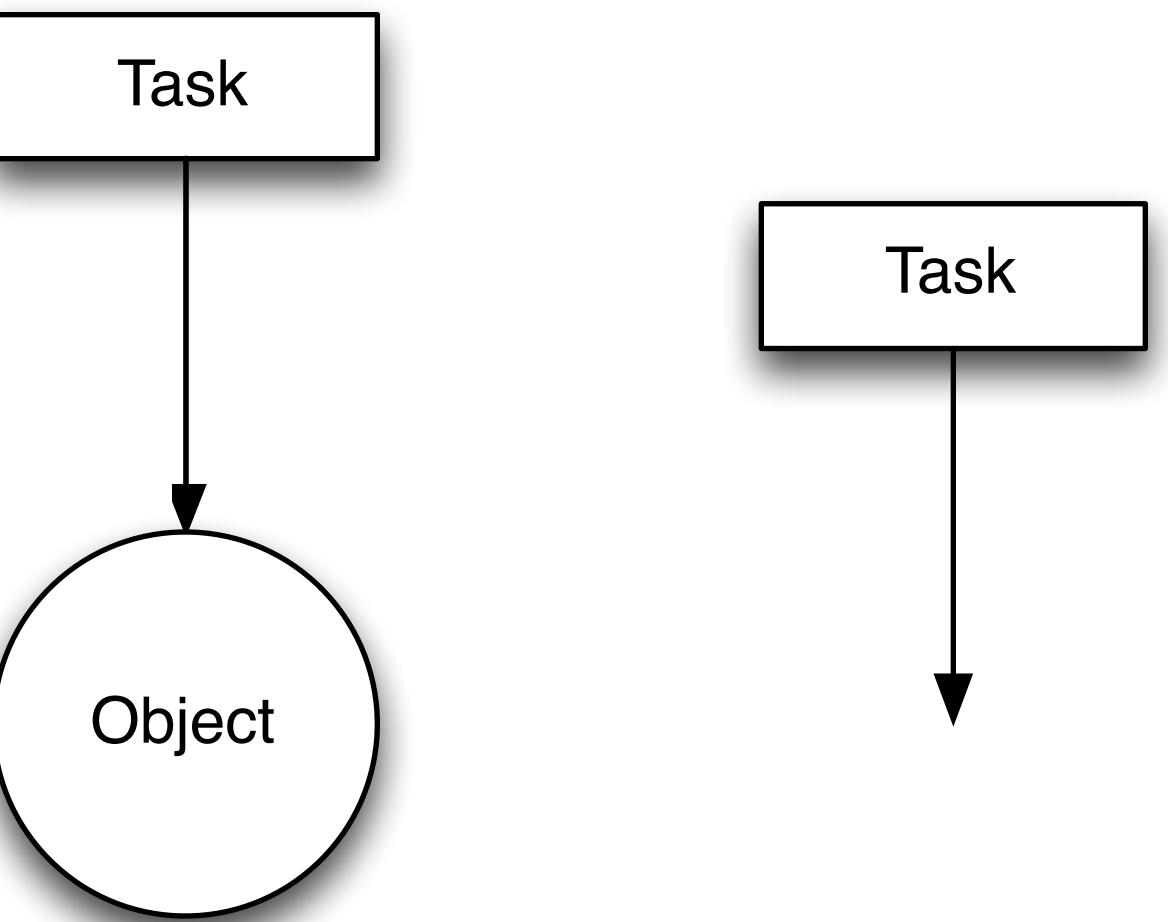
No Raw Synchronization Primitives



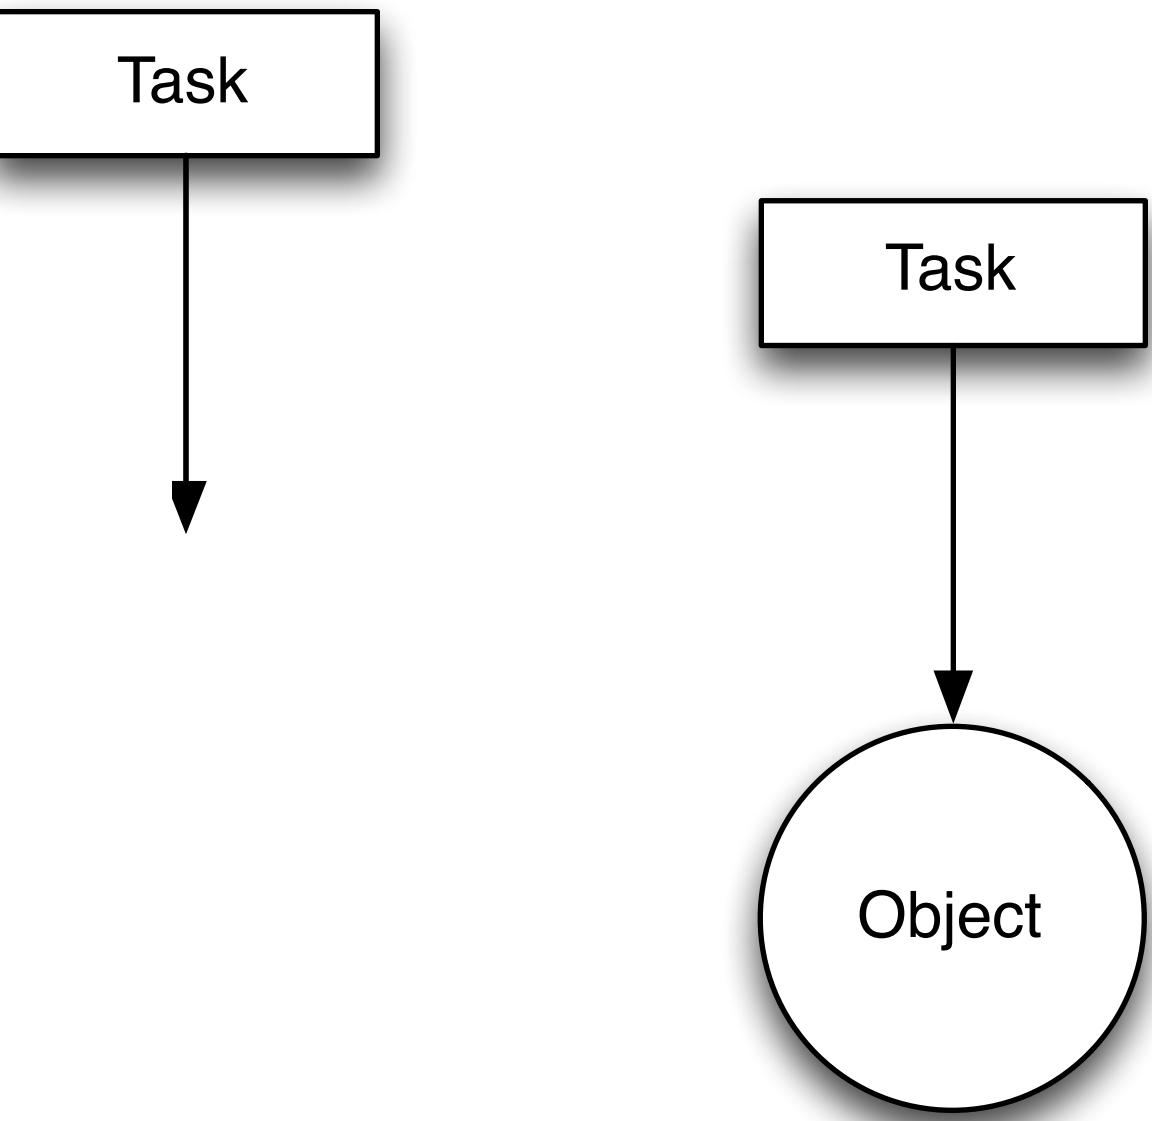
No Raw Synchronization Primitives



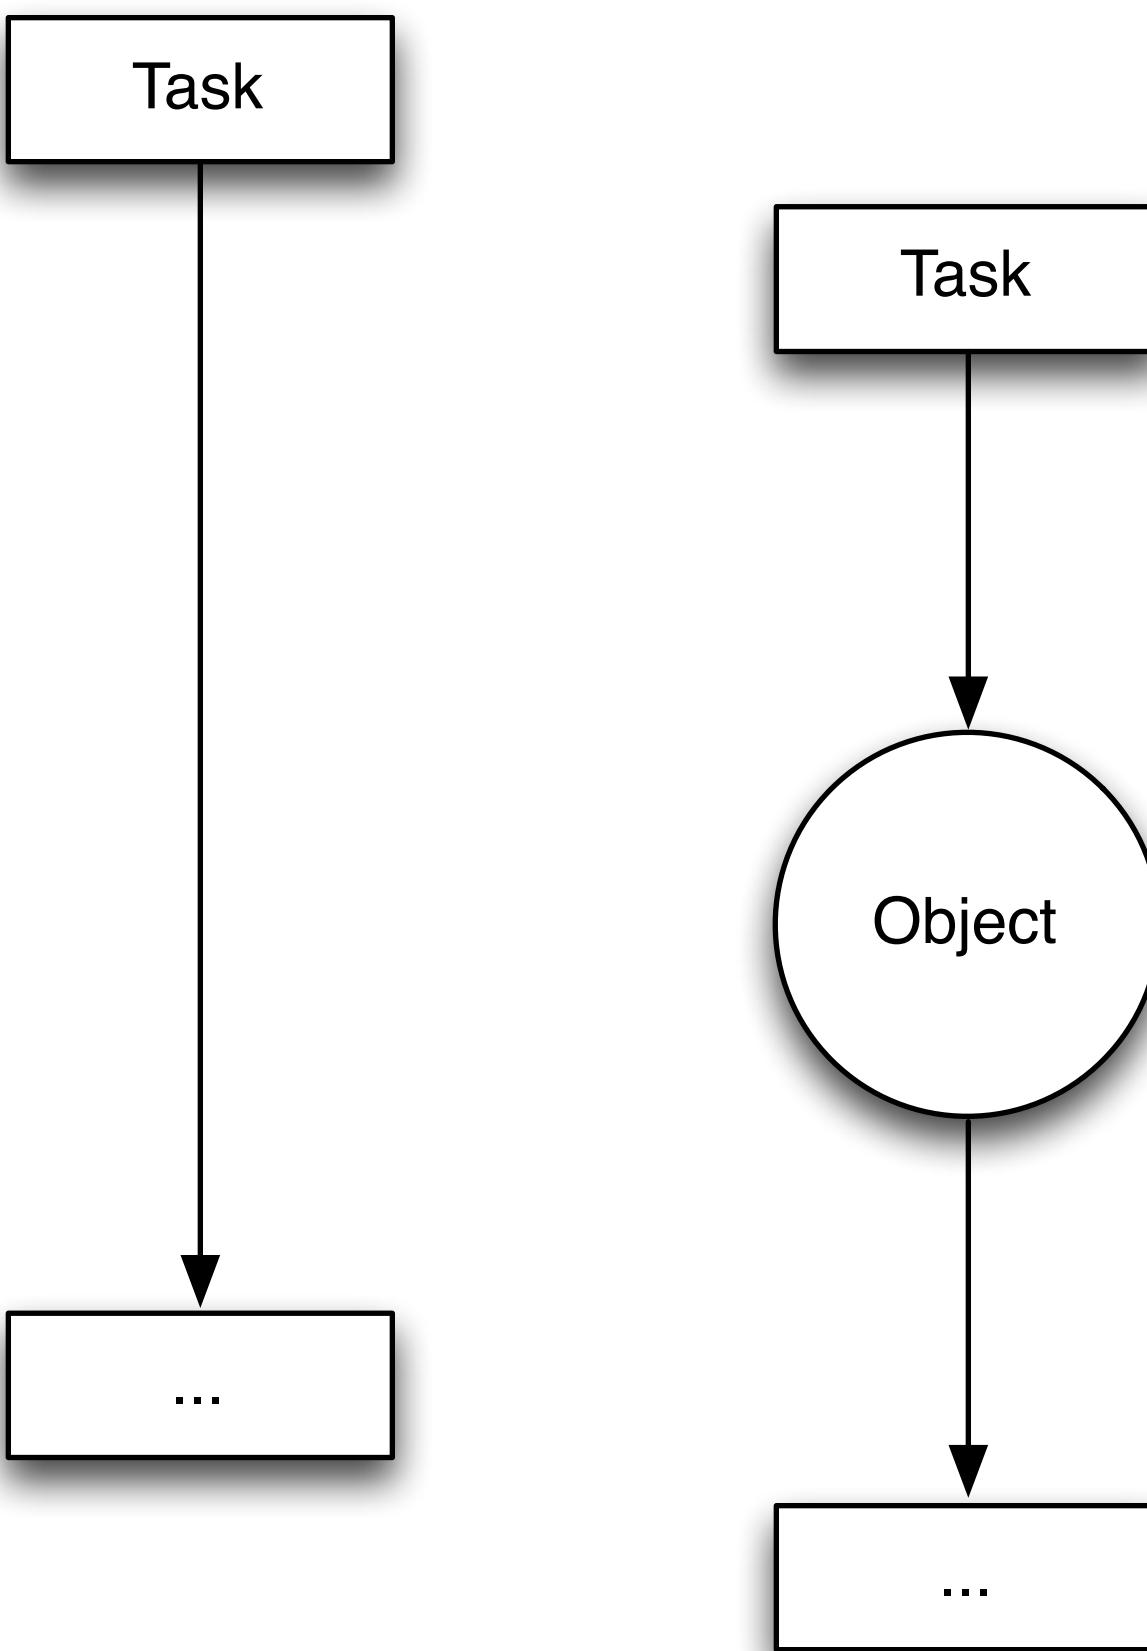
No Raw Synchronization Primitives



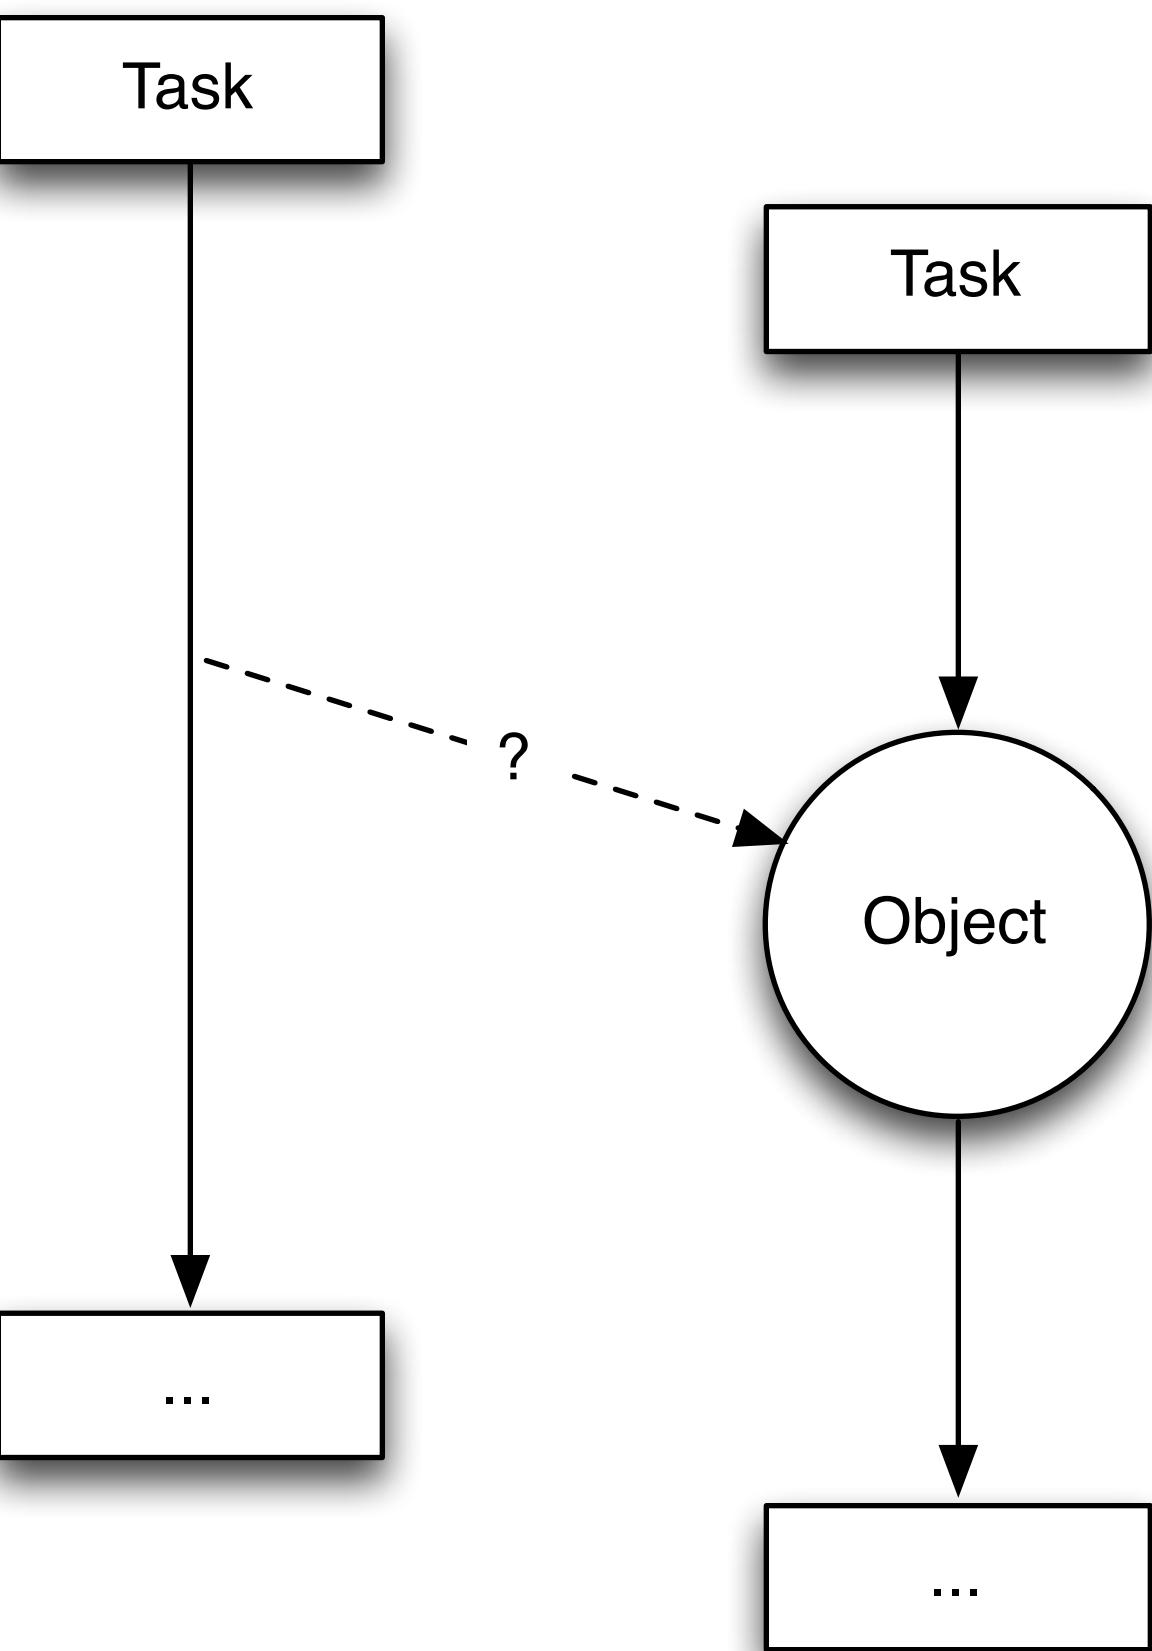
No Raw Synchronization Primitives



No Raw Synchronization Primitives



No Raw Synchronization Primitives



Futures

```
future<cpp_int> x = async( []{ return fibonacci<cpp_int>(1'000'000); });

// Do Something

cout << x.get() << endl;
```

Futures

```
future<cpp_int> x = async([]{ return fibonacci<cpp_int>(1'000'000); });

// Do Something

cout << x.get() << endl;
```

- Fibonacci is often used as an example for parallel algorithms
 - Please stop...

Public Service Announcement - How to Write Fibonacci

```
template <typename T, typename N, typename O>
T power(T x, N n, O op)
{
    if (n == 0) return identity_element(op);

    while ((n & 1) == 0) {
        n >>= 1;
        x = op(x, x);
    }

    T result = x;
    n >>= 1;
    while (n != 0) {
        x = op(x, x);
        if ((n & 1) != 0) result = op(result, x);
        n >>= 1;
    }
    return result;
}
```

Public Service Announcement - How to Write Fibonacci

```
template <typename T, typename N, typename O>
T power(T x, N n, O op)
{
    if (n == 0) return identity_element(op);

    while ((n & 1) == 0) {
        n >>= 1;
        x = op(x, x);
    }

    T result = x;
    n >>= 1;
    while (n != 0) {
        x = op(x, x);
        if ((n & 1) != 0) result = op(result, x);
        n >>= 1;
    }
    return result;
}
```

Egyptian Multiplication (Russian Peasant Algorithm)

See "From Mathematics to Generic Programming" - Alex Stepanov and Dan Rose

Public Service Announcement - How to Write Fibonacci

```
template <typename N>
struct multiply_2x2 {
    array<N, 4> operator()(const array<N, 4>& x, const array<N, 4>& y)
{
    return { x[0] * y[0] + x[1] * y[2], x[0] * y[1] + x[1] * y[3],
             x[2] * y[0] + x[3] * y[2], x[2] * y[1] + x[3] * y[3] };
}
};

template <typename N>
array<N, 4> identity_element(const multiply_2x2<N>&) { return { N(1), N(0), N(0), N(1) }; }

template <typename R, typename N>
R fibonacci(N n) {
    if (n == 0) return R(0);
    return power(array<R, 4>{ 1, 1, 1, 0 }, N(n - 1), multiply_2x2<R>())[0];
}
```

Public Service Announcement - How to Write Fibonacci

```
template <typename N>
struct multiply_2x2 {
    array<N, 4> operator()(const array<N, 4>& x, const array<N, 4>& y)
{
    return { x[0] * y[0] + x[1] * y[2], x[0] * y[1] + x[1] * y[3],
             x[2] * y[0] + x[3] * y[2], x[2] * y[1] + x[3] * y[3] };
}
};

template <typename N>
array<N, 4> identity_element(const multiply_2x2<N>&) { return { N(1), N(0), N(0), N(1) }; }

template <typename R, typename N>
R fibonacci(N n) {
    if (n == 0) return R(0);
    return power(array<R, 4>{ 1, 1, 1, 0 }, N(n - 1), multiply_2x2<R>())[0];
}
```

$$\begin{bmatrix} 1 & 1 \\ 1 & 0 \end{bmatrix}^n = \begin{bmatrix} F_{n+1} & F_n \\ F_n & F_{n-1} \end{bmatrix}$$

Futures

19532821287077577316320149475962563324435429965918733969534051945716252578870156947666419876341501461288795
24335220236084625510912019560233744015438115196636156919962125642894303370113827800638002767411527927466669
86557837931882283206127149758323033485489348957259923072291290192820926433162752173086146001791258204269965
9936020959339202005184862028402447343139811367418720203868480175318538621128781082406177413832935545616876
06454065125954718029126547942894036981659206361019359291352135410376799082940320155702716115395031975973247
78216295763162965335669477766328506234524559346064757502593581344345781676764625878859011372729907372947851
1448089572456191503507025589529116868550088020132334587472177947814475467920160901706425856293597475465327
57575740077432034913428785189795354304734560307765078938767286539166799232817449361991523768149557632085371
04785970618843873153058239562756087906310781900497516959470973671389174570455520213512335079440336071203050
41446852210415650373210679322756258647511914611417360349681217380234224786080292021093192496490409832397066
83247054441763512526732455275419501683845206023007394959854279298297831204382115757645787692495583351402522
1527206624418090032593807536284917966809529711850719137983367887377045991363933955814212036990261617972113
2509184002305532760710431647819097430043464779336328760146999612802392582947155731668894339455429292871877
48774789204296166356536610796023919702109728472966709427334586344798048633944635211654971507261342768205479
32093175079888010130416027982506354182344034558742236701282666356934611294613123128389060036547327660245693
15151850018328483150645480029978935985161237074046158229354440701748339514575869547491750264542126364262224
72060048855462589961190475892101224280542898621594646662478564373572217775549876087685912030118551635668902
01034463998397732663888903650784161807091545252992759735213957415477729146008794314339156060445825107823511
66271892637923313014643880597879468444879060576786297460989627426663569682474293386740207436559426057944790
71193052258931590719386545525880429139747140181849169733838138446154843063123649290835584278078456131936457
55911722136946338180311600307896211668652895953778436464402382516362449718197385444149563131714002850338928
22274134603018094224837216321854717270452813824078425638747365249141118080783866506339945376239206700513391
87333107136069698189628284763245423299306272870457991293245741167533902274499963096566680922262516468582544
55785134982414412726124015815753818098466667145006988839178551800894370189025721992485208742915560261917752
28124660628996787166529678487268484905041328497297712688011639978376434280202452251550102240354169885185375
01584673881194047619720619603126534496759917893244478170702904446589571950228809157793897642423751814020998
99958161231477902295781100168670186738619861797138139854666281969548553740707356228616165539428076418408092
12047932816683005984504787929406356318097479755152035094682765918741610907637506902765294367561539803261388
9019448594100452922754188094573562079542189966296341346396955980991375010053760259440113617219176881147364

1579015589283310034567384624310467690000936756893803676769777642059716492347060997973282994459039755683869
10568541105888505197986232161807165960864316652383369579251545877324797429523572491518310013505994095431367
Futures
2345441853967639642257048786844333673556851153585056517249014177233018072390350689838662532338266203548476
87722321662223383305226882245421258277211223435986491973881404168406609216954760818955479619408040043497601
35646408461148077885537891122888139618703907906033147416881433658136276942006644505679690480702792206520855
12245086839375655196861305232092138041808273198852928058246964575561801618520046644949262341864859342928965
21378574554544426221453176445385228867960454072522804961741905198550911362542849130027243353553345377968558
4978019597663651629059845721904348982135822120685692412113931313713213486574144089267000366555632446499775
56853514681289887391700907057970839124191923062570547772748610990924519168225326823578140721238189631411471
29610287340041050015549547086272721534936510345705849389706515725684266079756708385889612130516276472992631
59674474594901199950849178952149715987731953191759591623424021718579696778102054496598766846143959650647332
21985323521378108187030642875506951890343587181633604126397675020909133548480151135951824112432636080497447
37395896608759569909256138919905403404664655310556021101996525724843421071082933739200159651403373870955680
75656822683537933983982488022723770319785461480932302347255796621173892988541730741484707211664044157057536
0458225614322429985978068323969654385523783781413866750792868372058020433472254190336846843017198934115689
96526838242546875

0.72s to calculate
208,988 digits

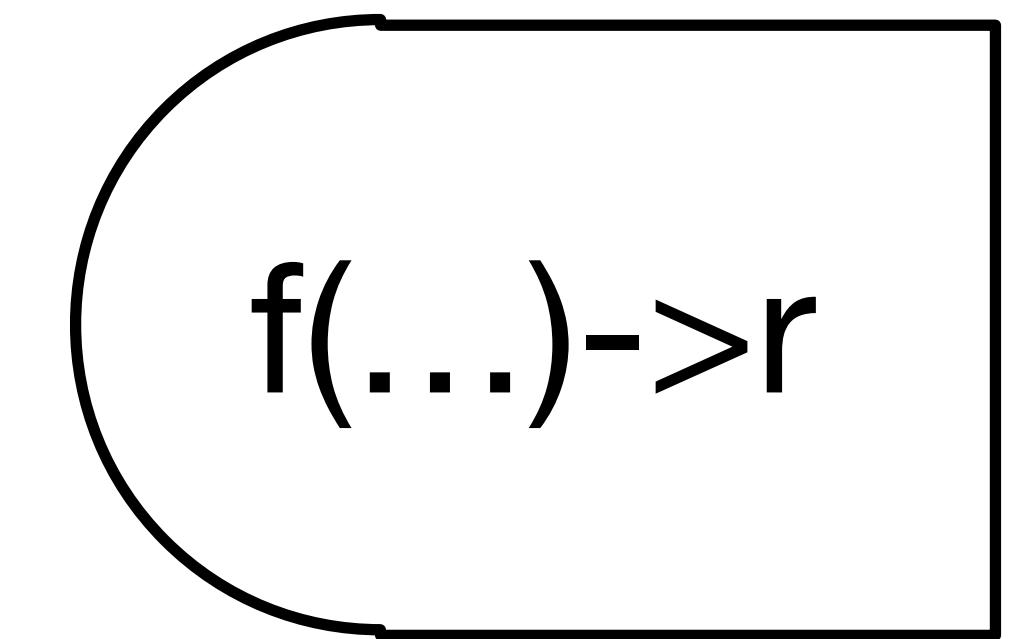
Futures

```
future<cpp_int> x = async( []{ return fibonacci<cpp_int>(1'000'000); });

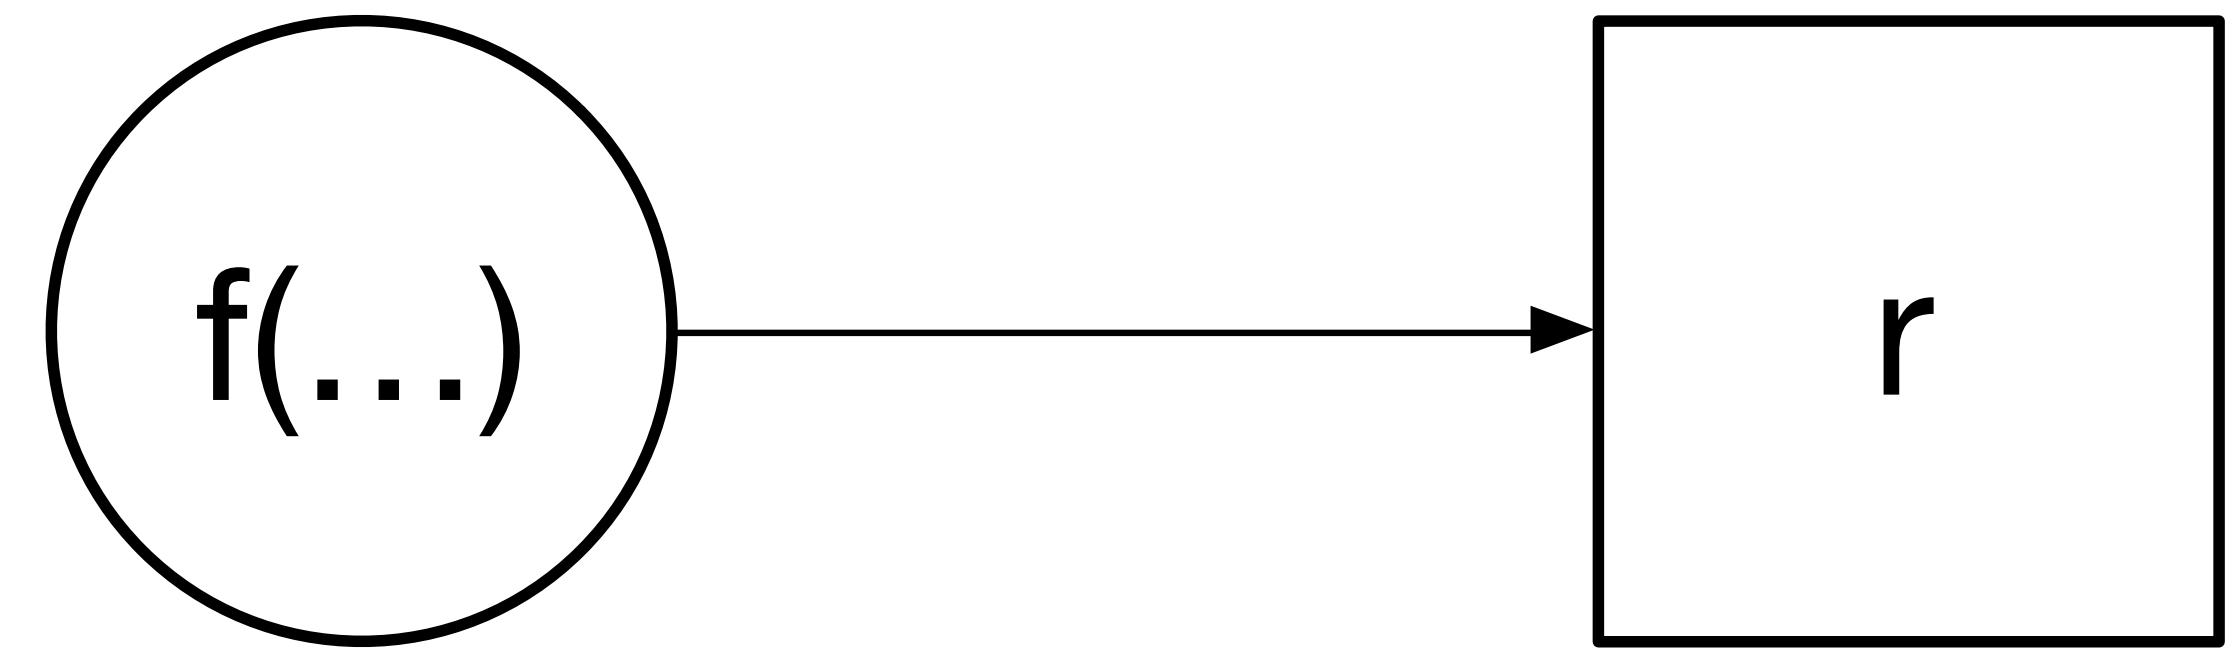
// Do Something

cout << x.get() << endl;
```

Futures



Futures



Futures

- Futures allow minimal code transformations to express dependencies

Exception Marshalling

```
future<cpp_int> x = async([]{
    throw runtime_error("failure");
    return fibonacci<cpp_int>(1'000'000);
});

// Do Something

try {
    cout << x.get() << endl;
} catch (const runtime_error& error) {
    cout << error.what() << endl;
}
```

Exception Marshalling

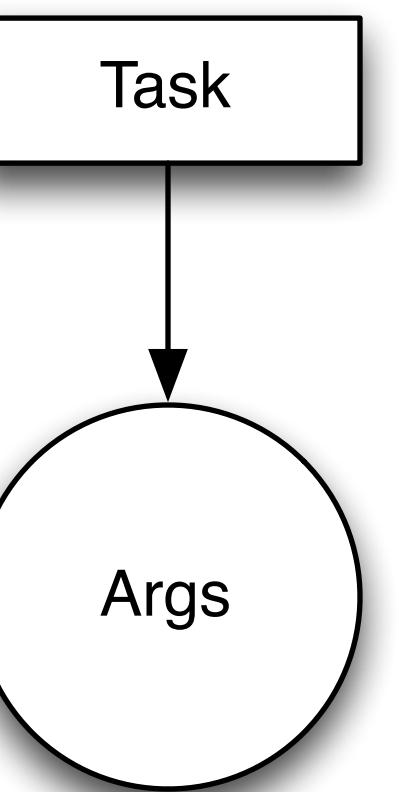
```
future<cpp_int> x = async([]{
    throw runtime_error("failure");
    return fibonacci<cpp_int>(1'000'000);
});

// Do Something

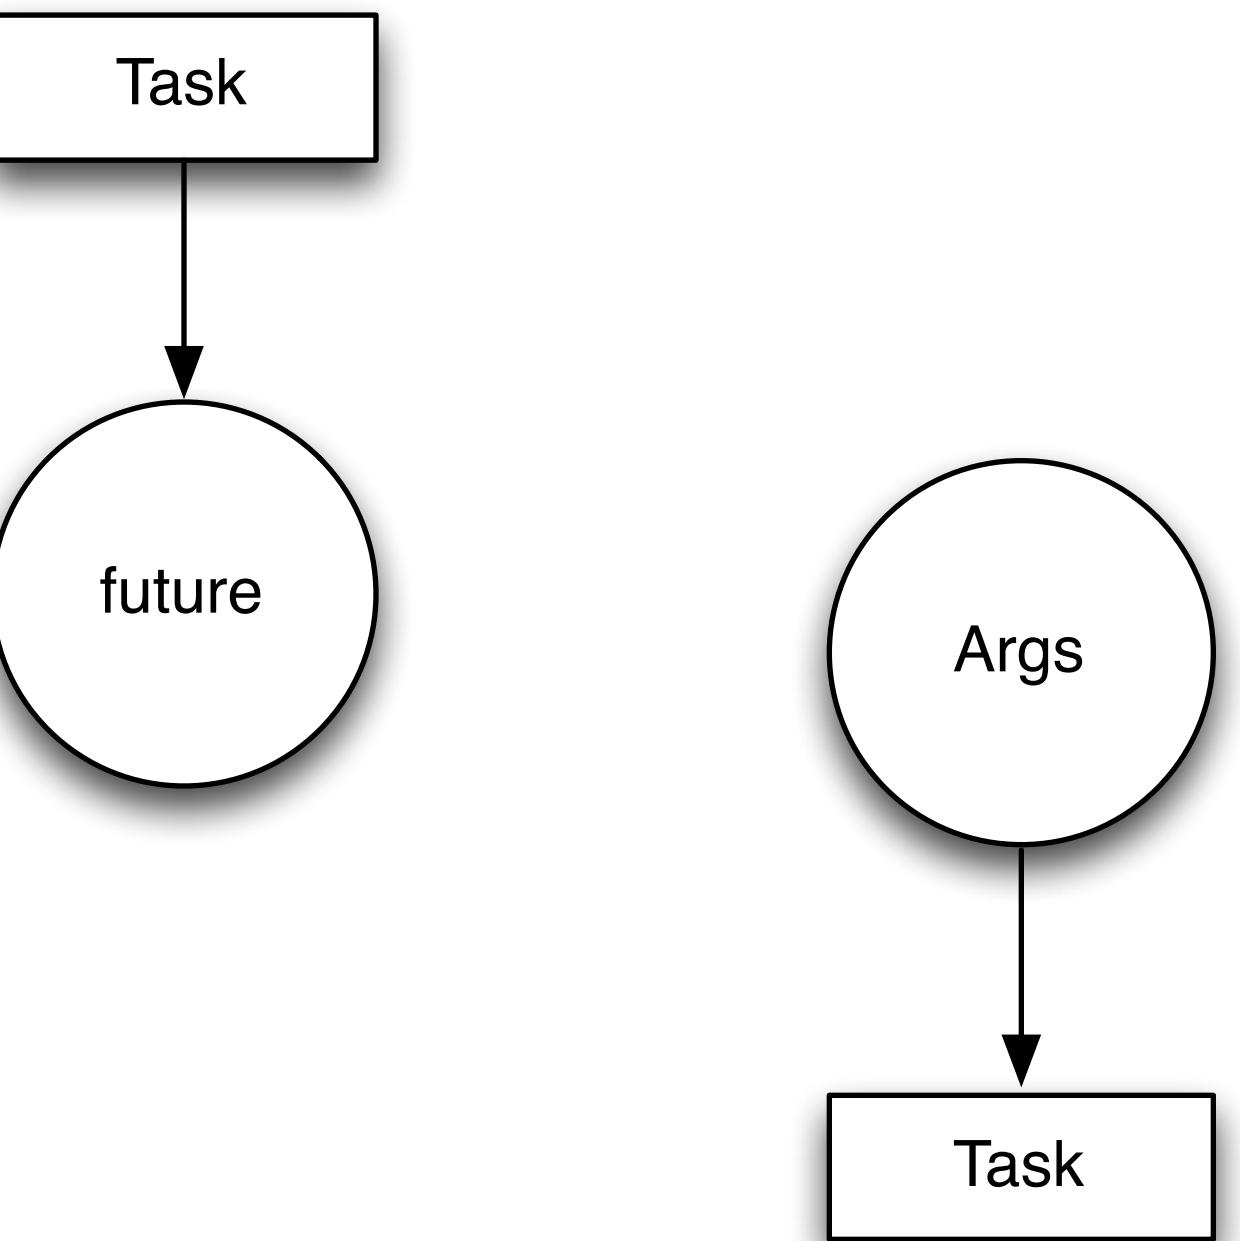
try {
    cout << x.get() << endl;
} catch (const runtime_error& error) {
    cout << error.what() << endl;
}
```

failure

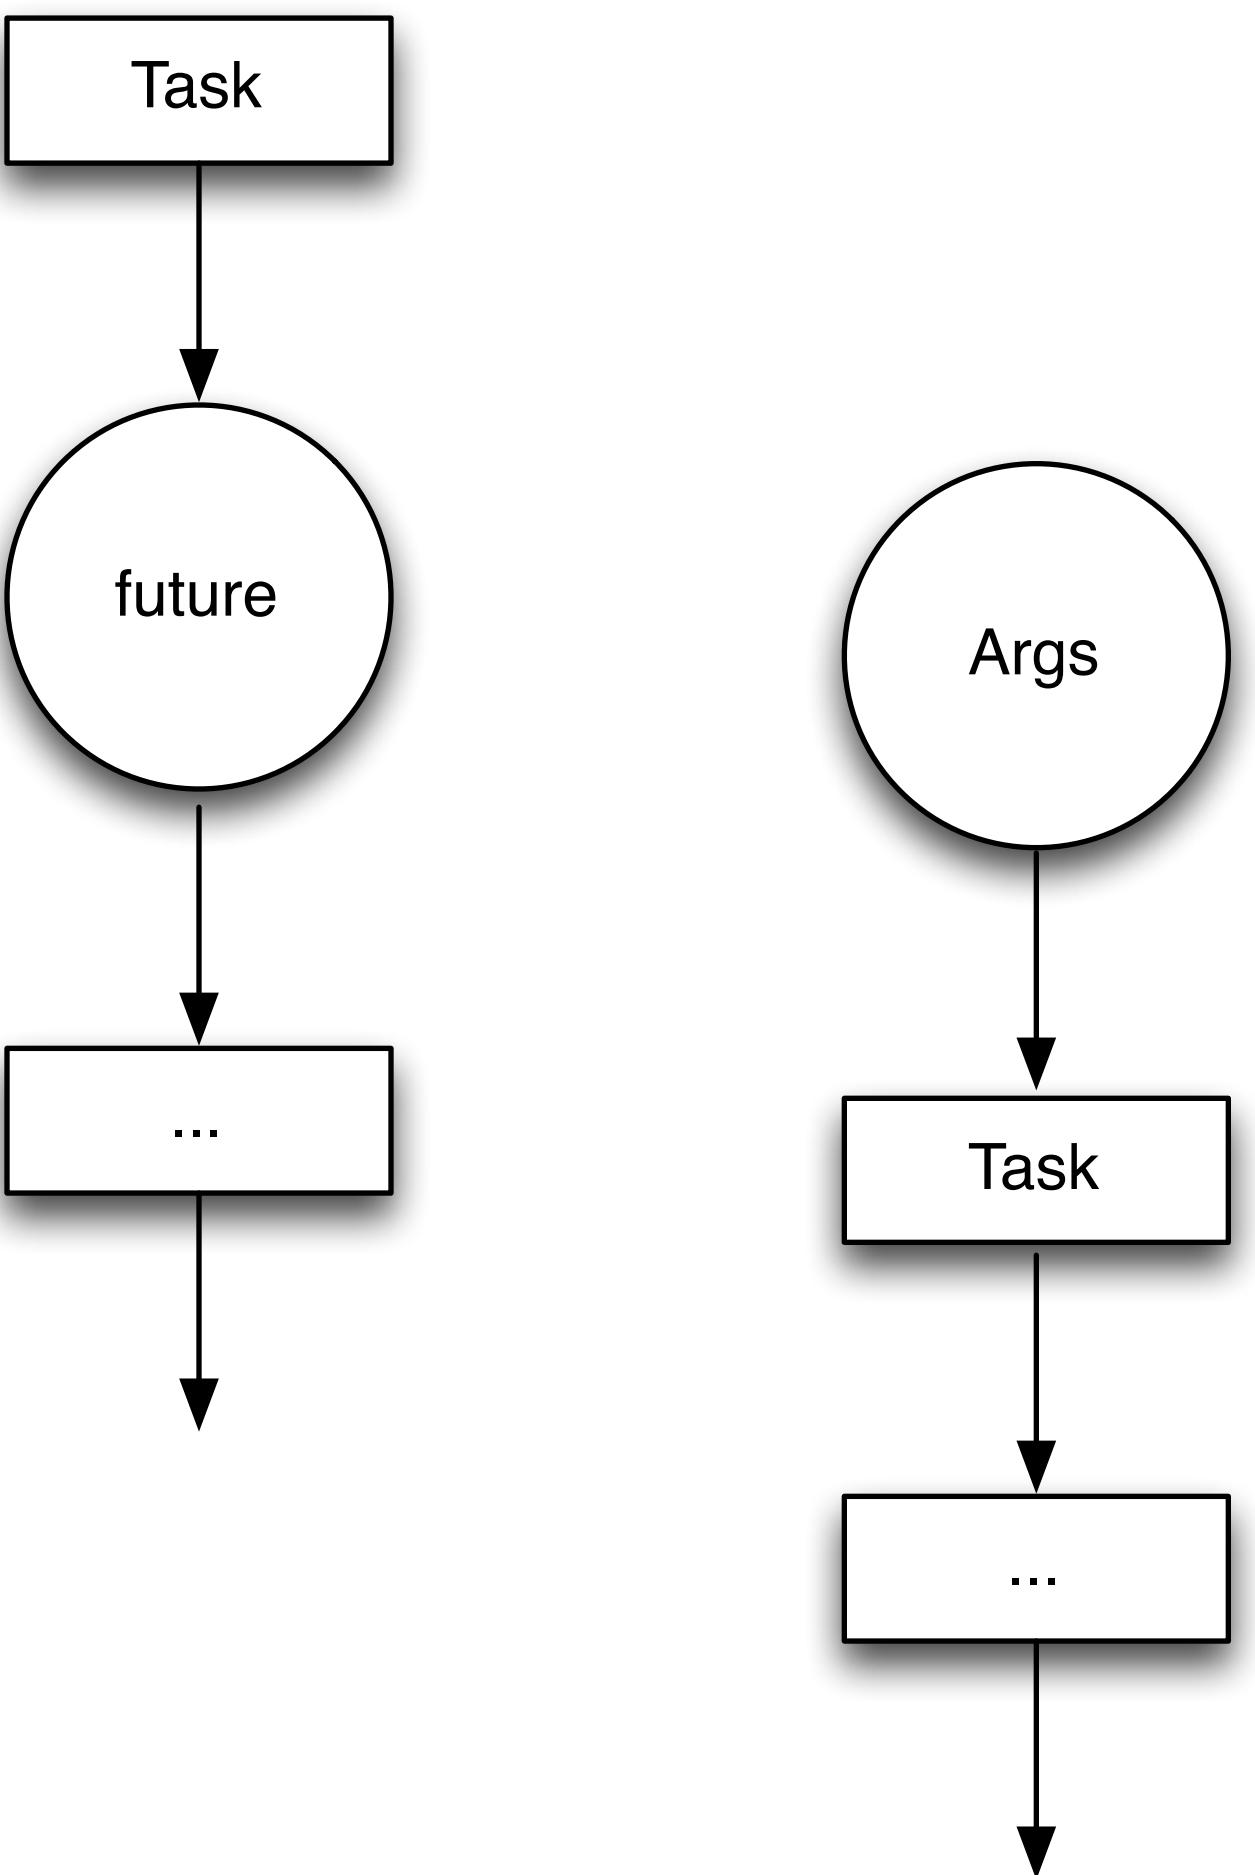
No Raw Synchronization Primitives



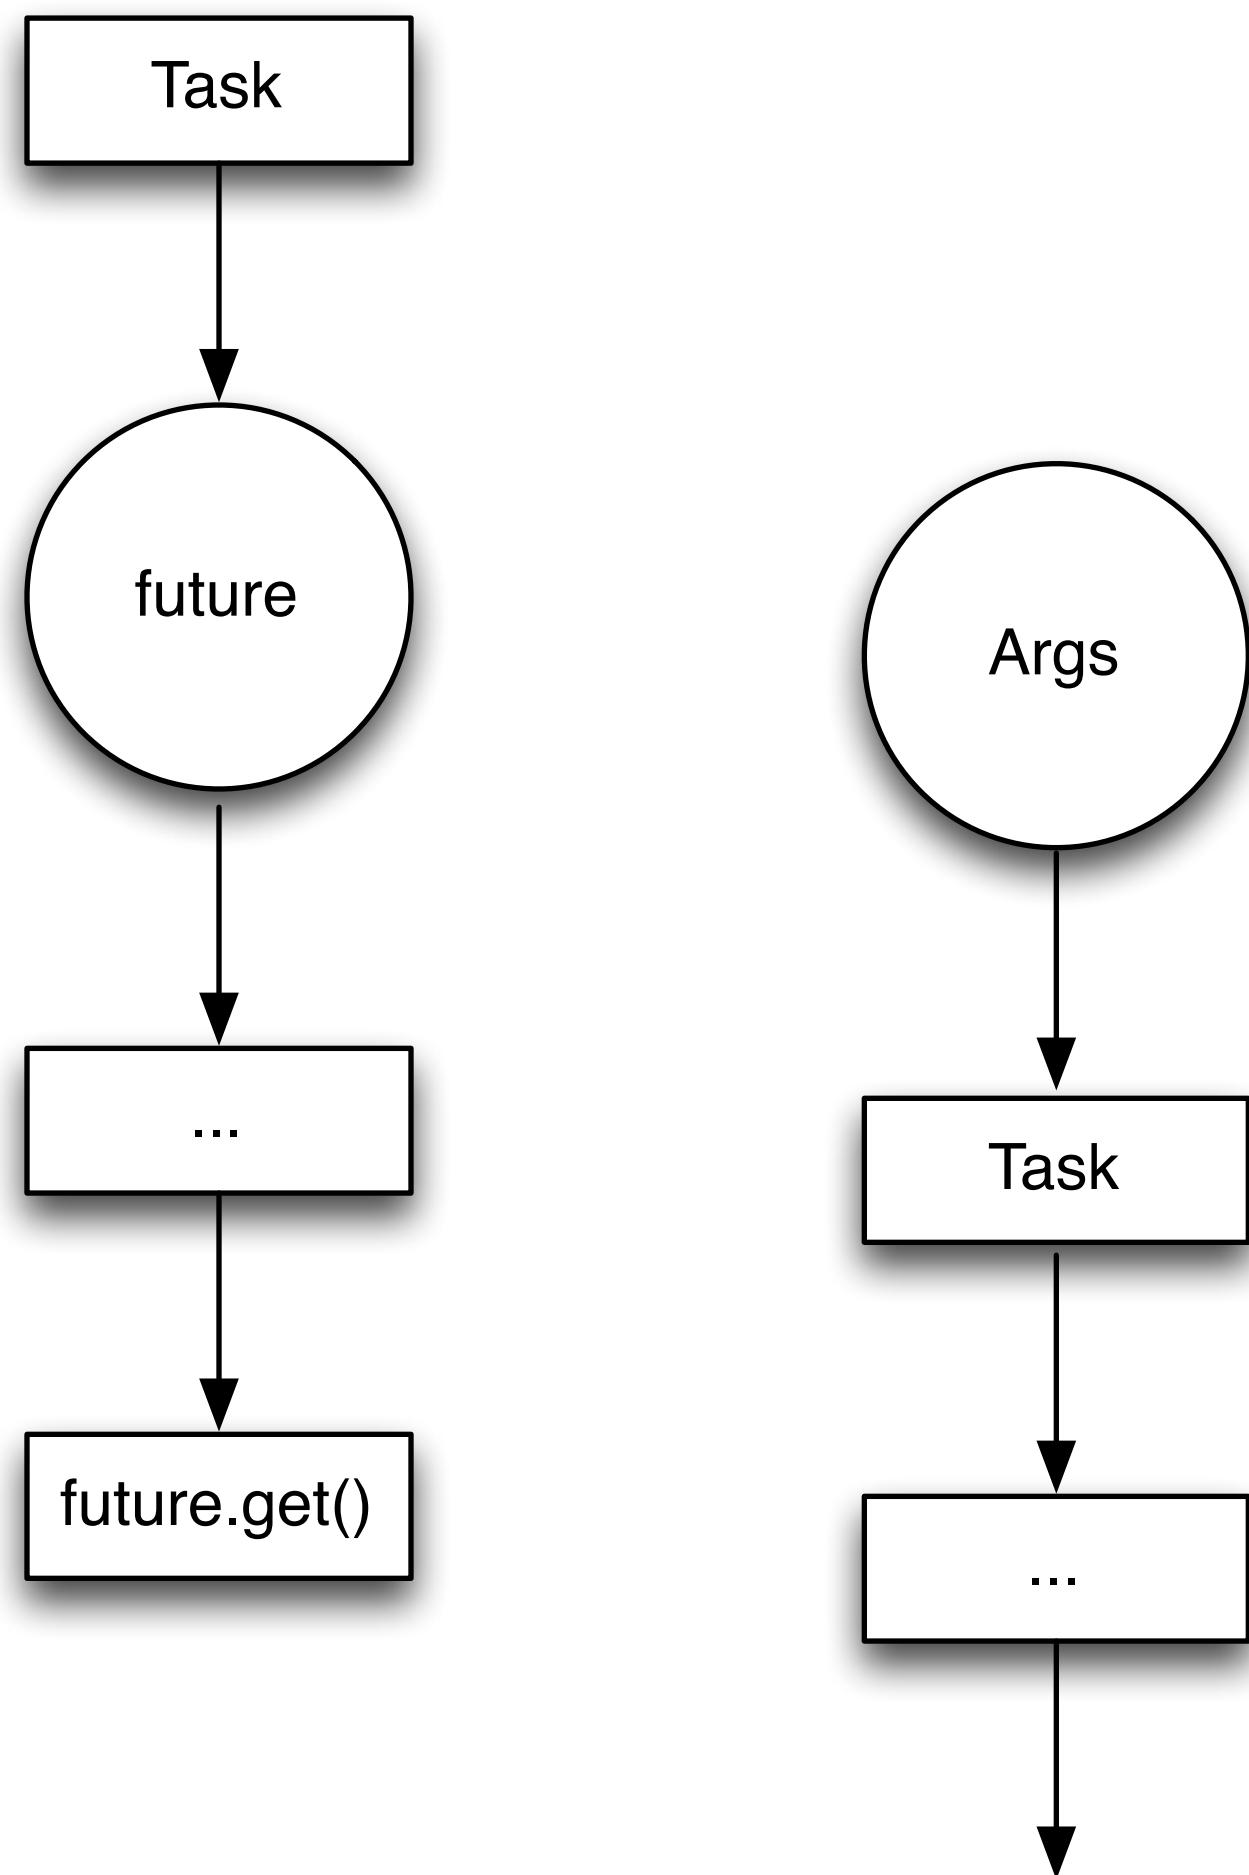
No Raw Synchronization Primitives



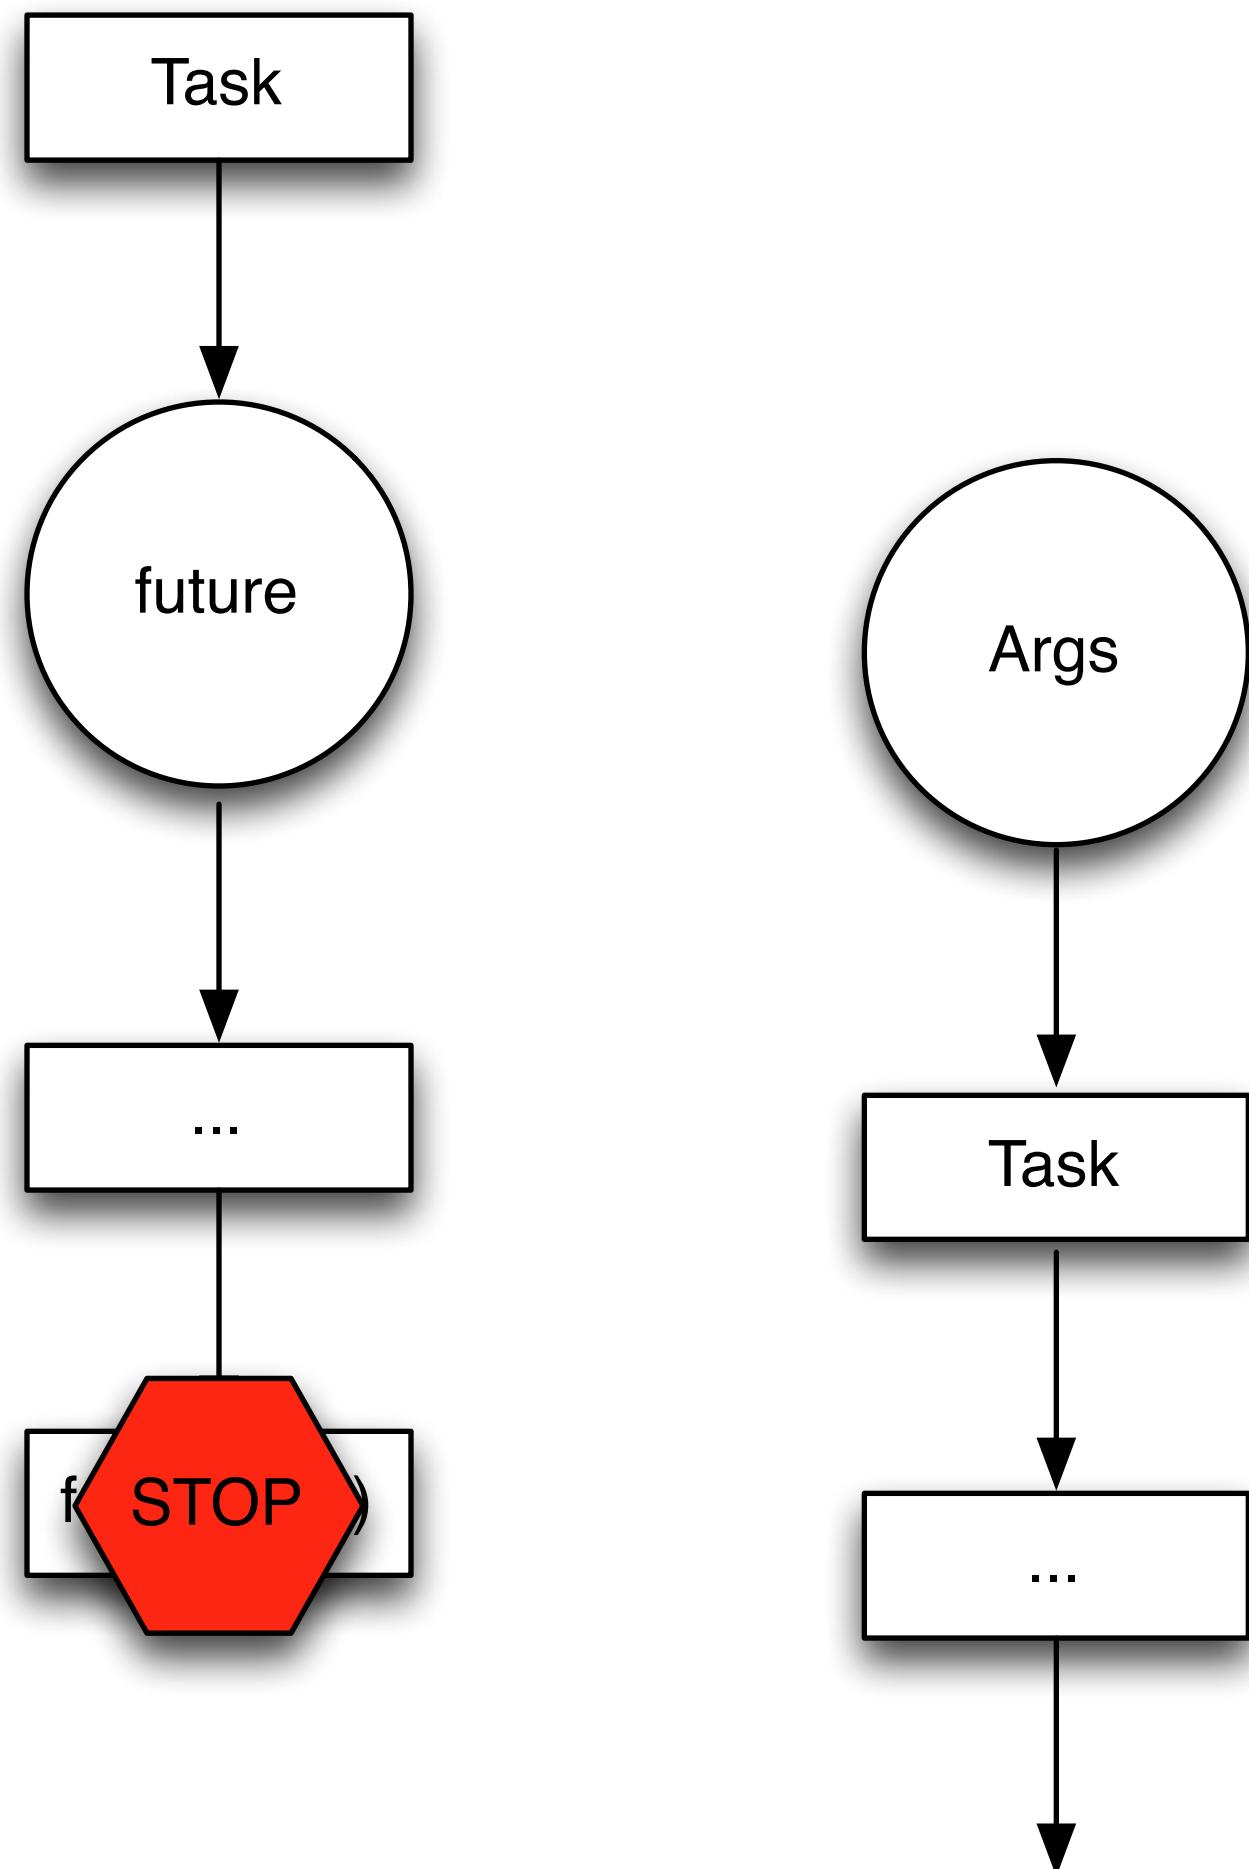
No Raw Synchronization Primitives



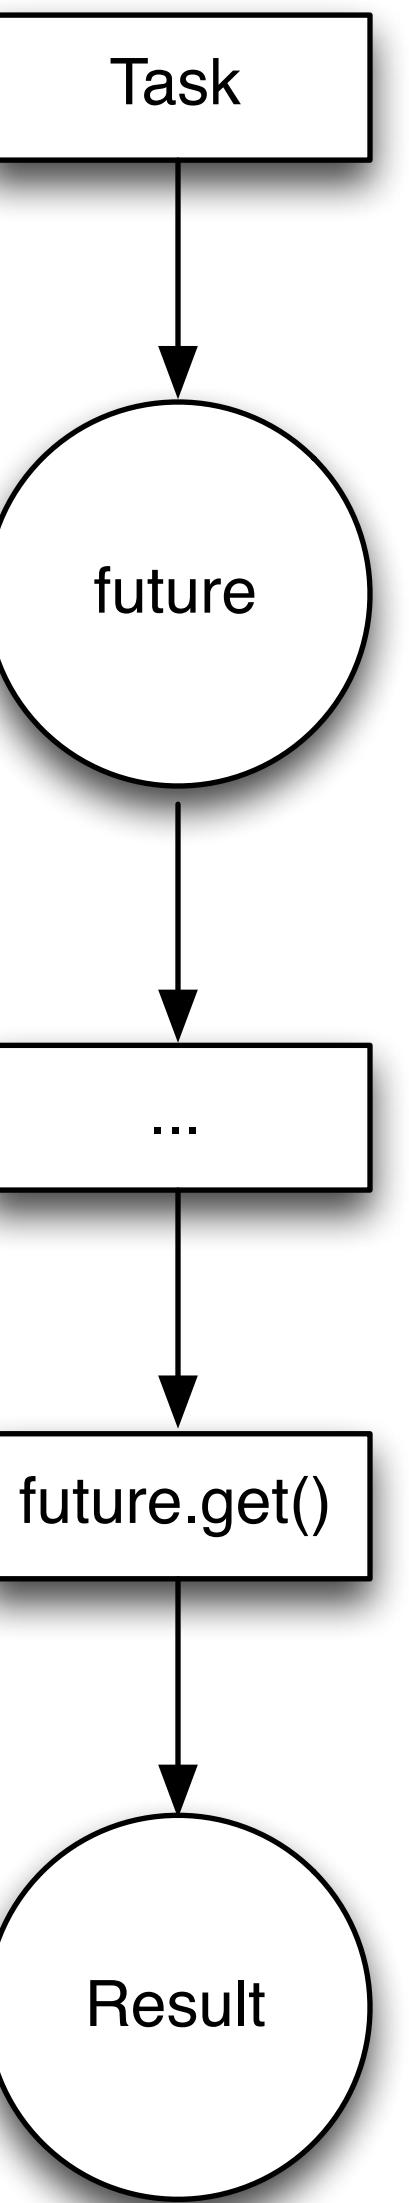
No Raw Synchronization Primitives



No Raw Synchronization Primitives



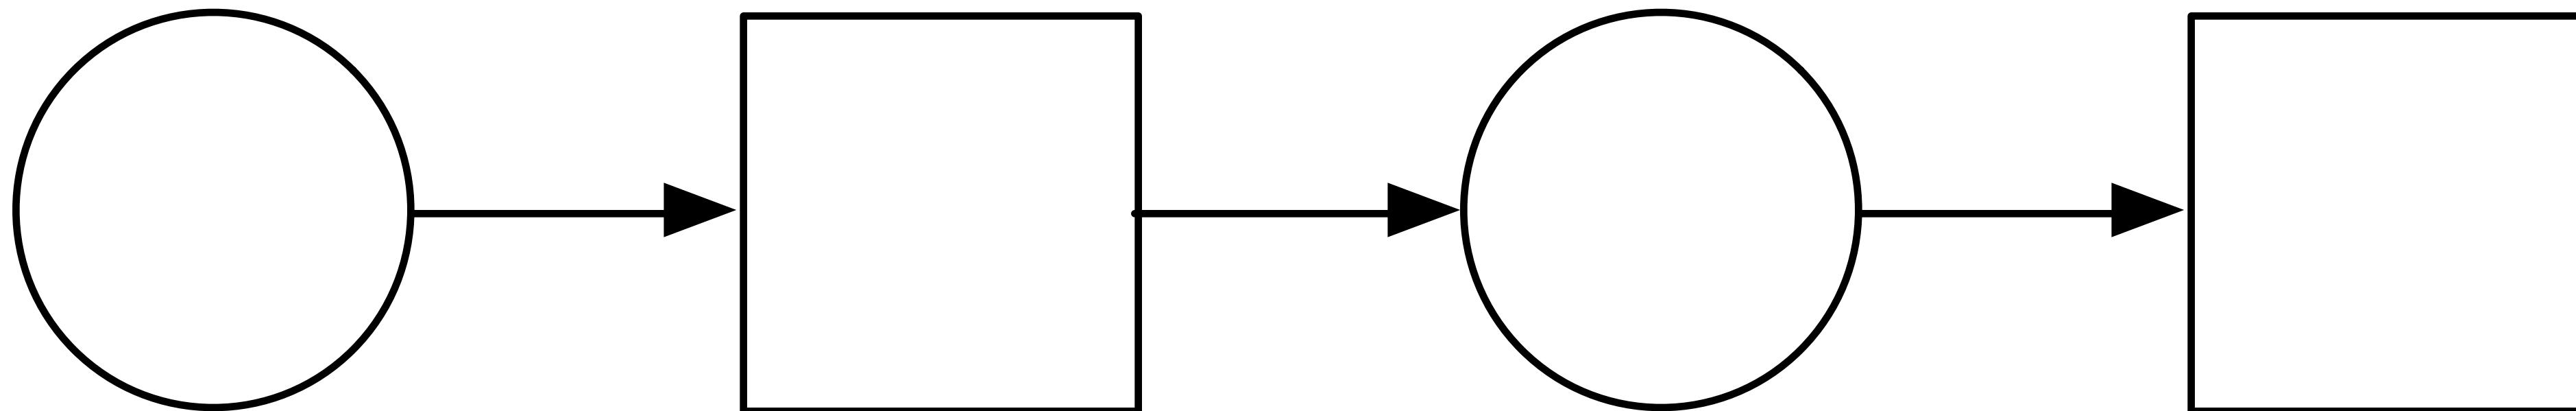
No Raw Synchronization Primitives



Futures: What year is this?

- C++14 futures lack:
 - Continuations - `.then()`
 - Joins - `when_all()`
 - Split
 - Cancellation
 - Progress Monitoring (Except Ready)
- And C++14 futures don't compose (easily) to add these features

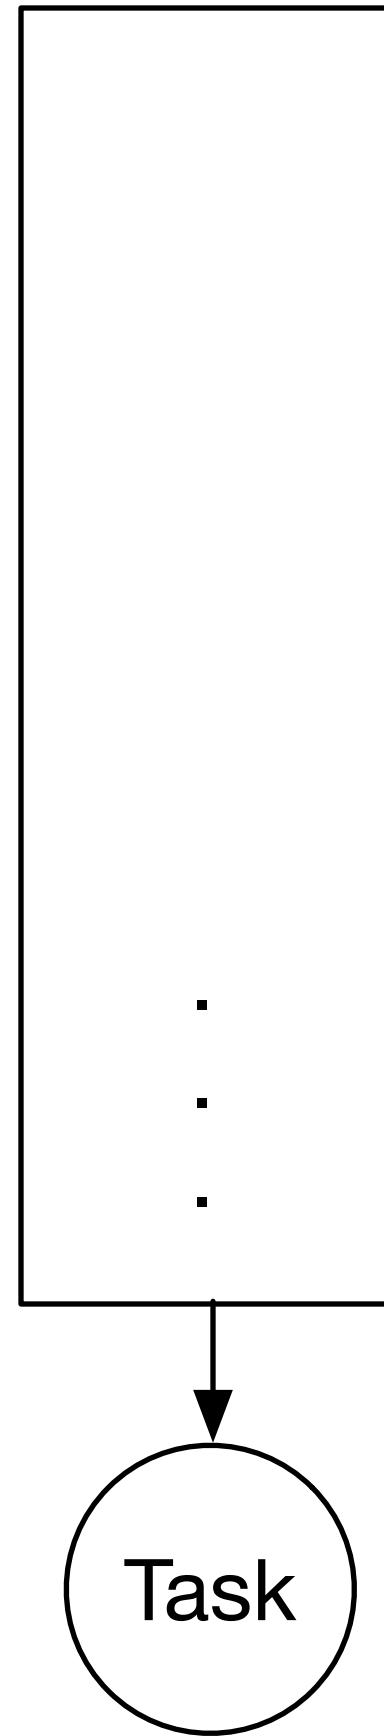
Futures: Continuations



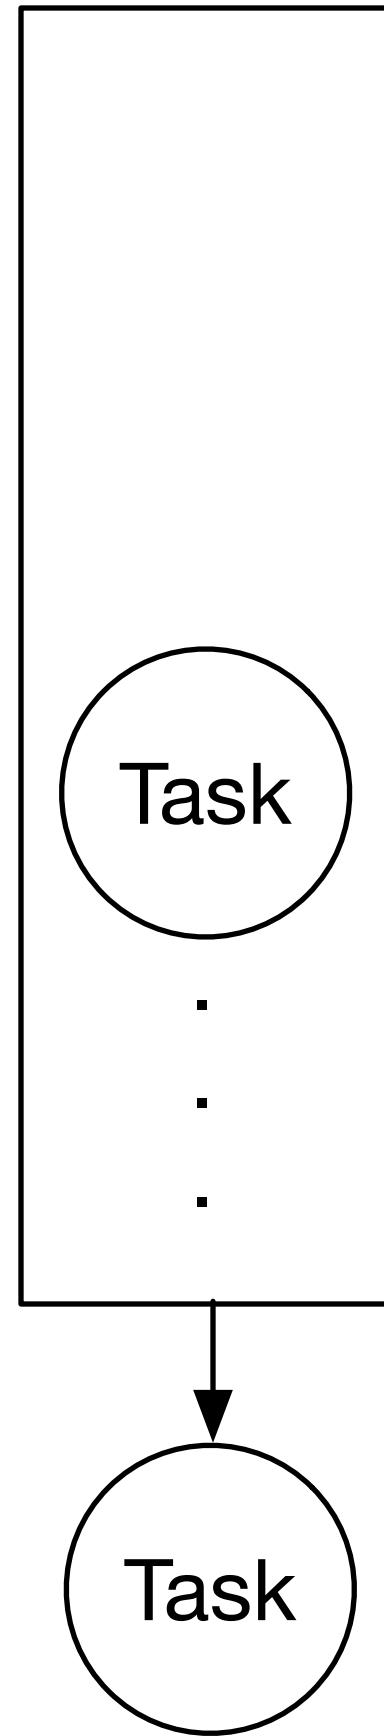
Futures: Continuations

- Blocking on `std::future.get()` has two problems
 - One thread resource is consumed, increasing contention
 - Possibly causing a deadlock in our tasking system!
 - Any subsequent non-dependent calculations on the task are also blocked
- C++14 doesn't have continuations
 - GCD has serialized queues and groups
 - PPL has chained tasks
 - TBB has flow graphs
 - TS Concurrency will have `.then()`
 - Boost futures have them now

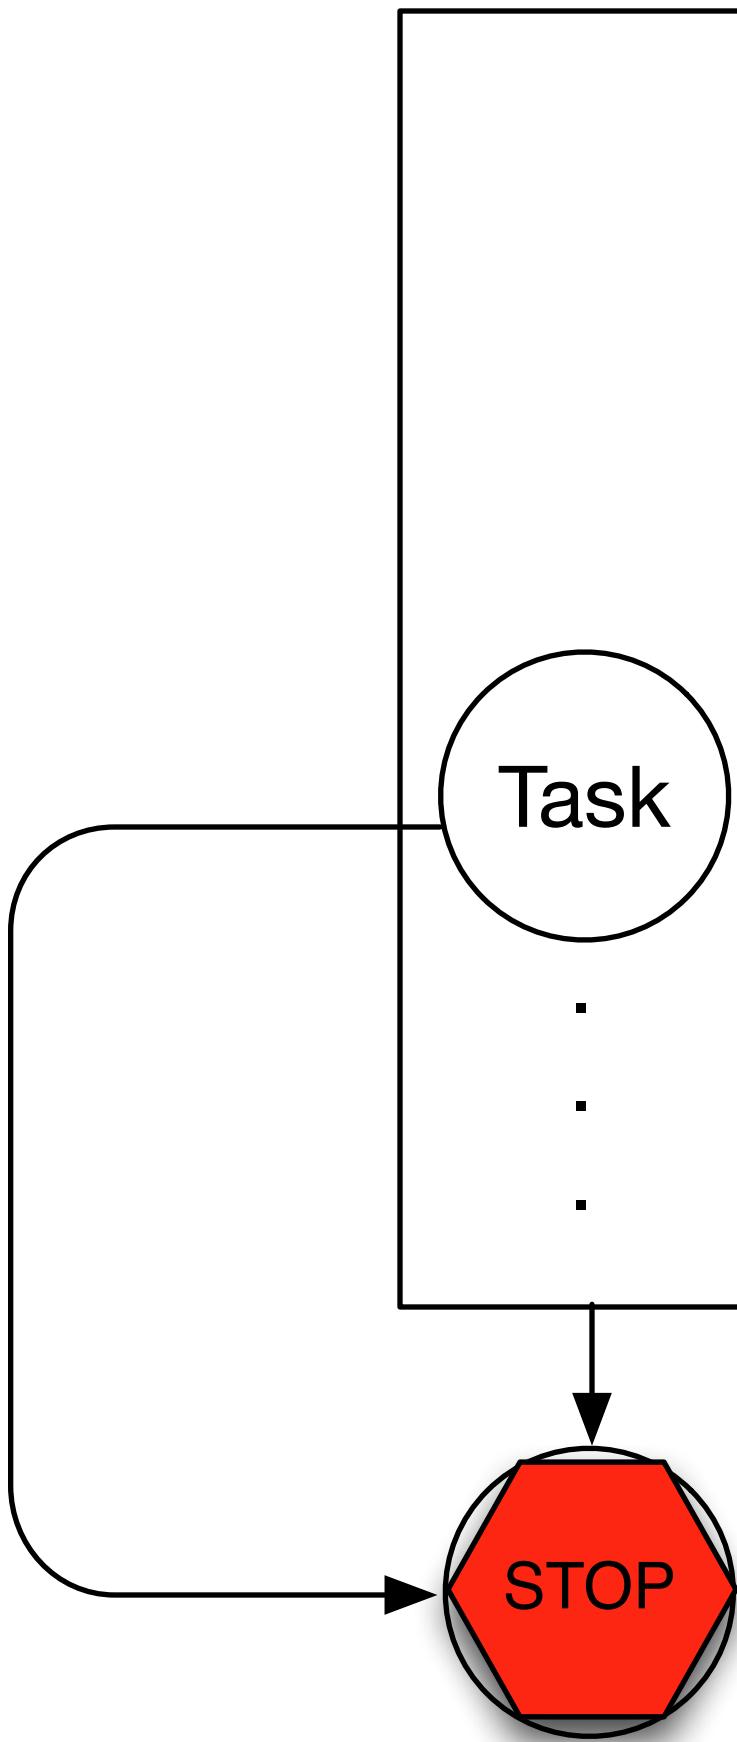
Futures: get() deadlock



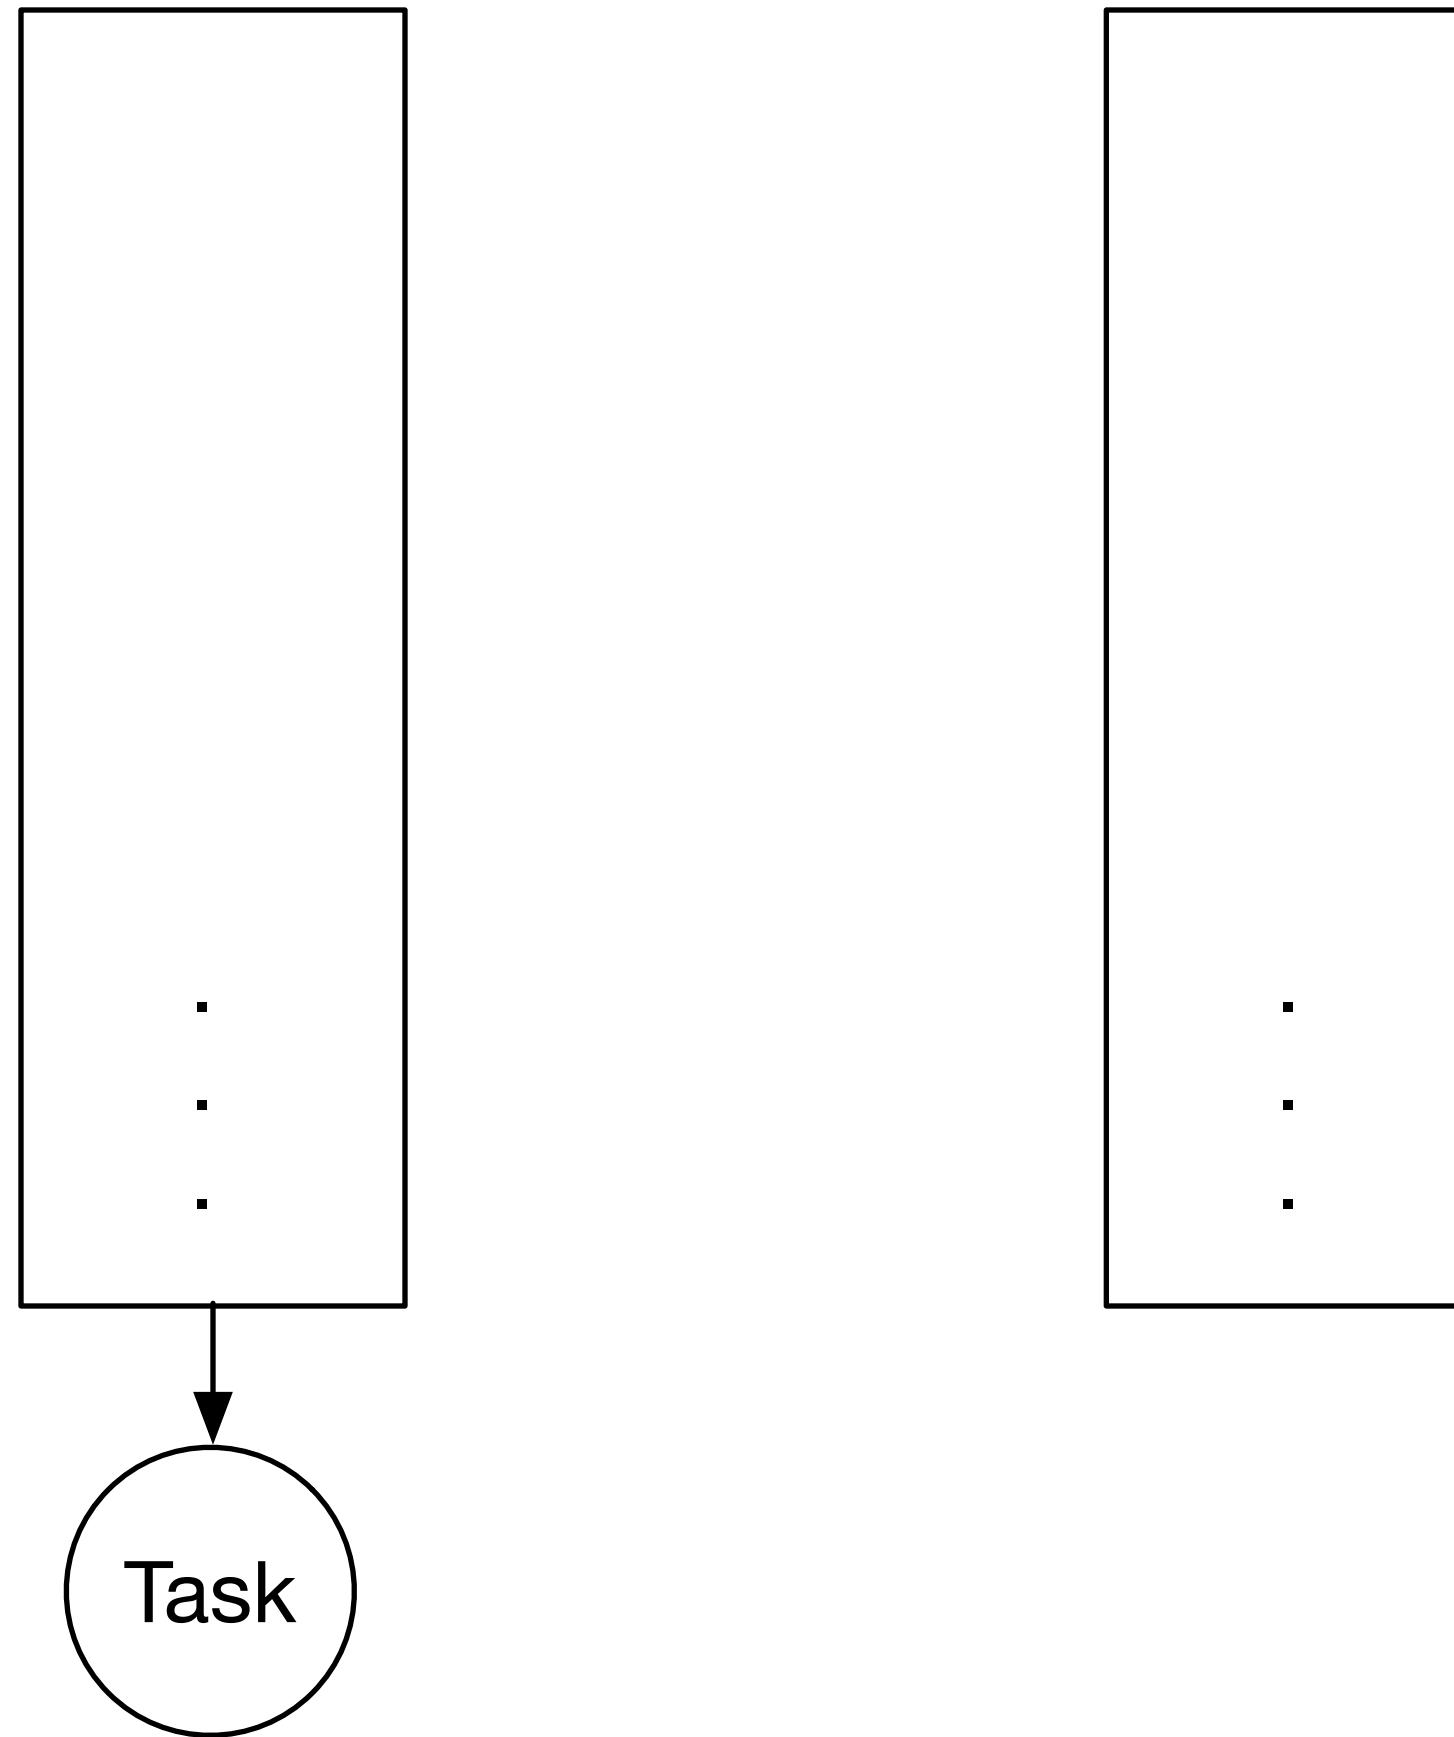
Futures: get() deadlock



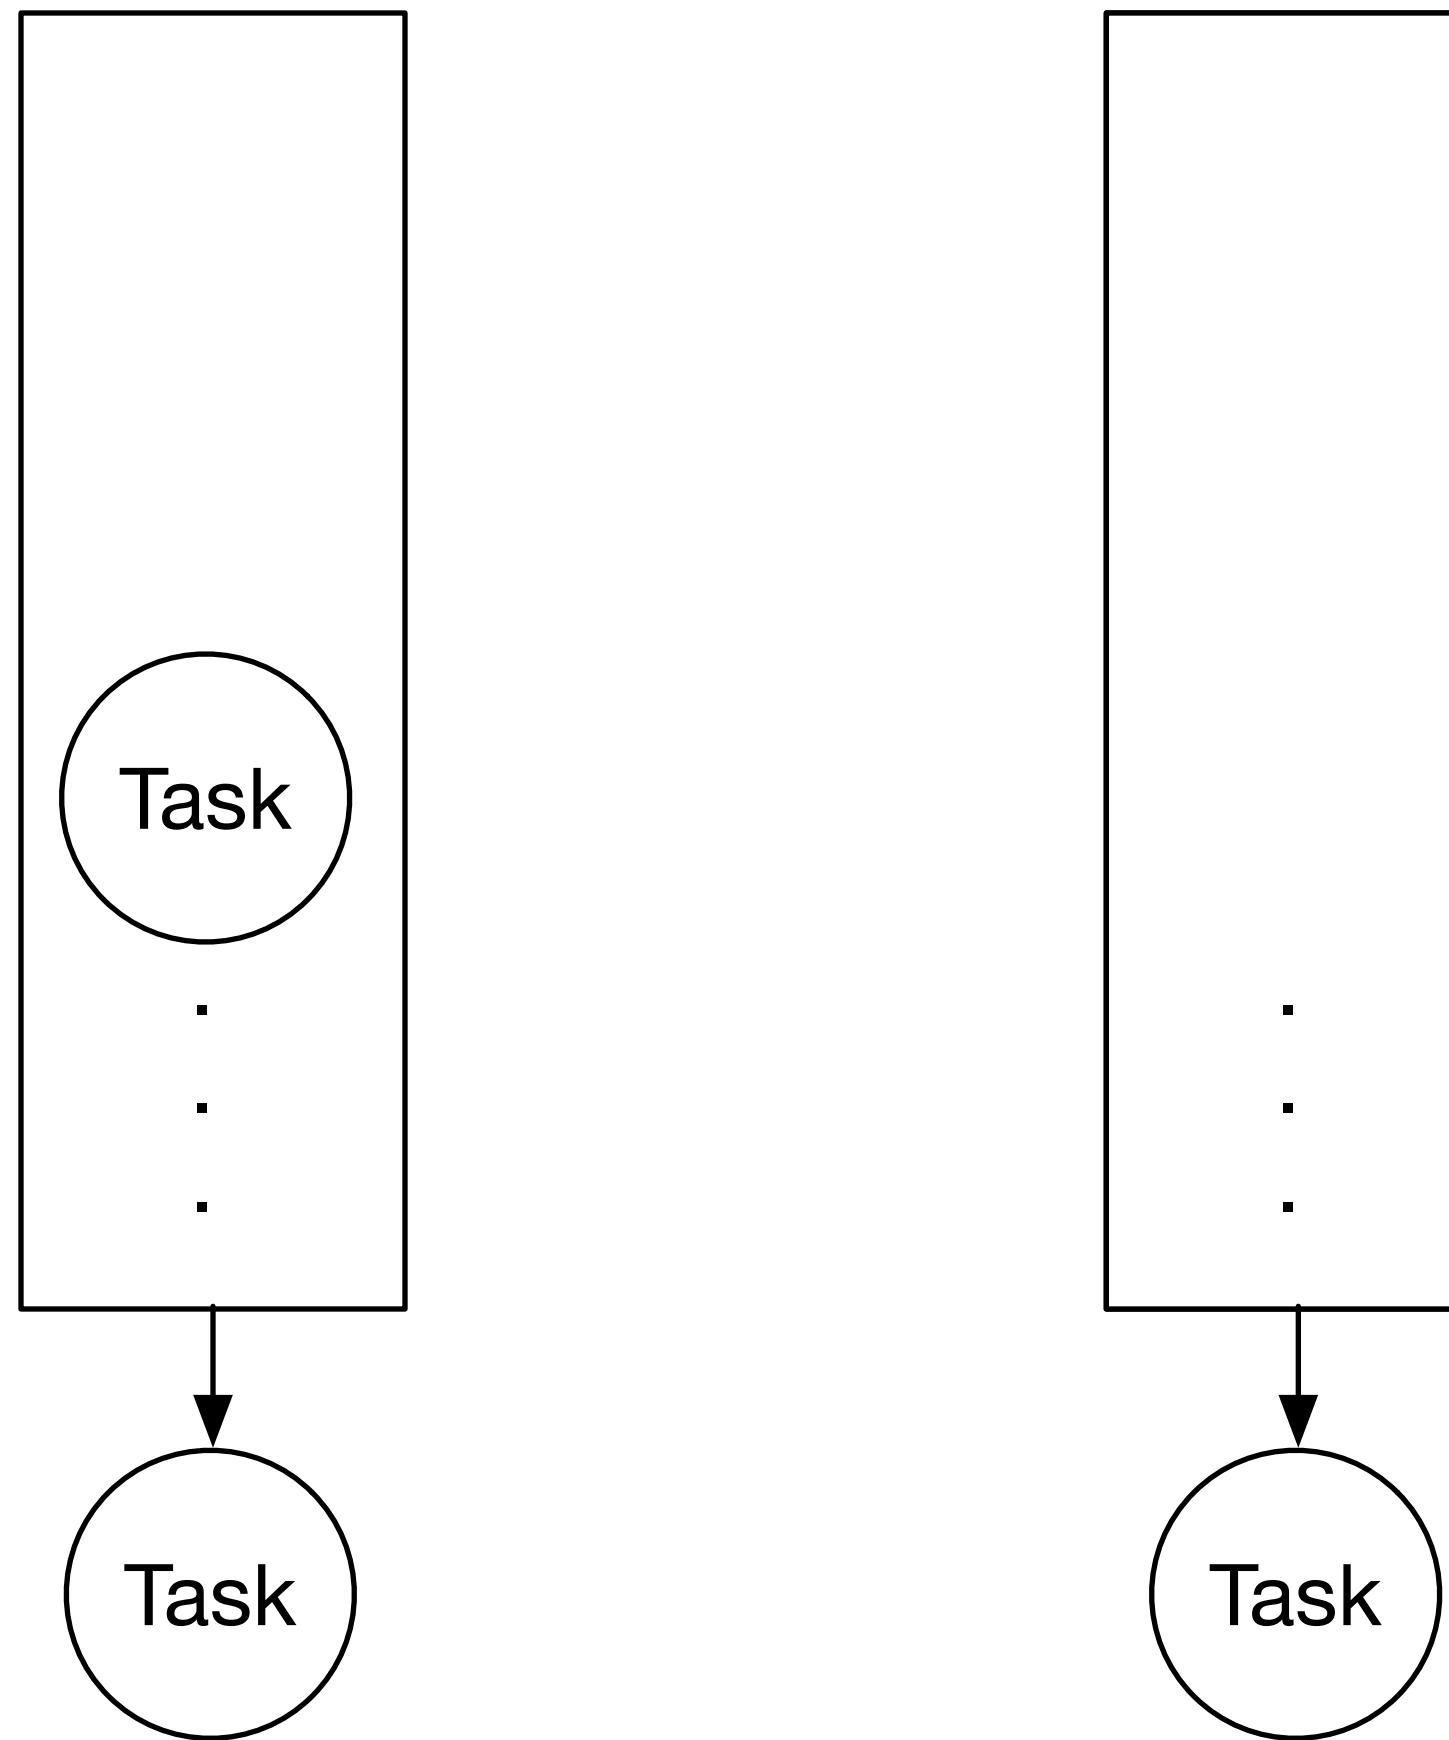
Futures: get() deadlock



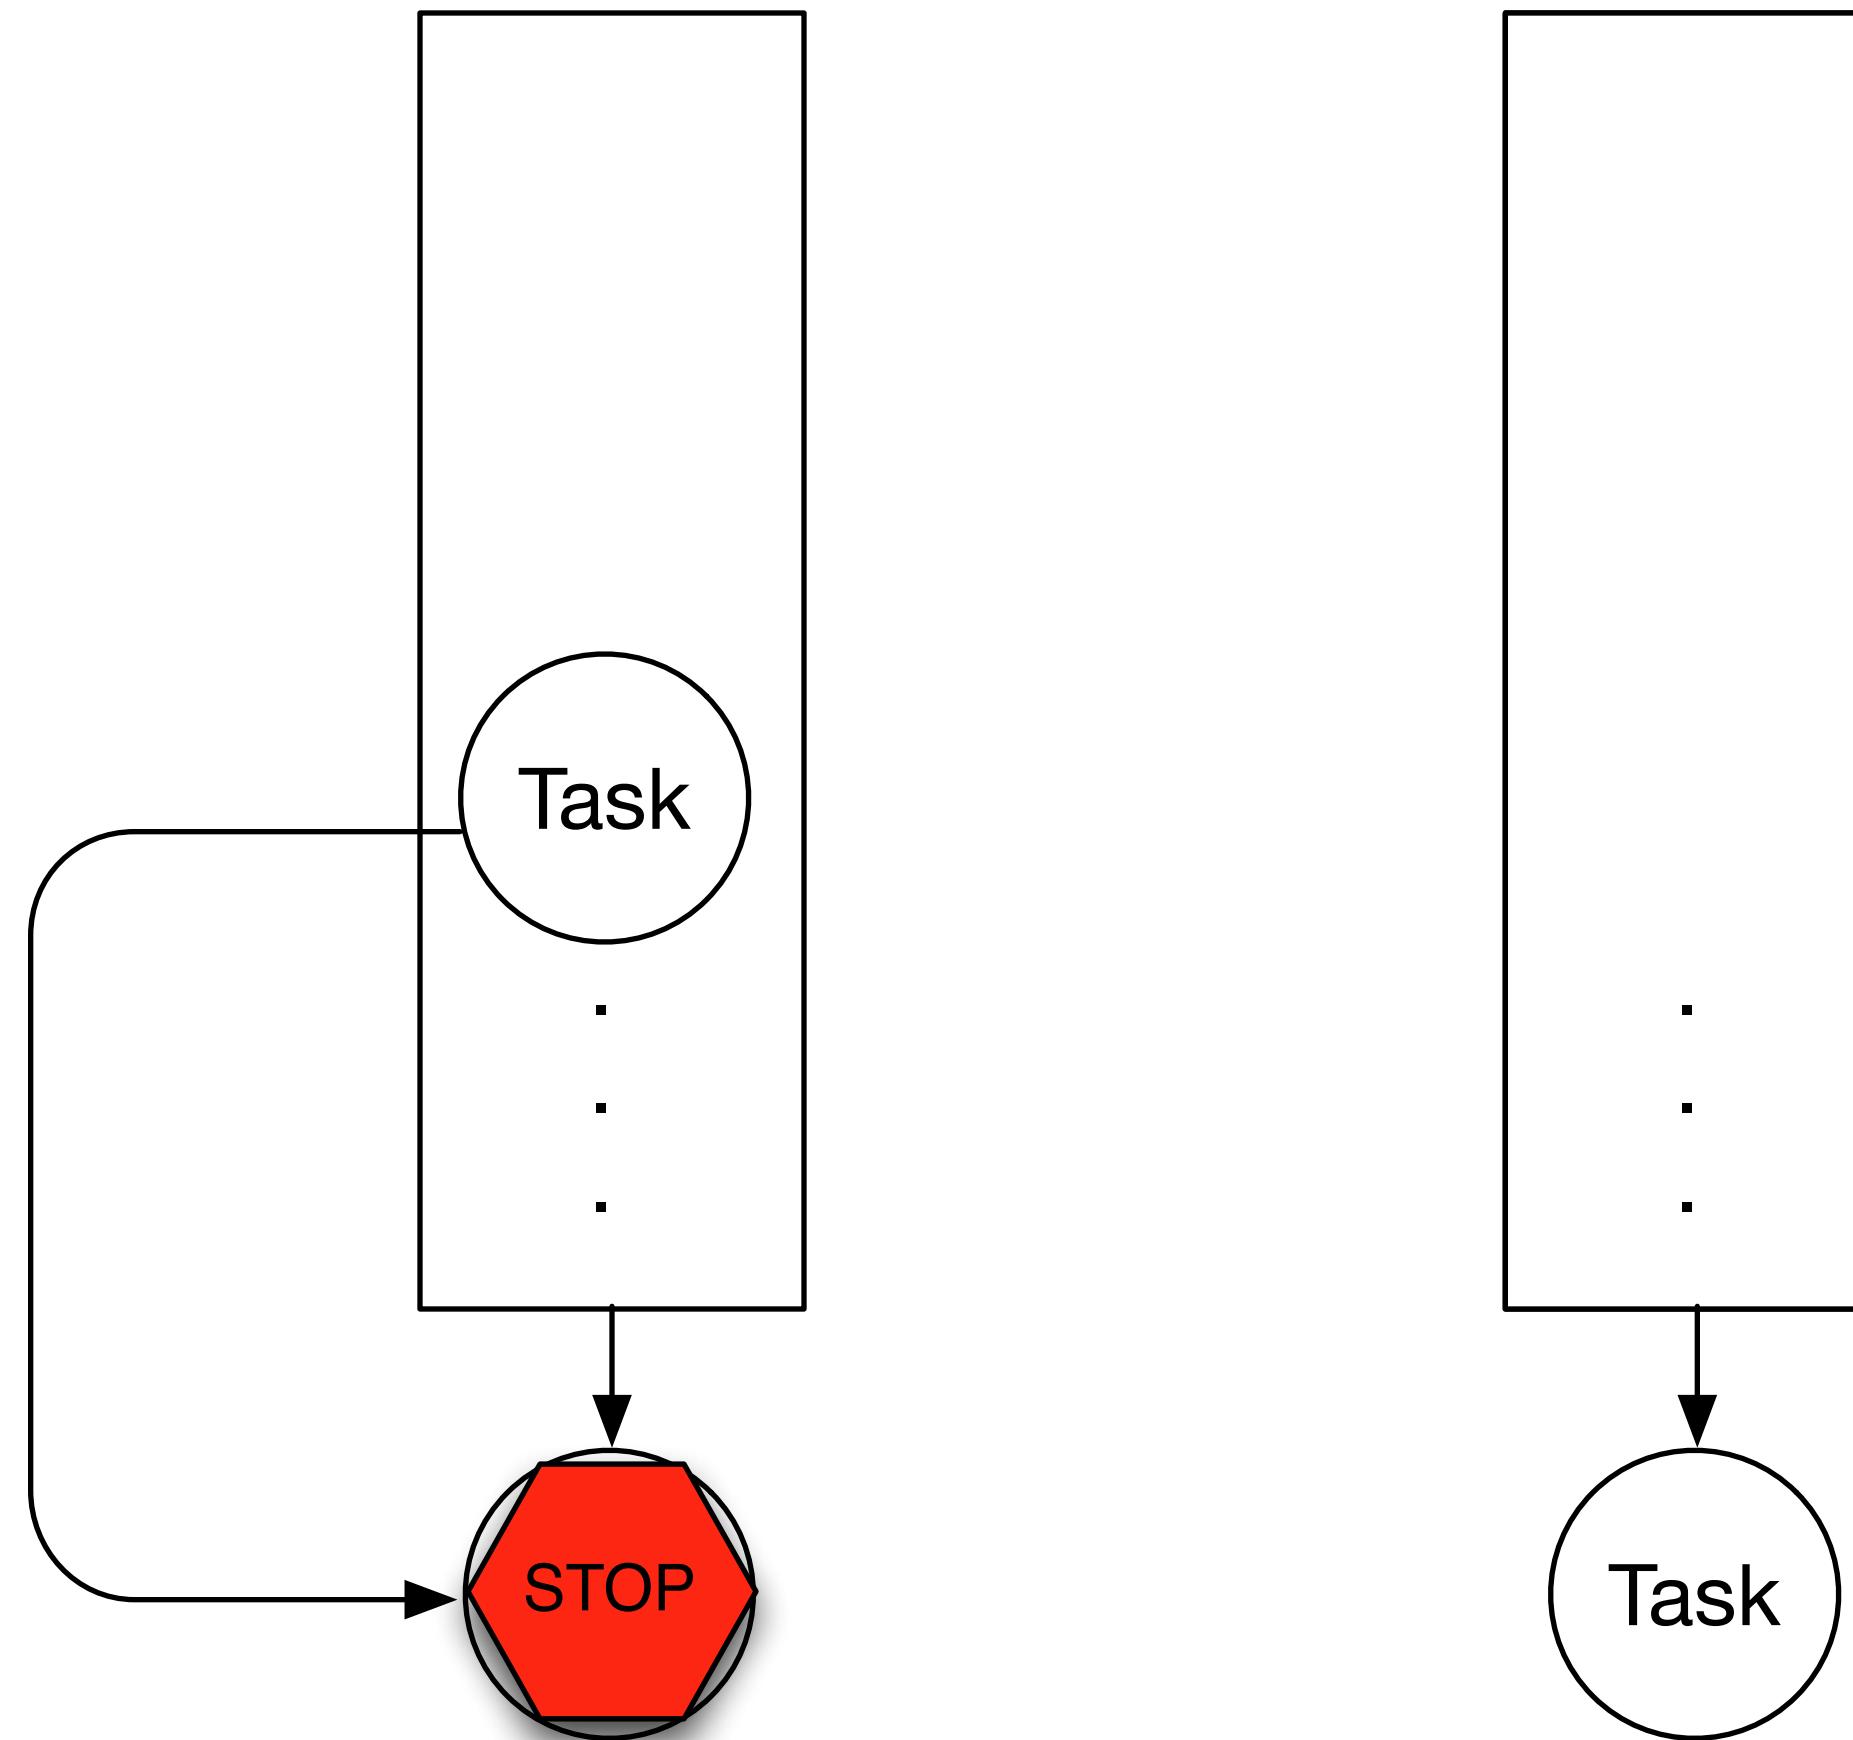
Futures: get() deadlock



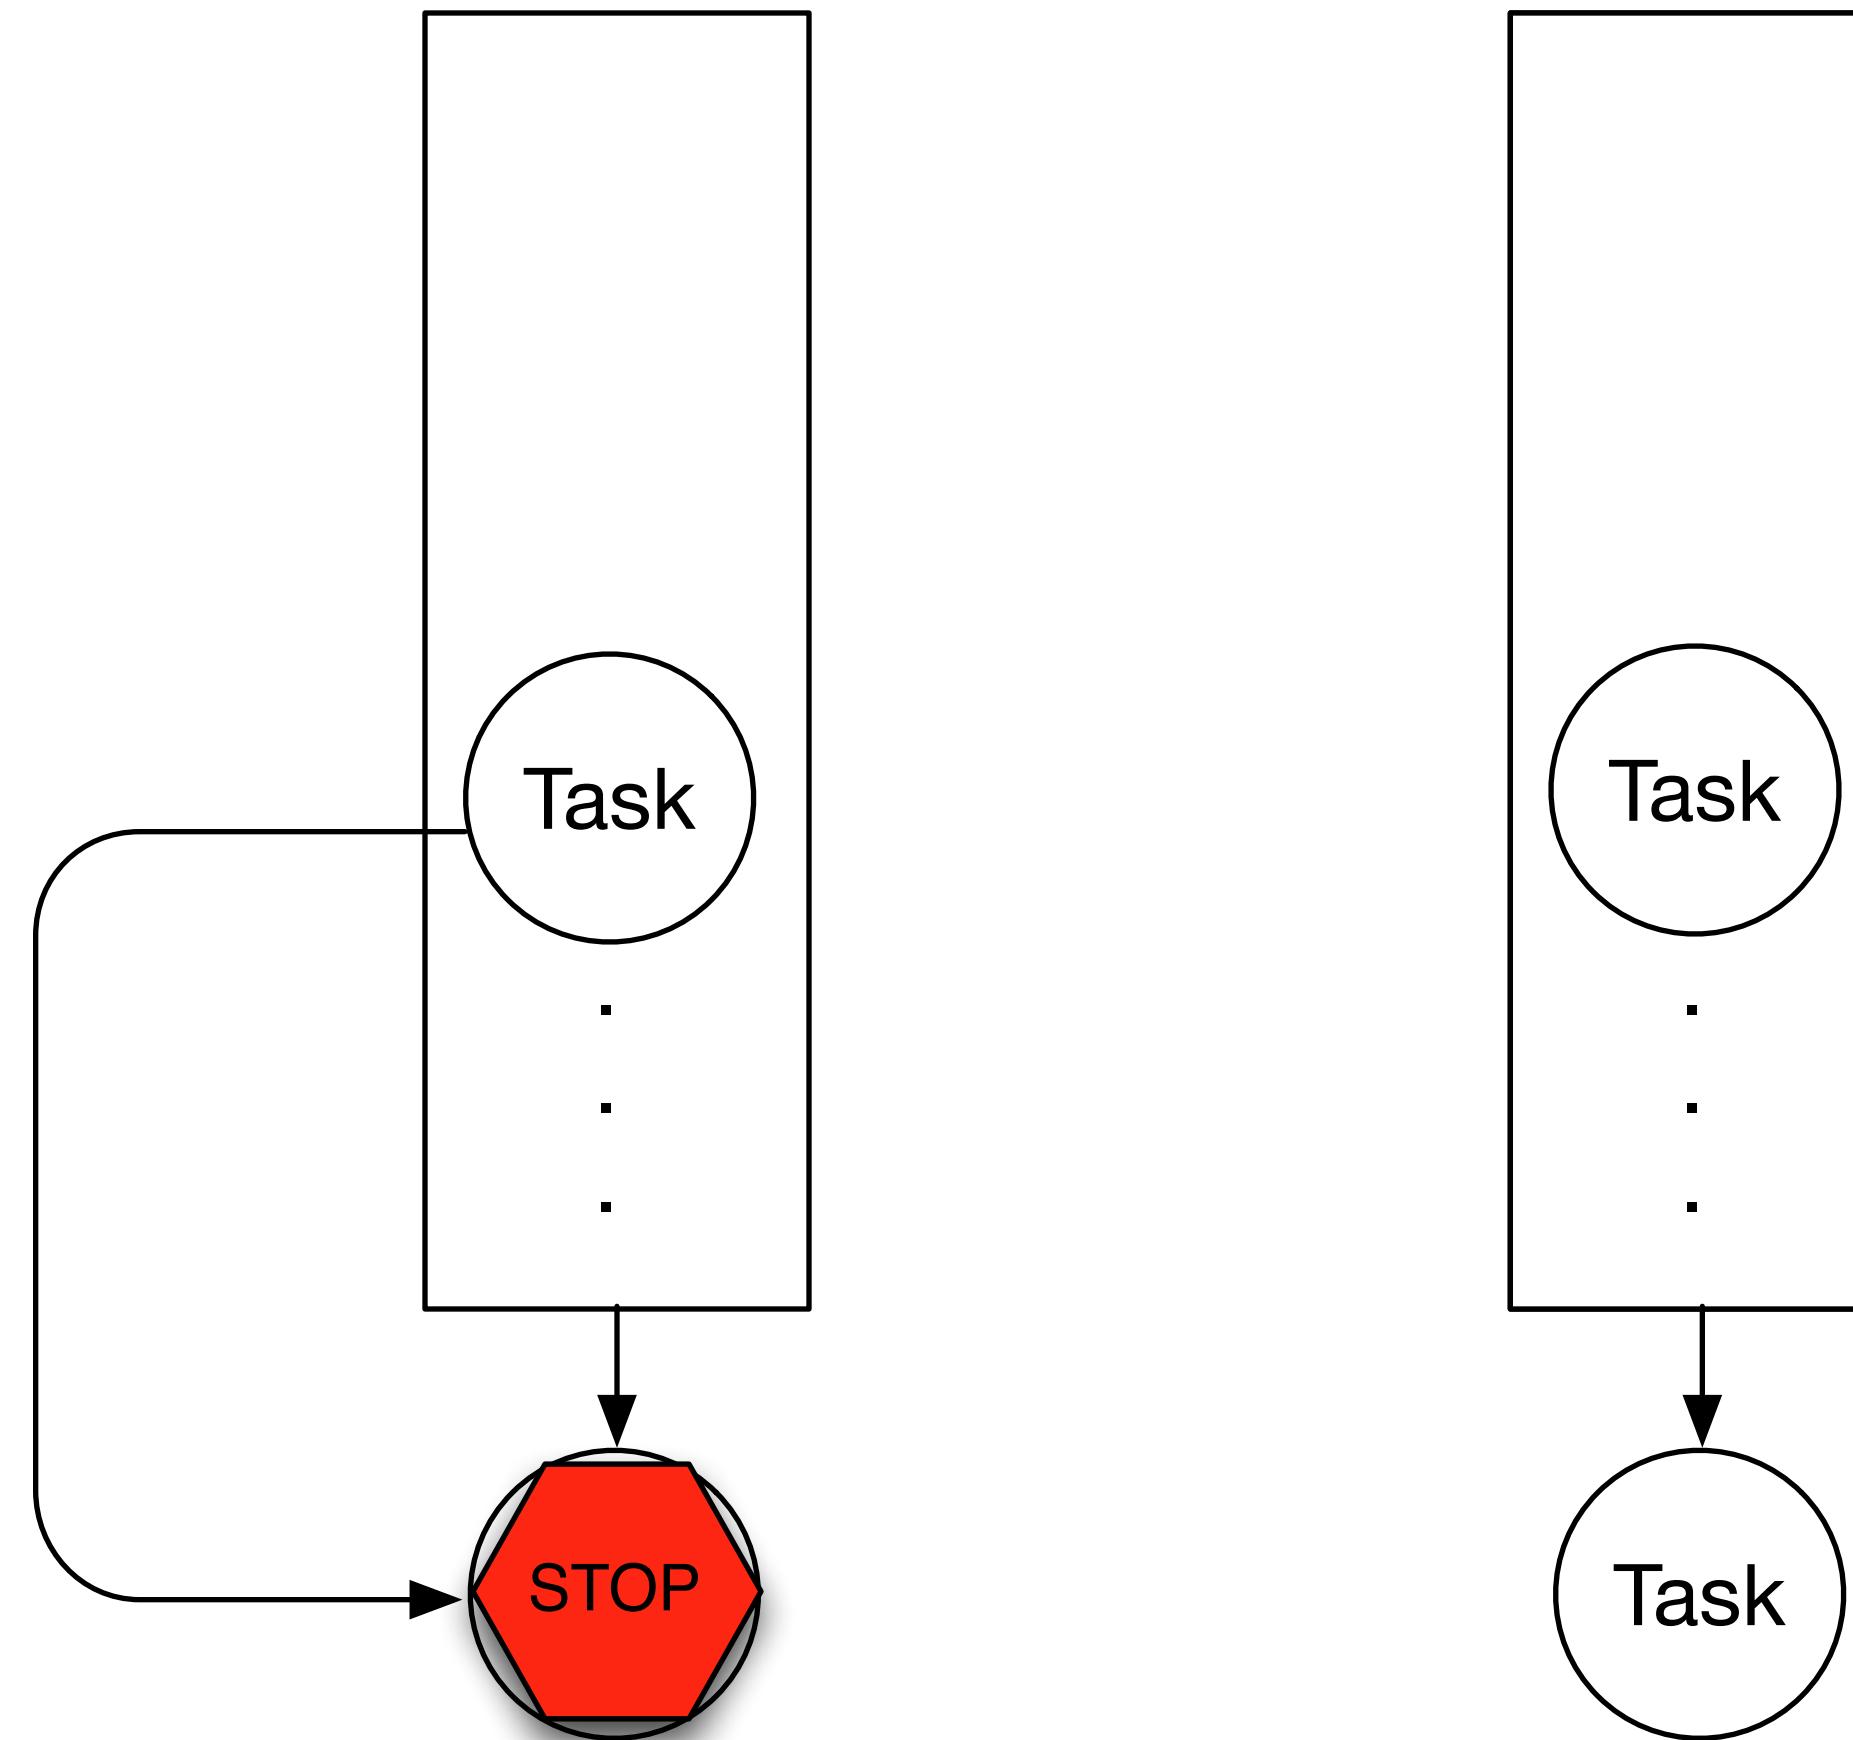
Futures: get() deadlock



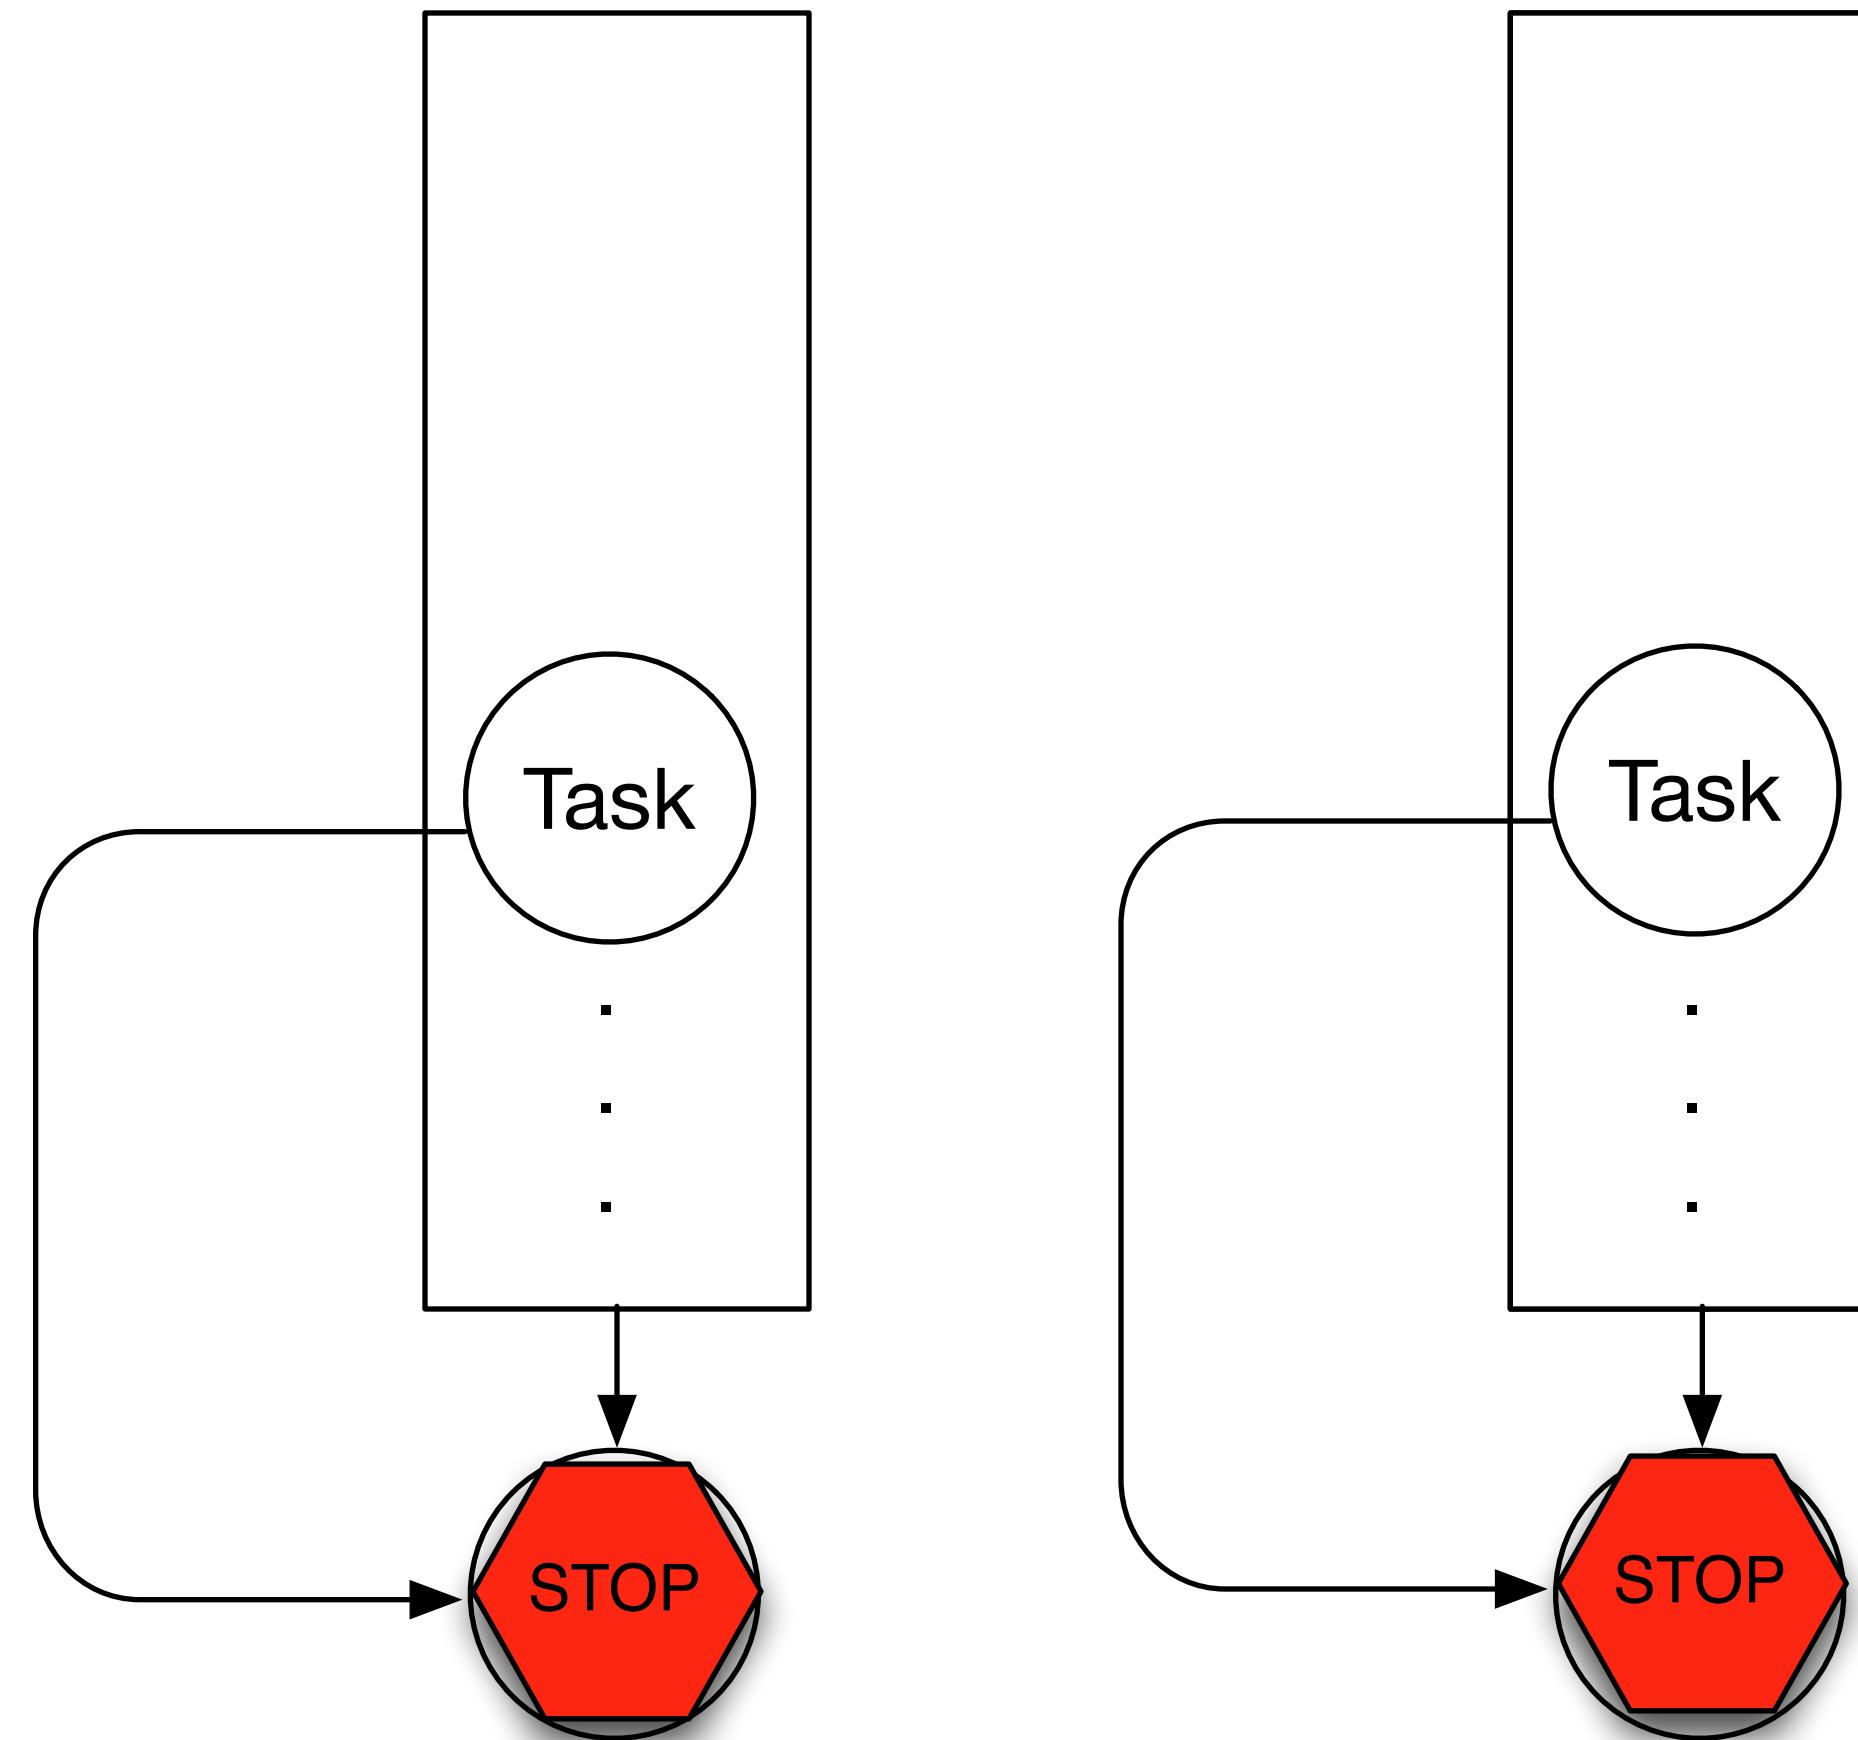
Futures: get() deadlock



Futures: get() deadlock



Futures: get() deadlock



Futures: Continuations

- Blocking on `std::future.get()`
 - Very difficult to use safely with a thread pool
 - C++14 allows `std::async()` to use a thread pool
- Not just `get()` - *any* blocking (condition variables, `wait`, ...) is problematic with a task system

Futures: Continuations

```
future<cpp_int> x = async([]{ return fibonacci<cpp_int>(1'000); });
future<void> y = x.then([](future<cpp_int> x){ cout << x.get() << endl; });
// Do something
y.wait();
```

Futures: Continuations

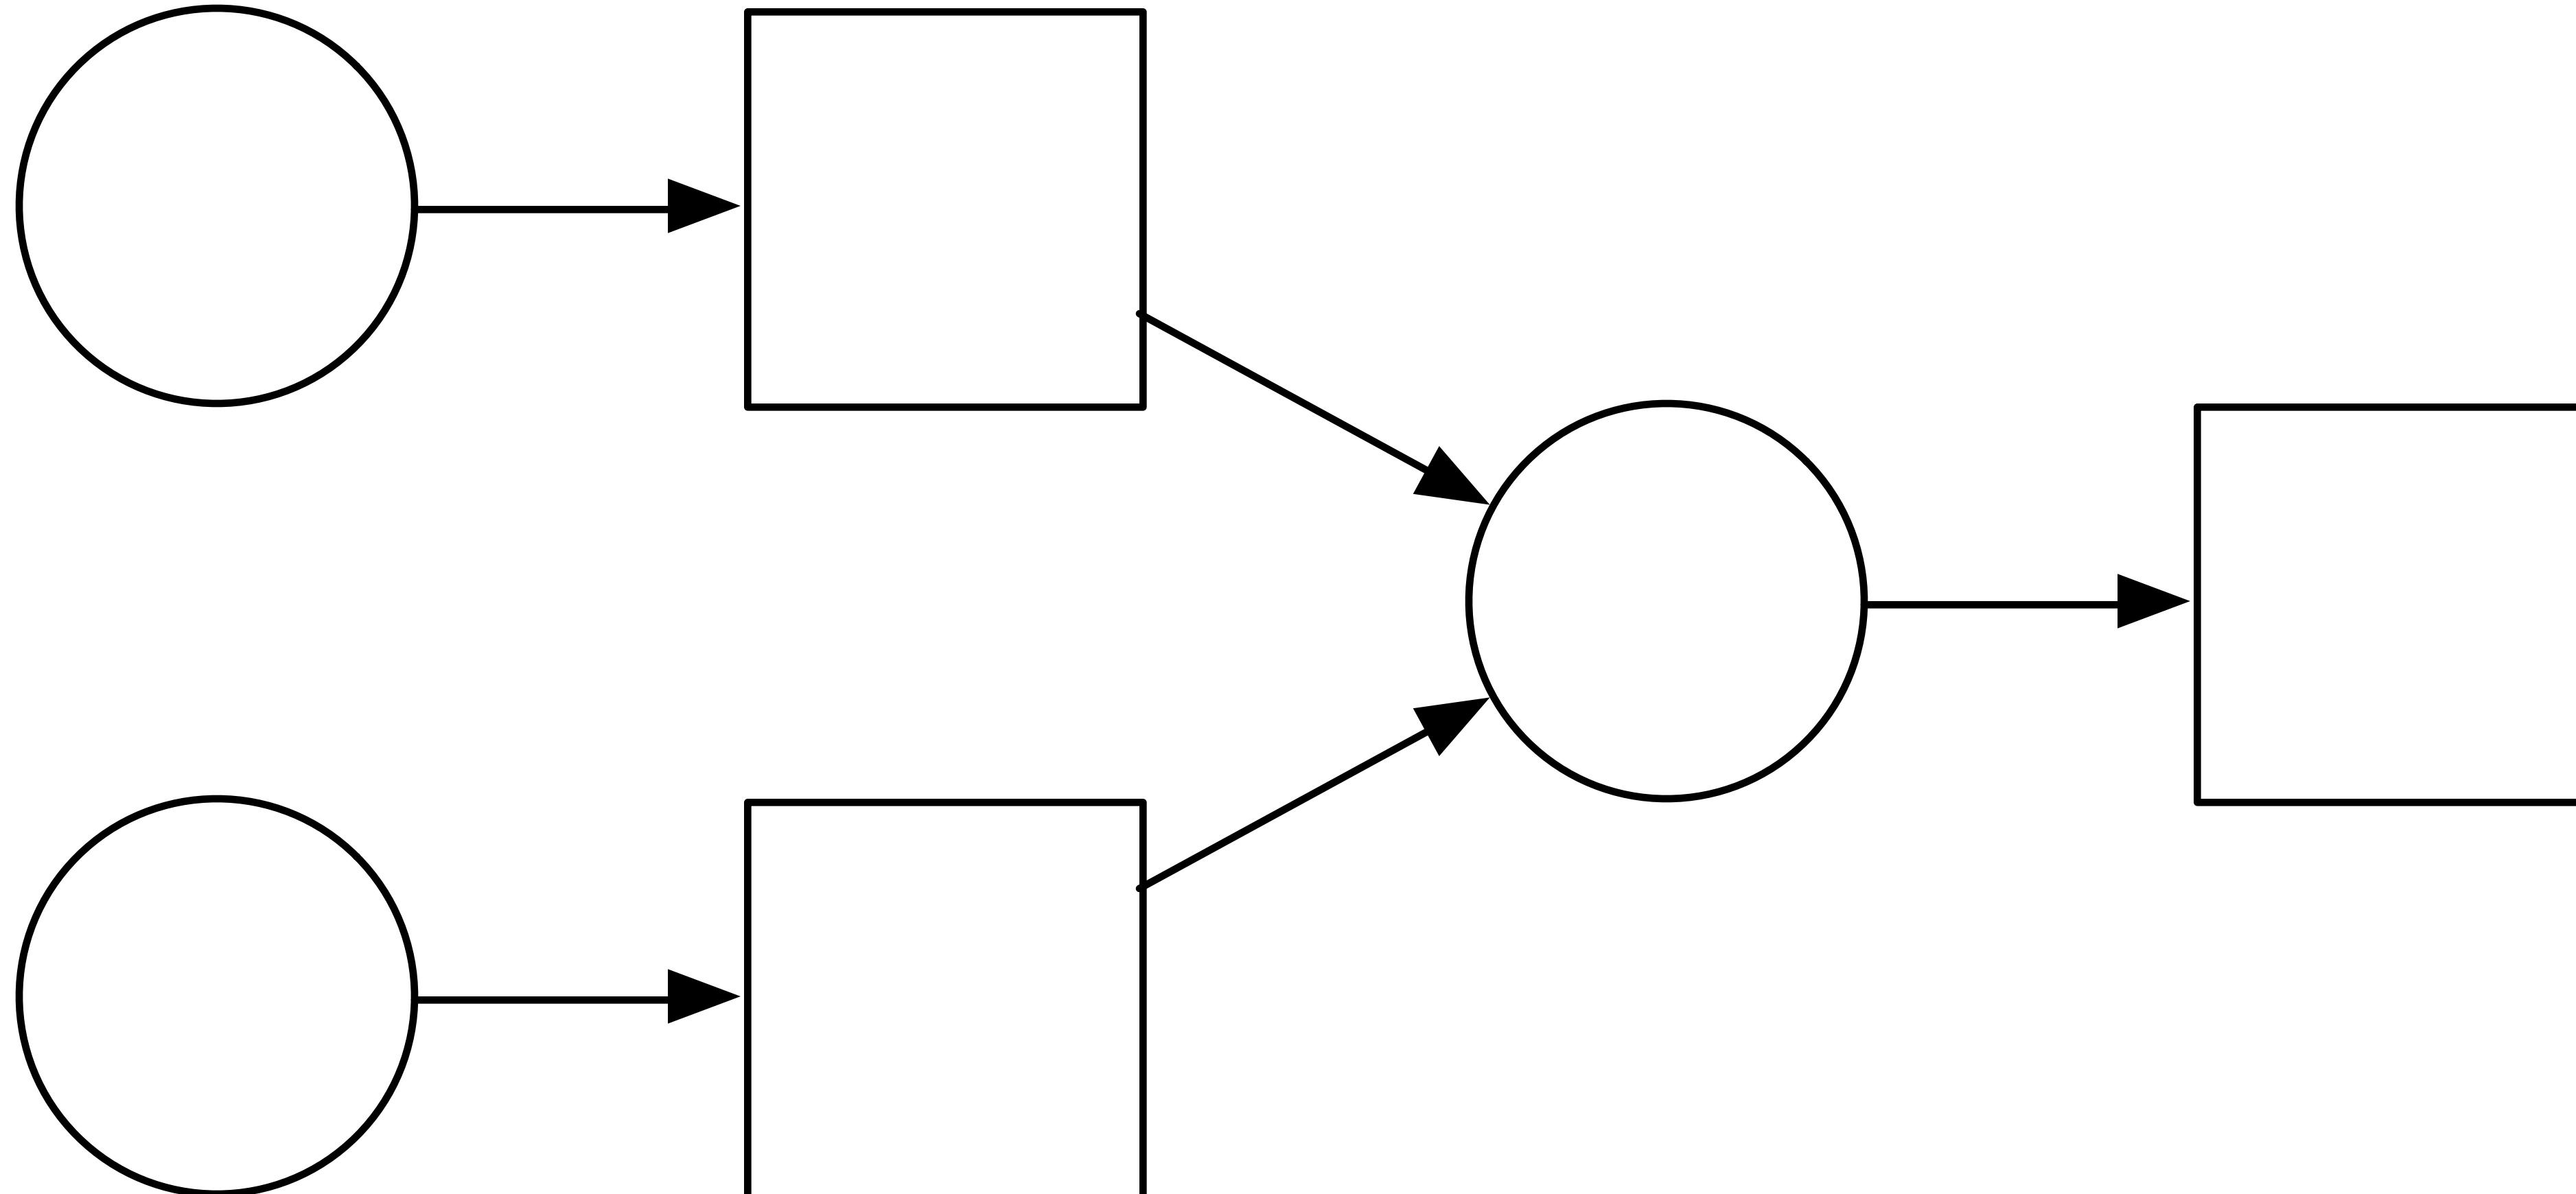
```
future<cpp_int> x = async([]{ return fibonacci<cpp_int>(1'000); });
future<void> y = x.then([](future<cpp_int> x){ cout << x.get() << endl; });
// Do something
y.wait();
```

**4346655768693745643568852767504062580256466051737178040248172908953655417949051890403879840079255169295922593080322634775209
689623239873322471161642996440906533187938298969649928516003704476137795166849228875**

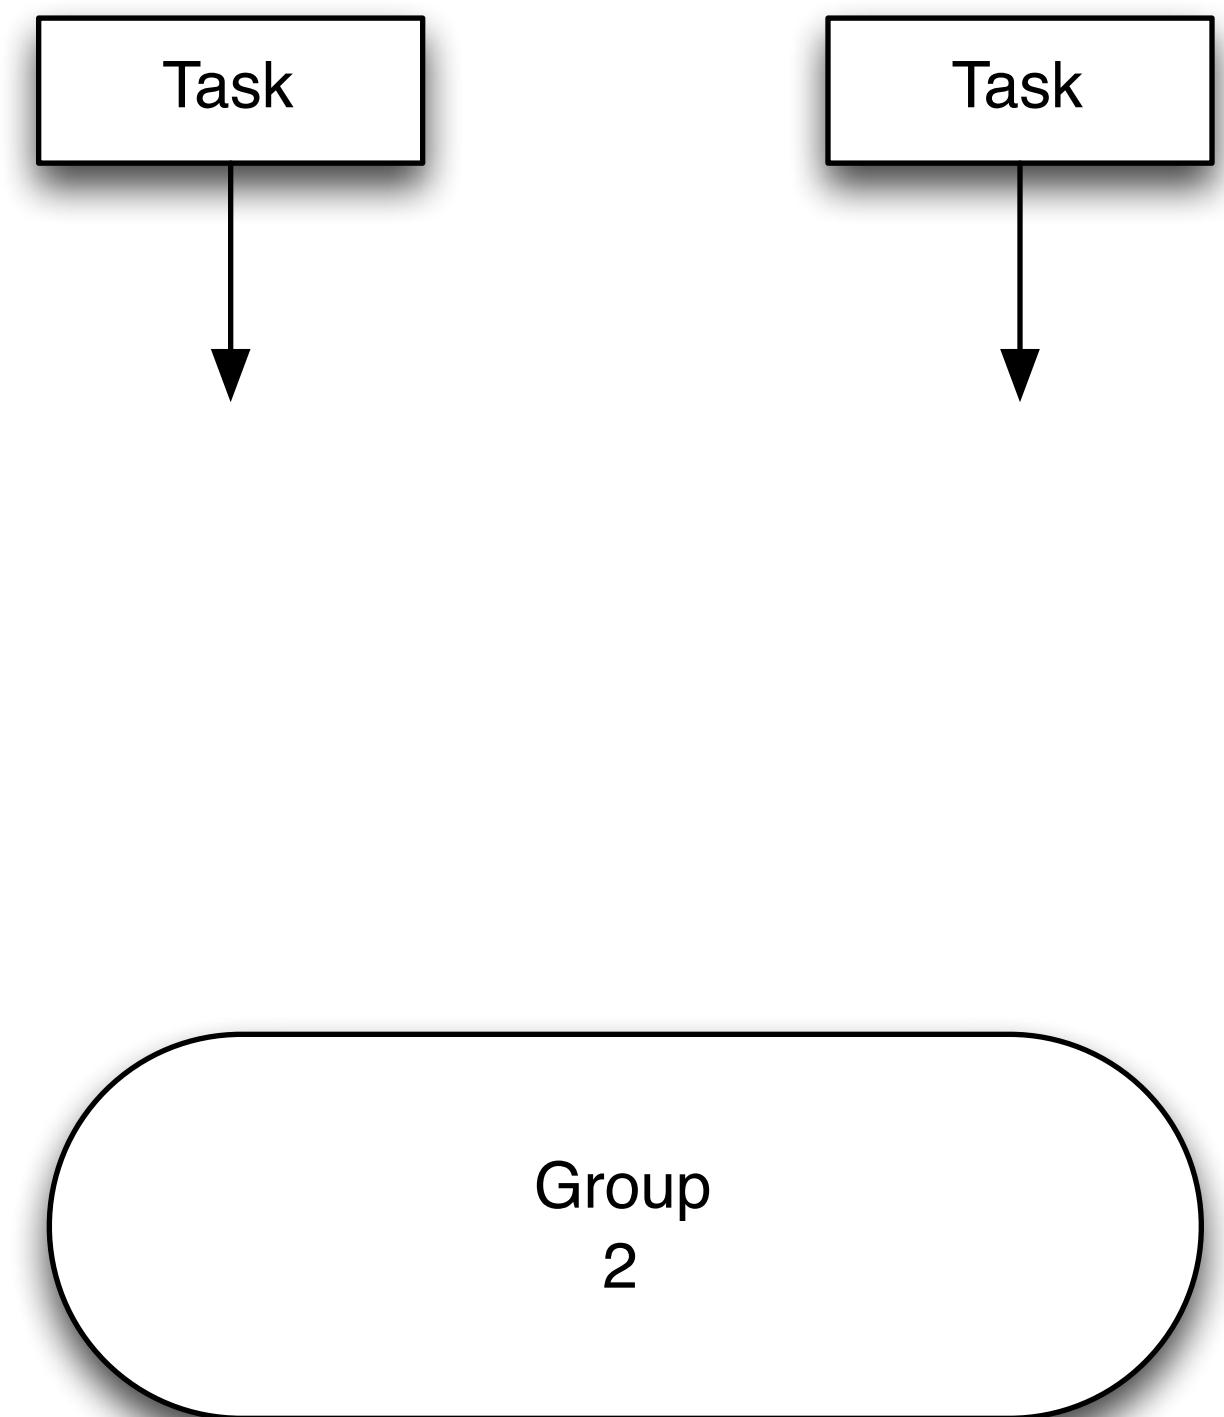
Futures vs Completion Handlers

- Completion handlers are callbacks, they must be known prior to the call
 - No need to synchronize between invoking and setting the continuation
- Futures allow setting the continuation after the sending call is in flight
 - Simpler to compose
 - Require synchronization between invoking and setting the continuation

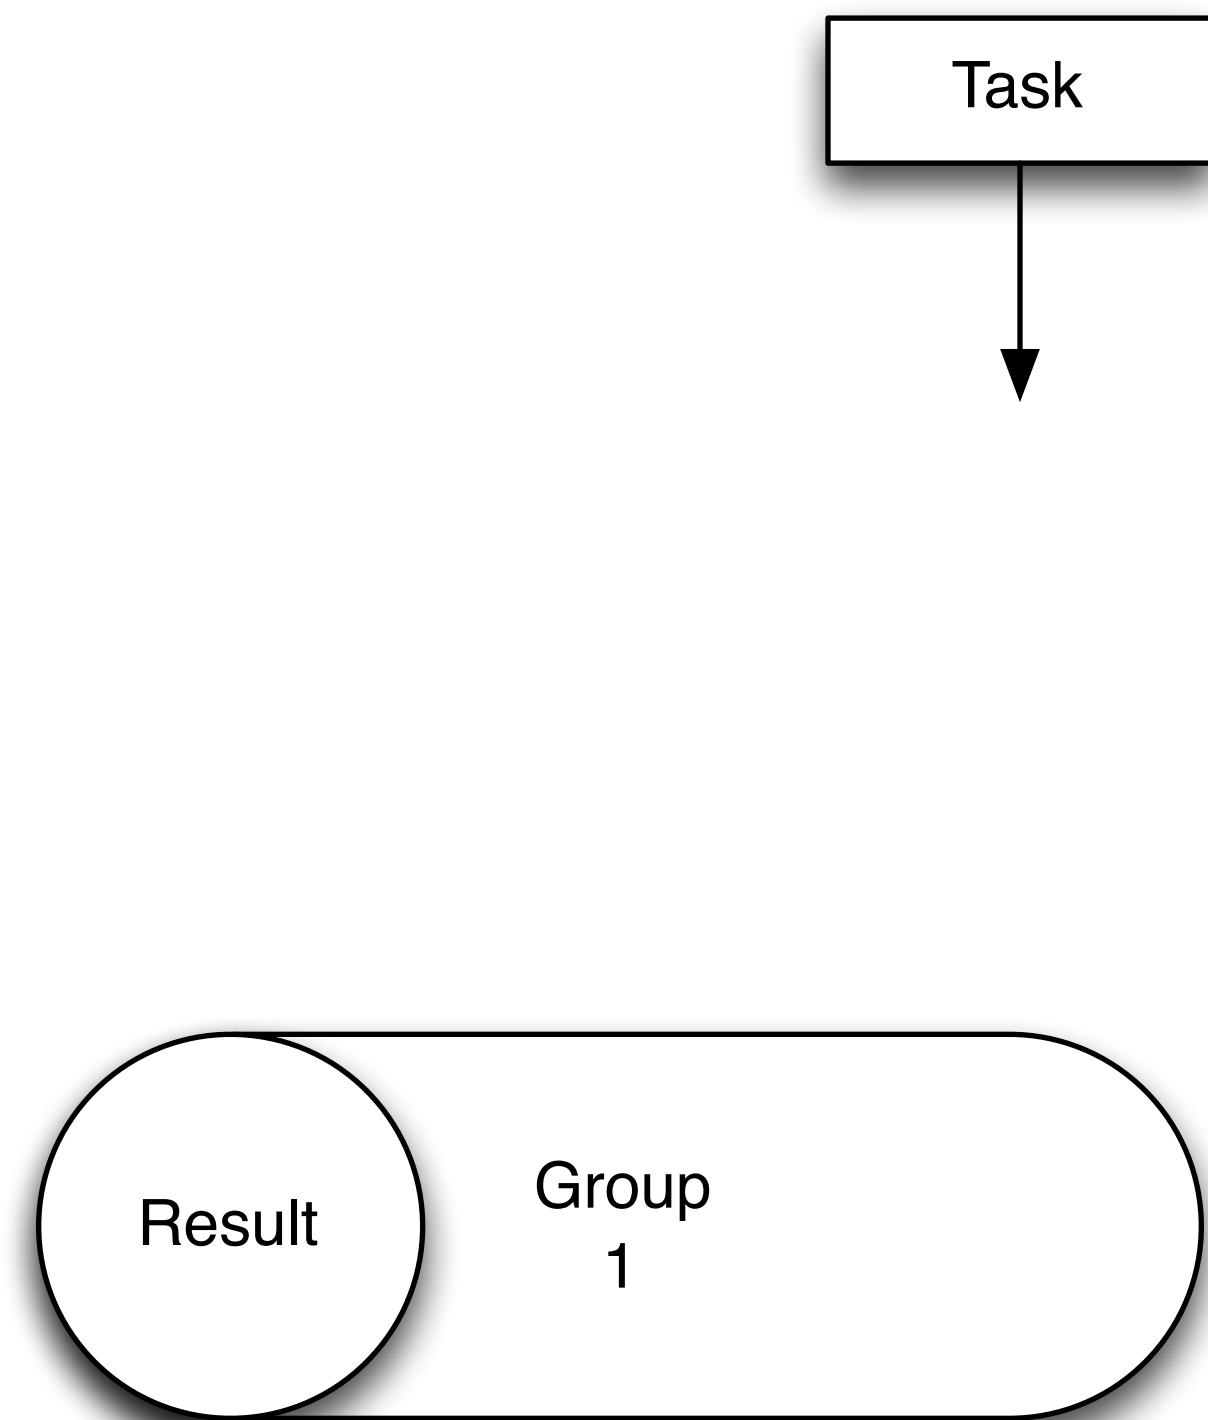
Futures: Joins



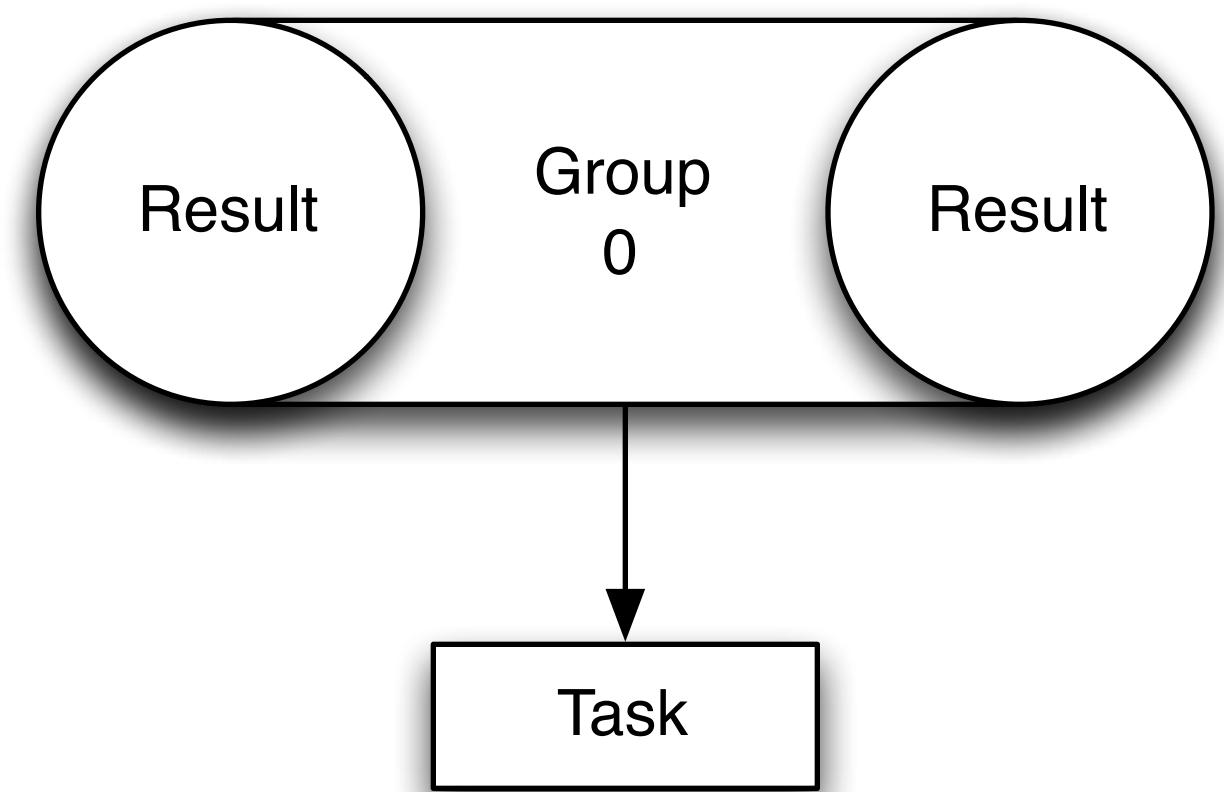
Task Systems



Task Systems



Task Systems



Futures: Continuations

```
auto x = async([]{ return fibonacci<cpp_int>(1'000'000); });
auto y = async([]{ return fibonacci<cpp_int>(2'000'000); });

auto z = when_all(std::move(x), std::move(y)).then([](auto f){
    auto t = f.get();
    return cpp_int(get<0>(t).get() * get<1>(t).get());
});

cout << z.get() << endl;
```

Futures: Continuations

```
auto x = async([]{ return fibonacci<cpp_int>(1'000'000); });
auto y = async([]{ return fibonacci<cpp_int>(2'000'000); });

auto z = when_all(std::move(x), std::move(y)).then([](auto f){
    auto t = f.get();
    return cpp_int(get<0>(t).get() * get<1>(t).get());
});

cout << z.get() << endl;
```

f is a future tuple of futures

Futures: Continuations

```
auto x = async([]{ return fibonacci<cpp_int>(1'000'000); });
auto y = async([]{ return fibonacci<cpp_int>(2'000'000); });

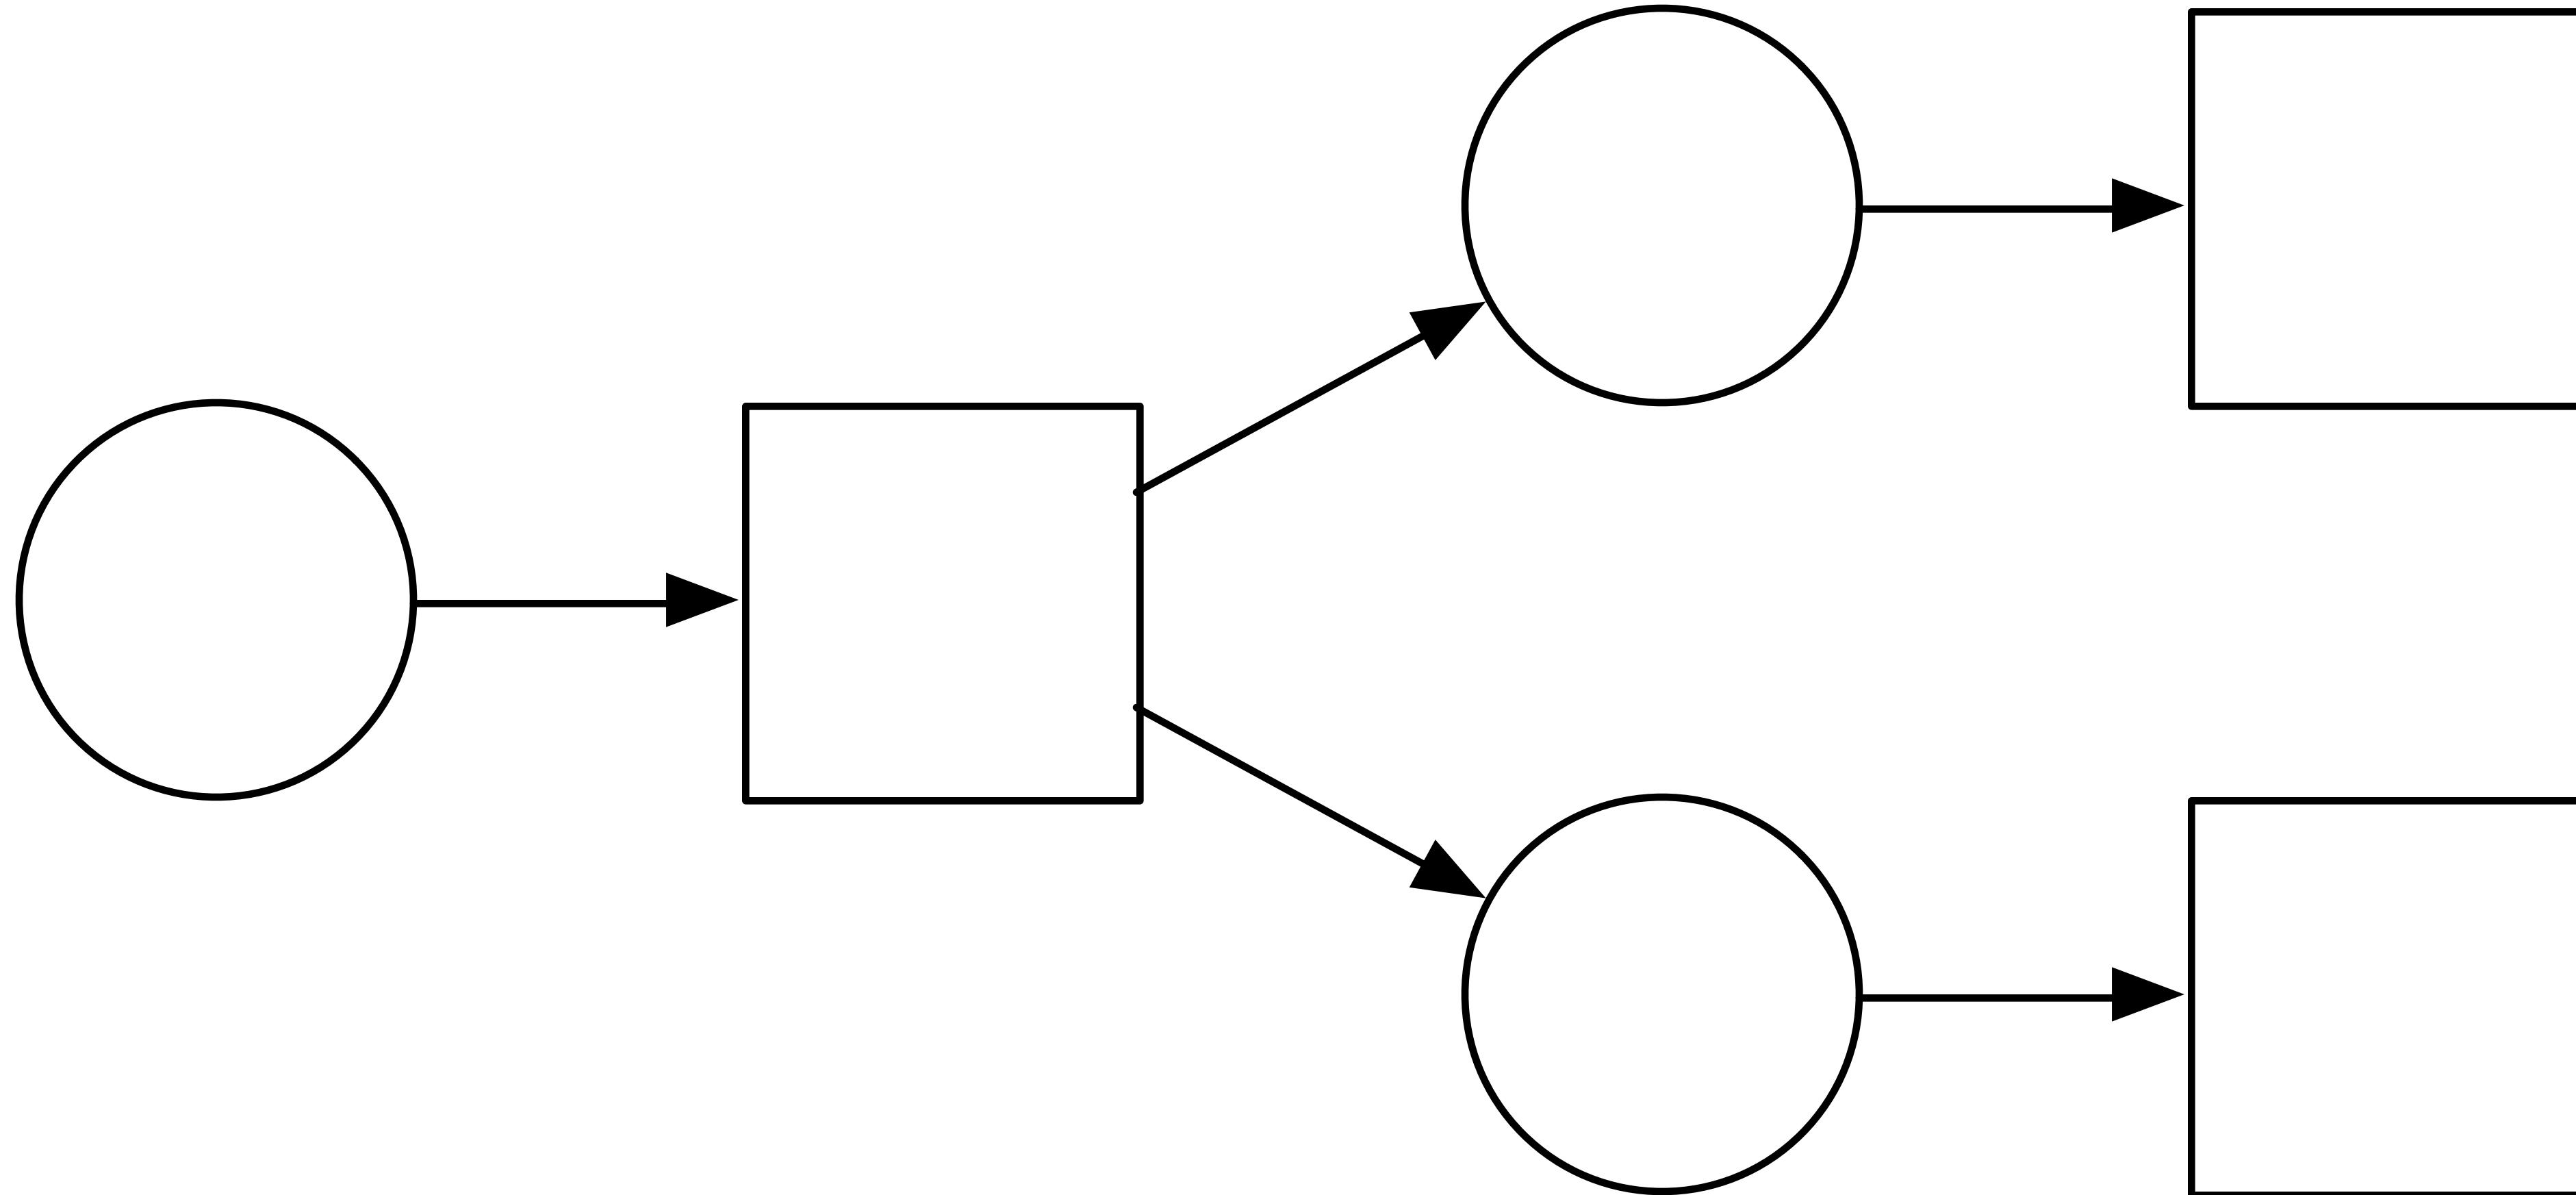
auto z = when_all(std::move(x), std::move(y)).then([](auto f){
    auto t = f.get();
    return cpp_int(get<0>(t).get() * get<1>(t).get());
});

cout << z.get() << endl;
```

f is a future tuple of futures

result is 626,964 digits

Futures: Split



Futures: Continuations

```
future<cpp_int> x = async([]{ return fibonacci<cpp_int>(100); });

future<cpp_int> y = x.then([](future<cpp_int> x){ return cpp_int(x.get() * 2); });
future<cpp_int> z = x.then([](future<cpp_int> x){ return cpp_int(x.get() / 15); });
```

Futures: Continuations

```
future<cpp_int> x = async([]{ return fibonacci<cpp_int>(100); });

future<cpp_int> y = x.then([](future<cpp_int> x){ return cpp_int(x.get() * 2); });

future<cpp_int> z = x.then([](future<cpp_int> x){ return cpp_int(x.get() / 15); });
```

Thread 1: signal SIGABRT

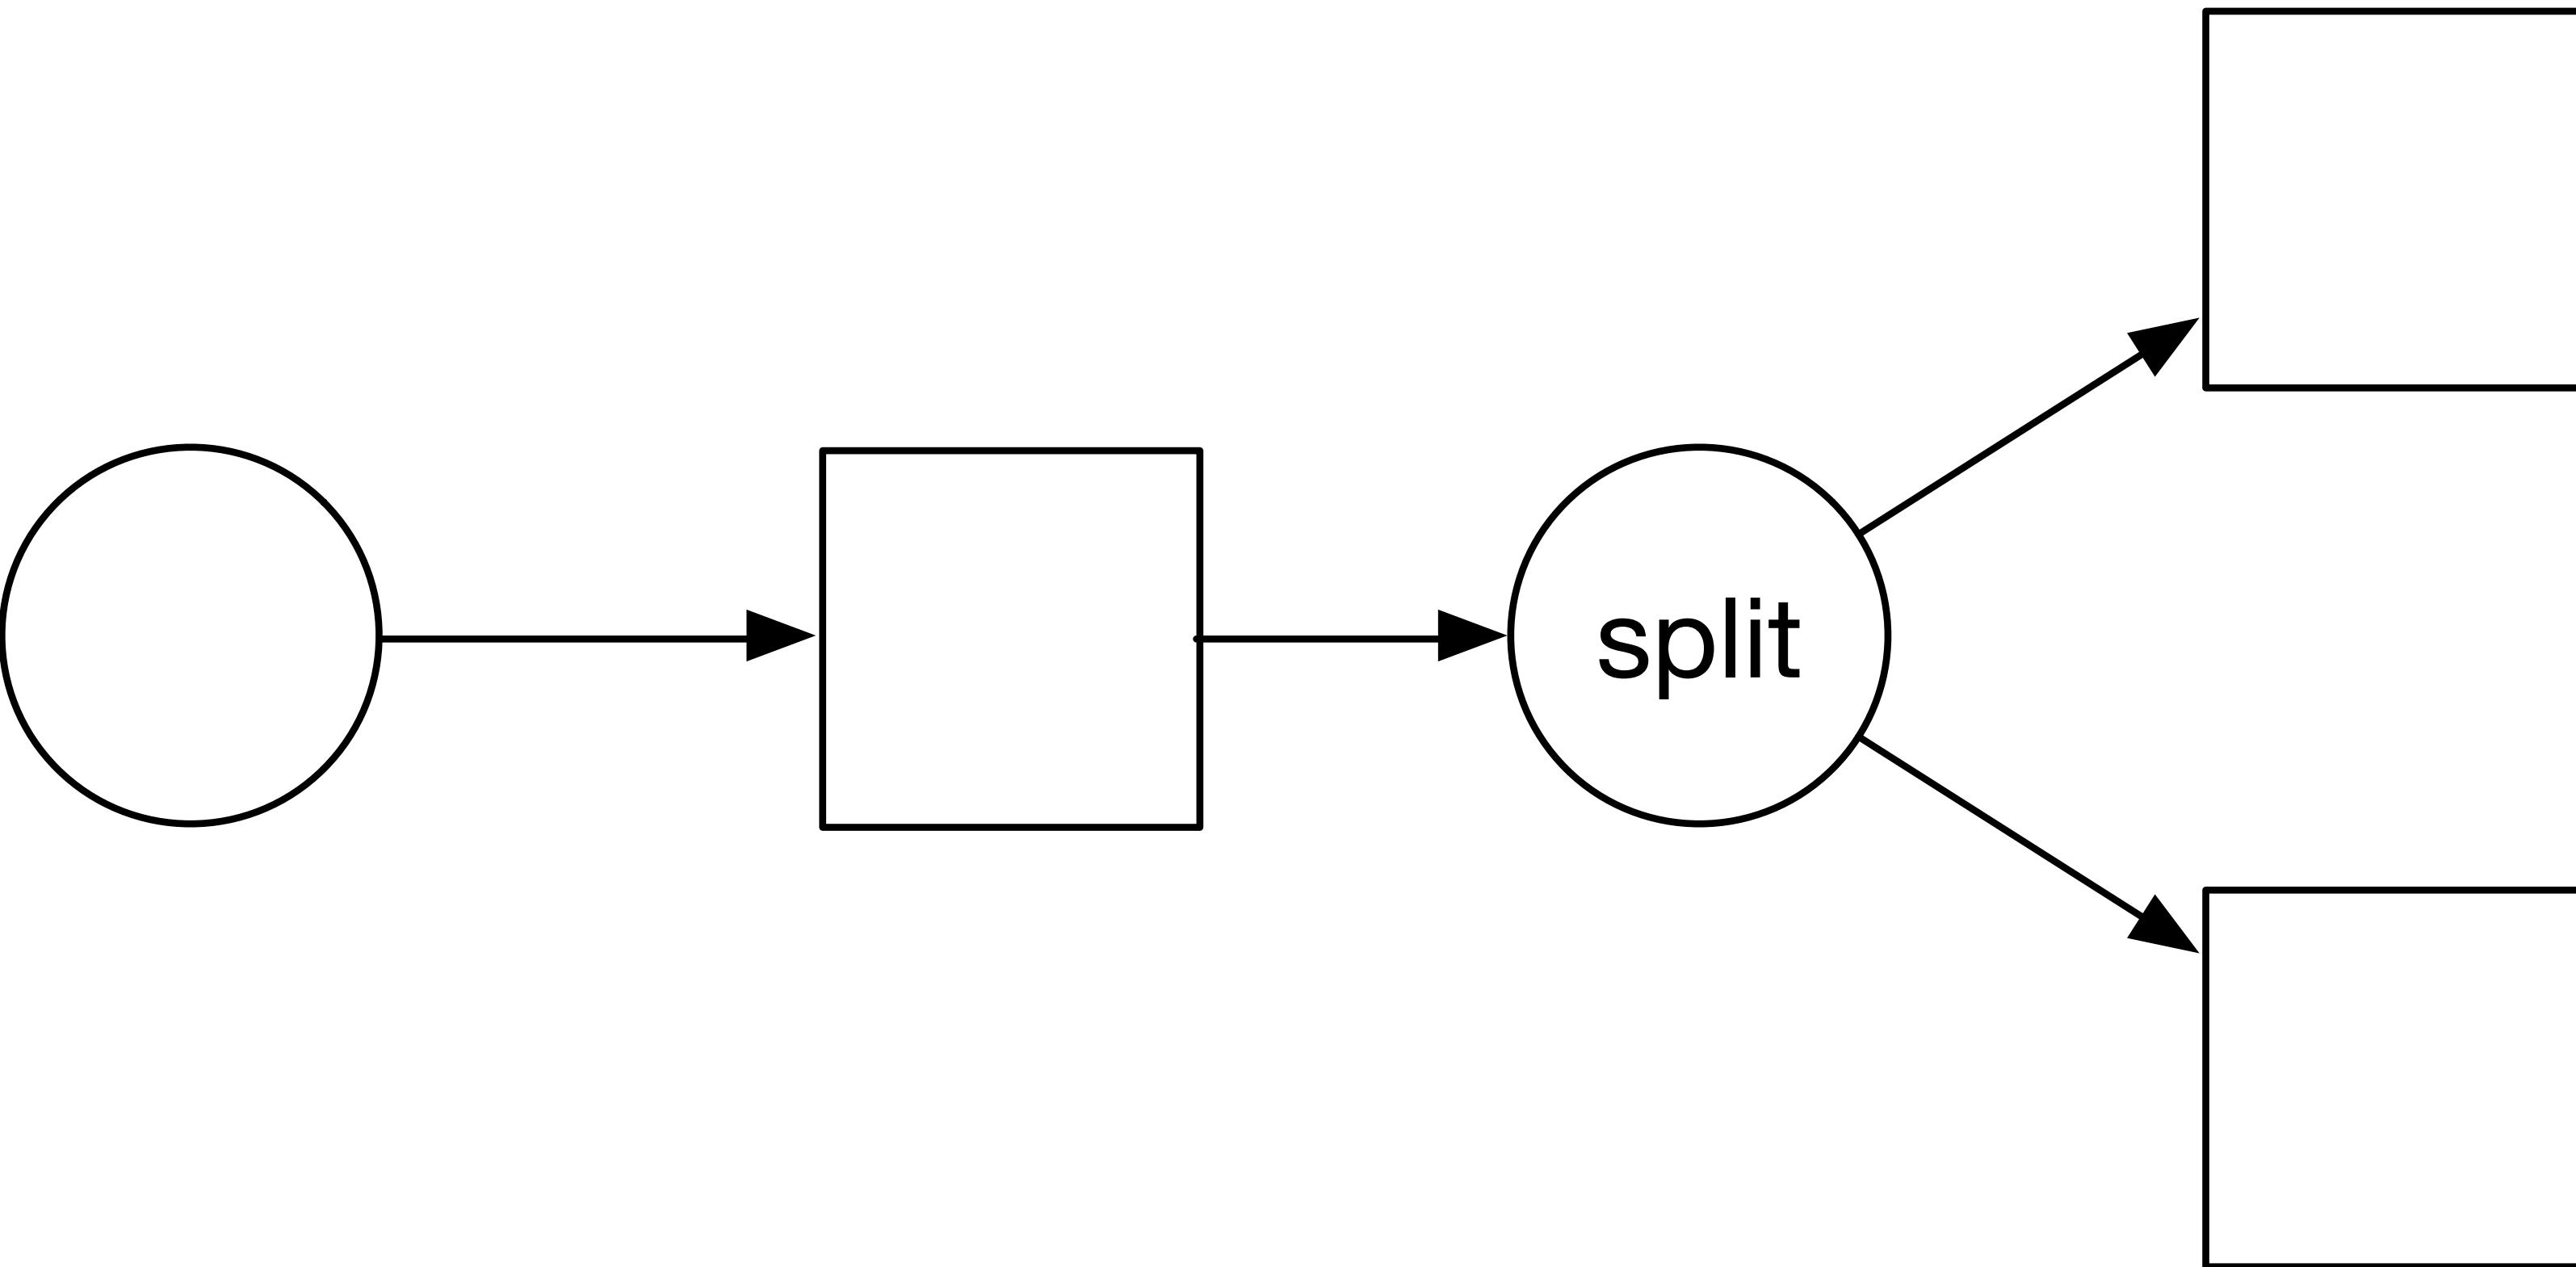
Assertion failed: (px != 0), function operator->, file shared_ptr.hpp, line 648.

Continuations

- Desired behavior
 - A future should behave as a *regular type* - a token for the actual value
 - `shared_futures` let me “copy” them around and do multiple `get()` operations
 - But not multiple continuations

Continuations

- We can write a pseudo-copy, split().



Futures: Continuations

Futures: Continuations

```
future<cpp_int> x = async([]{ return fibonacci<cpp_int>(100); });

future<cpp_int> y = split(x).then([](future<cpp_int> x){ return cpp_int(x.get() * 2); });
future<cpp_int> z = x.then([](future<cpp_int> x){ return cpp_int(x.get() / 15); });
```

Futures: Continuations

```
future<cpp_int> x = async([]{ return fibonacci<cpp_int>(100); });

future<cpp_int> y = split(x).then([](future<cpp_int> x){ return cpp_int(x.get() * 2); });
future<cpp_int> z = x.then([](future<cpp_int> x){ return cpp_int(x.get() / 15); });

future<void> done = when_all(std::move(y), std::move(z)).then([](auto f){
    auto t = f.get();
    cout << get<0>(t).get() << endl;
    cout << get<1>(t).get() << endl;
});

done.wait();
```

Futures: Continuations

```
future<cpp_int> x = async([]{ return fibonacci<cpp_int>(100); });

future<cpp_int> y = split(x).then([](future<cpp_int> x){ return cpp_int(x.get() * 2); });
future<cpp_int> z = x.then([](future<cpp_int> x){ return cpp_int(x.get() / 15); });

future<void> done = when_all(std::move(y), std::move(z)).then([](auto f){
    auto t = f.get();
    cout << get<0>(t).get() << endl;
    cout << get<1>(t).get() << endl;
});

done.wait();
```

**708449696358523830150
23614989878617461005**

Building Blocks

- Promise is the sending side of a future
- Promises are packaged with a function to formed a packaged task
 - Packaged tasks handle the exception marshalling through a promise

Promise

```
promise<int> x;
future<int> y = x.get_future();

x.set_value(42);
cout << y.get() << endl;
```

Promise

```
promise<int> x;
future<int> y = x.get_future();

x.set_value(42);
cout << y.get() << endl;
```

42

Futures: Split

```
template <typename T>
auto split(future<T>& x) {

    auto tmp = std::move(x);

    promise<T> p;
    x = p.get_future(); // replace x with new future

    return tmp.then([_p = move(p)](auto _tmp) mutable {
        auto value = _tmp.get();
        _p.set_value(value); // assign to new "x" future
        return value; // return value through future result
    });
}
```

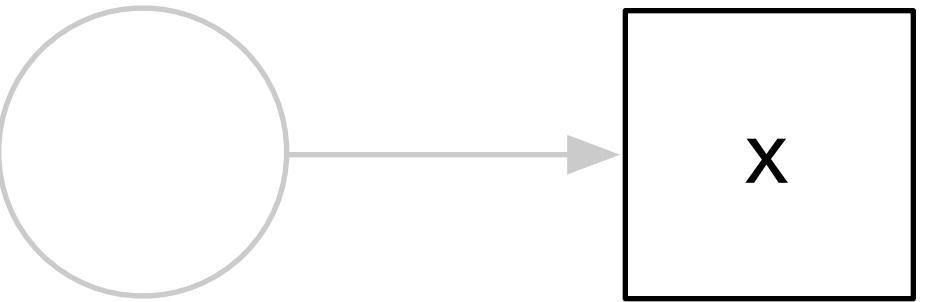
Futures: Split

```
template <typename T>
auto split(future<T>& x) {

    auto tmp = std::move(x);

    promise<T> p;
    x = p.get_future(); // replace x with new future

    return tmp.then([_p = move(p)](auto _tmp) mutable {
        auto value = _tmp.get();
        _p.set_value(value); // assign to new "x" future
        return value; // return value through future result
    });
}
```

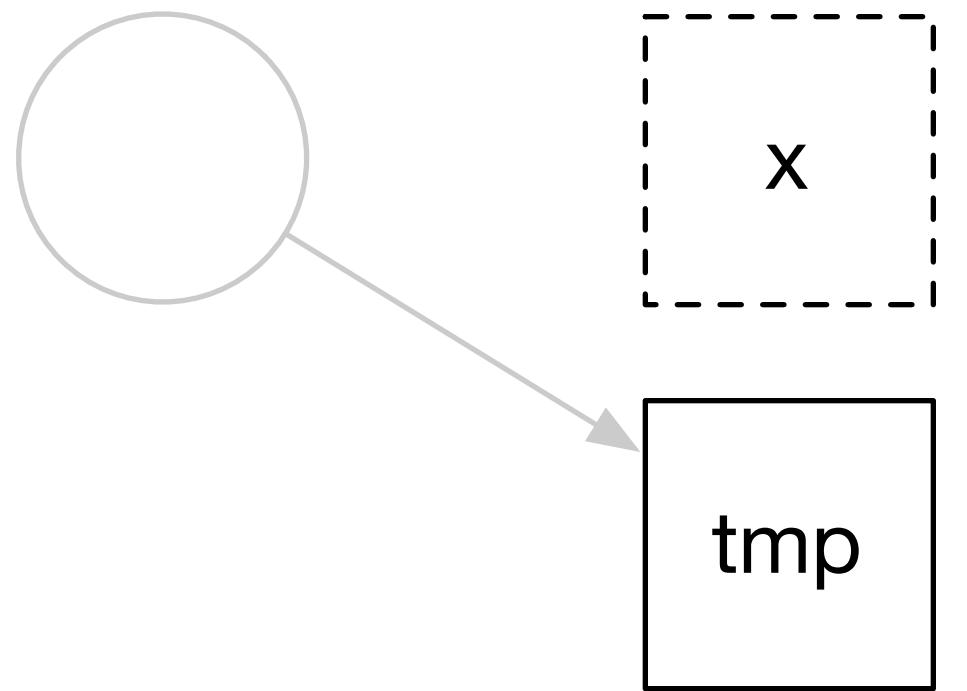


Futures: Split

```
template <typename T>
auto split(future<T>& x) {
    auto tmp = std::move(x);

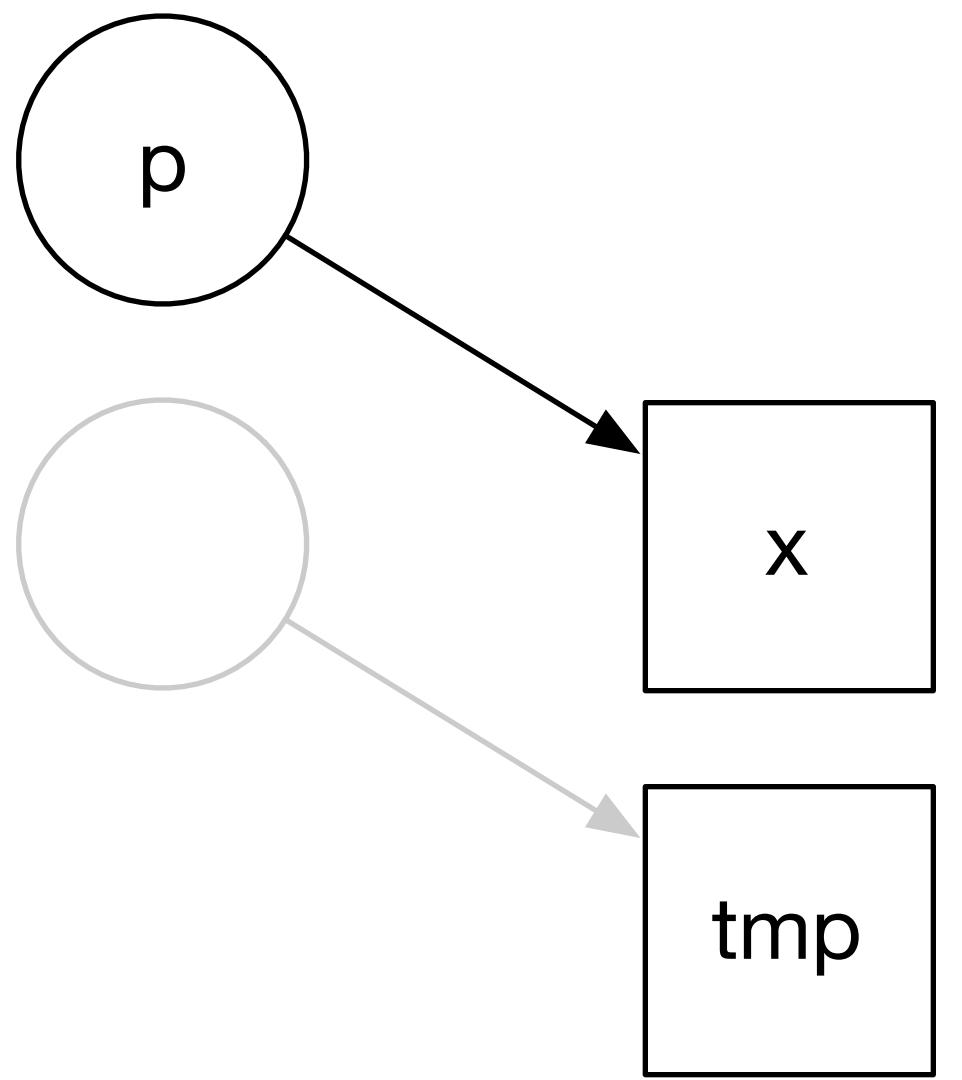
    promise<T> p;
    x = p.get_future(); // replace x with new future

    return tmp.then([_p = move(p)](auto _tmp) mutable {
        auto value = _tmp.get();
        _p.set_value(value); // assign to new "x" future
        return value; // return value through future result
    });
}
```



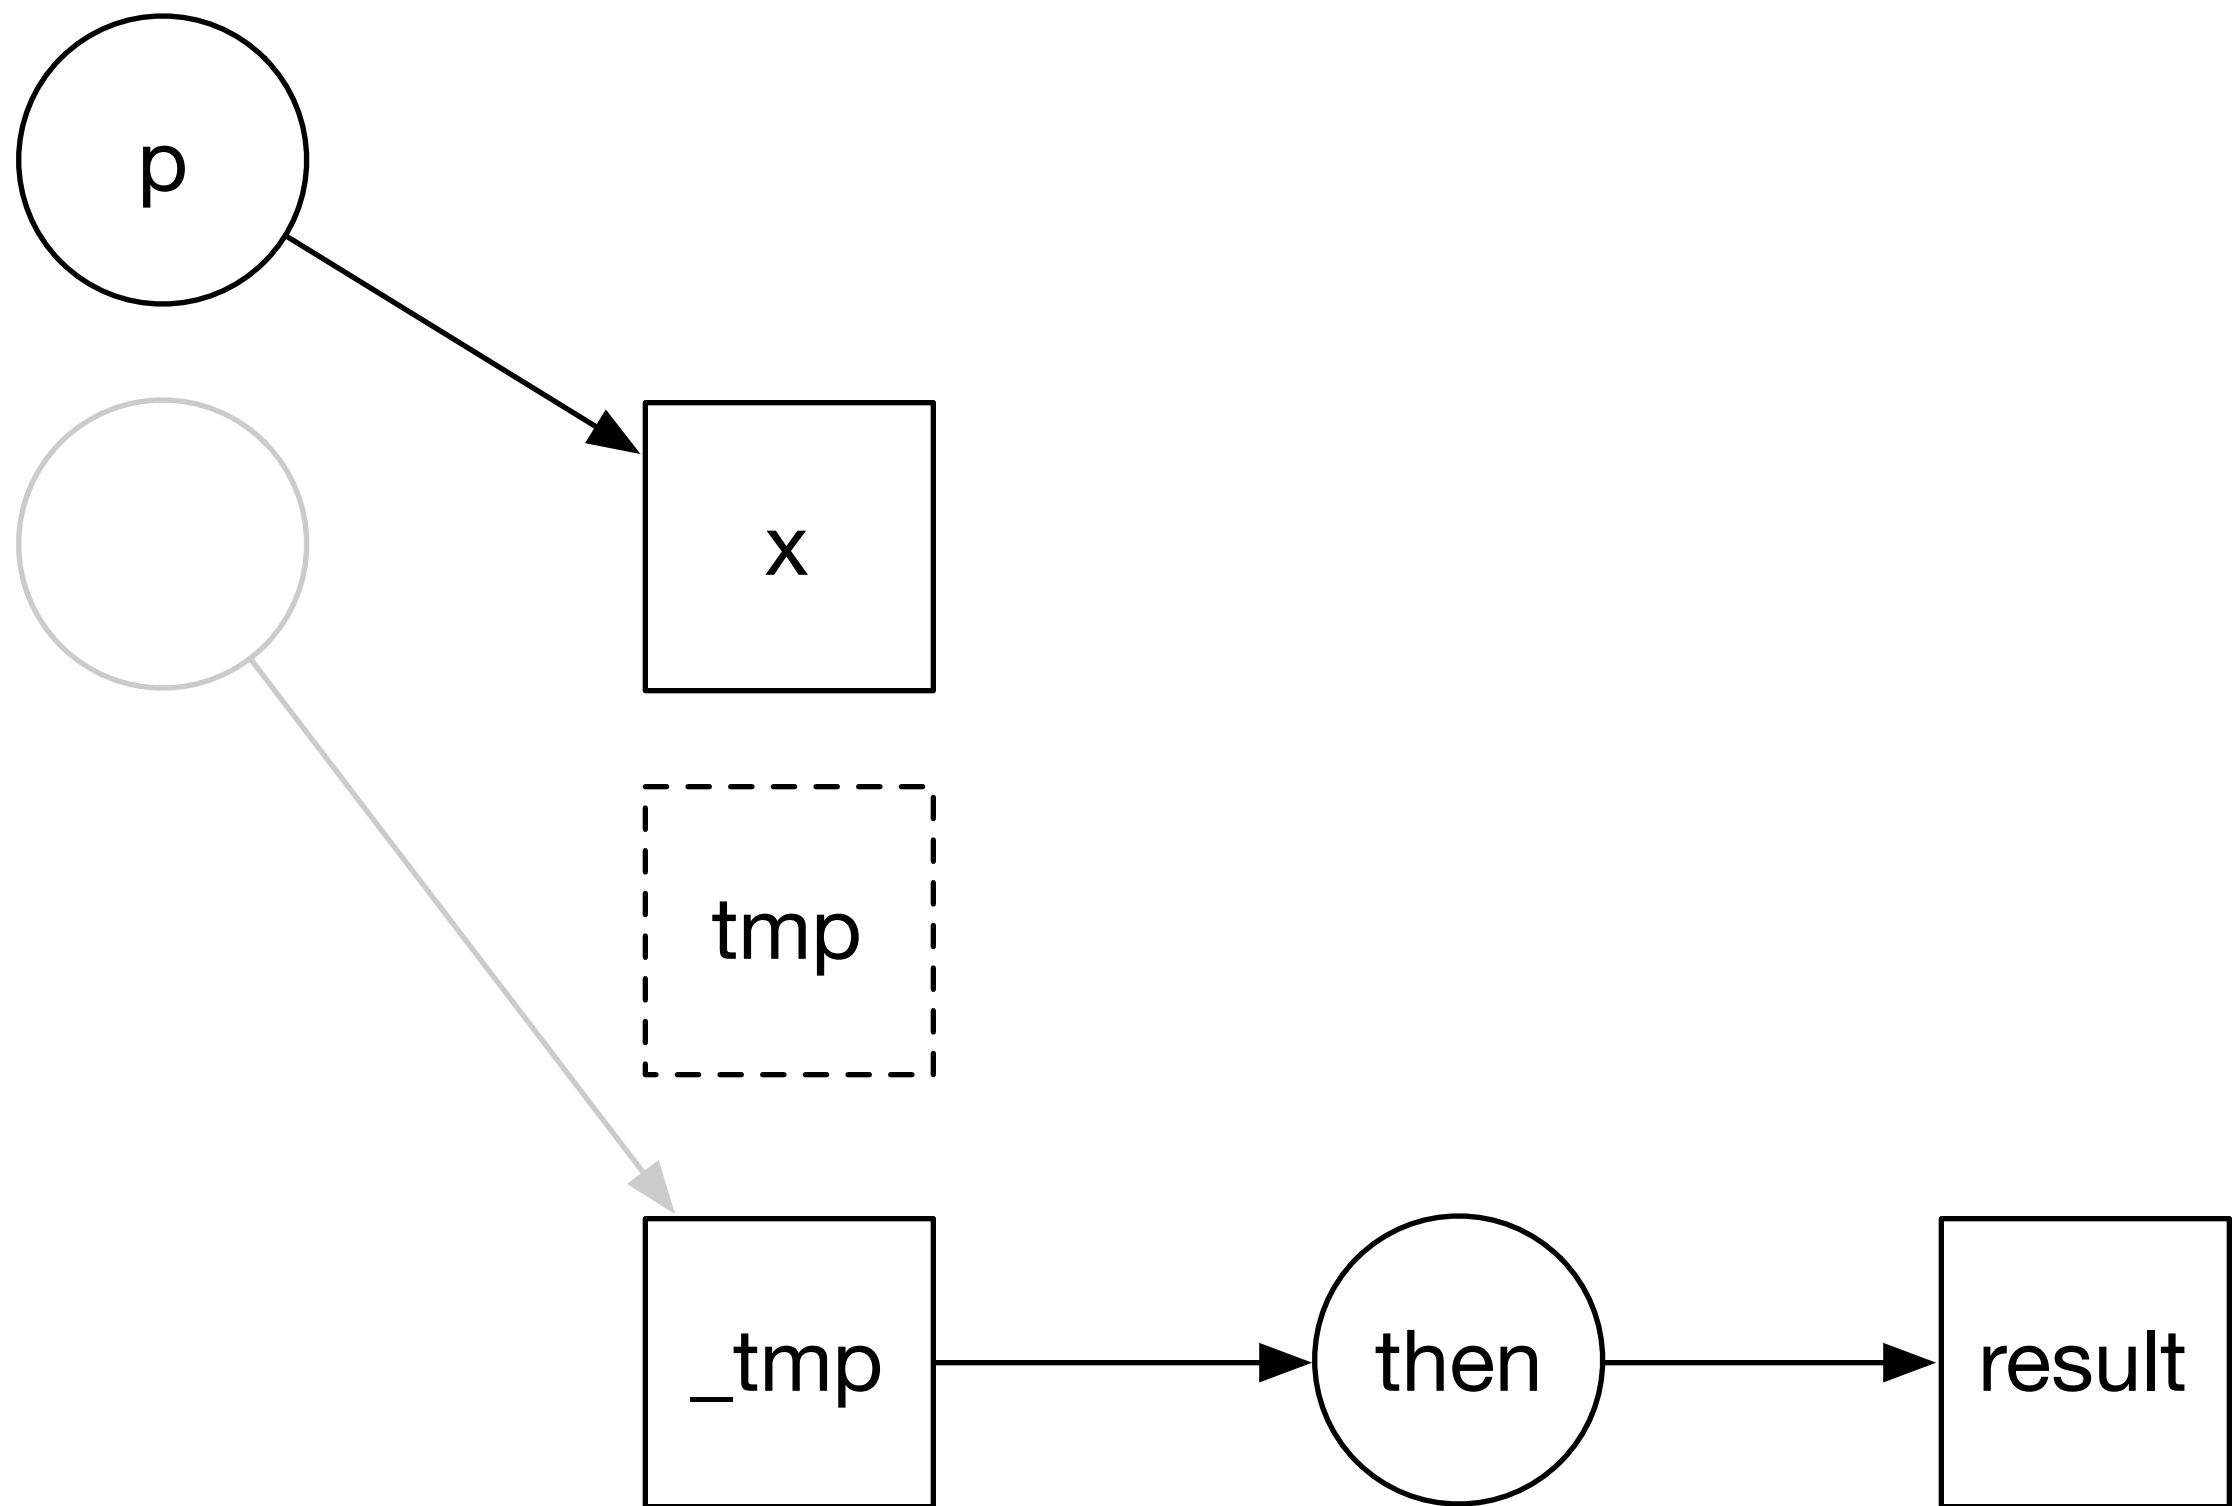
Futures: Split

```
template <typename T>
auto split(future<T>& x) {
    auto tmp = std::move(x);
    promise<T> p;
    x = p.get_future(); // replace x with new future
    return tmp.then([_p = move(p)](auto _tmp) mutable {
        auto value = _tmp.get();
        _p.set_value(value); // assign to new "x" future
        return value; // return value through future result
    });
}
```



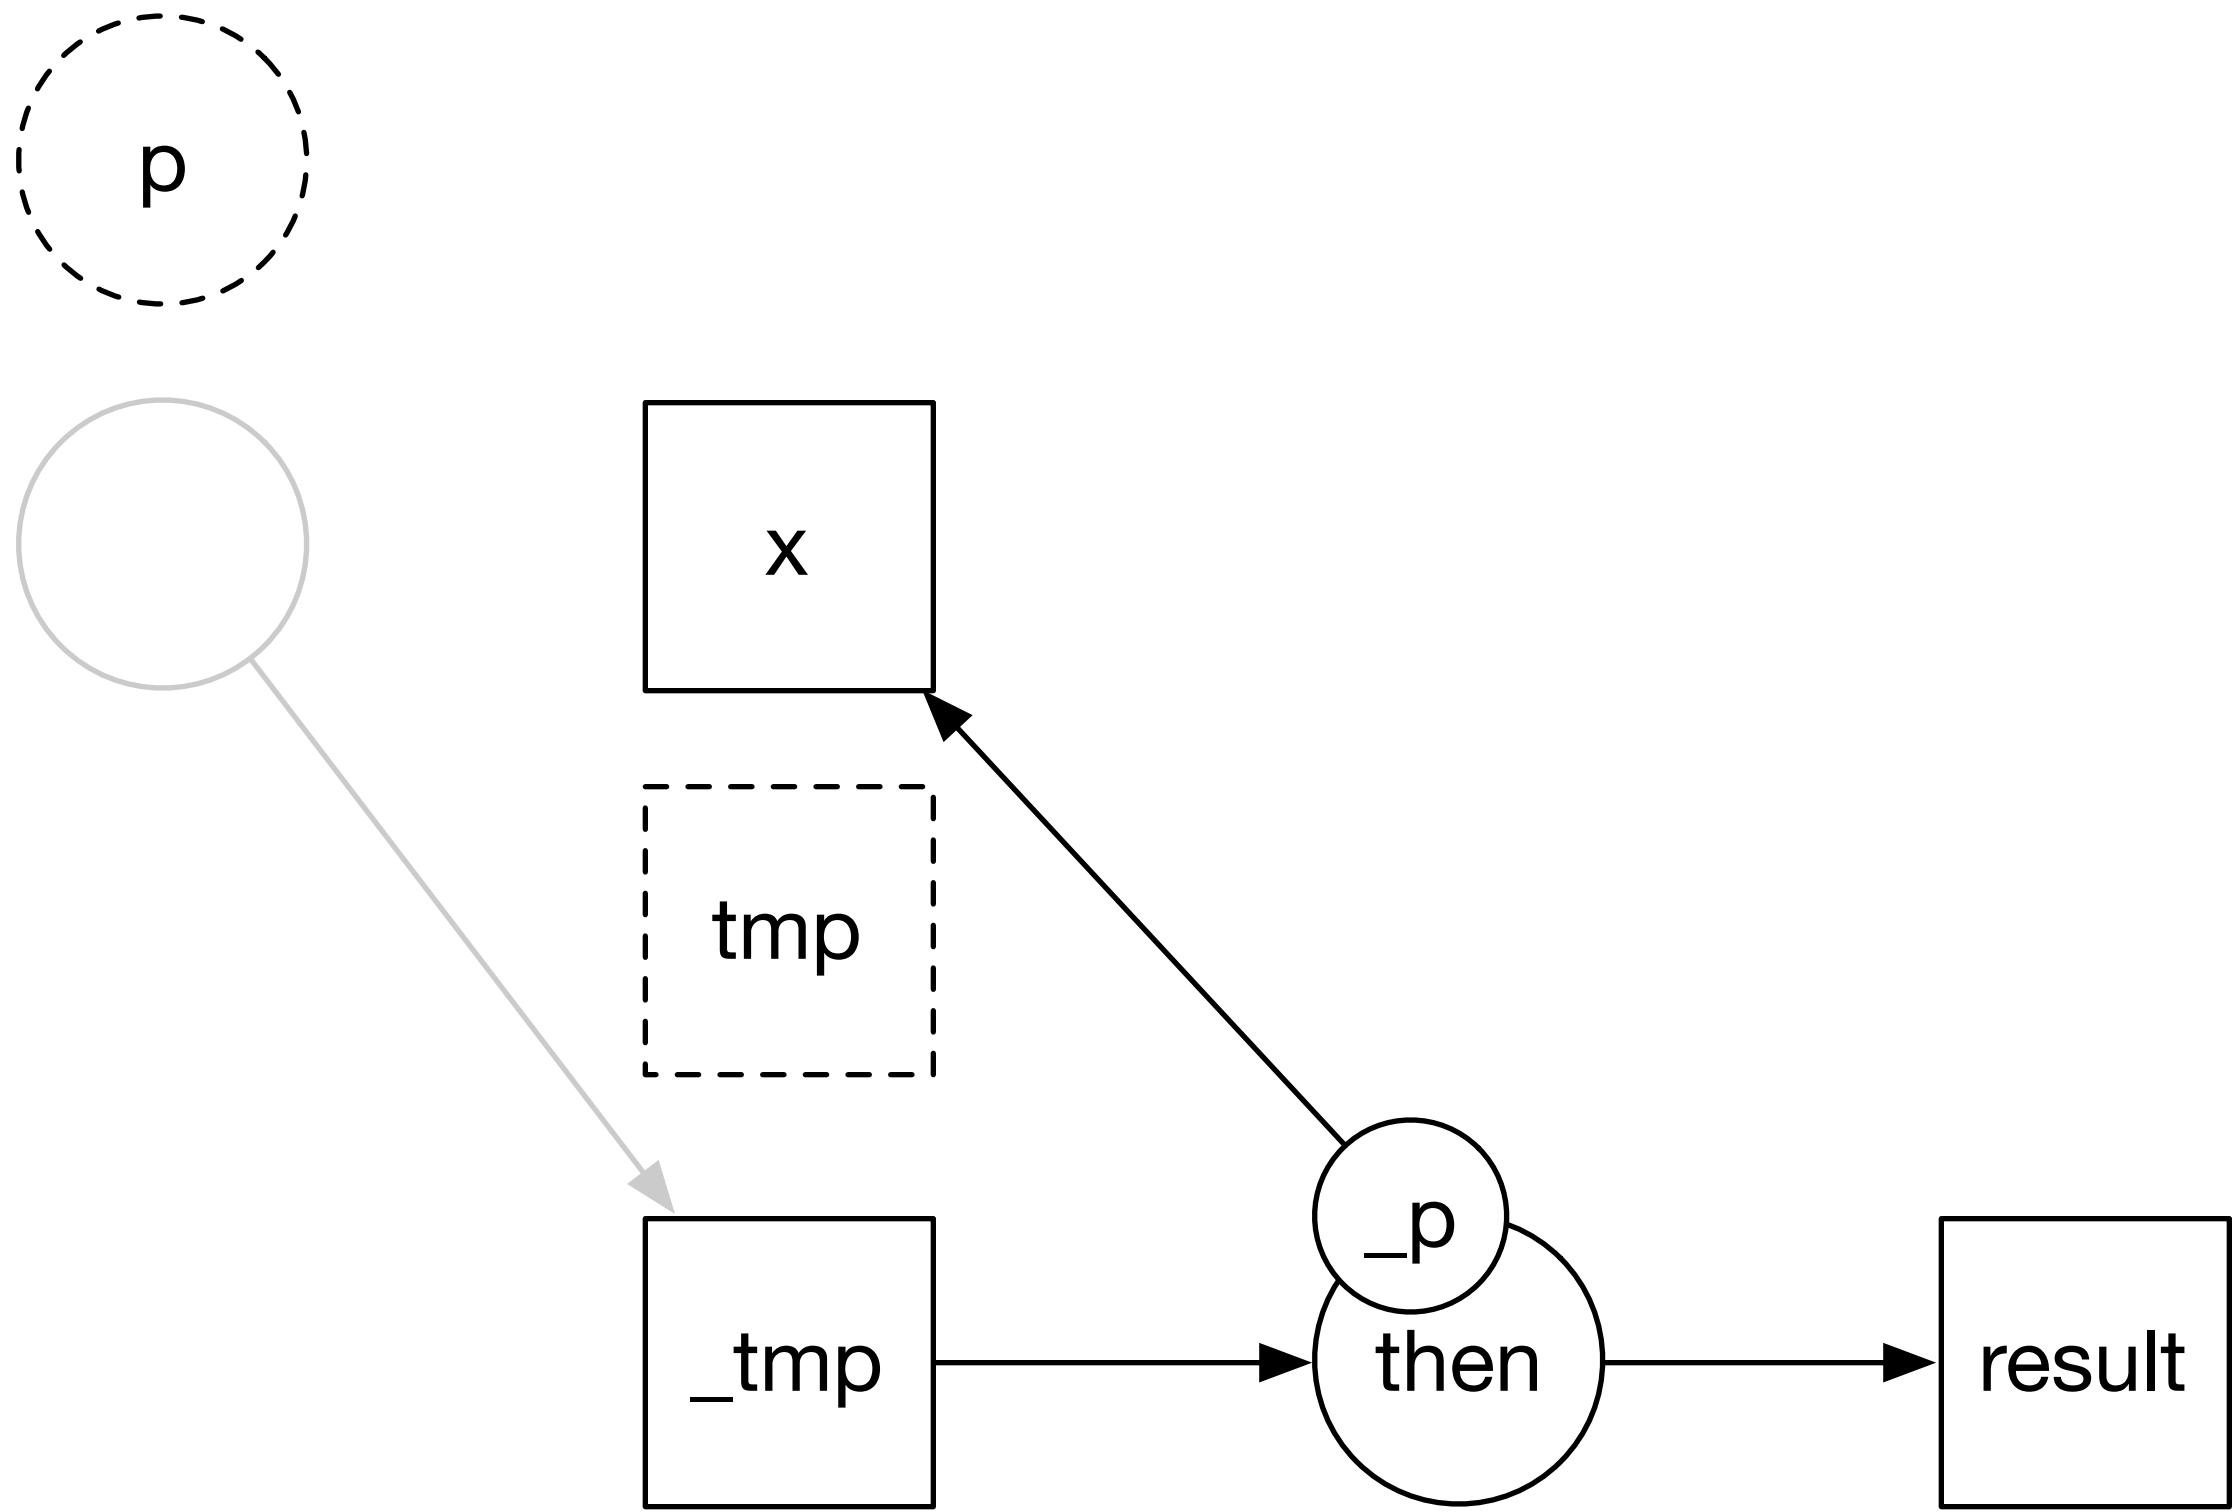
Futures: Split

```
template <typename T>
auto split(future<T>& x) {
    auto tmp = std::move(x);
    promise<T> p;
    x = p.get_future(); // replace x with new future
    return tmp.then([_p = move(p)](auto _tmp) mutable {
        auto value = _tmp.get();
        _p.set_value(value); // assign to new "x" future
        return value; // return value through future result
    });
}
```



Futures: Split

```
template <typename T>
auto split(future<T>& x) {
    auto tmp = std::move(x);
    promise<T> p;
    x = p.get_future(); // replace x with new future
    return tmp.then([_p = move(p)](auto _tmp) mutable {
        auto value = _tmp.get();
        _p.set_value(value); // assign to new "x" future
        return value; // return value through future result
    });
}
```



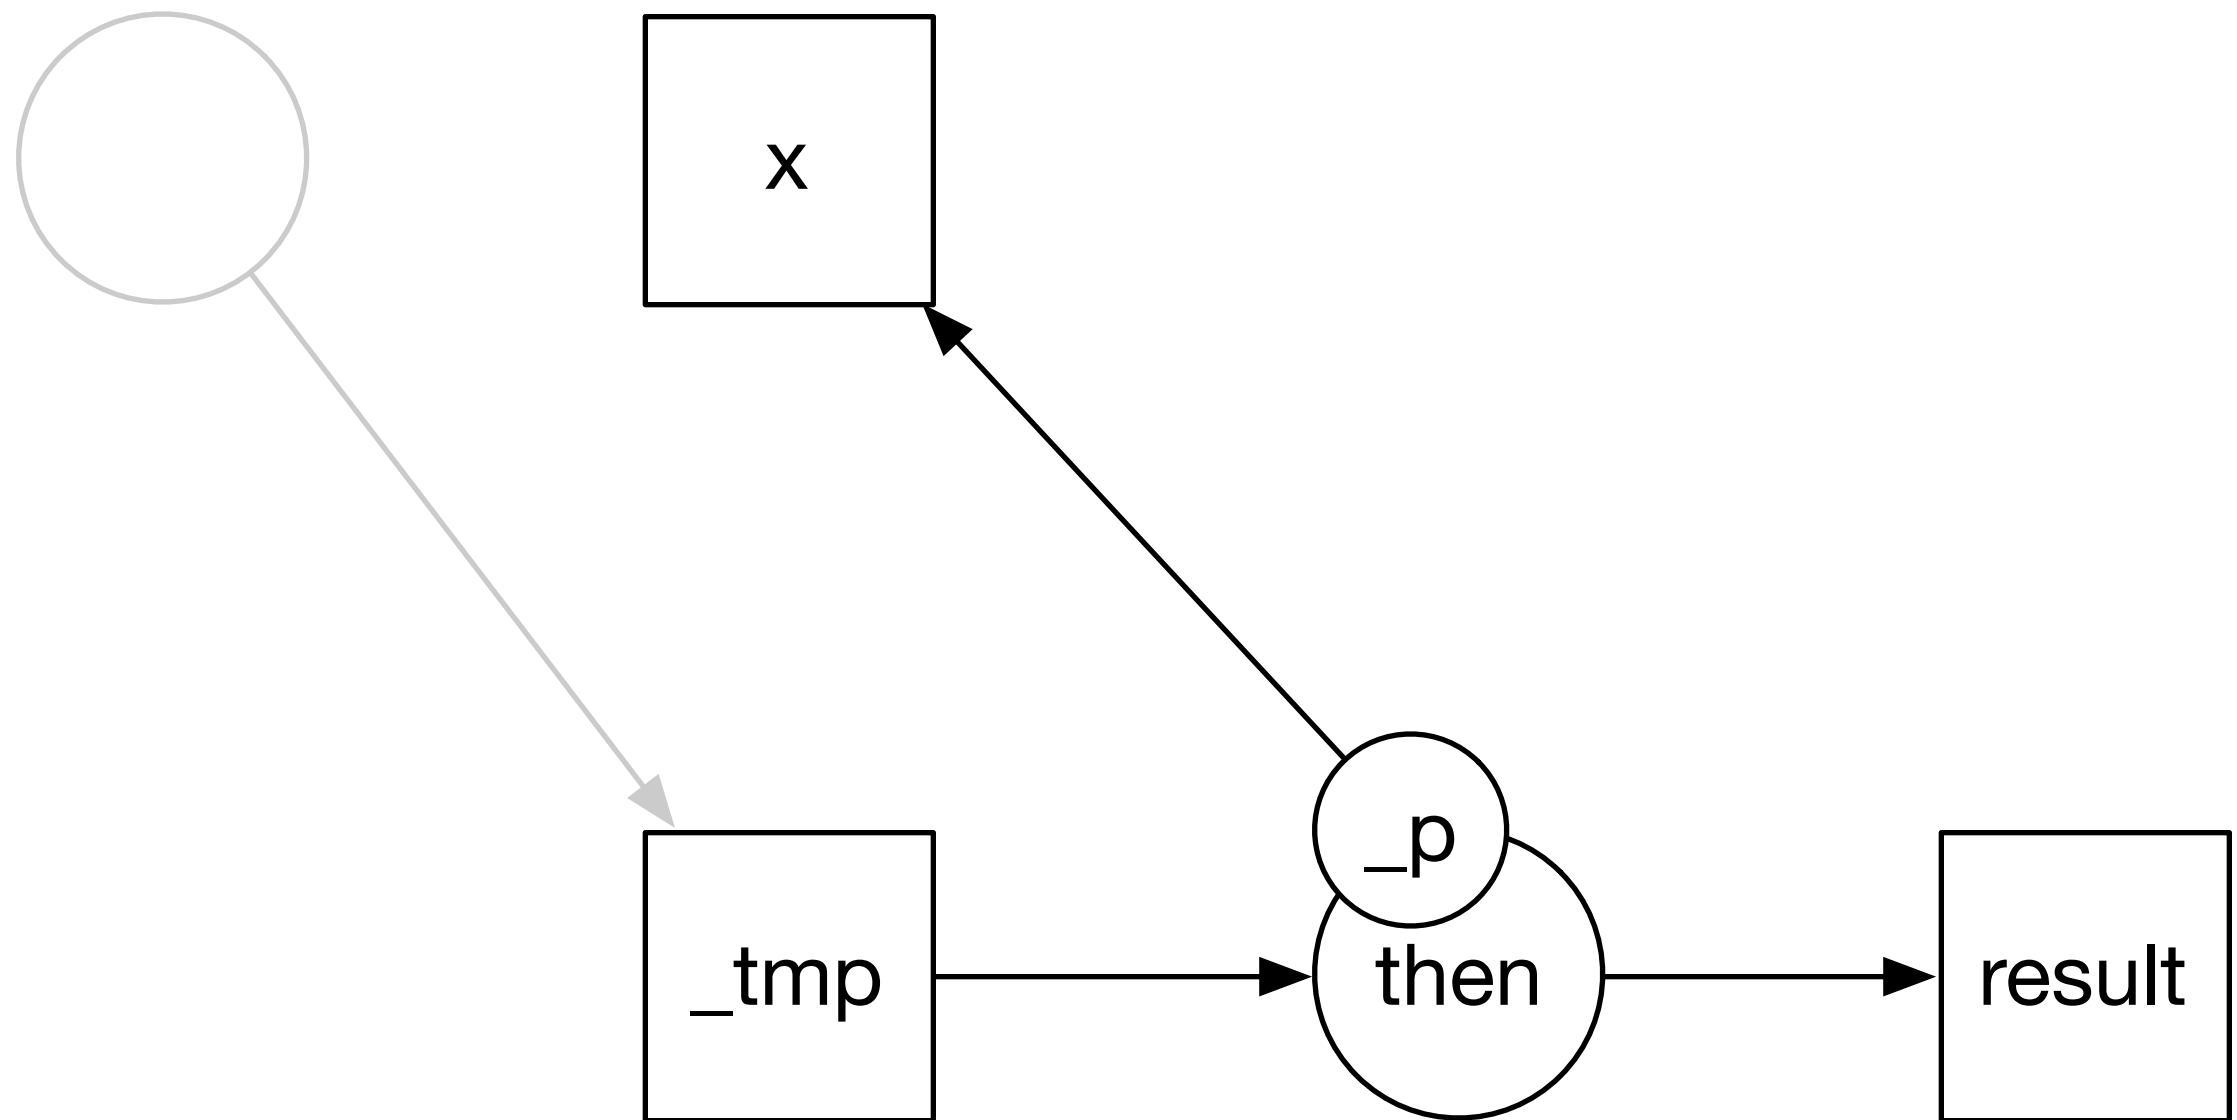
Futures: Split

```
template <typename T>
auto split(future<T>& x) {

    auto tmp = std::move(x);

    promise<T> p;
    x = p.get_future(); // replace x with new future

    return tmp.then([_p = move(p)](auto _tmp) mutable {
        auto value = _tmp.get();
        _p.set_value(value); // assign to new "x" future
        return value; // return value through future result
    });
}
```



Futures: Split

```
template <typename T>
auto split(future<T>& x) {

    auto tmp = std::move(x);

    promise<T> p;
    x = p.get_future(); // replace x with new future

    return tmp.then([_p = std::move(p)](auto _tmp) mutable {
        if (_tmp.has_exception()) {
            auto error = _tmp.get_exception_ptr();
            _p.set_exception(error);
            rethrow_exception(error);
        }

        auto value = _tmp.get();
        _p.set_value(value); // assign to new "x" future
        return value; // return value through future result
    });
}
```

Futures: Continuations

```
future<cpp_int> x = async([]{ return fibonacci<cpp_int>(100); });

future<cpp_int> y = split(x).then([](future<cpp_int> x){ return cpp_int(x.get() * 2); });
future<cpp_int> z = x.then([](future<cpp_int> x){ return cpp_int(x.get() / 15); });

future<void> done = when_all(std::move(y), std::move(z)).then([](auto f){
    auto t = f.get();
    cout << get<0>(t).get() << endl;
    cout << get<1>(t).get() << endl;
});

done.wait();
```

**708449696358523830150
23614989878617461005**

Cancelation

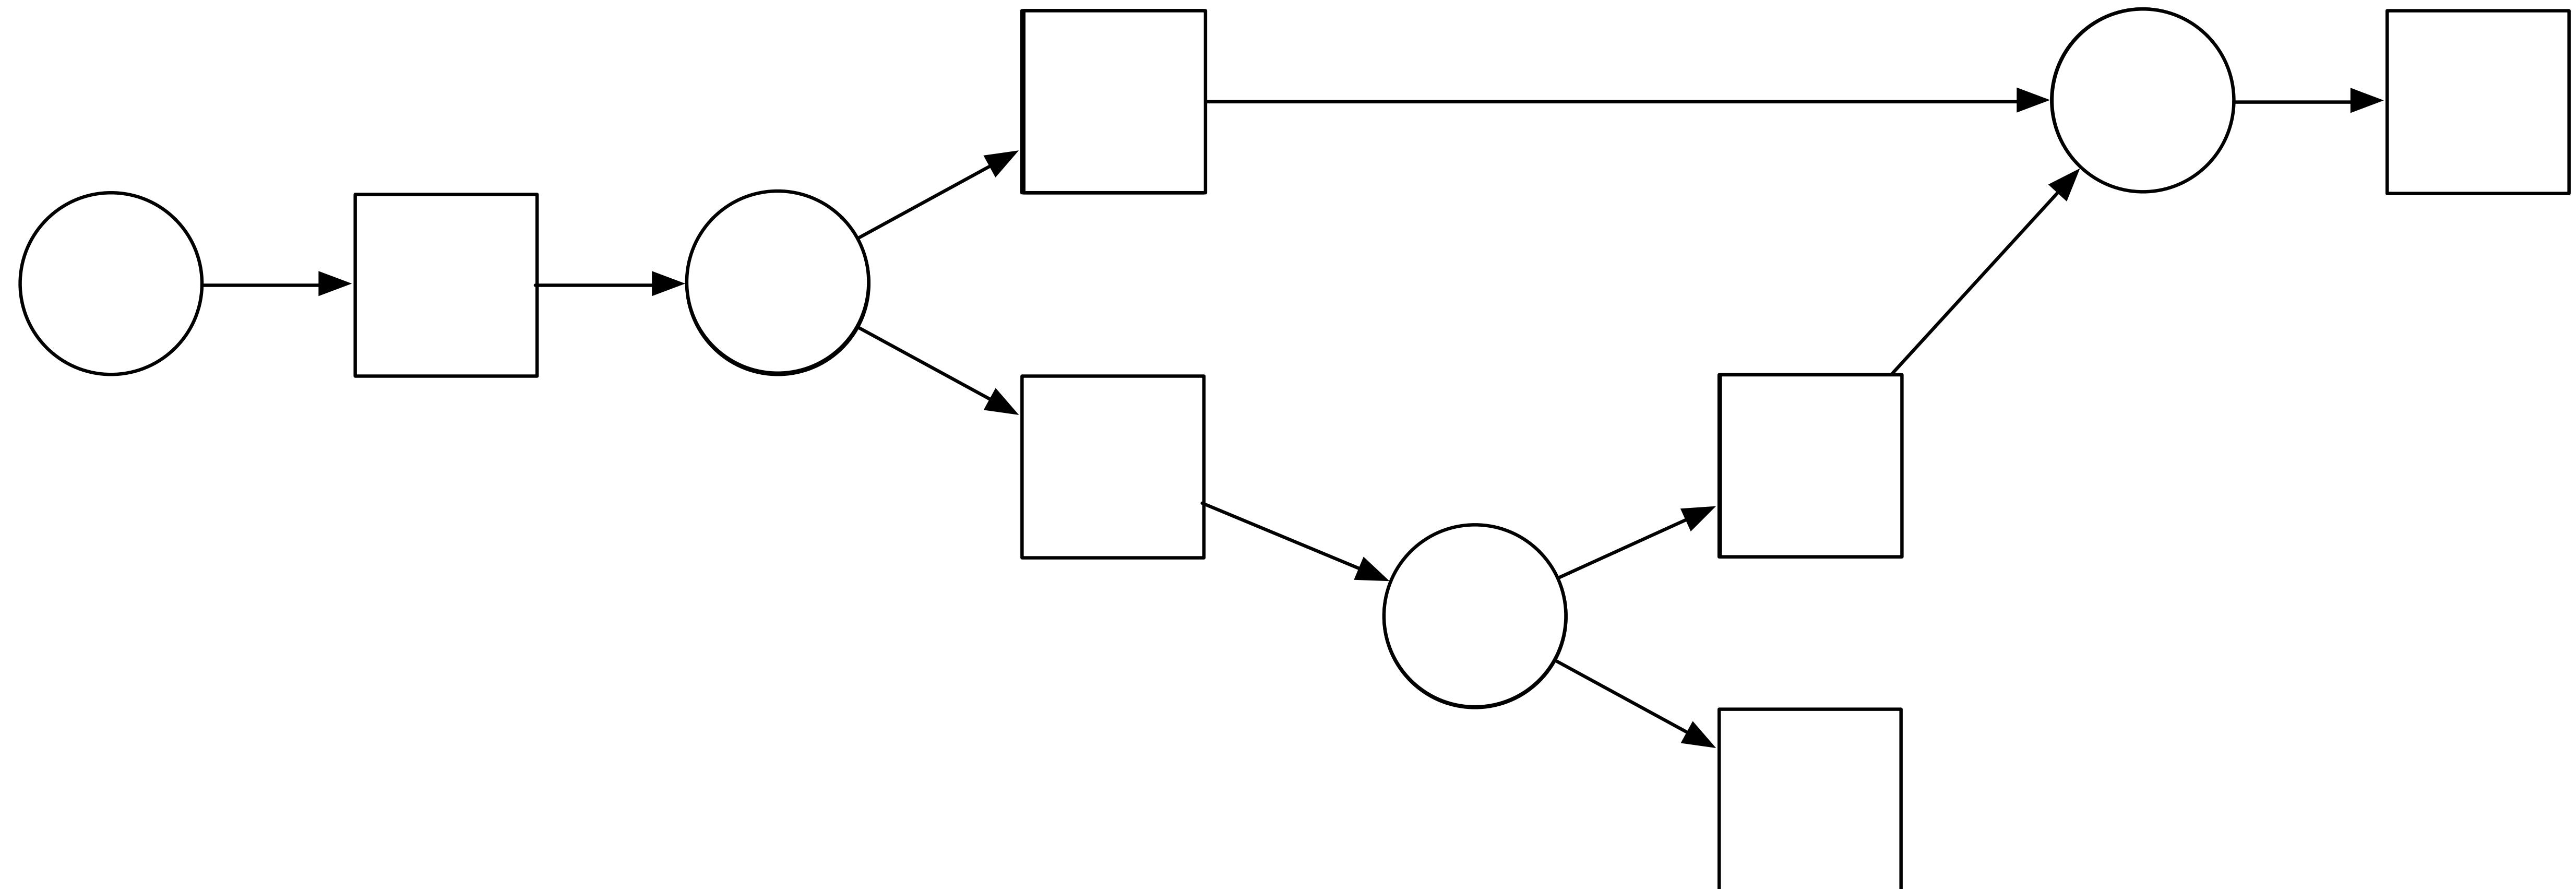
Cancelation

- When the (last) future destructs
 - The associated task that has not started, should not execute (NOP)
 - The resource held by that task should be released
 - Since that task may hold futures for other tasks, the system unravels

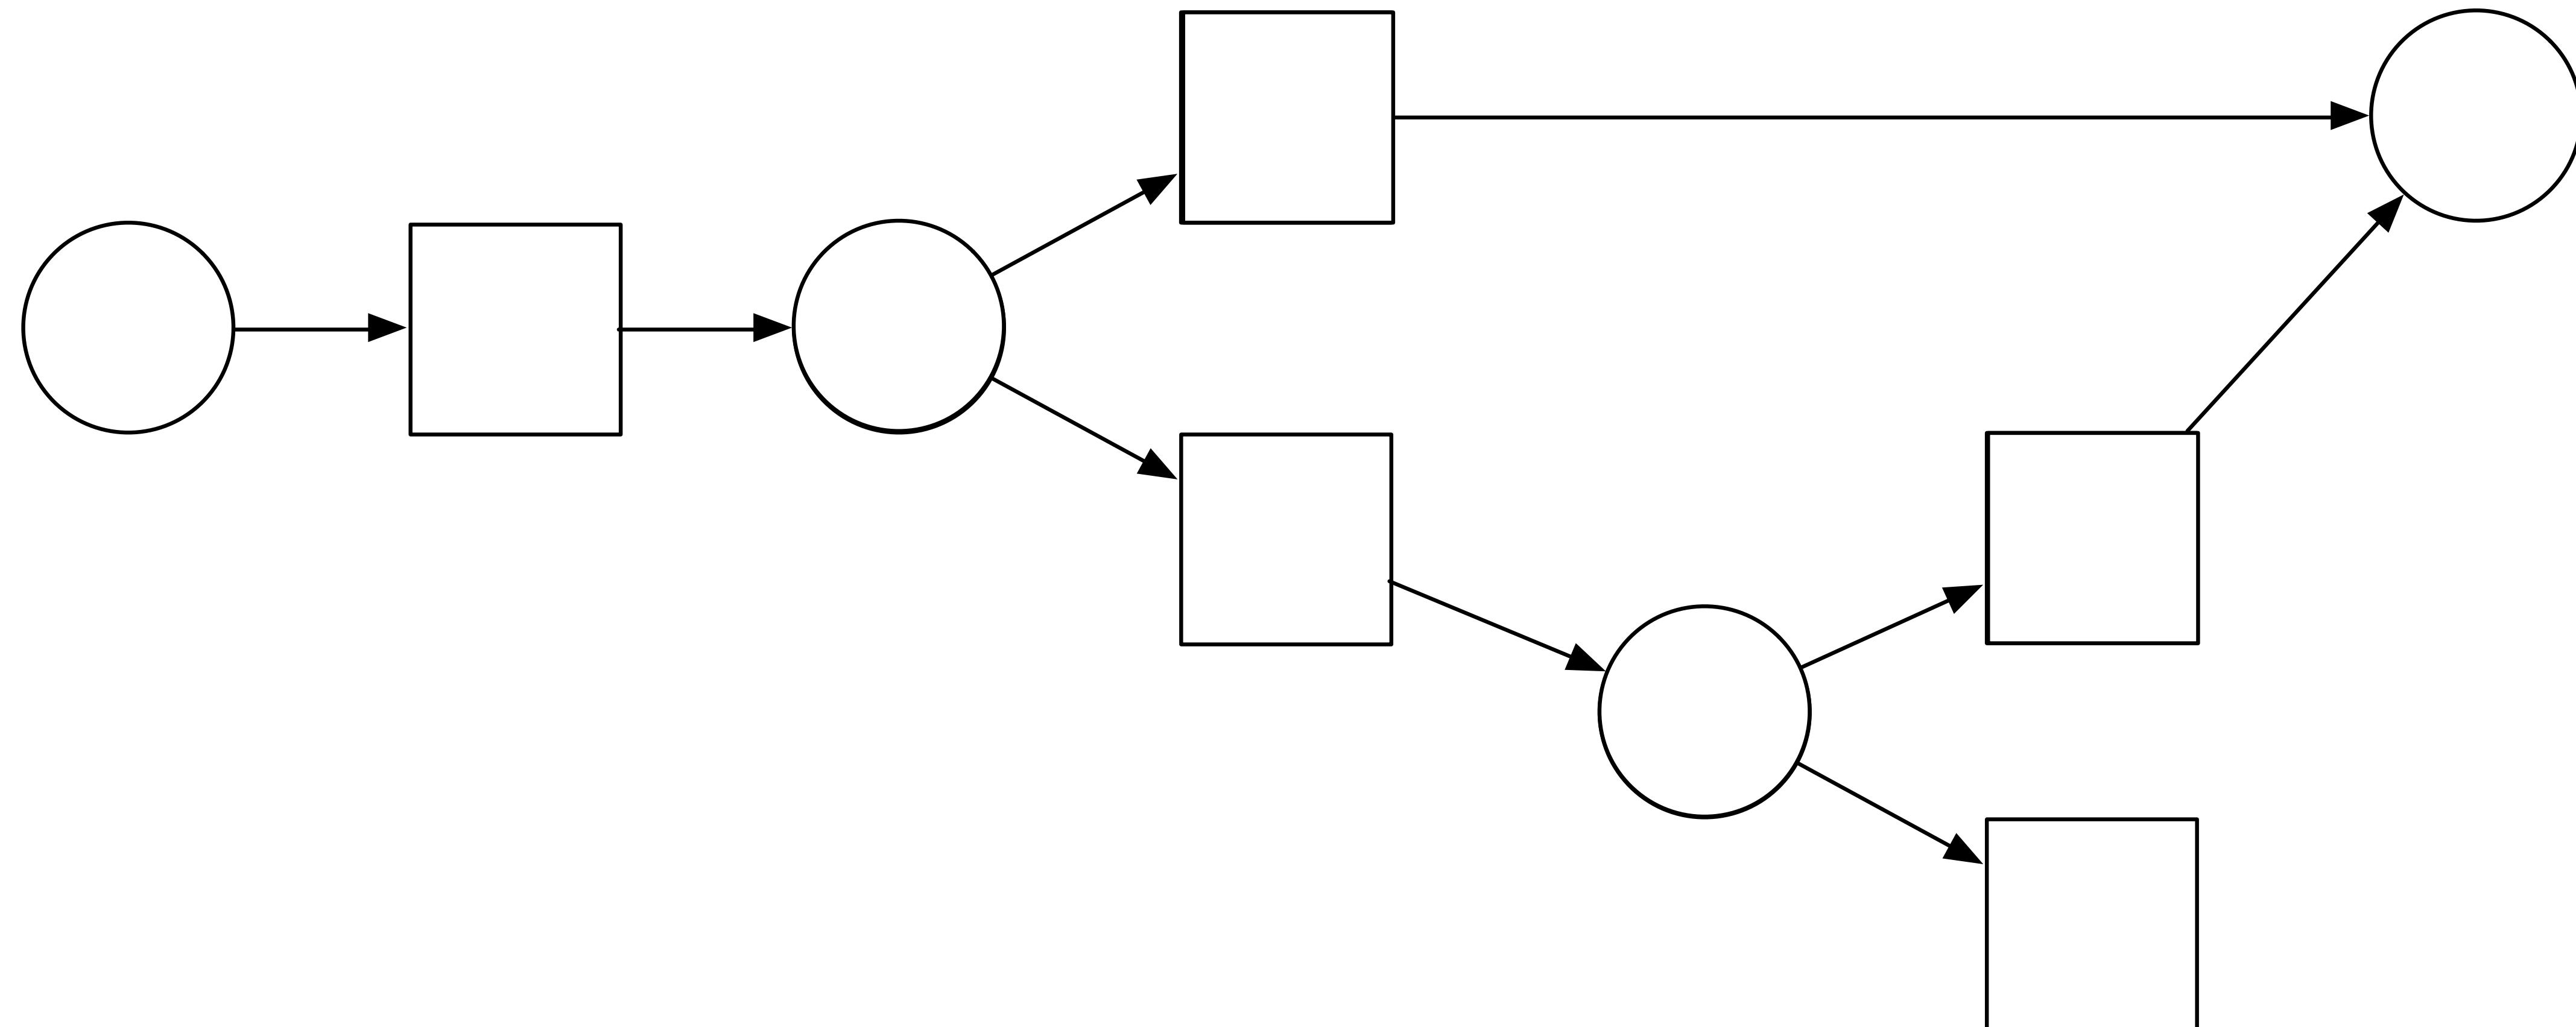
Cancelation

- When the (last) future destructs
 - The associated task that has not started, should not execute (NOP)
 - The resource held by that task should be released
 - Since that task may hold futures for other tasks, the system unravels
- I do not know of a good way to compose such cancelation with current futures
 - Except to create something more complex than re-implementing futures

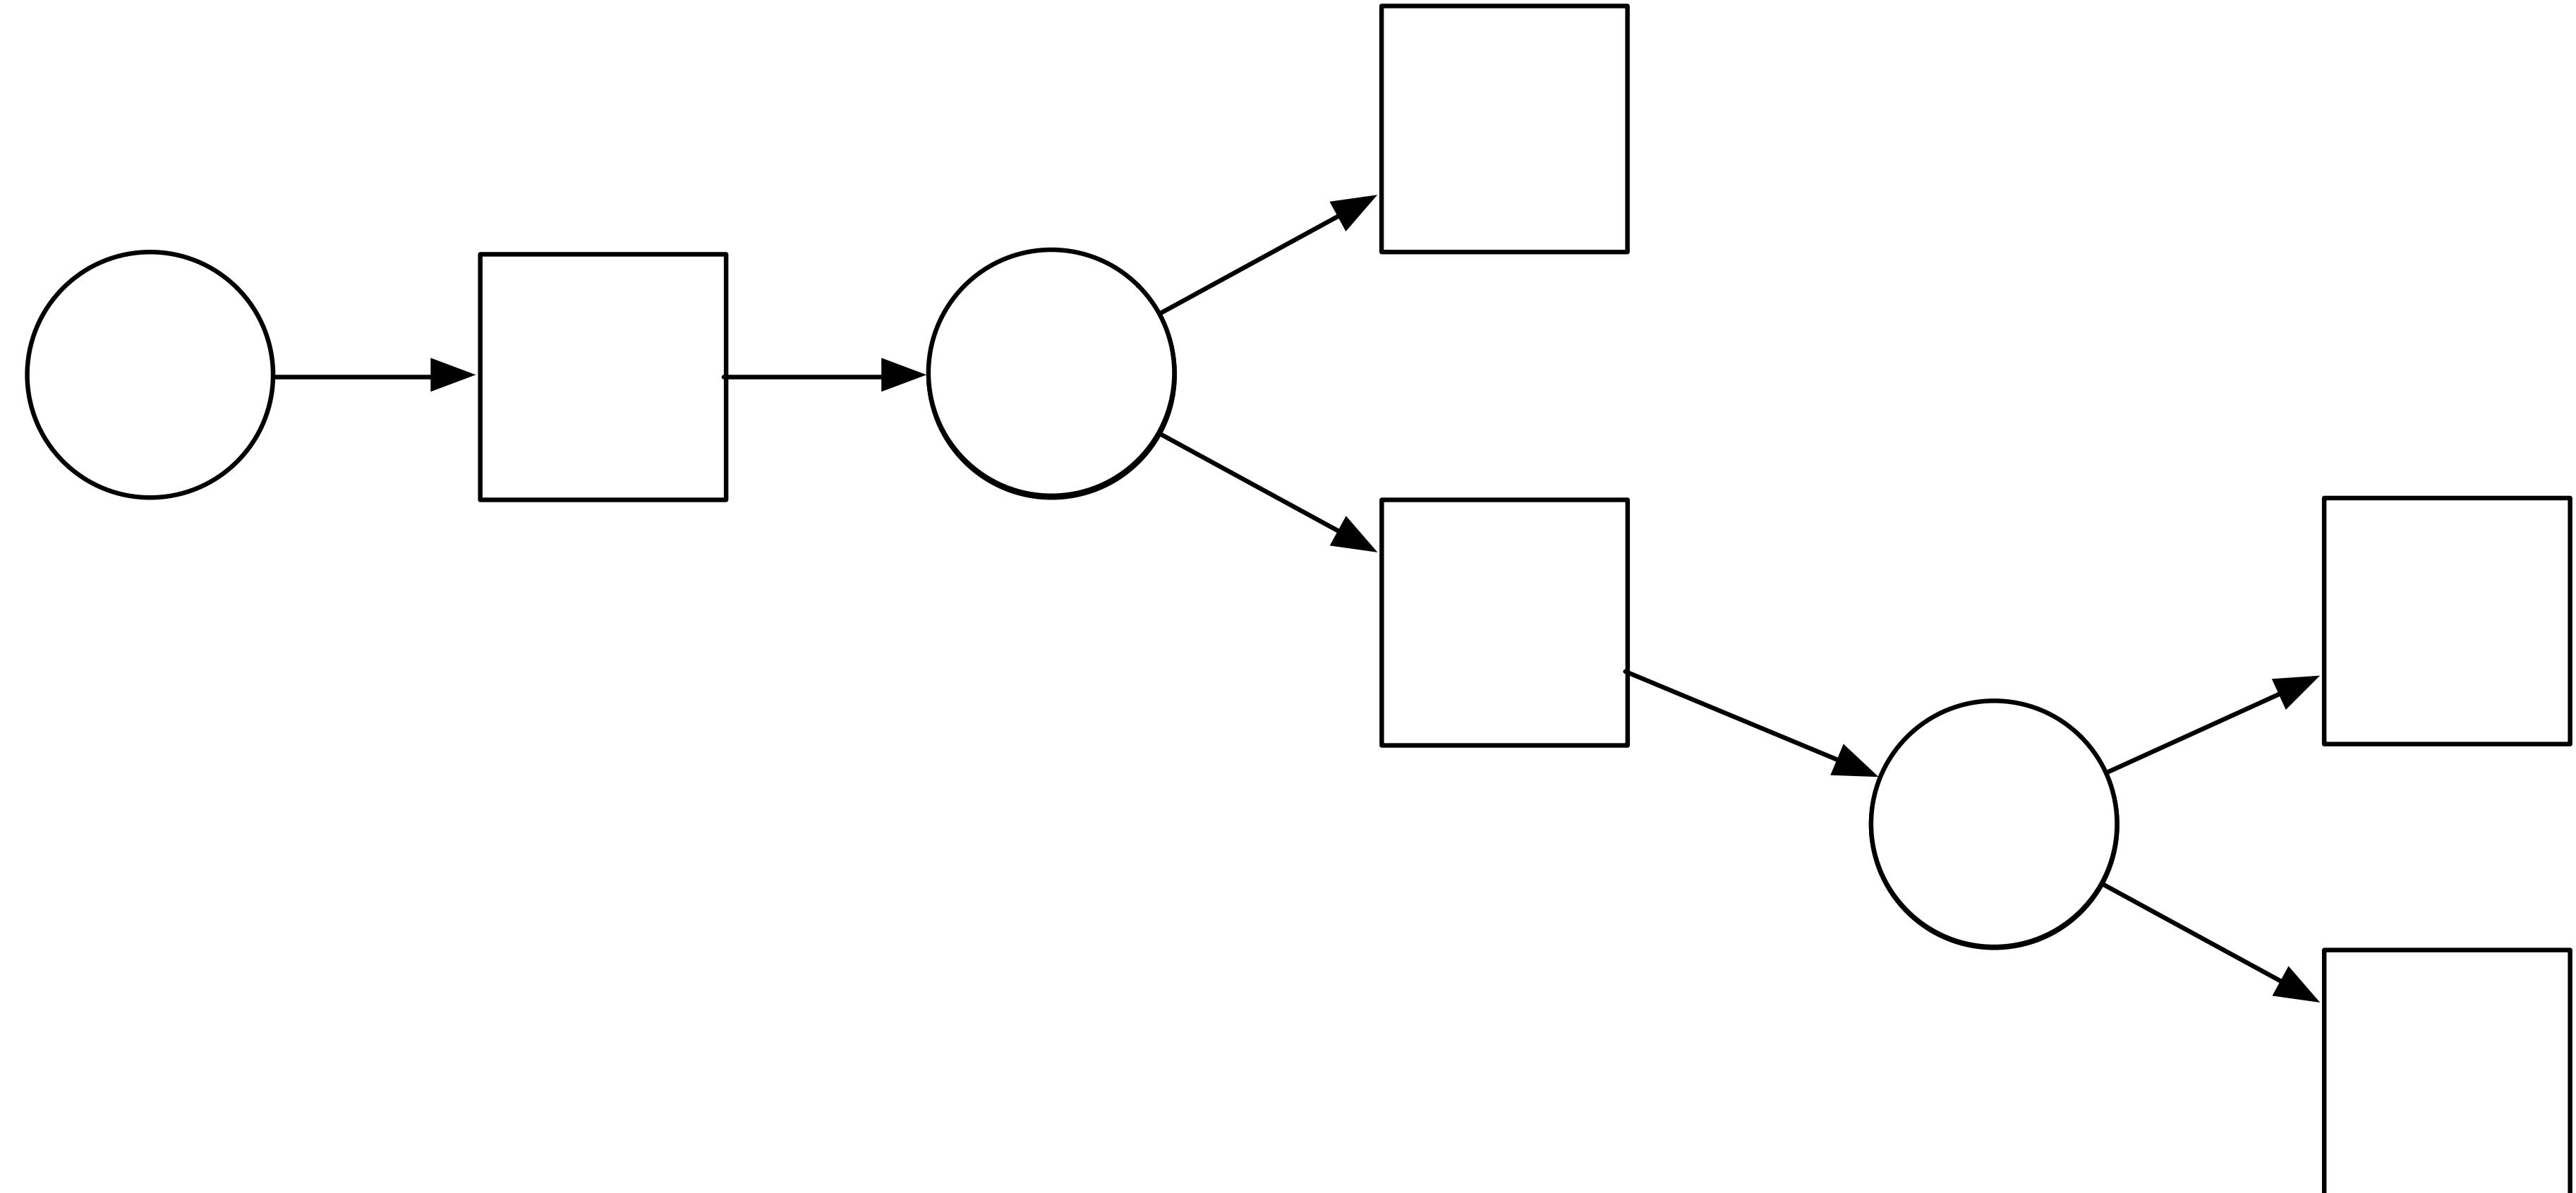
Cancelation



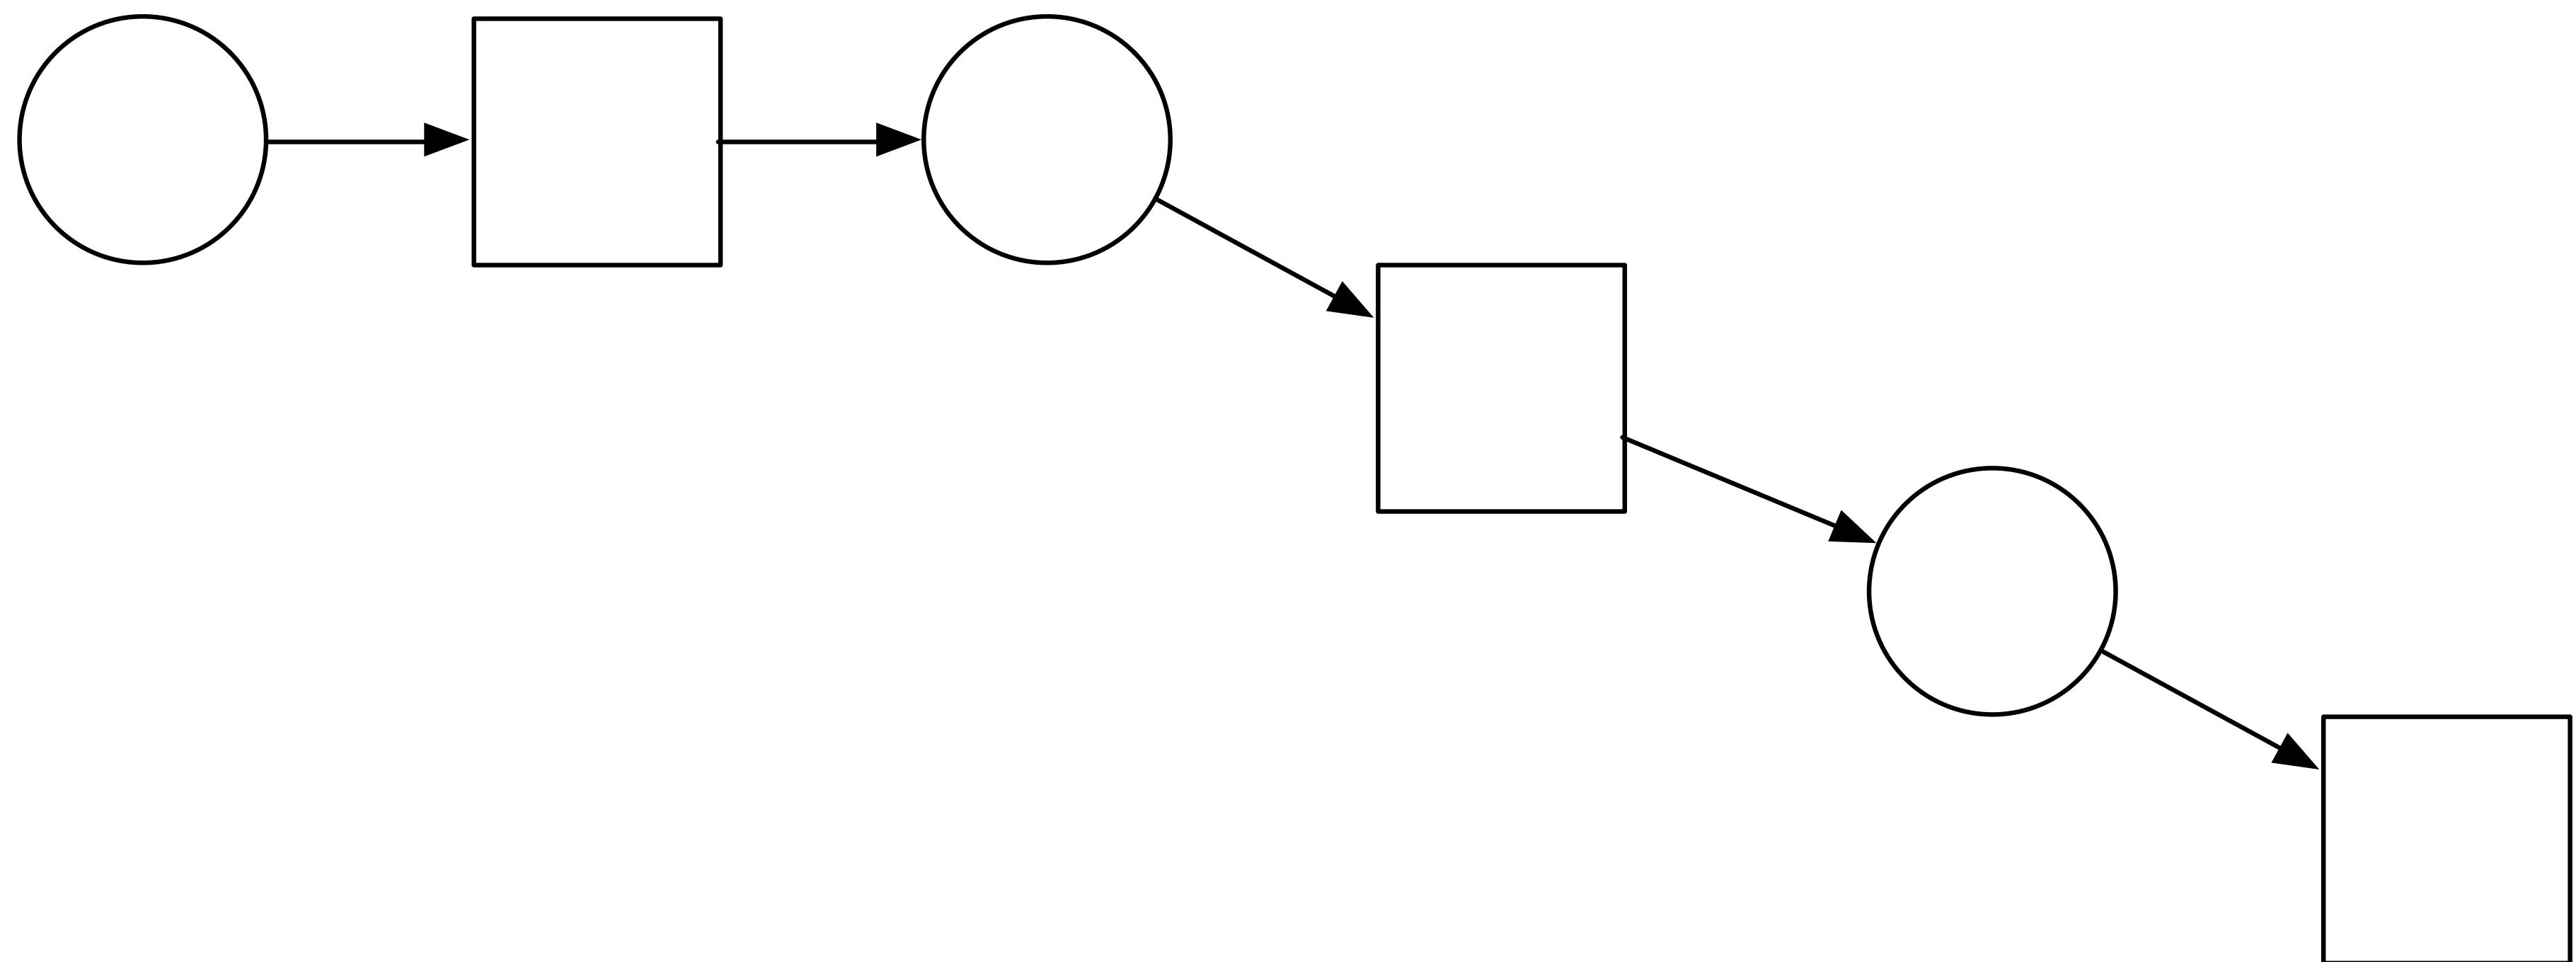
Cancelation



Cancelation

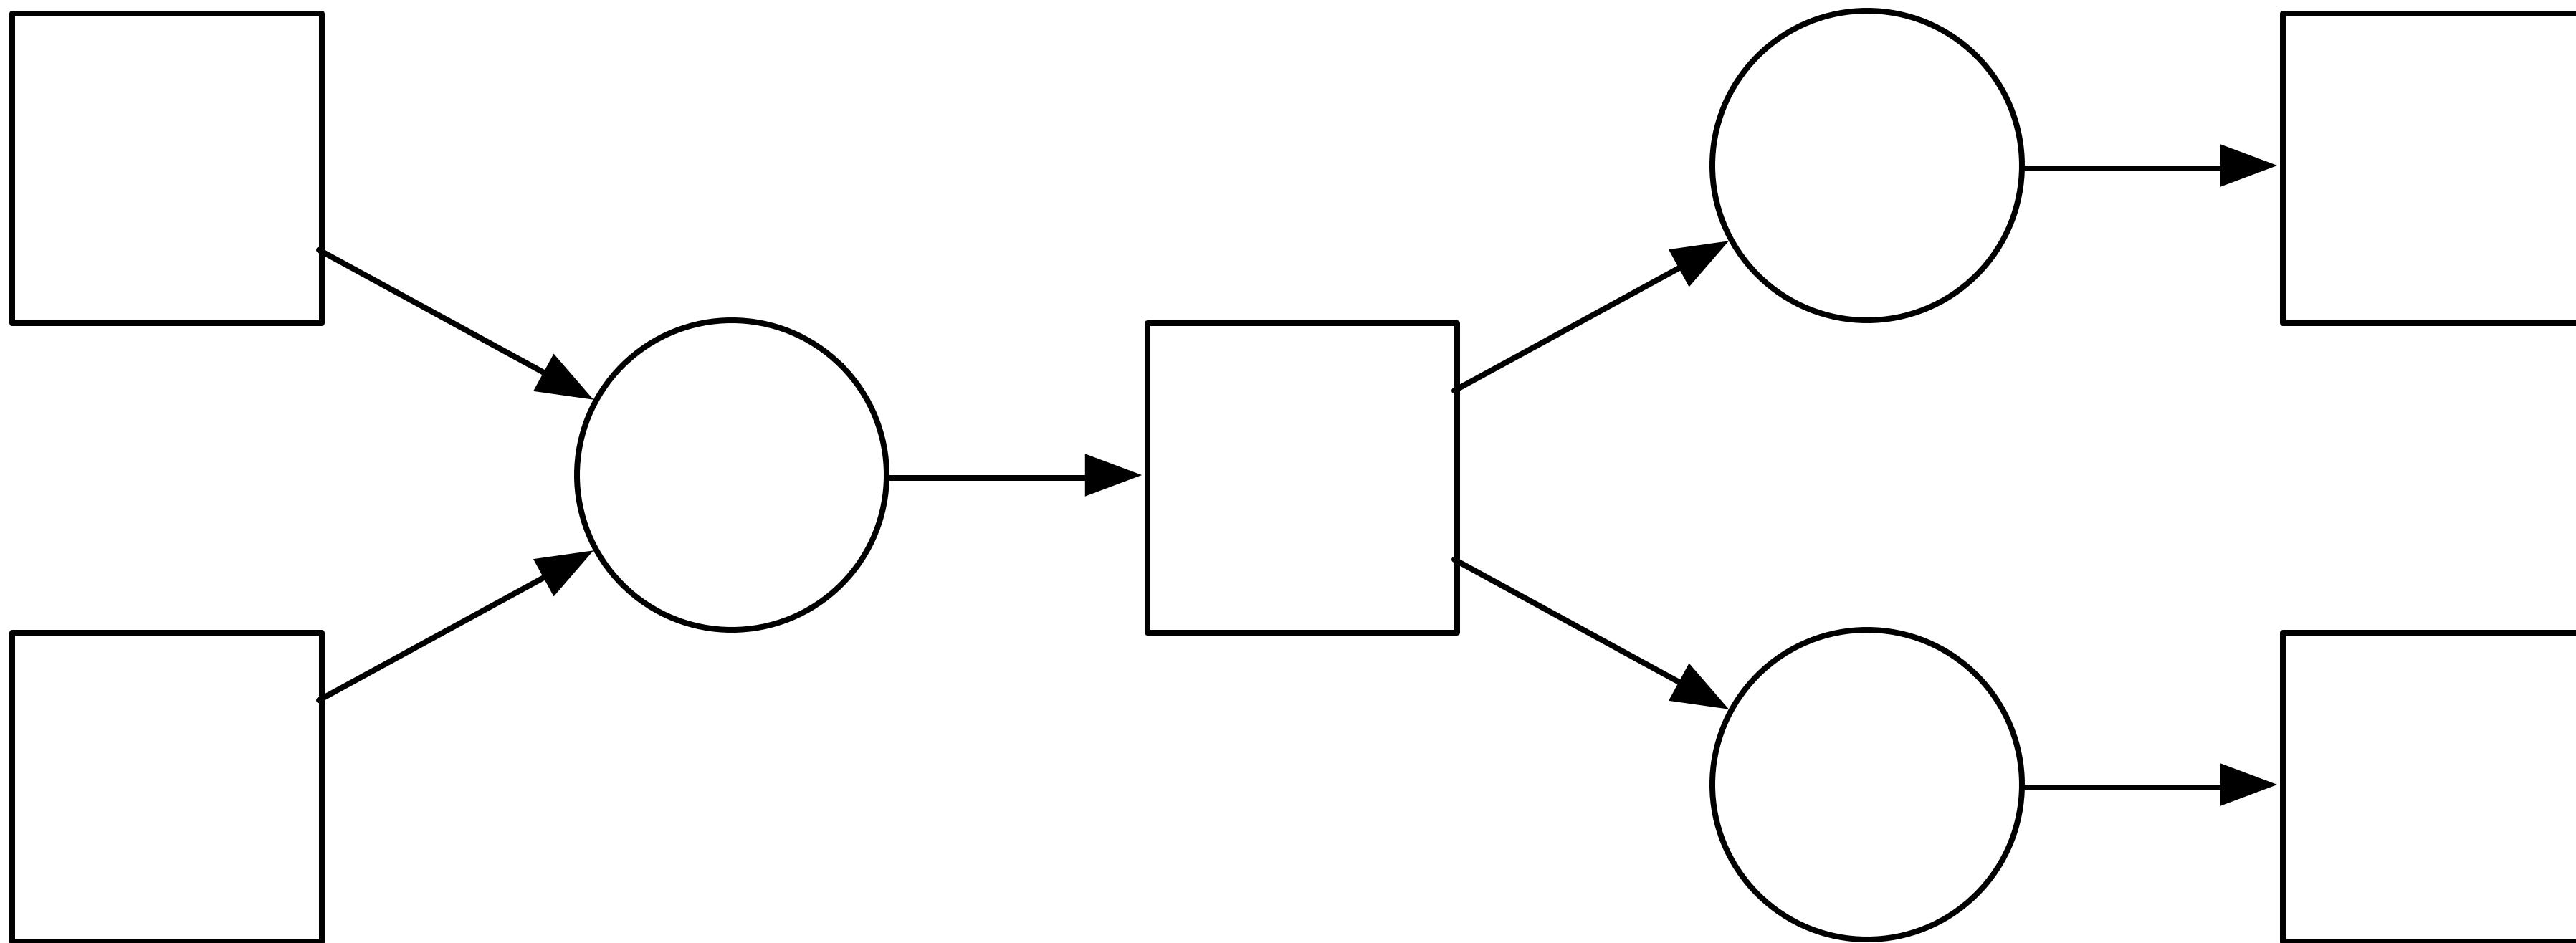


Cancelation



Channels

What if we persist the graph?



What if we persist the graph?

- Allow multiple invocations of the tasks by setting the source values
- Each change triggers a notification to the sink values
 - This is a reactive programming model and futures are known as *behaviors* or *channels*

Accumulators and Generator

- Each operation does not have to be a 1:1 mapping of input to output
- Coroutines are one way to write n:m functions

Futures: Continuations

```
channel<int> send;

auto hold = send
| [](const receiver<int>& r) {
    int sum = 0;
    while(auto v = co_await r) {
        sum += v.get();
    }
    return sum;
}
| [](int x){ cout << x << '\n'; };

send(1);
send(2);
send(3);
send.close();
```

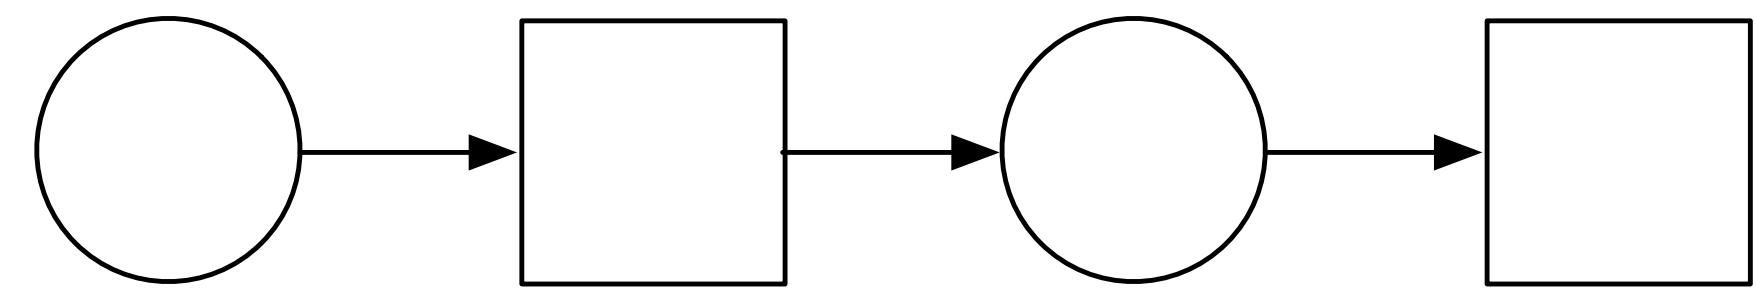
Futures: Continuations

```
channel<int> send;

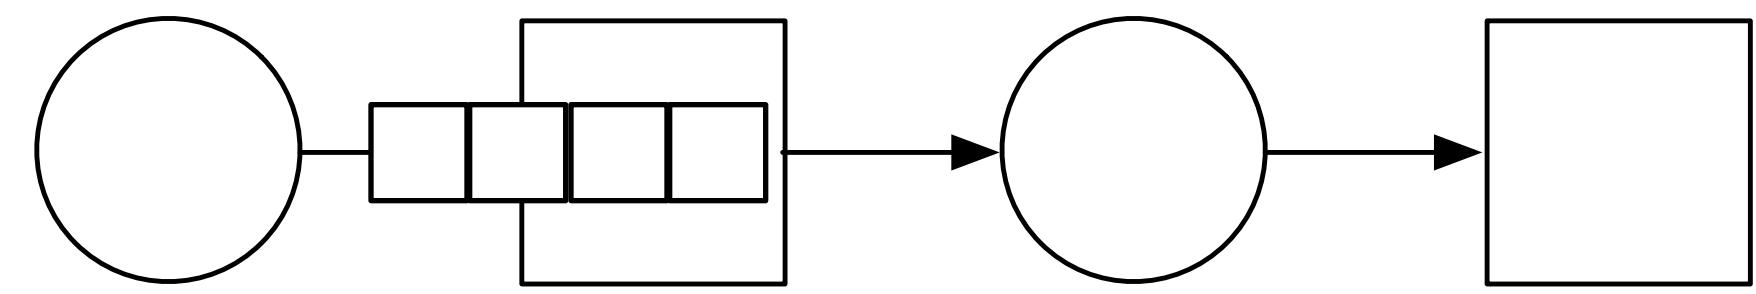
auto hold = send
| [](const receiver<int>& r) {
    int sum = 0;
    while(auto v = co_await r) {
        sum += v.get();
    }
    return sum;
}
| [](int x){ cout << x << '\n'; };

send(1);
send(2);
send(3);
send.close();
```

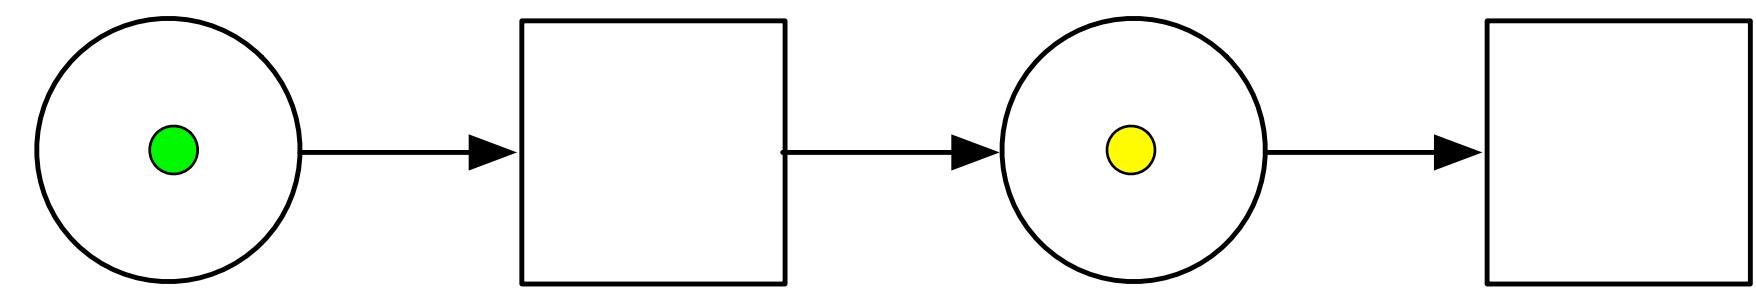
Flow Control



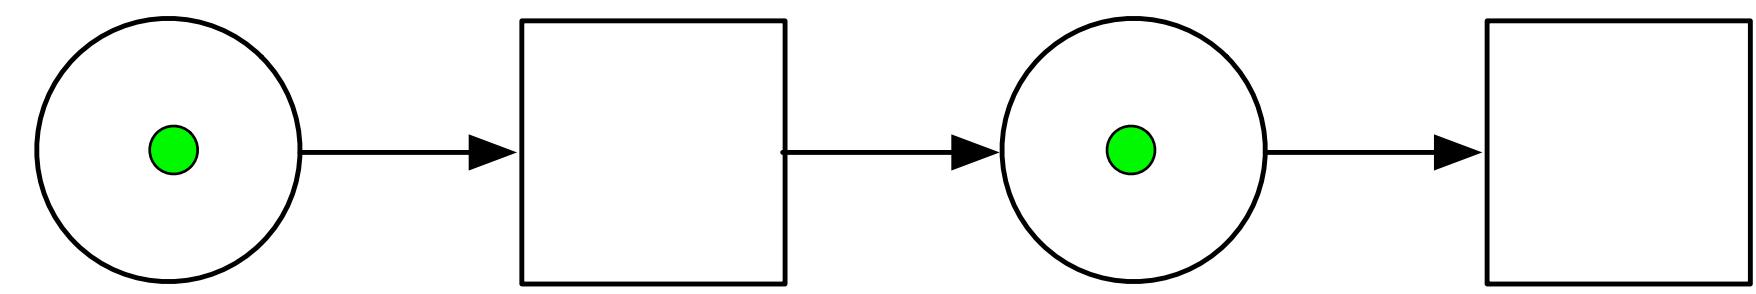
Flow Control



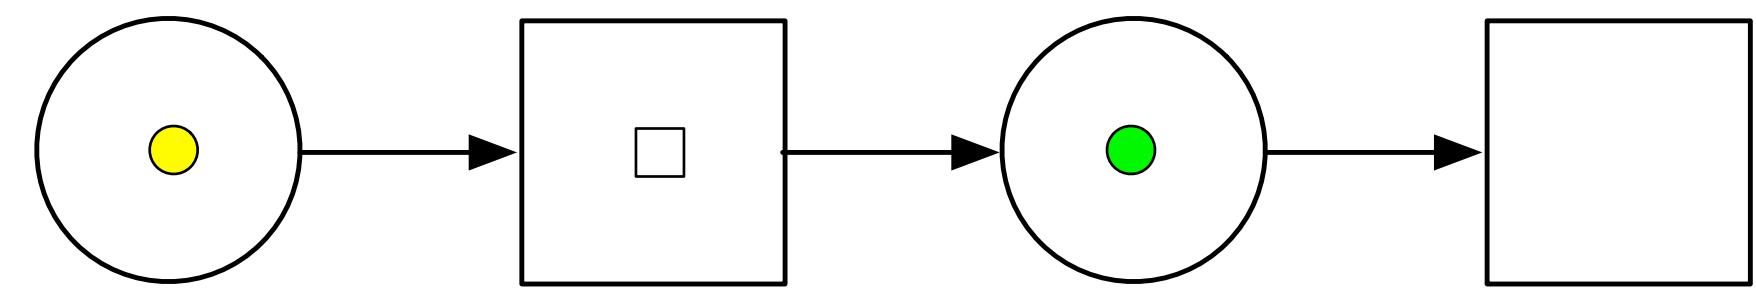
Flow Control



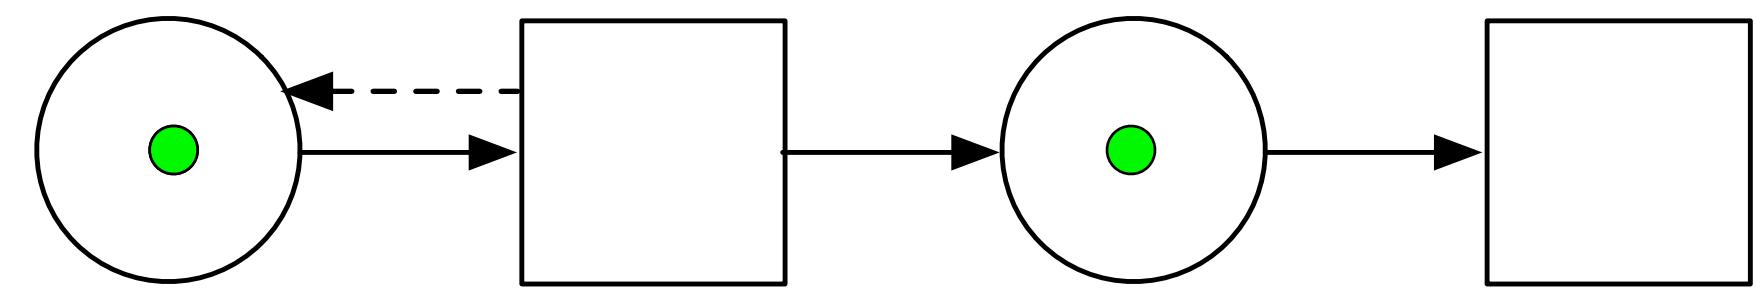
Flow Control



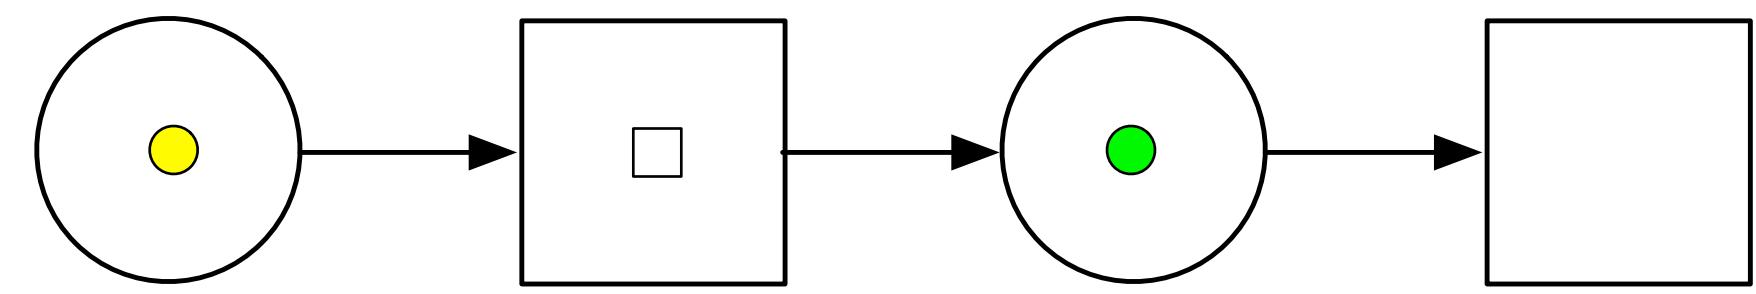
Flow Control



Flow Control



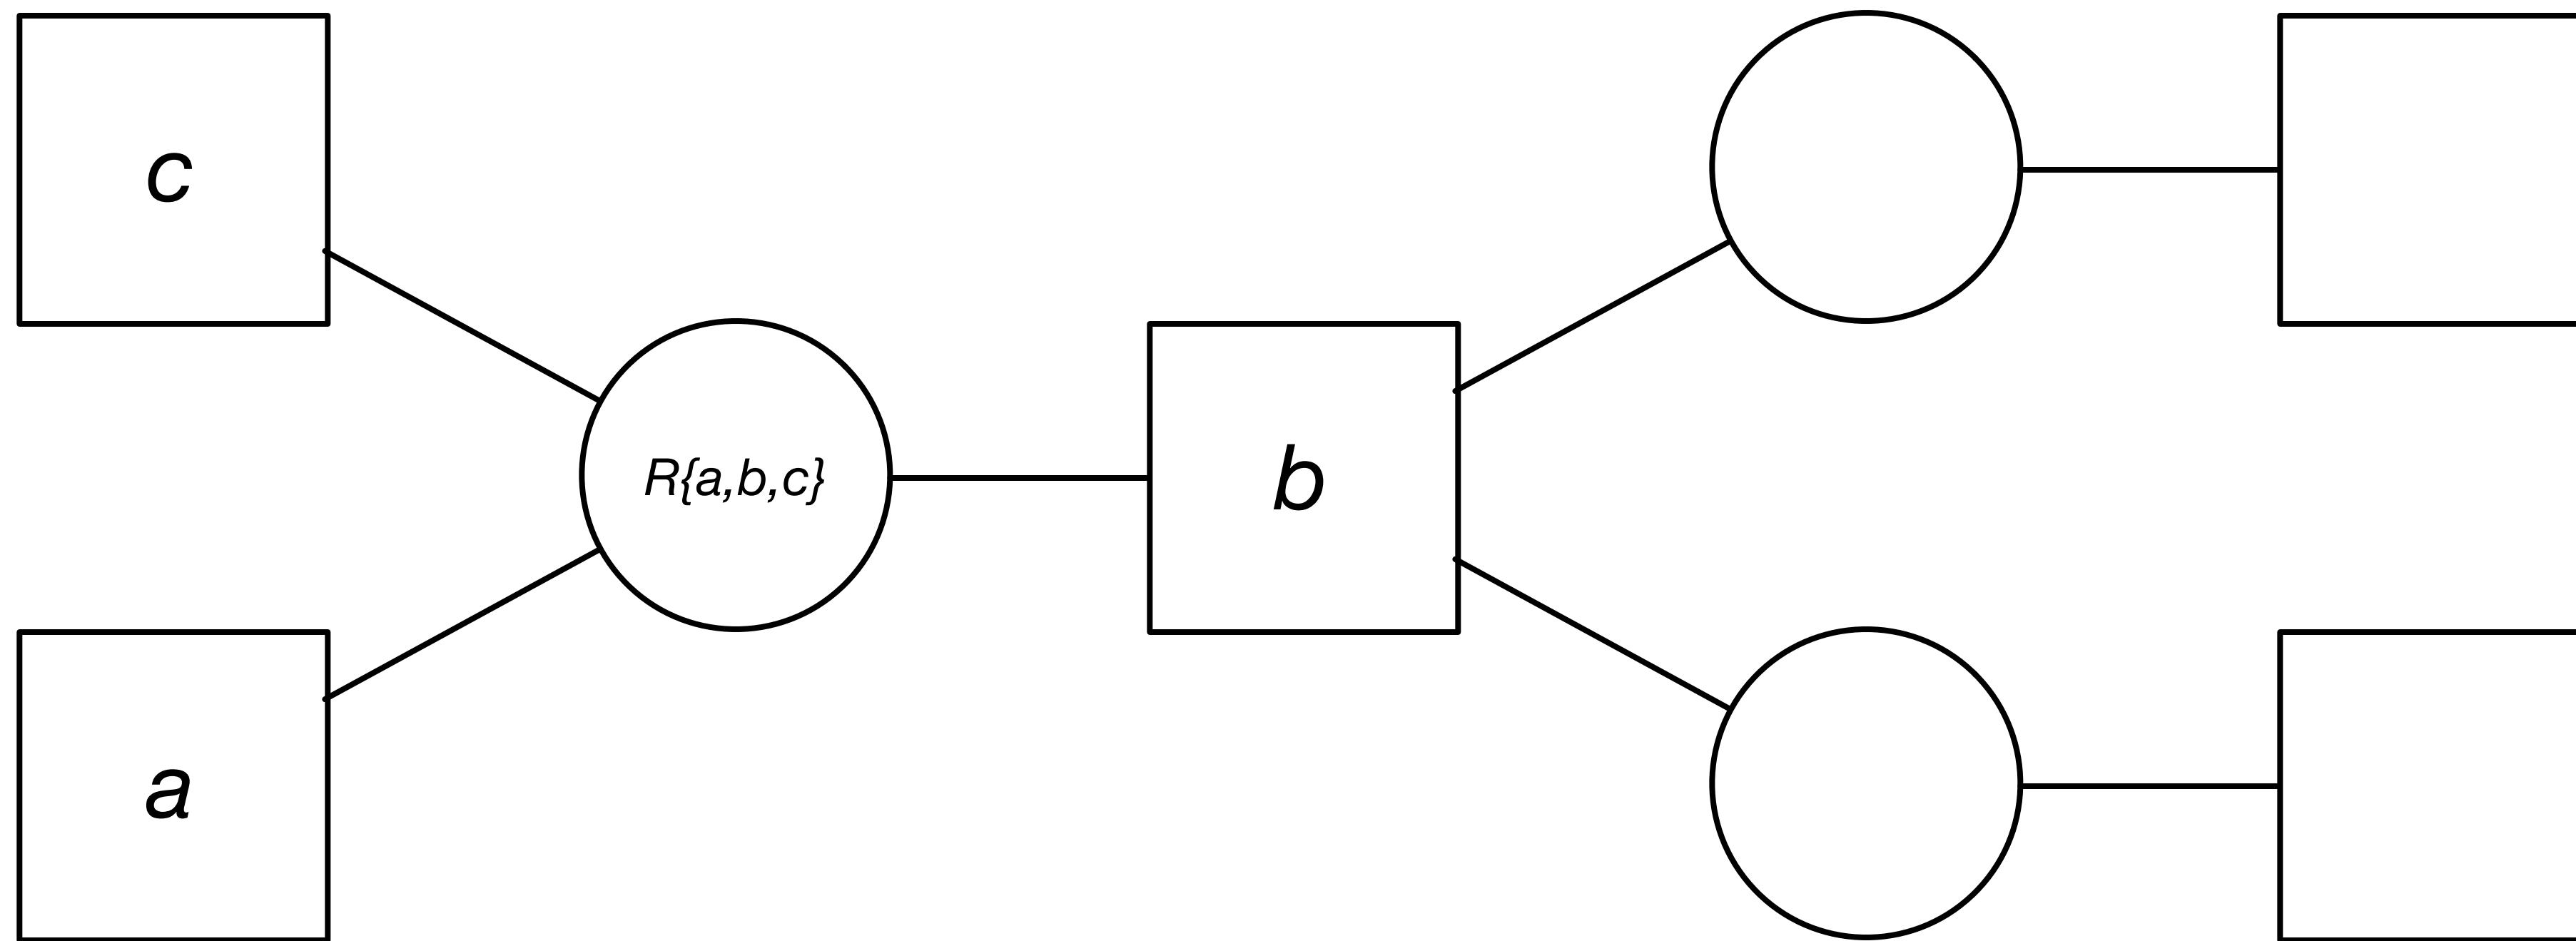
Flow Control



Property Models

How do the graphs change during execution?

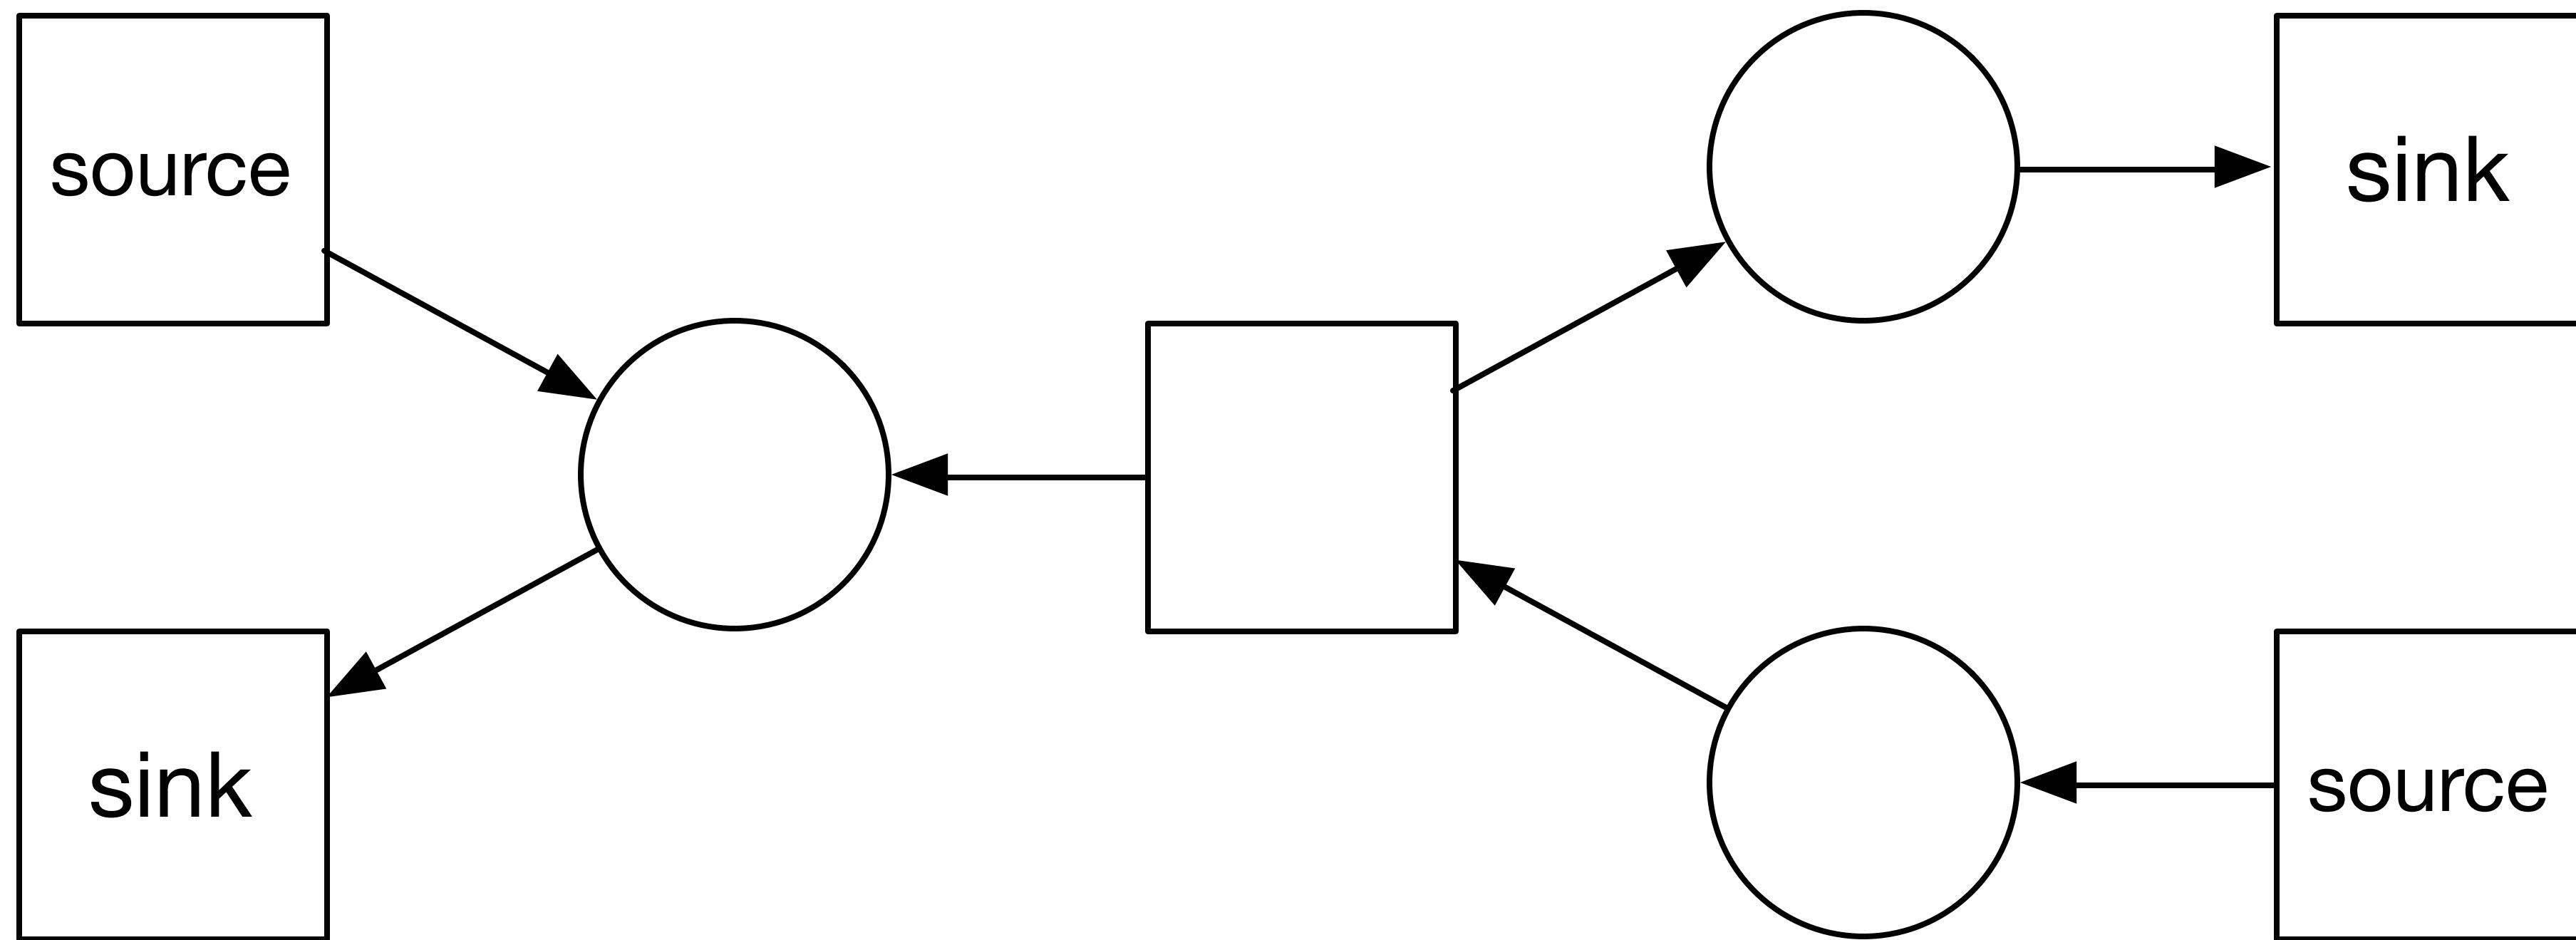
Property Model



A function is a directed relationship

- We can remove the arrows by providing a package of functions to represent the relationship
 - $a = b * c$
 - $b = a / c$
 - $c = a / b$
- This forms a type of constraint system called a *property model*
- Flow is determined by value, or *cell*, priority
- Relationships can be conditional, so long as predicate can be determined regardless of flow
- Cells can only have one in-edge for a given flow or the system is over constrained

Property Model

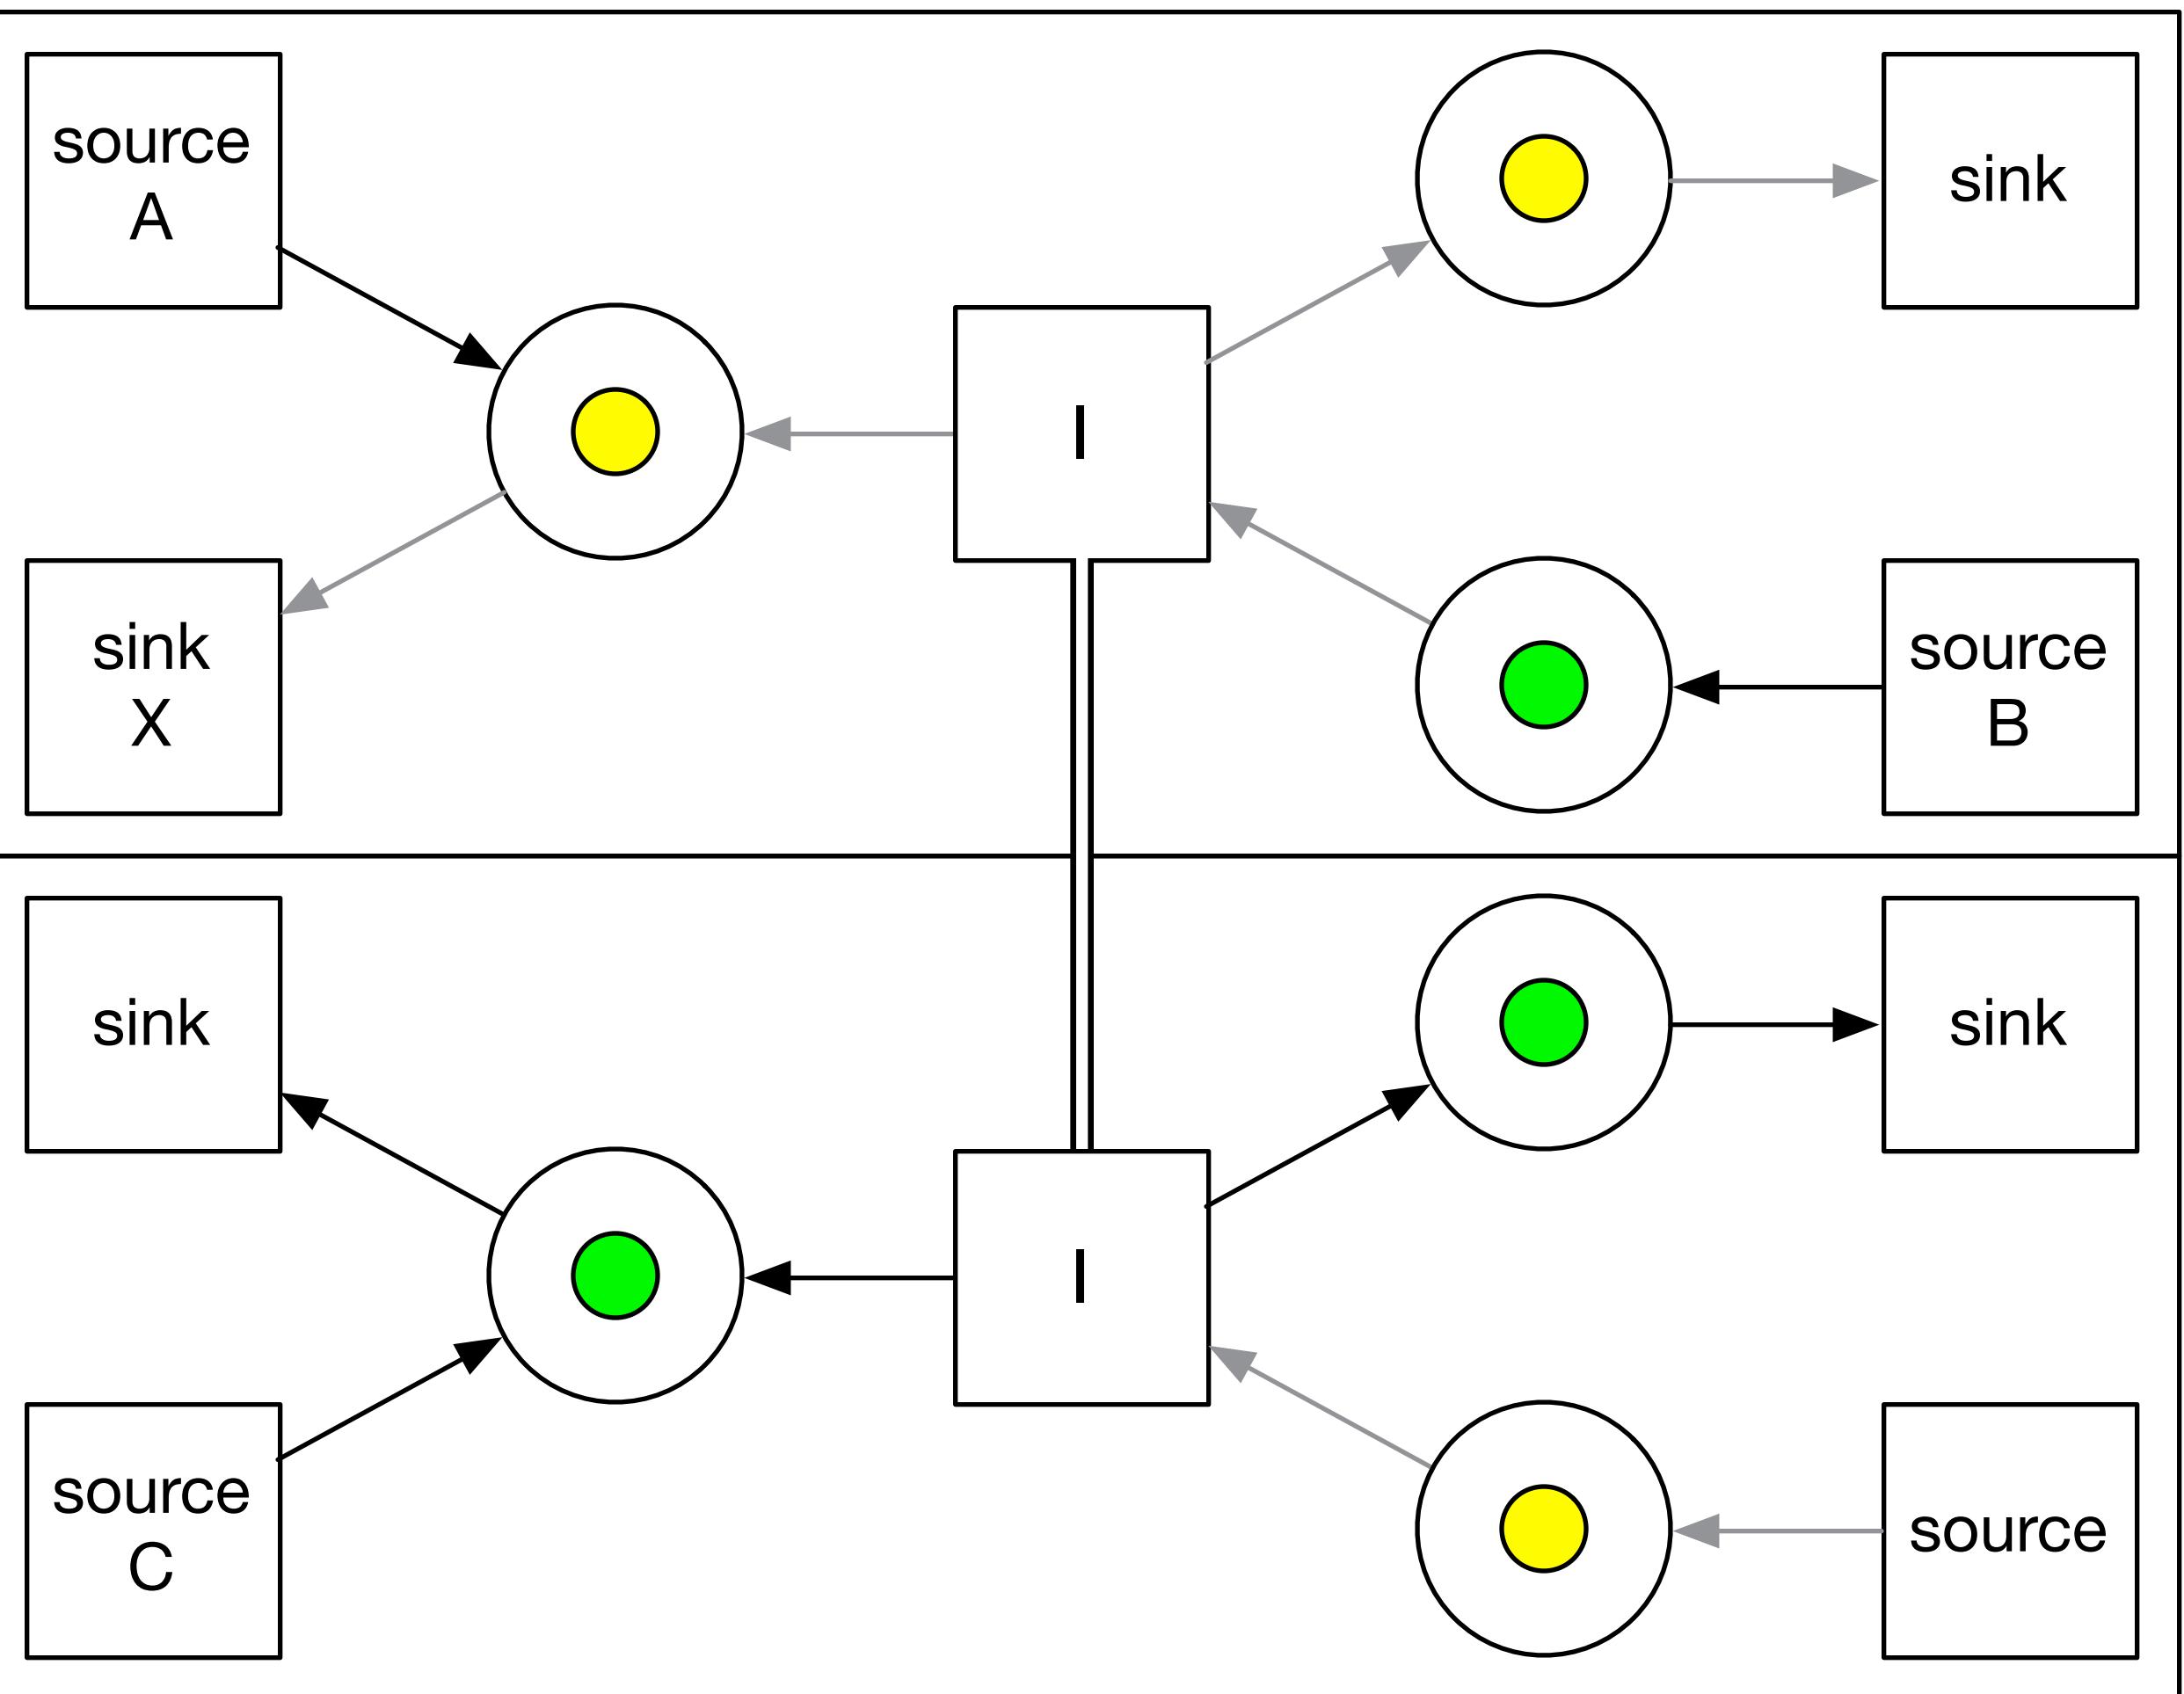


Property Models

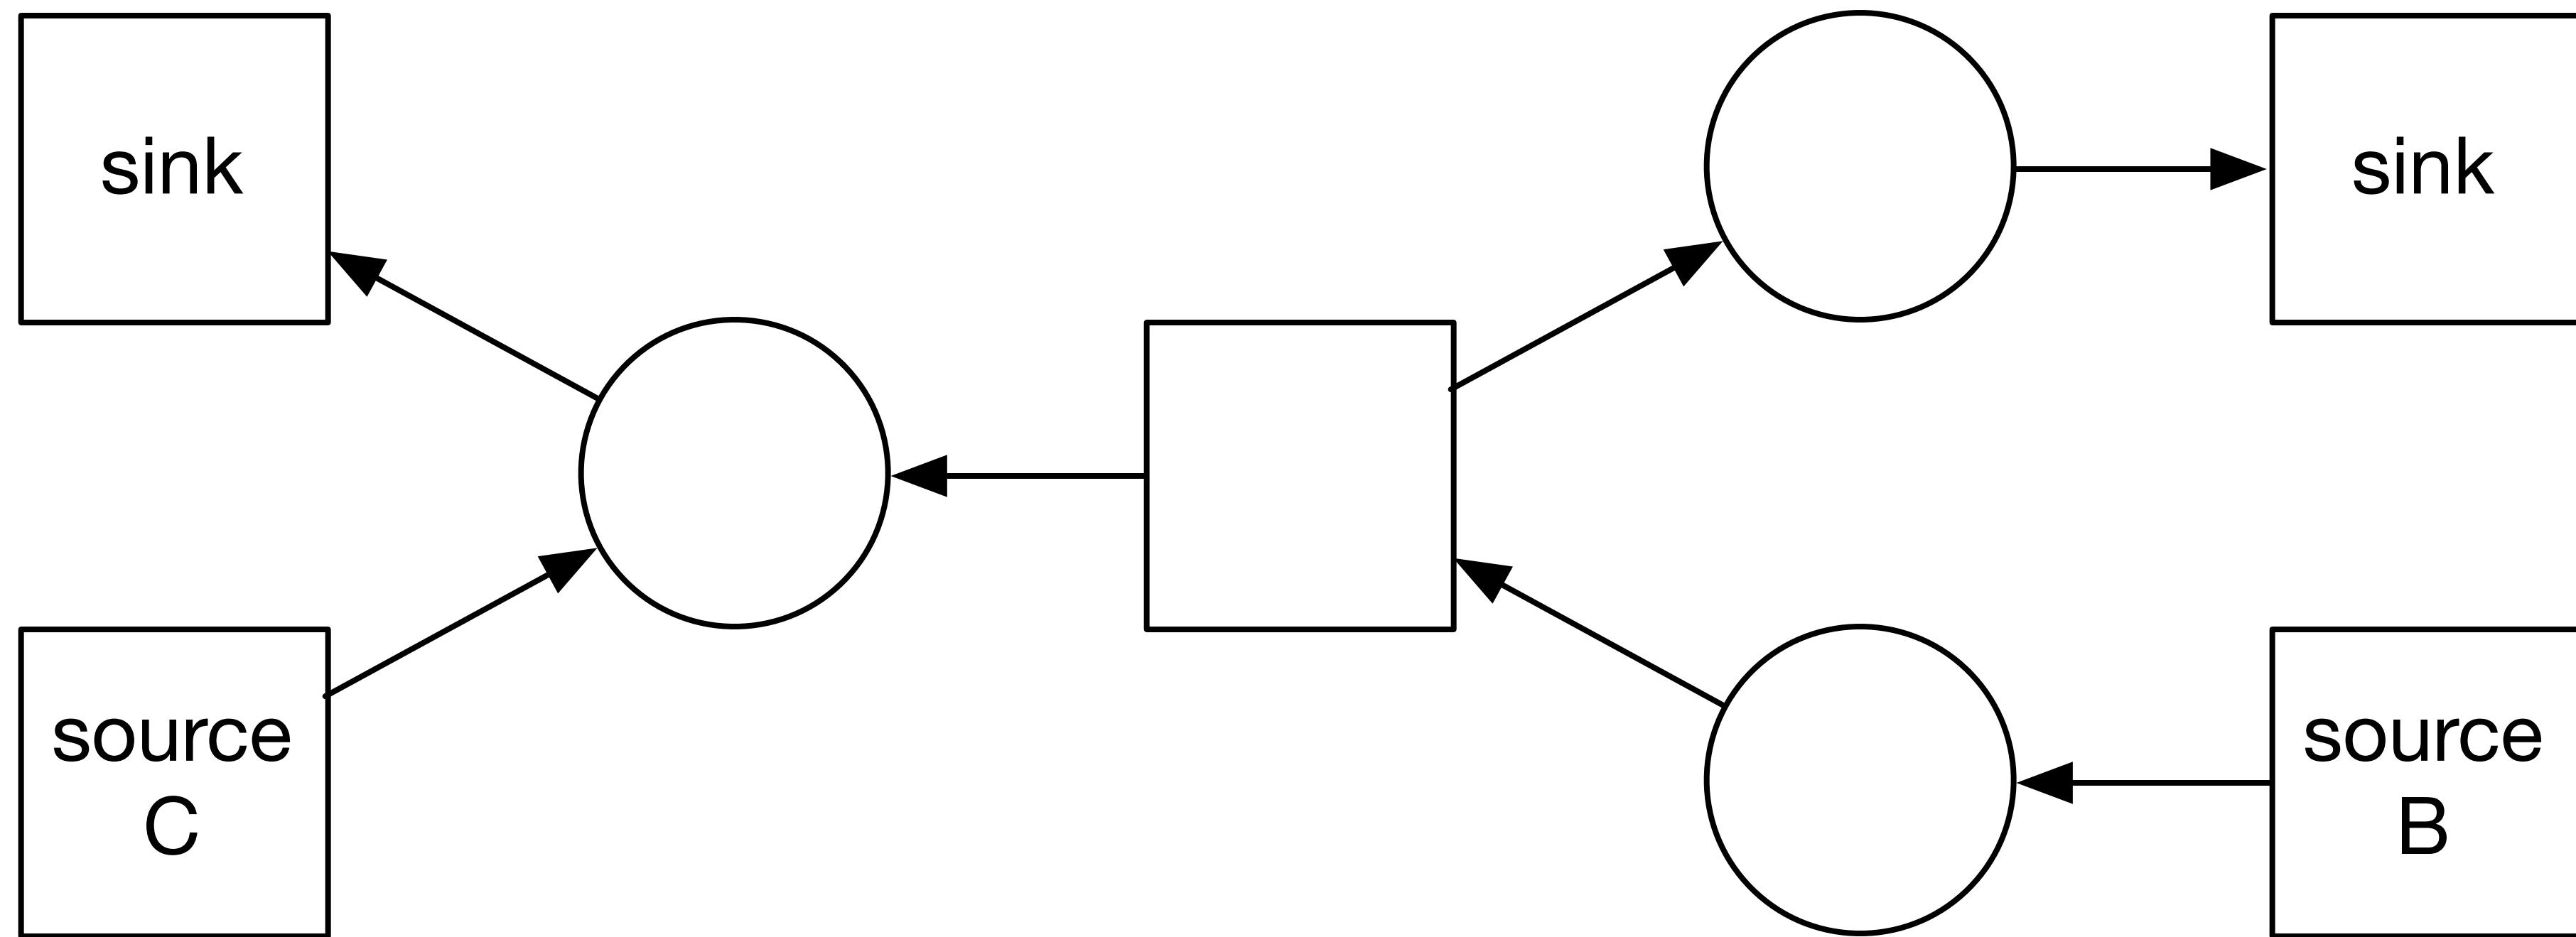
- Reflowing a property model doesn't require all relationships to be resolved
 - The task representing them can still be executing concurrently
- This creates a single dependency graph that is appended to for each new flow and is pruned and *unravels* as tasks are complete

Property Model

- Value set in source A then B
- Operation connected to B is high latency
- Value is set in source C causing reflow
- Sink X is no longer needed, pending operation is canceled
- Source A is discarded
- Intermediate value I is shared between flows, once determined
- Final values determined by source B (via I) and source C



Property Model



Property Models

- Very useful for UI behavior
- Significant information is in the graph
 - Source / Derived values form a partition set
 - Easily model checked
 - Equal result regardless of source order
 - Form an operational transform, useful for collaborative editing
 - A value is *implied* by the current state it only has in-edge
 - Source values determine *intent*
 - Values disconnected from result (sink) are *don't care*

Demo

Final Thoughts

- Perhaps representing such systems *as if* it were imperative code is not the correct approach
- Instead the a graph description can be compiled and statically validated
- Slides and code from talk:
- <http://sean-parent.stlab.cc/papers-and-presentations>
- Experimental future library:
- <https://github.com/stlab/libraries/tree/develop>

No raw synchronization primitives

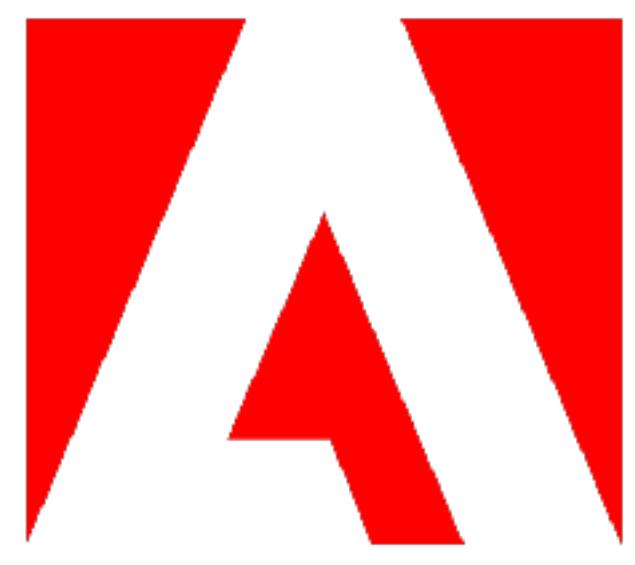
No raw synchronization primitives

Communicating Sequential Tasks

No raw synchronization primitives

Communicating Sequential Tasks

Better Code



Adobe